

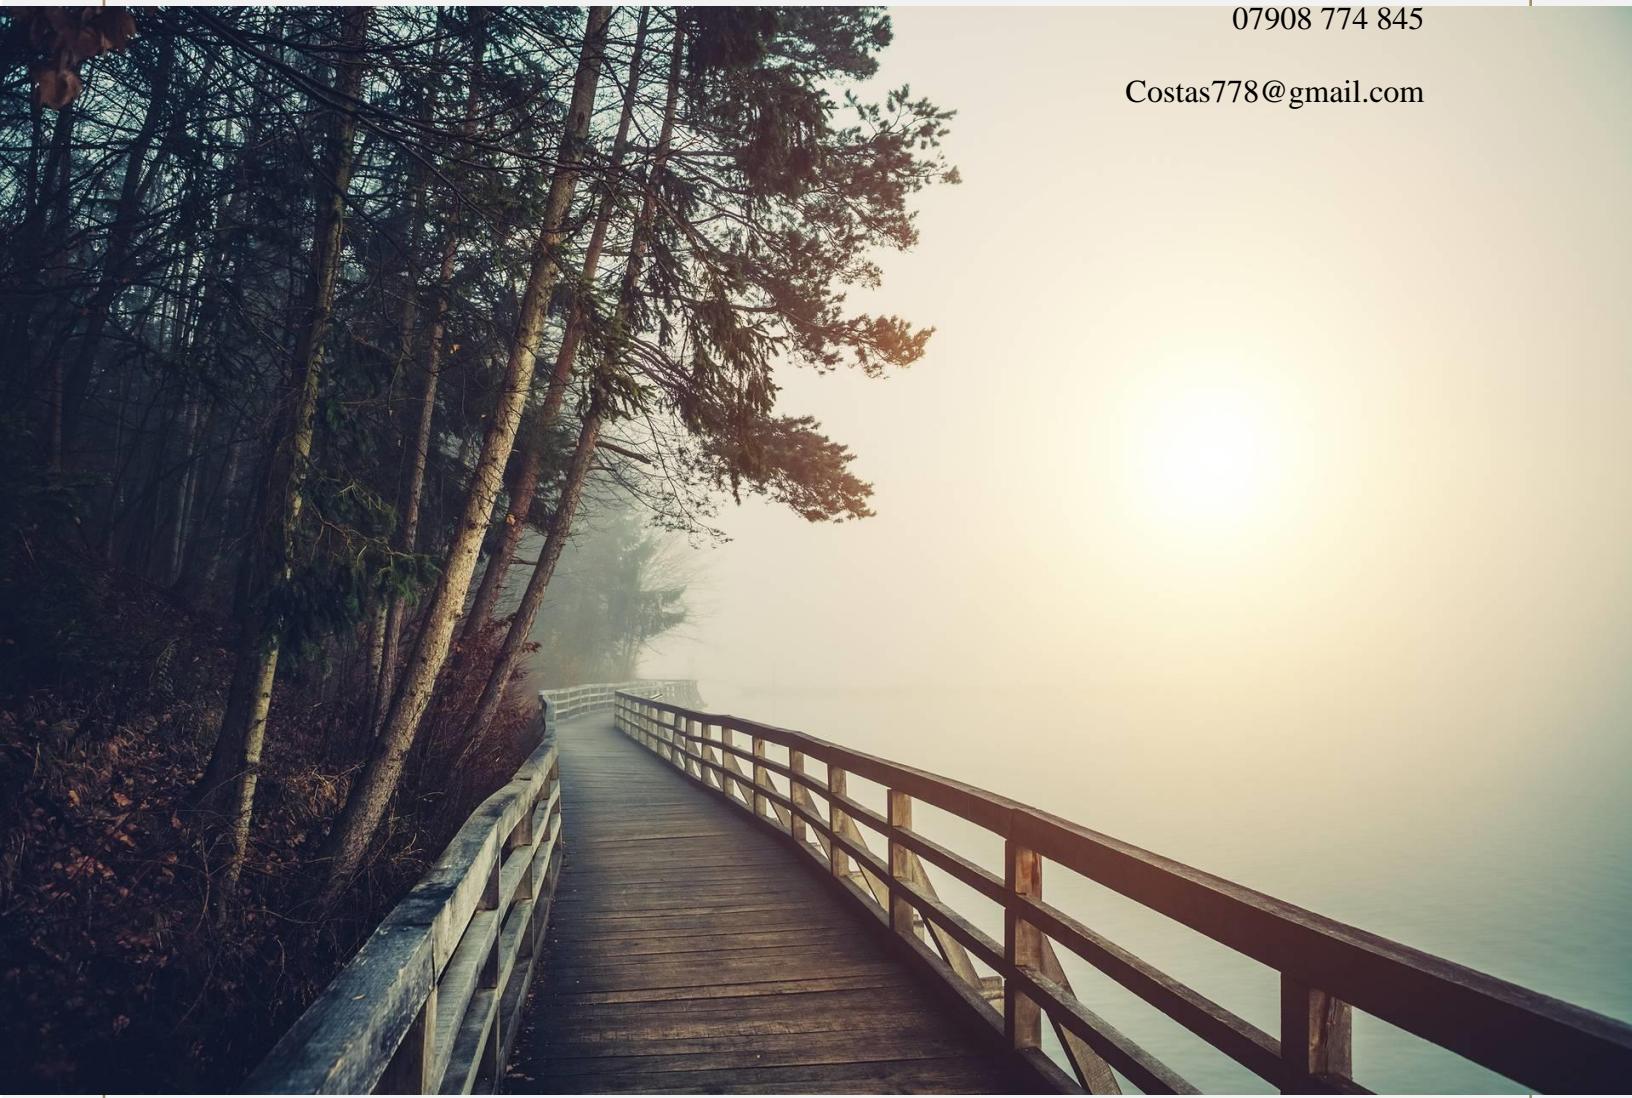
DEVOPS Projects

Costas Constantinou MSc

Version 2.0

07908 774 845

Costas778@gmail.com



DEVOPS PROJECTS 2.0

What is this document for?

Having spent over 10 years in the IT Industry, I have noticed that there are, from time to time, new “flavors of the month”. This leads to shortages of skilled professionals. DEVOPS is no exception.

What never changes are the dilemma of sifting through hundreds (and even thousands) of applications claiming they can add value to either your organisation or corporate entity.

HR professionals and recruiters must sift through many applications that have all the ‘right words’ on their CV or covering letter. However, there are more than a few applicants that slip through that turn out to be “economical with the truth”. In other words, they cannot support, either by knowledge or experience, what is written by them in their documents.

Now it is true that many of these people are screened and rejected by a technical test. However, even this is no guarantee that they, genuinely, know their subject matter considering all the resources on the internet to help you prep in various technical disciplines!

As a result, I decided to come up with this DEVOPS projects document that demonstrates, practically, what I can do with the various DEVOPS tools. Within this document you will find various projects I have done. It will be a growing list and one that will be updated as technology never sleeps.

If you are a HR or recruiter simply attach this to my application and pass it on to others that are best placed to make an assessment.

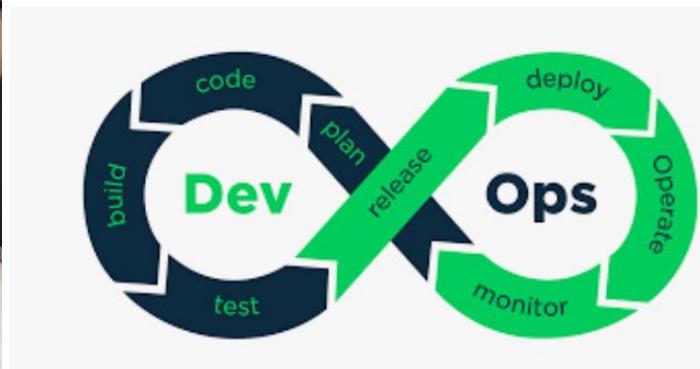


Table of Contents

[EKS Projects]

[Use Terraform to Create an EKS Deployment](#)

[Deploying a Kubernetes Cluster Using Elastic Kubernetes Service](#)

[Deploy Springboot Microservices App into Amazon EKS Cluster using Jenkins Pipeline and](#)

[Kubectl CLI Plug-in | Containerize Springboot App and Deploy into EKS Cluster using Jenkins](#)

[Pipeline](#)

[AKS Projects].....

[Create an AKS Cluster in Azure with Terraform](#)

[Deploying and Accessing an Application to an AKS Cluster](#)

[Jenkins Projects].....

[Deploying Jenkins Master and Worker Nodes in AWS Behind an ALB Using Terraform and Ansible](#)

[Deploying a Docker Container with Jenkins Pipelines](#)

[Deploy Python App into Kubernetes Cluster using Jenkins Pipeline | Containerize Python App and Deploy into AKS Cluster](#)

[Azure Pipeline].....

[Create and Test an ASP.NET Core App in Azure Pipelines](#)

[Deploy a Python App to an AKS Cluster Using Azure Pipelines](#)

[Docker].....

Load Balancing Containers (Docker Swarm and Docker Compose)

[Prometheus].....

Setting Up Prometheus and Adding Endpoints

Monitoring in Kubernetes with Prometheus and Grafana

How to setup monitoring on an EKS Kubernetes Cluster using both Prometheus and Grafana

[Ansible].....

Getting Started with Ansible

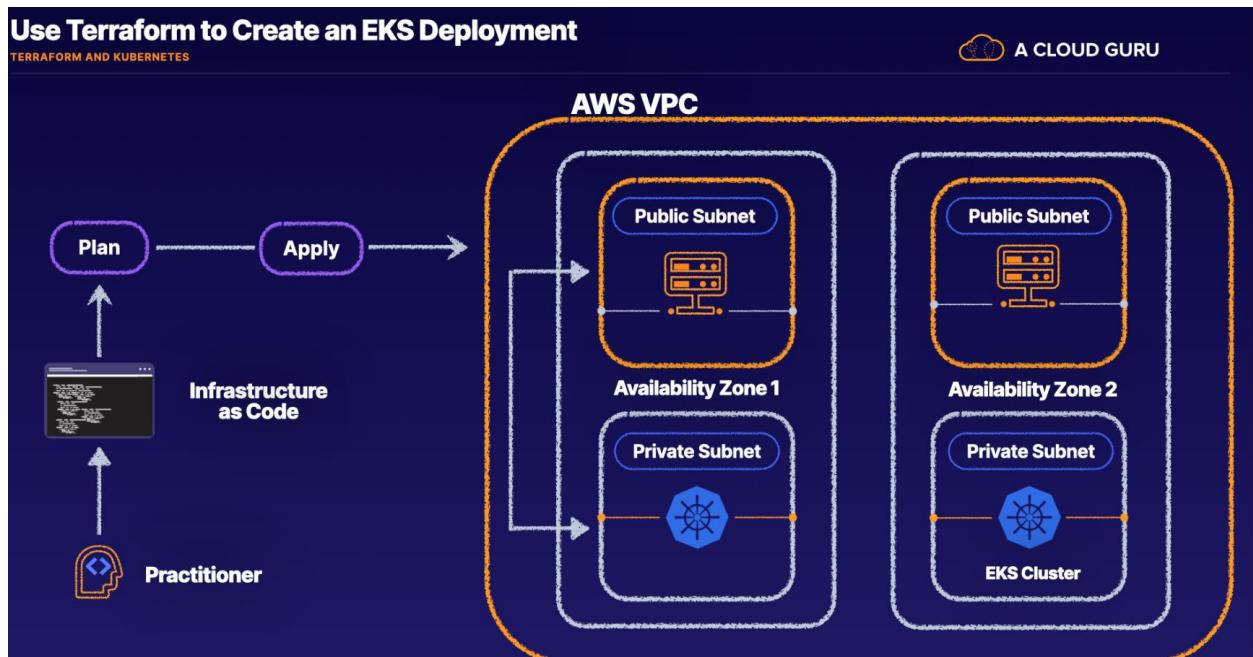
Ansible Playbooks: The Basics

Writing Your First Ansible Playbook

Deploying Services Using Ansible

Ad-Hoc Ansible Commands

[EKS Projects] — Use Terraform to Create an EKS Deployment



Solution:

Log in to the server using the credentials provided:

```
Host '3.80.152.91' is not in the trusted hosts file.
(sshd-ed25519 fingerprint shall!! 58:fd:99:ca:9a:f3:7a:6e:2a:32:9b:2c:37:0e:35:8c:63:0c:e8:f2)
Do you want to continue connecting? (y/n) y
cloud_user@3.80.152.91's password:
Last login: Fri Sep 22 09:45:08 2023
```

```
      _\   _\_)  Amazon Linux 2 AMI
     _\|_|_|
https://aws.amazon.com/amazon-linux-2/
```

ssh cloud_user@<PUBLIC_IP_ADDRESS>

Configure your AWS CLI:

aws configure

When prompted for your AWS Access Key ID, copy, and paste in the **Access Key**.

When prompted for your AWS Secret Access Key, copy, and paste in the **Secret Access Key**.

Press **Enter** to accept the **default** region.

Press **Enter** to accept the **default** output.

NOTE: if you don't have the above details, go to **IAM**, and generate the above.

```
[cloud_user@ip-10-0-1-201 ~]$ aws configure
AWS Access Key ID [None]: AKIAZXMQHZEGXNLZ67HY
AWS Secret Access Key [None]: 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
Default region name [us-east-1]:
Default output format [None]:
[cloud_user@ip-10-0-1-201 ~]$ █
```

Review the Contents of the Configuration Files in the lab-terraform-eks Directory:

View the contents of the directory:

Ls

Go into the lab-terraform-eks-2 directory:

cd lab-terraform-eks-2/

View the contents of the lab-terraform-eks-2 directory:

ls

```
[cloud_user@ip-10-0-1-201 ~]$ ls
lab-terraform-eks-2  lab-terraform-eks-2.zip  resource_ids.txt
[cloud_user@ip-10-0-1-201 ~]$ cd lab-terraform-eks-2/
[cloud_user@ip-10-0-1-201 lab-terraform-eks-2]$ ls
main.tf  outputs.tf  terraform.tf  variables.tf
[cloud_user@ip-10-0-1-201 lab-terraform-eks-2]$ vi terraform.tf █
```

Open the terraform.tf file:

vim terraform.tf

```
terraform {
  required_providers {
    aws = {
      source = "hashicorp/aws"
      version = "~> 4.47.0"
    }
    kubernetes = {
      source = "hashicorp/kubernetes"
      version = ">= 2.16.1"
    }
    random = {
      source = "hashicorp/random"
      version = "~> 3.4.3"
    }
    tls = {
```

```

source = "hashicorp/tls"
version = "~> 4.0.4"
}
}
required_version = "~> 1.3"
}

```

Review the contents of the file, and then enter :q! to exit the file.

Open the main.tf file:

vim main.tf

```

provider "aws" {
  region = var.region
}
data "aws_availability_zones" "available" {}
locals {
  cluster_name = "guru-eks-${random_string.suffix.result}"
}
resource "random_string" "suffix" {
  length = 8
  special = false
}
module "vpc" {
  source = "terraform-aws-modules/vpc/aws"
  version = "3.19.0"
  name = "guru-vpc"
  cidr = "10.0.0.0/16"
  azs = slice(data.aws_availability_zones.available.names, 0, 3)
  private_subnets = ["10.0.1.0/24", "10.0.2.0/24", "10.0.3.0/24"]
  public_subnets = ["10.0.4.0/24", "10.0.5.0/24", "10.0.6.0/24"]
  enable_nat_gateway = true
  single_nat_gateway = true
  enable_dns_hostnames = true
  public_subnet_tags = {
    "kubernetes.io/cluster/${local.cluster_name}" = "shared"
    "kubernetes.io/role/elb" = 1
  }
  private_subnet_tags = {
    "kubernetes.io/cluster/${local.cluster_name}" = "shared"
    "kubernetes.io/role/internal-elb" = 1
  }
}

```

Review the contents of the file, and then enter :q! to exit the file.

Open the variables.tf file:

vim variables.tf

```
variable "region" {
  description = "AWS region"
  type       = string
  default    = "us-east-1"
}
```

Review the contents of the file, and then enter :q! to exit the file.

Open the outputs.tf file:

vim outputs.tf

```
output "cluster_endpoint" {
  description = "Endpoint for EKS control plane"
  value      = module.eks.cluster_endpoint
}
output "cluster_security_group_id" {
  description = "Security group ids attached to the cluster control plane"
  value      = module.eks.cluster_security_group_id
}
output "region" {
  description = "AWS region"
  value      = var.region
}
output "cluster_name" {
  description = "Kubernetes Cluster Name"
  value      = module.eks.cluster_name
}
```

Review the contents of the file, and then enter :q! to exit the file.

NOTE: to recreate this Terraform project use the above code to create the files required.

```
[cloud_user@ip-10-0-1-201 ~]$ ls
lab-terraform-eks-2  lab-terraform-eks-2.zip  resource_ids.txt
[cloud_user@ip-10-0-1-201 ~]$ cd lab-terraform-eks-2/
[cloud_user@ip-10-0-1-201 lab-terraform-eks-2]$ ls
main.tf  outputs.tf  terraform.tf  variables.tf
[cloud_user@ip-10-0-1-201 lab-terraform-eks-2]$ vi terraform.tf
[cloud_user@ip-10-0-1-201 lab-terraform-eks-2]$ vi main.tf
[cloud_user@ip-10-0-1-201 lab-terraform-eks-2]$ vi variables.tf
[cloud_user@ip-10-0-1-201 lab-terraform-eks-2]$ vi outputs.tf
[cloud_user@ip-10-0-1-201 lab-terraform-eks-2]$ 
```

Deploy the EKS Cluster

Initialize the working directory (and download the appropriate plugins):

terraform init

Review the actions that will be performed and check for any potential errors:

```
terraform has been successfully initialized!

You may now begin working with Terraform. Try running "terraform plan" to see
any changes that are required for your infrastructure. All Terraform commands
should now work.

If you ever set or change modules or backend configuration for Terraform,
run this command to reinitialize your working directory. If you forget, other
commands will detect it and remind you to do so if necessary.
cloud_user@ip-10-0-1-201 lab-terraform-eks-2]$ 
```

Check to see if the provisioning plan of what you wish to deploy is acceptable to Terraform with the following command.

terraform plan

Now apply the configuration and deploy your cluster:

terraform apply

Type "**yes**" and hit **Enter** to confirm. Allow the configuration to complete successfully and create your resources, including the EKS cluster.

```
aws_eks_addon.ebs-csi: Still creating... [14m40s elapsed]
aws_eks_addon.ebs-csi: Still creating... [14m50s elapsed]
aws_eks_addon.ebs-csi: Still creating... [15m0s elapsed]
aws_eks_addon.ebs-csi: Creation complete after 15m8s [id=guru-eks-jxyiEf3V:aws-ebs-csi-driver]

Apply complete! Resources: 58 added, 0 changed, 0 destroyed.

Outputs:

cluster_endpoint = "https://3830E9EF67E279F402DEEAAD9AC100DA.gr7.us-east-1.eks.amazonaws.com"
cluster_name = "guru-eks-jxyiEf3V"
cluster_security_group_id = "sg-0657edeb6b12b8238"
region = "us-east-1"
[cloud_user@ip-10-0-1-201 lab-terraform-eks-2]$ 
```

Note: This can take around 10 minutes or more to complete so please be patient.

The **Output.tf** gives us the following:

```
cluster_endpoint = "https://3830E9EF67E279F402DEEAAD9AC100DA.gr7.us-east-1.eks.amazonaws.com"
cluster_name = "guru-eks-jxyiEf3V"
cluster_security_group_id = "sg-0657edeb6b12b8238"
region = "us-east-1"
[cloud_user@ip-10-0-1-201 lab-terraform-eks-2]$ 
```

Configure kubectl to interact with the cluster:

```
aws eks --region ${terraform output -raw region} update-kubeconfig --name ${terraform output -raw cluster_name}
```

```
[cloud_user@ip-10-0-1-201 lab-terraform-eks-2]$ aws eks --region ${terraform output -raw region} update-kubeconfig --name ${terraform output -raw cluster_name}
[cloud_user@ip-10-0-1-201 lab-terraform-eks-2]$
```

Confirm that kubectl was configured properly and that the cluster was successfully deployed:

kubectl get cs

All the components should be up and running with a status of Healthy as below.

| NAME | STATUS | MESSAGE | ERROR |
|--------------------|---------|-------------------------------|-------|
| etcd-1 | Healthy | {"health":"true","reason":""} | |
| etcd-0 | Healthy | {"health":"true","reason":""} | |
| etcd-2 | Healthy | {"health":"true","reason":""} | |
| controller-manager | Healthy | ok | |
| scheduler | Healthy | ok | |

```
[cloud_user@ip-10-0-1-201 lab-terraform-eks-2]$ kubectl get cs
Warning: v1 ComponentStatus is deprecated in v1.19+
NAME      STATUS  MESSAGE           ERROR
etcd-1    Healthy {"health": "true", "reason": ""}
etcd-0    Healthy {"health": "true", "reason": ""}
etcd-2    Healthy {"health": "true", "reason": ""}
controller-manager  Healthy   ok
scheduler   Healthy   ok
[cloud_user@ip-10-0-1-201 lab-terraform-eks-2]$
```

Deploy the NGINX Pods

Now download the **nginx.tf** file that contains the configuration for the NGINX deployments from the GitHub repo:

```
wget https://raw.githubusercontent.com/ACloudGuru-Resources/content-terraform-2021/main/nginx.tf
```

Confirm that the file was downloaded to the directory successfully:

ls

Open the nginx.tf file:

vim nginx.tf

```
data "terraform_remote_state" "eks" {
  backend = "local"

  config = {
    path = "../lab-terraform-eks-2/terraform.tfstate"
  }
}
```

```

data "aws_eks_cluster" "cluster" {
  name = data.terraform_remote_state.eks.outputs.cluster_name
}

provider "kubernetes" {
  host           = data.aws_eks_cluster.cluster.endpoint
  cluster_ca_certificate = base64decode(data.aws_eks_cluster.cluster.certificate_authority.0.data)
  exec {
    api_version = "client.authentication.k8s.io/v1beta1"
    command     = "aws"
    args = [
      "eks",
      "get-token",
      "--cluster-name",
      data.aws_eks_cluster.cluster.name
    ]
  }
}

resource "kubernetes_deployment" "nginx" {
  metadata {
    name = "long-live-the-bat"
    labels = {
      App = "LongLiveTheBat"
    }
  }

  spec {
    replicas = 2
    selector {
      match_labels = {
        App = "LongLiveTheBat"
      }
    }
    template {
      metadata {
        labels = {
          App = "LongLiveTheBat"
        }
      }
      spec {
        container {
          image = "nginx:1.7.8"
          name = "batman"

          port {
            container_port = 80
          }
        }
      }
    }
  }
}

```

```
resources {
    limits = {
        cpu   = "0.5"
        memory = "512Mi"
    }
    requests = {
        cpu   = "250m"
        memory = "50Mi"
    }
}
}
}
}
}
```

Review the contents of the file, and then enter :q! to exit the file.

```
[cloud_user@ip-10-0-1-201 lab-terraform-eks-2]$ wget https://raw.githubusercontent.com/ACloudGuru-Resources/content-terraform-2021/main/nginx.tf
--2023-09-22 11:09:19-- https://raw.githubusercontent.com/ACloudGuru-Resources/content-terraform-2021/main/nginx.tf
Resolving raw.githubusercontent.com (raw.githubusercontent.com)... 185.199.111.133, 185.199.108.133, 185.199.109.133, ...
Connecting to raw.githubusercontent.com (raw.githubusercontent.com)|185.199.111.133|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 1401 (1.4K) [text/plain]
Saving to: 'nginx.tf'

100%[=====] 1,401

2023-09-22 11:09:19 (22.5 MB/s) - 'nginx.tf' saved [1401/1401]

[cloud_user@ip-10-0-1-201 lab-terraform-eks-2]$ ls
main.tf  nginx.tf  outputs.tf  terraform.tf  terraform.tfstate  variables.tf
[cloud_user@ip-10-0-1-201 lab-terraform-eks-2]$ vi nginx.tf
[cloud_user@ip-10-0-1-201 lab-terraform-eks-2]$ terraform plan
```

Review the actions that will be performed and check for any potential errors:

Since we are updating the infrastructure, we need to run the same commands again.
terraform plan

Apply the configuration and deploy your NGINX pods to your EKS cluster:

terraform apply

Type "yes" and hit Enter to confirm. Allow the configuration to complete successfully and create your resources, including the NGINX pods.

```
Apply complete! Resources: 1 added, 0 changed, 0 destroyed.

Outputs:

cluster_endpoint = "https://3830E9EF67E279F402DEEAAD9AC100DA.gr7.us-east-1.eks.amazonaws.com"
cluster_name = "guru-eks-jxyiEf3V"
cluster_security_group_id = "sg-0657edeb6b12b8238"
region = "us-east-1"
[cloud_user@ip-10-0-1-201 lab-terraform-eks-2]$ █
```

Confirm that the NGINX pods were successfully deployed:

kubectl get deployments

The two long-live-the-bat pods should both be up and running.

```
[cloud_user@ip-10-0-1-201 lab-terraform-eks-2]$ kubectl get deployments
NAME           READY   UP-TO-DATE   AVAILABLE   AGE
long-live-the-bat   2/2      2          2          111s
[cloud_user@ip-10-0-1-201 lab-terraform-eks-2]$ █
```

Destroy Your Cluster

Before you destroy the infrastructure, view the resources Terraform is managing:

terraform state list

A whole list of resources will appear.

Destroy all of the resources you just created:

terraform destroy

Type "yes" and hit Enter to confirm.

```
random_string.suffix: Destroying... [id=jxyiEf3V]
random_string.suffix: Destruction complete after 0s
module.vpc.aws_vpc.this[0]: Destroying... [id=vpc-0
module.vpc.aws_vpc.this[0]: Destruction complete af

Destroy complete! Resources: 59 destroyed.
[cloud_user@ip-10-0-1-201 lab-terraform-eks-2]$ █
```

Confirm the resources were all destroyed:

terraform state list

Nothing should be returned, confirming that all resources were properly destroyed.

Addendum

To recreate the project we went through we will need to take the following pre requisite steps beforehand.

Create and connect to a new instance on AWS (hence we will need an AWS account).

Create an Admin User Role within IAM

Upgrade AWS CLI to version 2 (for AMI distros)

Install the AWS CLI into that directory

Install and configure Terraform on an EC2 instance

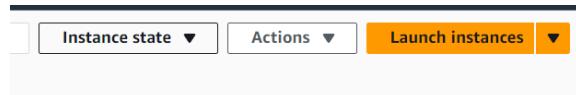
Install Kubectl

Create and Connect to New Instance

In the search bar on top, enter **EC2** and select EC2 from the search results to access EC2.

Select **Instances (running)**.

Click the **Launch instances** button.



On the Launch an instance button, set the following parameters:

Name and tags: Enter **eks-admin**.

Application and OS Images (Amazon Machine Image): Select **Amazon Linux**.

Key pair name: Select **Create new key pair**.

On the Create key pair page, under Key pair name, enter **eks-key**.

Click the **Create key pair** button to create the key pair.

Create key pair

Key pair name
Key pairs allow you to connect to your instance securely.
eks-key

The name can include up to 255 ASCII characters. It can't include leading or trailing spaces.

Key pair type

- RSA RSA encrypted private and public key pair
- ED25519 ED25519 encrypted private and public key pair

Private key file format

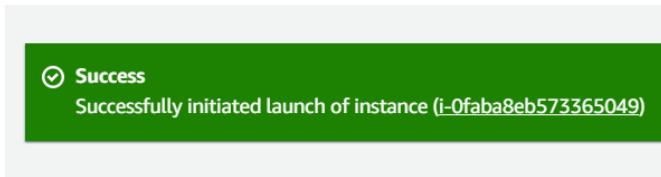
- .pem For use with OpenSSH
- .ppk For use with PuTTY

⚠️ When prompted, store the private key in a secure and accessible location on your computer. You will need it later to connect to your instance. [Learn more](#)

Cancel **Create key pair**

Click the **Launch instance** button.

Under **Success**, select the link to the instance to see its status. It can take a little while for the instance to launch successfully and for the Instance state to change from Pending to Running.

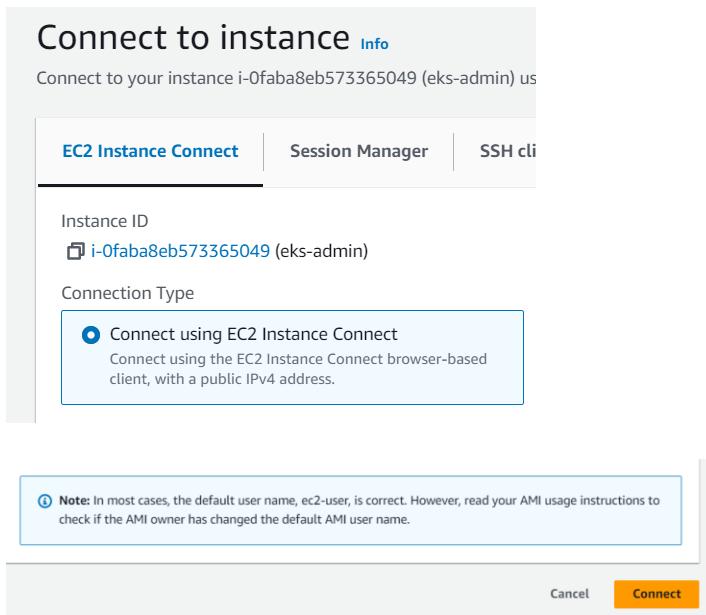


Click the checkbox next to the eks-admin instance to select it.

Click the **Connect** button.

| Instances (1/1) Info | | C | Connect |
|--|---------------------|---|-------------------------|
| <input type="text"/> Find instance by attribute or tag (case-sensitive) | | | |
| <input type="text"/> Instance ID = i-0faba8eb573365049 X Clear filters | | | |
| Name | Instance ID | Instance state | |
| <input checked="" type="checkbox"/> eks-admin | i-0faba8eb573365049 | Running Details | |

On the Connect to instance page, click the **Connect** button.



Create an Admin User Role within IAM

Under Favourites, select **IAM**. If you don't see a Favourites section, you can enter **IAM** in the search bar on top of the console and select **IAM** from the search results to access the **IAM** dashboard.

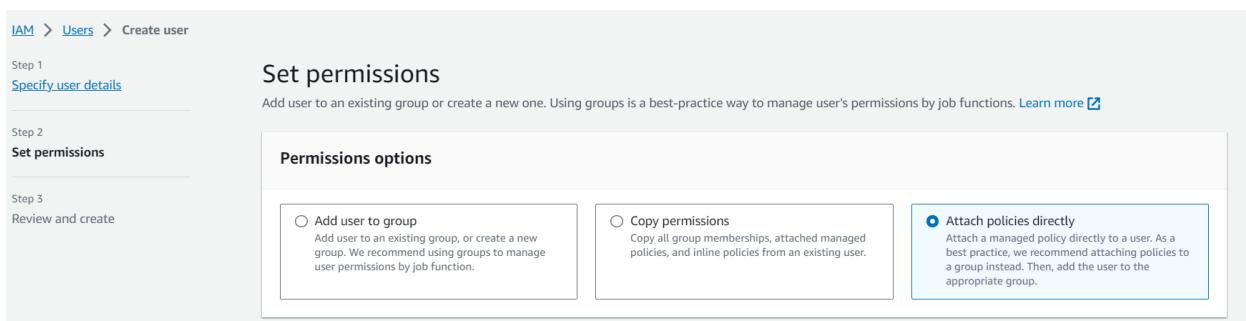
In the lefthand navigation menu, select **Users**.

Click the **Add Users** button.

On the Specify user details page, under Username, enter **eks-user**.

Click the **Next** button.

On the Set permissions page, under Permissions options, select **Attach policies directly**.



Under Permissions policies, click the checkbox next to **AdministratorAccess** to grant admin access to the role being created.

| Policy name | Type |
|--|----------------------------|
| <input type="checkbox"/> <input checked="" type="checkbox"/>  AccessAnalyzerServiceRolePolicy | AWS managed |
| <input checked="" type="checkbox"/> <input type="checkbox"/>  AdministratorAccess | AWS managed - job function |
| <input type="checkbox"/> <input type="checkbox"/>  AdministratorAccess | AWS managed |

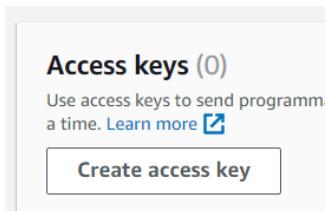
Click the **Next** button.

On the Review and create page, click the **Create user** button.

Once the user is created, select **eks-user**.

Select the **Security credentials** tab.

Under **Access keys**, click the **Create access key** button.



On the Access key best practices & alternatives page, select **Command Line Interface (CLI)**.

Use case

Command Line Interface (CLI)

You plan to use this access key to enable the AWS CLI to access your AWS account.

Click the checkbox next to **I understand the recommendation and want to proceed to create an access key**.

Click the **Next** button.

On the Set description tag - optional page, click the **Create access key** button.

Click the **Download .csv** file.

[Upgrade AWS CLI to version 2](#)

[Download AWS CLI 2](#)

```
curl "https://awscli.amazonaws.com/awscli-exe-linux-x86\_64.zip" -o "awscliv2.zip"
```

[Unzip the downloaded CLI 2 file:](#)

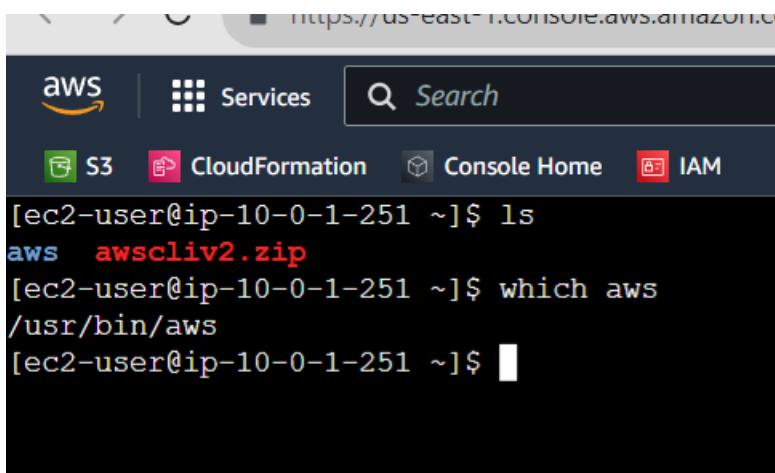
unzip awscli2.zip

This can take a few minutes to unzip and create everything needed.

Check the location of our AWS files:

which aws

You should see they are located in **usr/bin/aws**.



```
[ec2-user@ip-10-0-1-251 ~]$ ls
aws awscli2.zip
[ec2-user@ip-10-0-1-251 ~]$ which aws
/usr/bin/aws
[ec2-user@ip-10-0-1-251 ~]$ 
```

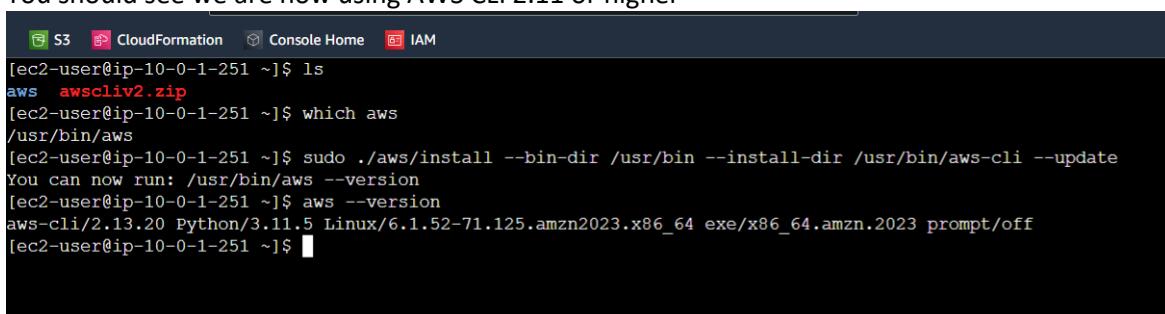
Install the AWS CLI into that directory:

sudo ./aws/install --bin-dir /usr/bin --install-dir /usr/bin/aws-cli --update

Re-check which version of AWS is being used:

aws --version

You should see we are now using AWS CLI 2.11 or higher



```
[ec2-user@ip-10-0-1-251 ~]$ ls
aws awscli2.zip
[ec2-user@ip-10-0-1-251 ~]$ which aws
/usr/bin/aws
[ec2-user@ip-10-0-1-251 ~]$ sudo ./aws/install --bin-dir /usr/bin --install-dir /usr/bin/aws-cli --update
You can now run: /usr/bin/aws --version
[ec2-user@ip-10-0-1-251 ~]$ aws --version
aws-cli/2.13.20 Python/3.11.5 Linux/6.1.52-71.125.amzn2023.x86_64 exe/x86_64.amzn.2023 prompt/off
[ec2-user@ip-10-0-1-251 ~]$ 
```

Install and configure Terraform on an EC2 instance:

```
[cloud_user@ip-10-0-1-201 ~]$ terraform version
Terraform v1.5.7
on linux_amd64
[cloud_user@ip-10-0-1-201 ~]$ █
```

sudo yum update -y

wget https://releases.hashicorp.com/terraform/1.5.7/terraform_1.5.7_linux_amd64.zip

unzip terraform_1.5.7_linux_amd64.zip

sudo mv terraform /usr/local/bin/

terraform --version

rm terraform_1.5.7_linux_amd64.zip

Install Kubectl

curl -LO "<https://dl.k8s.io/release/v1.21.2/bin/linux/amd64/kubectl>"

chmod +x kubectl

sudo mv kubectl /usr/local/bin/

kubectl version --client

rm kubectl

[cloud_user@ip-10-0-1-201 ~]\$ kubectl version

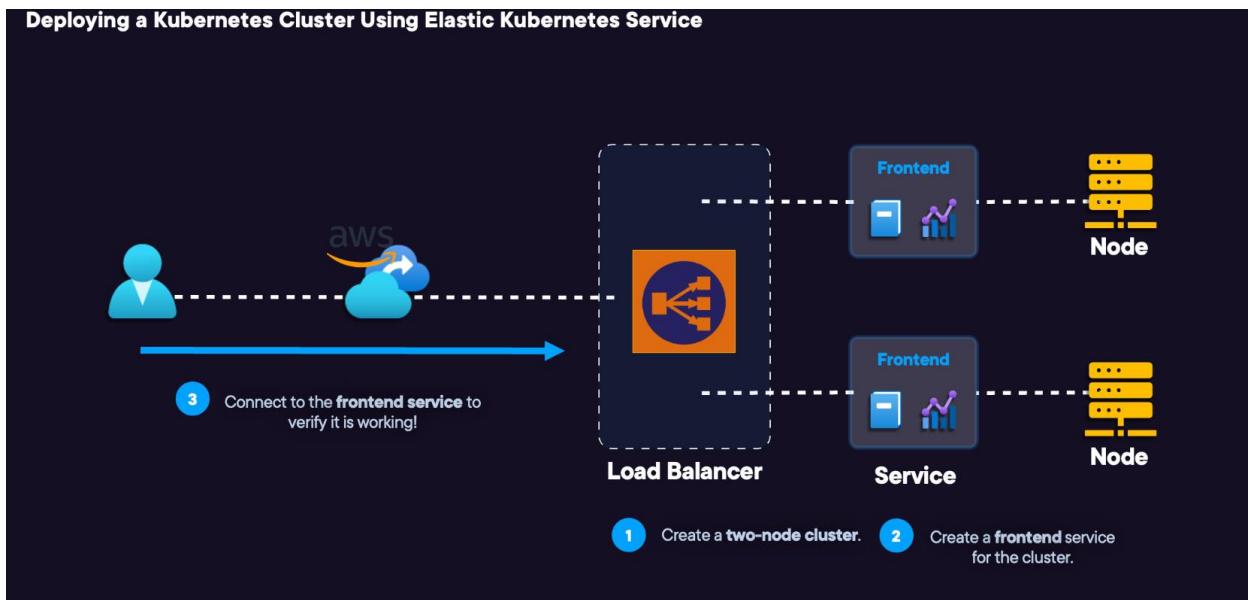
Client Version: version.Info{Major:"1", Minor:"21+", GitVersion:"v1.21.2-13+d2965f0db10712",
 GitCommit:"d2965f0db1071203c6f5bc662c2827c71fc8b20d", GitTreeState:"clean", BuildDate:"2021-06-
 26T01:02:11Z", GoVersion:"go1.16.5", Compiler:"gc", Platform:"linux/amd64"}

The connection to the server localhost:8080 was refused -

did you specify the right host or port?

[cloud_user@ip-10-0-1-201 ~]\$

Deploying a Kubernetes Cluster Using Elastic Kubernetes Service



Solution

Log into the AWS Management Console.

Note: you will need to have an AWS account prior to starting.

Note: Make sure you're using the us-east-1 Region.

Create an EKS Cluster in the us-east-1 Region:

Create an Admin User Role:

Under Favourites, select **IAM**. If you don't see a Favourites section, you can enter **IAM** in the search bar on top of the console and select **IAM** from the search results to access the **IAM** dashboard.

In the lefthand navigation menu, select **Users**.

Click the **Add Users** button.

On the Specify user details page, under Username, enter **eks-user**.

Click the **Next** button.

On the **Set permissions** page, under Permissions options, select **Attach policies directly**.

IAM > Users > Create user

Step 1: Specify user details

Step 2: Set permissions

Step 3: Review and create

Set permissions

Add user to an existing group or create a new one. Using groups is a best-practice way to manage user's permissions by job functions. [Learn more](#)

Permissions options

- Add user to group Add user to an existing group, or create a new group. We recommend using groups to manage user permissions by job function.
- Copy permissions Copy all group memberships, attached managed policies, and inline policies from an existing user.
- Attach policies directly Attach a managed policy directly to a user. As a best practice, we recommend attaching policies to a group instead. Then, add the user to the appropriate group.

Under Permissions policies, click the checkbox next to **AdministratorAccess** to grant admin access to the role being created.

| Policy name | Type |
|--|----------------------------|
| <input type="checkbox"/> AccessAnalyzerServiceRolePolicy | AWS managed |
| <input checked="" type="checkbox"/> AdministratorAccess | AWS managed - job function |

Click the **Next** button.

On the Review and create page, click the **Create user** button.

Once the user is created, select **eks-user**.

Select the **Security credentials** tab.

Under **Access keys**, click the **Create access key** button.

Access keys (0)

Use access keys to send programmatic requests at any time. [Learn more](#)

Create access key

On the Access key best practices & alternatives page, select **Command Line Interface (CLI)**.

Use case

- Command Line Interface (CLI)
You plan to use this access key to enable the AWS CLI to access your AWS account.

Click the checkbox next to **I understand the recommendation and want to proceed to create an access key**.

Click the **Next** button.

On the Set description tag - optional page, click the **Create access key** button.

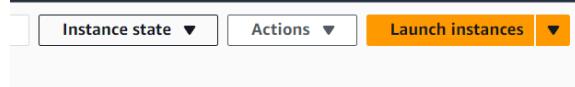
Click the **Download .csv** file.

Create and Connect to New Instance:

In the search bar on top, enter **EC2** and select **EC2** from the search results to access EC2.

Select **Instances (running)**.

Click the **Launch instances** button.



On the Launch an instance button, set the following parameters:

Name and tags: Enter **eks-admin**.

Application and OS Images (Amazon Machine Image): Select **Amazon Linux**.

Key pair name: Select **Create new key pair**.

On the Create key pair page, under Key pair name, enter **eks-key**.

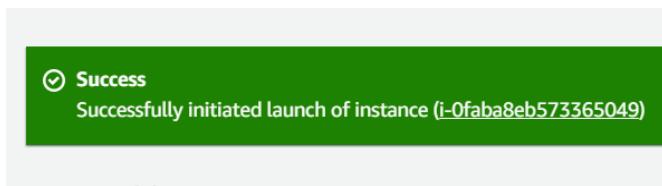
Click the **Create key pair** button to create the key pair.

 A screenshot of the 'Create key pair' wizard. Step 1: Set key pair name. It shows a text input field with 'eks-key' typed into it. Below the input field is a note: 'The name can include up to 255 ASCII characters. It can't include leading or trailing spaces.' Step 2: Set key pair type. It shows two options: 'RSA' (selected) and 'ED25519'. Step 3: Set private key file format. It shows two options: '.pem' (selected) and '.ppk'. Step 4: A warning message: '⚠ When prompted, store the private key in a secure and accessible location on your computer. You will need it later to connect to your instance. [Learn more](#)' Step 5: 'Create key pair' button at the bottom.

| | | |
|--|---|--|
| Key pair name | eks-key | |
| Key pair type | <input checked="" type="radio"/> RSA RSA encrypted private and public key pair | <input type="radio"/> ED25519 ED25519 encrypted private and public key pair |
| Private key file format | <input checked="" type="radio"/> .pem For use with OpenSSH | |
| | <input type="radio"/> .ppk For use with PuTTY | |
| ⚠ When prompted, store the private key in a secure and accessible location on your computer. You will need it later to connect to your instance. Learn more | | |

Click the **Launch instance** button.

Under **Success**, select the link to the instance to see its status. It can take a little while for the instance to launch successfully and for the Instance state to change from Pending to Running.



Click the checkbox next to the eks-admin instance to select it.

Click the **Connect** button.

A screenshot of the AWS EC2 Instances page. It shows a single instance listed: "eks-admin" with Instance ID "i-0faba8eb573365049". The "Name" column has a checked checkbox. The "Instance state" column shows "Running". There are "Connect" and "Info" buttons at the top right.

On the Connect to instance page, click the **Connect** button.

A screenshot of the "Connect to instance" page. It shows the instance "i-0faba8eb573365049 (eks-admin)" selected. Under "Connection Type", the radio button for "Connect using EC2 Instance Connect" is selected. A note at the bottom left says: "Note: In most cases, the default user name, ec2-user, is correct. However, read your AMI usage instructions to check if the AMI owner has changed the default AMI user name." At the bottom right are "Cancel" and "Connect" buttons.

Create Cluster

Check the version of the CLI being used to verify whether it is NOT AWS CLI 2:

aws --version

You should see that the version being used is AWS CLI 1.18 or above.

The screenshot shows a terminal window within the AWS CloudShell interface. The title bar indicates the URL is <https://us-east-1.console.aws.amazon.com/ec2-instance-connect/ssh?connType=standard&instanceId=i-0faba8eb573365049&osUser=ec2-user>. The terminal displays a welcome message for Amazon Linux 2023, followed by the command `aws --version` which outputs the AWS CLI version 2.9.19. The user is also prompted to upgrade to version 3.9.16.

```

Amazon Linux 2023
https://aws.amazon.com/linux/amazon-linux-2023

[ec2-user@ip-10-0-1-251 ~]$ aws --version
aws-cli/2.9.19 Python/3.9.16 Linux/6.1.52-71.125.amzn2023.x86_64 source/x86_64.amzn.2023 prompt/off
[ec2-user@ip-10-0-1-251 ~]$ 

```

For this project we will need to upgrade our AWS CLI to version 2.

Note: if AWS CLI does NOT exist at all then follow these steps as well.

Download AWS CLI 2:

```
curl "https://awscli.amazonaws.com/awscli-exe-linux-x86_64.zip" -o "awscliv2.zip"
```

Unzip the downloaded CLI 2 file:

```
unzip awscliv2.zip
```

This can take a few minutes to unzip and create everything needed.

Check the location of our AWS files:

```
which aws
```

You should see they are located in **usr/bin/aws**.

The screenshot shows a terminal window within the AWS CloudShell interface. The title bar indicates the URL is <https://us-east-1.console.aws.amazon.com/ec2-instance-connect/ssh?connType=standard&instanceId=i-0faba8eb573365049&osUser=ec2-user>. The terminal lists the contents of the current directory, showing the `awscliv2.zip` file. It then runs the `which aws` command, which returns the path `/usr/bin/aws`.

```

[ec2-user@ip-10-0-1-251 ~]$ ls
aws  awscliv2.zip
[ec2-user@ip-10-0-1-251 ~]$ which aws
/usr/bin/aws
[ec2-user@ip-10-0-1-251 ~]$ 

```

Install the AWS CLI into the destination directory:

```
sudo ./aws/install --bin-dir /usr/bin --install-dir /usr/bin/aws-cli --update
```

Re-check which version of AWS is being used:

aws –version

You should see we are now using AWS CLI 2.11 or above.

```
S3 CloudFormation Console Home IAM
[ec2-user@ip-10-0-1-251 ~]$ ls
aws awscli2.zip
[ec2-user@ip-10-0-1-251 ~]$ which aws
/usr/bin/aws
[ec2-user@ip-10-0-1-251 ~]$ sudo ./aws/install --bin-dir /usr/bin --install-dir /usr/bin/aws-cli --update
You can now run: /usr/bin/aws --version
[ec2-user@ip-10-0-1-251 ~]$ aws --version
aws-cli/2.13.20 Python/3.11.5 Linux/6.1.52-71.125.amzn2023.x86_64 exe/x86_64.amzn.2023 prompt/off
[ec2-user@ip-10-0-1-251 ~]$
```

Log into our AWS environment:

aws configure

From the CSV file created earlier, copy and paste the **Access Key ID** when prompted.

From the same CSV file, copy and paste the **Secret Access Key** when prompted.

| A | B | C | D | E | F | G |
|----------------------|--|---|---|---|---|---|
| Access key ID | Secret access key | | | | | |
| AKIASXI2QSGGEJ4KZV4C | d/0AmSA/N/WPE1RSGebTE7wzFhSc1pPvpdgKQitU | | | | | |

When prompted for the default region name, enter **us-east-1**.

When prompted for the default output format, enter **json**.

```
https://us-east-1.console.aws.amazon.com/ec2-instance-connect/ssh?connTyp
aws Services Search
S3 CloudFormation Console Home IAM
[ec2-user@ip-10-0-1-251 ~]$ aws configure
AWS Access Key ID [None]: AKIASXI2QSGGEJ4KZV4C
AWS Secret Access Key [None]: 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
Default region name [None]: us-east-1
Default output format [None]: json
```

Note: In a few hours this infrastructure will be gone with the key pair credentials!

Download kubectl:

```
curl -O https://s3.us-west-2.amazonaws.com/amazon-eks/1.25.6/2023-01-30/bin/linux/amd64/kubectl
```

Add execution permissions to the file:

```
chmod +x ./kubectl
```

Create a new bin directory and copy kubectl to the bin directory and home directory:

```
mkdir -p $HOME/bin && cp ./kubectl $HOME/bin/kubectl && export PATH=$PATH:$HOME/bin
```

Download eksctl:

```
curl --silent --location "https://github.com/weaveworks/eksctl/releases/latest/download/eksctl_$(uname -s)_amd64.tar.gz" | tar xz -C /tmp
```

This will install the CLI tool you will use to create a cluster on Amazon EKS.

Move the files to usr/local/bin:

```
sudo mv /tmp/eksctl /usr/local/bin
```

Check the eksctl version:

```
eksctl version
```

This should show the version as 0.133.0 or higher!

```
[ec2-user@ip-10-0-1-251 ~]$ curl -O https://s3.us-west-2.amazonaws.com/amazon-eks/1.25.6/2023-01-30/bin/linux/amd64/kubectl
$ Total % Received % Xferd Average Speed Time Time Current
          Dload Upload Total Spent Left Speed
100 42.9M 100 42.9M 0 0 17.2M 0 0:00:02 0:00:02 --:-- 17.2M
[ec2-user@ip-10-0-1-251 ~]$ chmod +x ./kubectl
[ec2-user@ip-10-0-1-251 ~]$ mkdir -p $HOME/bin && cp ./kubectl $HOME/bin/kubectl && export PATH=$PATH:$HOME/bin
[ec2-user@ip-10-0-1-251 ~]$ curl --silent --location "https://github.com/weaveworks/eksctl/releases/latest/download/eksctl_$(uname -s)_amd64.tar.gz" | tar xz -C /tmp
[ec2-user@ip-10-0-1-251 ~]$ sudo mv /tmp/eksctl /usr/local/bin
[ec2-user@ip-10-0-1-251 ~]$ eksctl version
0.157.0
```

Create the cluster:

```
eksctl create cluster --name myeks --region us-east-1 --zones=us-east-1a,us-east-1b --nodegroup-name eks-workers --node-type t3.medium --nodes 2 --nodes-min 2 --nodes-max 4 --managed
```

This may take several minutes to finish creating the cluster.

```
[ec2-user@ip-10-0-1-251 ~]$ eksctl create cluster --name myeks --region us-east-1 --zones=us-east-1a,us-east-1b --nodegroup-name eks-workers --node-type t3.medium --nodes 2 --nodes-min 2 --nodes-max 4 --managed
2023-09-21 20:29:49 [i] eksctl version 1.25
2023-09-21 20:29:49 [i] using region us-east-1
2023-09-21 20:29:49 [i] subnets for us-east-1a - public:192.168.0.0/19 private:192.168.64.0/19
2023-09-21 20:29:49 [i] subnets for us-east-1b - public:192.168.32.0/19 private:192.168.96.0/19
2023-09-21 20:29:50 [i] nodegroup "eks-workers" will use "" (AmazonLinux/1.25)
2023-09-21 20:29:50 [i] using Kubernetes version 1.25
2023-09-21 20:29:50 [i] creating EKS cluster "myeks" in "us-east-1" region with managed nodes
2023-09-21 20:29:50 [i] Will create 2 separate CloudFormation stacks for cluster itself and the initial managed nodegroup
2023-09-21 20:29:50 [i] If you encounter any issues, check CloudFormation console or try 'eksctl utils describe-stacks --region=us-east-1 --cluster=myeks'
2023-09-21 20:29:50 [i] Kubernetes API endpoint access will use default of [publicAccess=true, privateAccess=false] for cluster "myeks" in "us-east-1"
2023-09-21 20:29:50 [i] CloudWatch logging will not be enabled for cluster "myeks" in "us-east-1"
2023-09-21 20:29:50 [i] You can enable it with 'eksctl utils update-cluster-logging --enable-types=(SPECIFY-YOUR-LOG-TYPES-HERE (e.g. all)) --region=us-east-1 --cluster=myeks'
2023-09-21 20:29:50 [i] sequential tasks: ( create cluster control plane "myeks",
2023-09-21 20:29:50 [i]   2 sequential sub-tasks: (
2023-09-21 20:29:50 [i]     wait for control plane to become ready,
2023-09-21 20:29:50 [i]     create managed nodegroup "eks-workers",
2023-09-21 20:29:50 [i]   )
2023-09-21 20:29:50 [i]   building cluster stack "eksctl-myeks-cluster"
2023-09-21 20:29:50 [i]   deploying stack "eksctl-myeks-cluster"
```

| Stack name | Status | Created time | Description |
|--|-----------------|------------------------------|--|
| eksctl-myeks-cluster | CREATE_COMPLETE | 2023-09-21 23:29:50 UTC+0300 | EKS cluster (dedicated VPC: true, dedicated IAM: true) [created and managed by eksctl] |
| cfst-4038-af246102f34a0759f74dd7d062a0fe4b | CREATE_COMPLETE | 2023-09-21 21:00:40 UTC+0300 | Creates a VPC with (3) public subnets allowing ssh |

Once the cluster is created, connect to the cluster:

aws eks update-kubeconfig --name myeks --region us-east-1

Verify that it is running eksctl:

eksctl get cluster

You should see that it is running.

Deploy a Test Application to Mimic the Application

Install the Git package:

sudo yum install git -y

Clone a repository from the EKS Basics course:

git clone https://github.com/costas778/Course_EKS-Basics

Change into the Course_EKS-Basics directory:

cd Course_EKS-Basics/

View the contents of the directory:

ls

You should see a license file, an nginx-deployment.yaml file, and an nginx-svc.yaml file.

```
[ec2-user@ip-10-0-1-251 ~]$ git clone https://github.com/costas778/Course_EKS-Basics
Cloning into 'Course_EKS-Basics'...
remote: Enumerating objects: 21, done.
remote: Counting objects: 100% (21/21), done.
remote: Compressing objects: 100% (20/20), done.
remote: Total 21 (delta 8), reused 0 (delta 0), pack-reused 0
Receiving objects: 100% (21/21), 5.37 KiB | 1.34 MiB/s, done.
Resolving deltas: 100% (8/8), done.
[ec2-user@ip-10-0-1-251 ~]$ cd Course_EKS-Basics/
[ec2-user@ip-10-0-1-251 Course_EKS-Basics]$ ls
LICENSE  nginx-deployment.yaml  nginx-svc.yaml
[ec2-user@ip-10-0-1-251 Course_EKS-Basics]$
```

View the nginx-deployment.yaml file:

cat nginx-deployment.yaml

View the nginx-svc.yaml file:

cat nginx-svc.yaml

Apply the nginx-svc.yaml file:

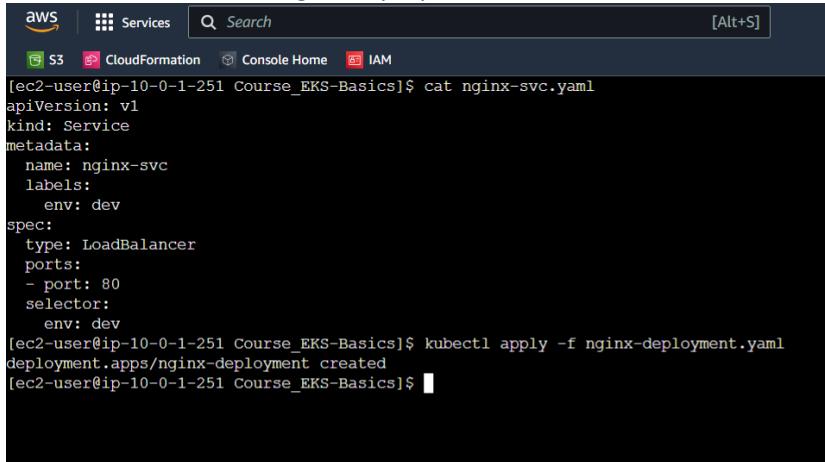
kubectl apply -f nginx-svc.yaml

This should create the Nginx service.

Apply the nginx-deployment.yaml file:

kubectl apply -f nginx-deployment.yaml

This should create the Nginx deployments.



```
[ec2-user@ip-10-0-1-251 Course_EKS-Basics]$ cat nginx-svc.yaml
apiVersion: v1
kind: Service
metadata:
  name: nginx-svc
  labels:
    env: dev
spec:
  type: LoadBalancer
  ports:
  - port: 80
    selector:
      env: dev
[ec2-user@ip-10-0-1-251 Course_EKS-Basics]$ kubectl apply -f nginx-deployment.yaml
deployment.apps/nginx-deployment created
[ec2-user@ip-10-0-1-251 Course_EKS-Basics]$
```

Verify that the deployments are available:

kubectl get deployment

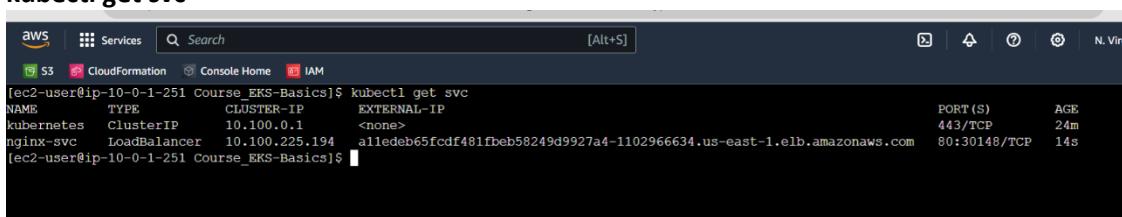
You should see that all three deployments are available.

Verify that all Pods are running:

kubectl get pods

Verify that the service is running:

kubectl get svc



```
[ec2-user@ip-10-0-1-251 Course_EKS-Basics]$ kubectl get svc
NAME      TYPE      CLUSTER-IP   EXTERNAL-IP
kubernetes ClusterIP  10.100.0.1   <none>
nginx-svc  LoadBalancer 10.100.225.194  a11edeb65fcdf481fbefb58249d9927a4-1102966634.us-east-1.elb.amazonaws.com
[ec2-user@ip-10-0-1-251 Course_EKS-Basics]$
```

You should see above that the load balancer has a DNS name.

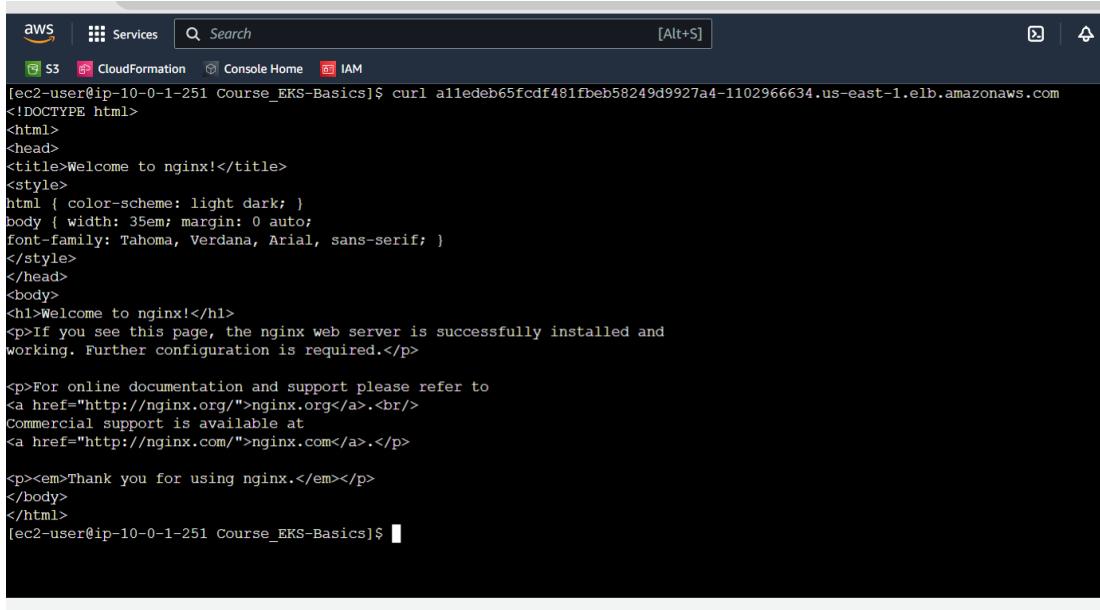
Use the DNS Name of the Load Balancer to Test the Cluster:

Copy the DNS name of the load balancer.

Use the DNS name to test that the cluster is working:

curl <PASTE DNS NAME HERE>

curl a11edeb65fcdf481fbeb58249d9927a4-1102966634.us-east-1.elb.amazonaws.com



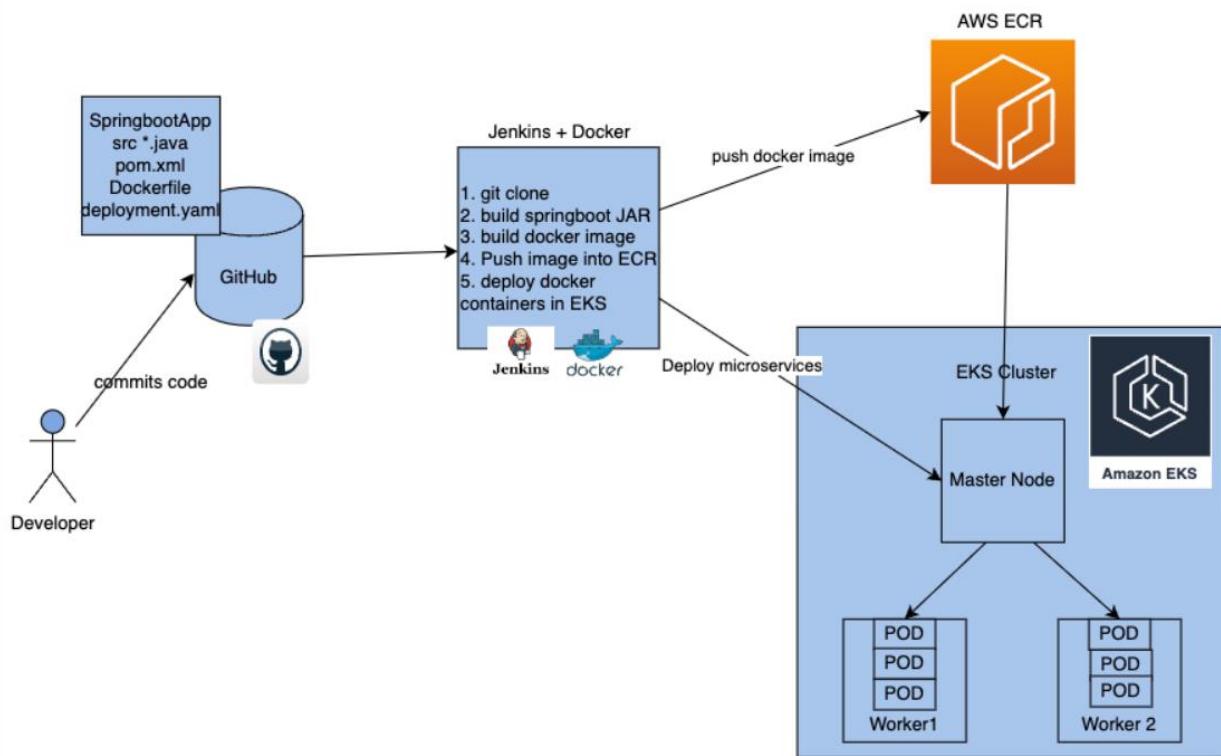
The screenshot shows a terminal window within the AWS CloudWatch interface. The user has run the command `curl a11edeb65fcdf481fbeb58249d9927a4-1102966634.us-east-1.elb.amazonaws.com`. The response is a standard Nginx welcome page, indicating successful installation and providing links for documentation and support.

```
[ec2-user@ip-10-0-1-251 Course_EKS-Basics]$ curl a11edeb65fcdf481fbeb58249d9927a4-1102966634.us-east-1.elb.amazonaws.com
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<style>
html { color-scheme: light dark; }
body { width: 35em; margin: 0 auto;
font-family: Tahoma, Verdana, Arial, sans-serif; }
</style>
</head>
<body>
<h1>Welcome to nginx!</h1>
<p>If you see this page, the nginx web server is successfully installed and
working. Further configuration is required.</p>
<p>For online documentation and support please refer to
<a href="http://nginx.org/">nginx.org</a>.<br/>
Commercial support is available at
<a href="http://nginx.com/">nginx.com</a>.</p>
<p><em>Thank you for using nginx.</em></p>
</body>
</html>
[ec2-user@ip-10-0-1-251 Course_EKS-Basics]$
```

You should see the test page for the Nginx service.

Deploy Springboot Microservices App into Amazon EKS Cluster using Jenkins Pipeline and Kubectl CLI Plug-in | Containerize Springboot App and Deploy into EKS Cluster using Jenkins Pipeline

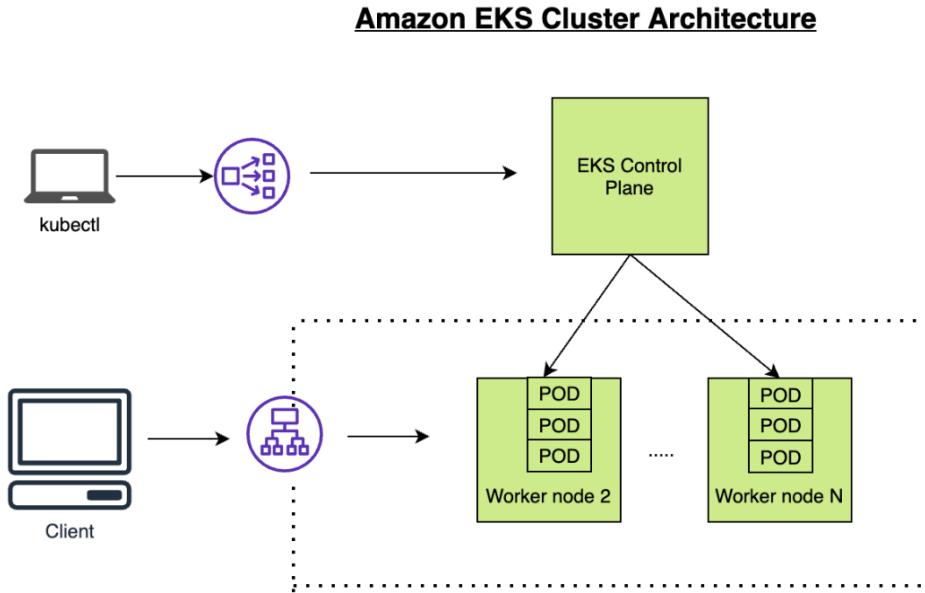
Springboot Microservices Deployment into EKS Cluster using Jenkins Pipeline



Pre-requisites:

1. Amazon EKS Cluster is setup and running.
2. Create ECR repo in AWS
3. Jenkins Master is up and running
4. Docker installed on Jenkins instance
5. Docker, Docker pipeline and Kubernetes CLI plug-ins are installed in Jenkins

Setup Amazon EKS Cluster



To help achieve this we need to first setup a Jenkins **EC2 instance**. This **EC2 instance**, that will be an **Ubuntu 22.4 LTS** server, needs to have following both installed and configured:

- Install AWS CLI** – Command line tools for working with AWS services, including Amazon EKS.
- Install eksctl** – A command line tool for working with EKS clusters that automates many individual tasks.
- Install kubectl** – A command line tool for working with Kubernetes clusters.

Create the Ubuntu instance.

Go to EC2 and launch an instance.

Name and tags [Info](#)

Name
Jenkins [Add additional tags](#)

Application and OS Images (Amazon Machine Image) [Info](#)

An AMI is a template that contains the software configuration (operating system, application server, and applications) required to launch your instance. Search or Browse for AMIs if you don't see what you are looking for below

Quick Start

[Amazon Linux](#) [macOS](#) [Ubuntu](#) [Windows](#) [Red Hat](#) [SUSE Linux](#) [SUSI](#) [Browse more AMIs](#) Including AMIs from AWS, Marketplace and the Community

Ubuntu Server 22.04 LTS (HVM), SSD Volume Type Free tier eligible

ami-053b0d53c279acc90 (64-bit (x86)) / ami-0a0c8eebddd6dcdb0 (64-bit (Arm))
Virtualization: hvm ENA enabled: true Root device type: ebs

Description
Canonical, Ubuntu, 22.04 LTS, amd64 jammy image build on 2023-05-16

Architecture 64-bit (x86) **AMI ID** ami-053b0d53c279acc90 **Verified provider**

Instance type [Info](#)

Instance type **t2.medium**

Family: t2 2 vCPU 4 GiB Memory Current generation: true
On-Demand Linux base pricing: 0.0464 USD per Hour
On-Demand RHEL base pricing: 0.1064 USD per Hour
On-Demand Windows base pricing: 0.0644 USD per Hour
On-Demand SUSE base pricing: 0.1464 USD per Hour

[All generations](#) [Compare instance types](#)

[Additional costs apply for AMIs with pre-installed software](#)

Give it the name **Jenkins**, choose Ubuntu under **Quick Start** and select a **t2.medium** instance type.

Create key pair

Key pair name
Key pairs allow you to connect to your instance securely.

The name can include up to 255 ASCII characters. It can't include leading or trailing spaces.

Key pair type

RSA
RSA encrypted private and public key pair

ED25519
ED25519 encrypted private and public key pair

Private key file format

.pem
For use with OpenSSH

.ppk
For use with PuTTY

⚠️ When prompted, store the private key in a secure and accessible location on your computer. You will need it later to connect to your instance. [Learn more ↗](#)

[Cancel](#) **Create key pair**

On the same page choose to create a key pair and click **Create key pair**.

Create the instance, then once created navigate to EC2 and locate the Jenkins instance you just created security group.

EC2 > Security Groups > sg-085eef18246ebc44b8 - launch-wizard-1 > Edit inbound rules

Edit inbound rules [Info](#)

Inbound rules control the incoming traffic that's allowed to reach the instance.

| Security group rule ID | Type Info | Protocol Info | Port range Info | Source Info | Description - optional Info |
|------------------------|---------------------------|-------------------------------|---------------------------------|-----------------------------|---|
| sgr-0654a505485c1a4e2 | SSH | TCP | 22 | Custom | <input type="text" value="0.0.0.0/0"/> Delete |
| - | Custom TCP | TCP | 8080 | Any... Info | <input type="text" value="0.0.0.0/0"/> Delete |

[Add rule](#)

[Cancel](#) [Preview changes](#) **Save rules**

Create a new inbound rule of port **8080** with **0.0.0.0/0** to give access to the Jenkins server from anywhere. Click **Save Rules**.

Click the **Connect** button within **EC2**.

Use the key pair to login into the Ubuntu instance:

 **Command copied**

 `ssh -i "Jenkins.pem" ubuntu@ec2-54-146-82-200.compute-1.amazonaws.com`

Copy the above link into a command prompt on your local workstation that has the **Jenkins.pem** key pair file.

Note: Ensure that you run the above command in the same directory as the **Jenkins.pem** file. Otherwise, you will have to alter the above command to reflect the path of the **Jenkins.pem** file.

Please follow the steps to install **Java**, **Jenkins**, and **Maven** on Ubuntu 22.0.4 instance. **Jenkins** and, **Maven** are Java based applications, so we need to install Java first.

Change Host Name to Jenkins

sudo hostnamectl set-hostname Jenkins

Note: if you have already named the server Jenkins you can skip this step.

Perform update first

sudo apt update

Install Java 11

sudo apt install default-jdk -y

```
ubuntu@ip-172-31-44-89:~$ sudo apt install default-jdk -y
Reading package lists... Done
Building dependency tree
Reading state information... Done
The following additional packages will be installed:
  at-spi2-core default-jdk-headless default-jre fonts-dejavu-extra libatk-bridge2.0-0
  libatk-wrapper-java libatk-wrapper-java-jni libatk1.0-0 libatk1.0-data libatspi2.0-0
  libdrm-amdgpu1 libdrm-intel1 libdrm-nouveau2 libdrm-radeon1 libfontenc1 libgif7 libgl1
  libgl1-mesa-dri libglapi-mesa libglvnd0 libglx-mesa0 libglx0 libice-dev libice6
  libllvm10 libpciaccess0 libpthread-stubs0-dev libsensors4 libsm-dev libsm6 libx11-dev
  libx11-doc libx11-xcb1 libxau-dev libxaw7 libxcb-dri2-0 libxcb-dri3-0 libxcb-glx0
  libxcb-present0 libxcb-shape0 libxcb-sync1 libxcb1-dev libcomposite1 libxdamage1
```

Once install java, enter the below command

Verify Java Version

java -version

```
ubuntu@ip-172-31-19-22:~/opt$ java -version
openjdk version "11.0.6" 2020-01-14
OpenJDK Runtime Environment (Build 11.0.6+10-post-Ubuntu-Ubuntu18.04.1)
OpenJDK 64-Bit Server VM (Build 11.0.6+10-post-Ubuntu-Ubuntu18.04.1, mixed mode, sharing)
ubuntu@ip-172-31-19-22:~/opt$
```

Maven Installation

You can install Maven by executing below command:

sudo apt install maven -y

you can type **mvn –version**

you should see the below output.

```
ubuntu@ip-172-31-34-93:~$ mvn --version
Apache Maven 3.6.0
Maven home: /usr/share/maven
Java version: 11.0.7, vendor: Ubuntu, runtime: /usr/lib/jvm/java-11-openjdk-amd64
Default locale: en, platform encoding: UTF-8
OS name: "linux", version: "4.15.0-1865-aws", arch: "amd64", family: "unix"
ubuntu@ip-172-31-34-93:~$
```

Note: you can skip the Java step and just install Maven. This is because Maven will do the job of installing Java for you.

Now let's start Jenkins installation.

Jenkins SetupAdd Repository key to the system

```
curl -fsSL https://pkg.jenkins.io/debian/jenkins.io-2023.key | sudo tee \
/usr/share/keyrings/jenkins-keyring.asc > /dev/null
```

Append debian package repo address to the system

```
echo deb [signed-by=/usr/share/keyrings/jenkins-keyring.asc] \
https://pkg.jenkins.io/debian binary/ | sudo tee \
/etc/apt/sources.list.d/jenkins.list > /dev/null
```

Update Ubuntu packagesudo apt update

```
ubuntu@ip-172-31-37-211:~$ sudo apt-get update
Hit:1 http://us-east-2.ec2.archive.ubuntu.com/ubuntu xenial InRelease
Hit:2 http://us-east-2.ec2.archive.ubuntu.com/ubuntu xenial-updates InRelease
Hit:3 http://us-east-2.ec2.archive.ubuntu.com/ubuntu xenial-backports InRelease
Hit:4 http://security.ubuntu.com/ubuntu xenial-security InRelease
Hit:5 http://ppa.launchpad.net/webupd8team/java/ubuntu xenial InRelease
Ign:6 http://pkg.jenkins.io/debian-stable binary/ InRelease
Get:7 http://pkg.jenkins.io/debian-stable binary/ Release [2,042 B]
Get:8 http://pkg.jenkins.io/debian-stable binary/ Release.gpg [181 B]
Get:9 http://pkg.jenkins.io/debian-stable binary/ Packages [13.2 kB]
Fetched 15.5 kB in 0s (17.8 kB/s)
Reading package lists... Done
```

Install Jenkins

sudo apt install jenkins -y

```
ubuntu@ip-172-31-37-211:~$ sudo apt-get install jenkins -y
Reading package lists... Done
Building dependency tree
Reading state information... Done
The following additional packages will be installed:
  daemon
The following NEW packages will be installed:
  daemon jenkins
0 upgraded, 2 newly installed, 0 to remove and 8 not upgraded.
Need to get 72.0 MB of archives.
After this operation, 75.1 MB of additional disk space will be used.
Get:1 http://us-east-2.ec2.archive.ubuntu.com/ubuntu xenial/universe amd64 daemon [71.9 MB]
Get:2 http://pkg.jenkins.io/debian-stable binary/ jenkins 2.121.3 [71.9 MB]
Fetched 72.0 MB in 5s (12.5 MB/s)
Selecting previously unselected package daemon.
(Reading database ... 51843 files and directories currently installed.)
Preparing to unpack .../daemon_0.6.4-1_amd64.deb ...
Unpacking daemon (0.6.4-1) ...
Selecting previously unselected package jenkins.
Preparing to unpack .../jenkins_2.121.3_all.deb ...
Unpacking jenkins (2.121.3) ...
Processing triggers for man-db (2.7.5-1) ...
Processing triggers for systemd (229-4ubuntu21.4) ...
Processing triggers for ureadahead (0.100.0-19) ...
Setting up daemon (0.6.4-1) ...
Setting up Jenkins (2.121.3) ...
Processing triggers for systemd (229-4ubuntu21.4) ...
Processing triggers for ureadahead (0.100.0-19) ...
ubuntu@ip-172-31-37-211:~$
```

The above screenshot should confirm that Jenkins is successfully installed.

Access Jenkins in web browser

Now Go to AWS console. Click on EC2 then click on the running instances link.
Connect to your instance using its Public IP Address:

The screenshot shows the AWS EC2 Instances page. A specific instance, **i-0fbfb784941fcf59b5**, is selected. The instance summary for this instance is displayed, showing details such as Instance ID (i-0fbfb784941fcf59b5 (Jenkins)), IPv6 address (none), Hostname type (IP name: ip-172-31-51-115.ec2.internal), Answer private resource DNS name (IPv4 (A)), and Auto-assigned IP address (54.146.82.200 [Public IP]).

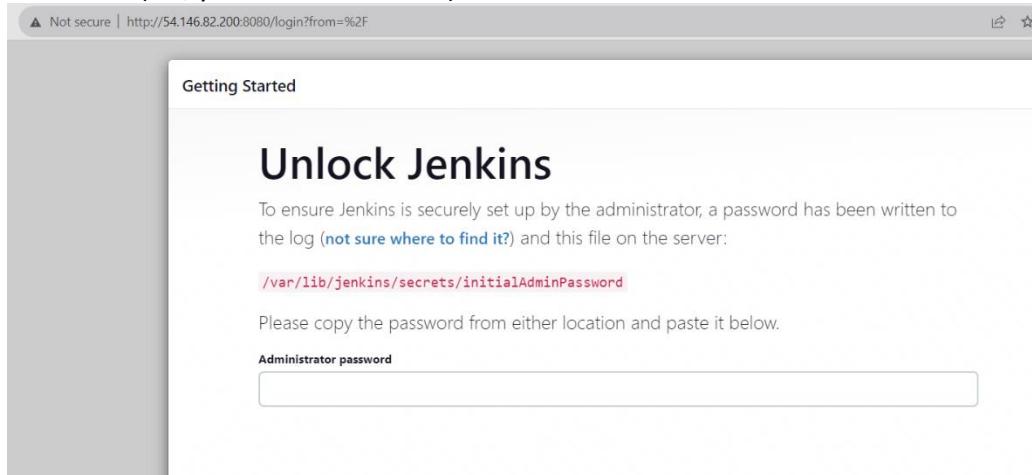
Now go to a browser and enter the **Auto-assigned (public) IP address** with port no **8080**.

http://Public IP Address:8080

It should load without issues after adding the inbound rule with the port 8080.

Unlock Jenkins

Since this is a fresh build, you will get the following screen. Enter the **cat** command below into the bash shell. (i.e., your Ubuntu console).



Get the initial password from the below file

```
sudo cat /var/lib/jenkins/secrets/initialAdminPassword
```

```
ubuntu@ip-172-31-37-211:~$ sudo cat /var/lib/jenkins/secrets/initialAdminPassword
```

Copy the password and paste it into the browser's text box under **Administrator password**. Move to the next page.

Click on **Install suggested plugins**.

Customize Jenkins

Plugins extend Jenkins with additional features to support many different needs.



Also create an admin user.

We choose for our non-production server to use both the username **admin** and password **admin**.

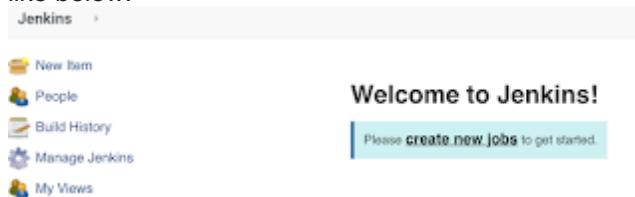
Create User

| | |
|-------------------|-----------------|
| Username: | admin |
| Password: | |
| Confirm password: | |
| Full name: | admin |
| E-mail address: | admin@admin.com |

Create User

Click **Create User**.

Click on **Save and Finish**. Click on **start using Jenkins**. Now you should see a screen like below:



Install the AWS CLI version 2 on Linux.

Follow these steps from the command line to install AWS CLI on Linux.

```
curl "https://awscli.amazonaws.com/awscli-exe-linux-x86_64.zip" -o "awscliv2.zip"
sudo apt install unzip
sudo unzip awscliv2.zip
sudo ./aws/install
aws --version
```

it should display the below output.

```
ubuntu@ip-172-31-24-217:~$ aws --version
aws-cli/2.2.32 Python/3.8.8 Linux/5.4.0-1045-aws exe/x86_64.ubuntu.18 prompt/off
ubuntu@ip-172-31-24-217:~$
```

Install eksctl on a Ubuntu Instance

eksctl is a command line tool for working with EKS clusters that automates many individual tasks. The **eksctl** tool uses **CloudFormation** under the hood, creating one stack for the EKS master control plane and another stack for the worker nodes.

In Amazon **Elastic Kubernetes Service (EKS)**, a **stack** typically refers to a collection of AWS **CloudFormation** resources that are used to provision and manage an EKS cluster and its associated infrastructure

Download and extract the latest release of **eksctl** with the following command.

```
curl --silent --location
```

```
"https://github.com/weaveworks/eksctl/releases/latest/download/eksctl_$(uname -
```

```
s)_amd64.tar.gz" | tar xz -C /tmp
```

Move the extracted binary to /usr/local/bin.

```
sudo mv /tmp/eksctl /usr/local/bin  
eksctl version
```

```
ubuntu@ip-172-31-17-55:~$ eksctl version  
0.92.0  
ubuntu@ip-172-31-17-55:~$
```

Install kubectl on Ubuntu Instance

Kubernetes uses a command line utility called **kubectl** for communicating with the cluster API server. It is tool for controlling **Kubernetes** clusters. **kubectl** looks for a file named config in the \$HOME directory.

Install kubectl

```
sudo curl --silent --location -o /usr/local/bin/kubectl https://s3.us-west-  
2.amazonaws.com/amazon-eks/1.22.6/2022-03-09/bin/linux/amd64/kubectl
```

```
sudo chmod +x /usr/local/bin/kubectl
```

Verify if kubectl got installed

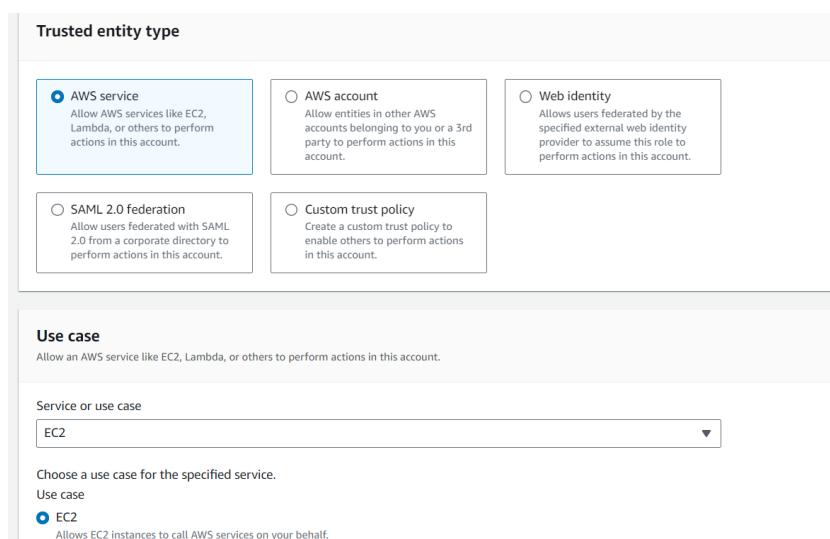
```
kubectl version --short --client
```

NOTE: You may get an error with Ubuntu with the above command. It may complain of deprecated command. If you find this, drop the **-short** switch and just use **-client**.

Its fine on AMI as of writing this!

Create an IAM Role with Administrator Access

You need to create an **IAM** role with **AdministratorAccess** policy.
Go to the AWS console, **IAM**, click on **Roles**. create a **role**.



Select AWS service.

Under **Use case**, select **EC2**.

Click **Next**

Add permissions Info

Permissions policies (1/879) Info

Choose one or more policies to attach to your new role.

Filter by Type

Search All types ▾

| Policy name | Type |
|---|----------------------------|
| <input checked="" type="checkbox"/> AdministratorAccess | AWS managed - job function |

Search for **AdministratorAccess**, select it and choose **Next**.

IAM > Roles > Create role

Step 1 Select trusted entity

Step 2 Add permissions

Step 3 Name, review, and create

Name, review, and create

Role details

Role name
Enter a meaningful name to identify this role.

Maximum 64 characters. Use alphanumeric and '+-=,.@-_' characters.

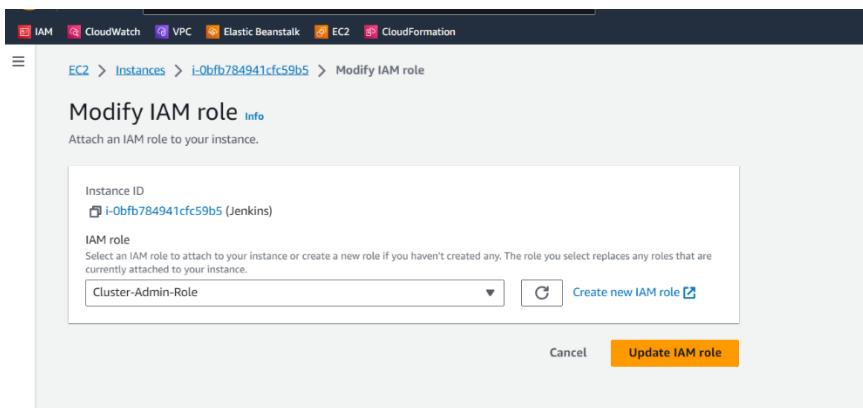
Give it a name such as **Cluster-Admin-Role** and click **Create Role**.

Assign the role to EC2 instance

Go to AWS console, click on **EC2**, select **EC2 instance**, Choose **Security**.

Click on **Modify IAM Role**

| Instance state | Instance type | Actions |
|----------------|---------------|--|
| Stopped | t2.small | Connect View details Manage instance state Instance settings ► Networking ► Security ► Image and templates ► Monitor and troubleshoot ► |
| Stopped | t2.medium | |
| Stopped | t2.small | Change security groups Get Windows password Modify IAM role |
| Stopped | t2.small | |



Choose the role you have created from the dropdown.

Click on **Update IAM role**.

Instances (1/1) [Info](#)

[Find instance by attribute or tag \(case-sensitive\)](#)

| Name | Instance ID |
|---------|--------------------|
| Jenkins | i-0fb784941fcf59b5 |

Instance: i-0fb784941fcf59b5 (Jenkins)

[Details](#) | **Security** | [Networking](#) | [Storage](#) | [Status checks](#) | [Monitoring](#)

Security details

IAM Role: [Cluster-Admin-Role](#)

Owner ID: 642660919026

You can see from above that the IAM role is attached to the Jenkins instance.
Let's got back to the Ubuntu Instance.

Switch to Jenkins user

sudo su - jenkins

Create EKS Cluster with two worker nodes using eksctl

eksctl create cluster --name demo-eks --region us-east-1 --nodegroup-name my-nodes --node-type t3.small --managed --nodes 2

Note: if we didn't explicitly declare the -nodes 2 it will create two nodes by default.

The above command should create an EKS cluster in AWS, it might take 15 to 20 mins.

```
Mac-Silver-Coach:~ devopscoach$ eksctl create cluster --name demo-eks --region us-east-2 --nodegroup-name my-nodes
--node-type t3.small --managed
[1] using region us-east-2
[1] setting availability zones as [us-east-2c us-east-2b us-east-2a]
[1] subnets for us-east-2c - public:192.168.0.0/19 private:192.168.96.0/19
[1] subnets for us-east-2b - public:192.168.32.0/19 private:192.168.128.0/19
[1] subnets for us-east-2a - public:192.168.64.0/19 private:192.168.168.0/19
[1] using Kubernetes version 1.17
[1] creating the cluster "demo-eks" in "us-east-2" region with managed nodes
[1] will create 2 separate CloudFormation stacks for cluster itself and the initial managed nodegroup
[1] if you encounter any issues, check CloudFormation console or try `eksctl utils describe-stacks --region=us-east-2 --cluster=demo-eks`
[1] CloudWatch Logging will not be enabled for cluster "demo-eks" in "us-east-2"
[1] you can enable it with `eksctl utils update-cluster-logging --enable-types=[SPECIFY-YOUR-LOG-TYPES-HERE] (e.g. --regions=us-east-2 --cluster=demo-eks)`
[1] Kubernetes API endpoint access will use default of {publicAccess=true, privateAccess=false} for cluster "demo-eks" in "us-east-2"
[1] 2 sequential tasks: [ 0 tasks, create managed nodegroup "my-nodes" ] 1
[1] building cluster stack "eksctl-demo-eks-cluster"
[1] deploying stack "eksctl-demo-eks-cluster"
[1] building managed nodegroup stack "eksctl-demo-eks-nodegroup-my-nodes"
[1] deploying stack "eksctl-demo-eks-nodegroup-my-nodes"
[1] waiting for the control plane availability...
[1] saved kubeconfig as "/Users/devopscoach/.kube/config"
[1] no tasks
[2] all EKS cluster resources for "demo-eks" have been created
[1] nodegroup "my-nodes" has 2 nodes(s)
[1] node "ip-192-168-3-147.us-east-2.compute.internal" is ready
[1] node "ip-192-168-3-233.us-east-2.compute.internal" is ready
[1] waiting for at least 2 node(s) to become ready in "my-nodes"
[1] nodegroup "my-nodes" has 2 nodes(s)
[1] node "ip-192-168-3-347.us-east-2.compute.internal" is ready
[1] node "ip-192-168-3-233.us-east-2.compute.internal" is ready
[1] kubelet command should work with /Users/devopscoach/.kube/config, try "kubectl get nodes"
[2] the cluster "demo-eks" is now healthy, flag is ready
```

After the cluster build run the following command:
eksctl get cluster --name demo-eks --region us-east-1

This should confirm that EKS cluster is up and running.

Update Kube config by entering below command:

aws eks update-kubeconfig --name demo-eks --region us-east-1

```
jenkins@ip-172-31-18-53:~ aws eks update-kubeconfig --name demo-eks --region us-east-2
added new context: arn:aws:eks:us-east-2:231223701595:cluster/demo-eks to /var/lib/jenkins/.kube/config
```

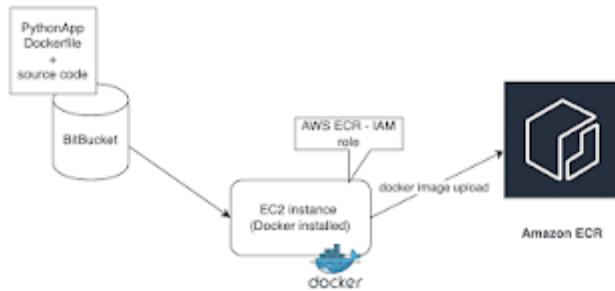
kubeconfig file be updated under the **/var/lib/jenkins/.kube** folder.
 You can view the kubeconfig file by entering the below command:

cat /var/lib/jenkins/.kube/config

Create ECR repo in AWS

Amazon Elastic Container Registry (ECR) uses Amazon **S3** for storage to make your container images highly available and accessible, allowing you to reliably deploy new containers for your applications. Amazon ECR transfers your container images over HTTPS and automatically encrypts your images at rest. Amazon **ECR** is integrated with Amazon **Elastic Container Service (ECS)**, simplifying your development to production workflow.

Create Docker image & Push into Amazon ECR



Pre-requisites:

- **EC2** instance up and running with **Docker** installed.
- Make sure you open port **8081**.
- Install **aws cli**

Create a repo in ECR

Go to AWS console and search for ECR.

The screenshot shows the AWS ECR 'Repositories' page. At the top, there's a 'Create a new repository' button. Below it, there's a search bar labeled 'Find repositories' and a table with columns for 'Repository name', 'URI', 'Created at', 'Tag immutability', and 'Size on path'. The table currently shows one repository entry.

Click on **Create Repository**

The screenshot shows the 'Create Repository' dialog box. It has fields for 'Repository name' (containing 'my-docker-repo'), 'URI', 'Created at', 'Tag immutability', and 'Size on path'. At the bottom right, there is a large red 'Create repository' button.

Enter a name for your repo - all lower case and click create repository.
We chose my-docker-repo

Repository name
287301745136.dkr.ecr.us-east-1.amazonaws.com/**my-docker-repo**
A namespace can be included with your repository name (e.g., namespace/repo-name).

Tag immutability
Enable tag immutability to prevent image tags from being overwritten by subsequent image pushes using the same tag. Disable tag immutability to allow image tags to be overwritten.
 Disabled

Image scan settings
Scan on push
Enable scan on push to have each image automatically scanned after being pushed to a repository. If disabled, each image scan must be manually started to get scan results.
 Disabled

Create repository

Once the repo is created, choose the repo and click on **View push commands**.

Note down the **account ID**
my-docker-repo

Successfully created repository my-docker-repo

Repositories (1 of 1)

| Repository name | URI | Created at |
|-----------------------|---|-------------|
| my-docker-repo | 287301745136.dkr.ecr.us-east-1.amazonaws.com/my-docker-repo | 05/21/20 AM |

View push commands

Note the URL below, as this will be used for tagging and pushing docker images into ECR.

642660919026.dkr.ecr.us-east-1.amazonaws.com/my-docker-repo
642660919026.dkr.ecr.us-east-1.amazonaws.com/my-python-repo

Push commands for mywebapp

[macOS / Linux](#) [Windows](#)

Make sure that you have the latest version of the AWS CLI and Docker installed. For more information, see [Started with Amazon ECR](#).

Use the following steps to authenticate and push an image to your repository. For additional registry authentication methods, including the Amazon ECR credential helper, see [Registry Authentication](#).

1. Retrieve an authentication token and authenticate your Docker client to your registry.
Use the AWS CLI:


```
aws ecr get-login-password --region us-east-2 | docker login --username AWS --password-stdin
```

Note: If you receive an error using the AWS CLI, make sure that you have the latest version of the AWS CLI and Docker installed.
2. Build your Docker image using the following command. For information on building a Docker file from scratch see the instructions [here](#). You can skip this step if your image is already built:


```
docker build -t mywebapp .
```
3. After the build completes, tag your image so you can push the image to this repository:


```
docker tag mywebapp:latest 642660919026.dkr.ecr.us-east-2.amazonaws.com/mywebapp
```
4. Run the following command to push this image to your newly created AWS repository:

Copy the above into a location for later reference.

Make sure that you have the latest version of AWS CLI and Docker installed. For more information, see [Getting Started with Amazon ECR](#).

Use the following steps to authenticate and push an image to your repository. For additional registry authentication methods, including the Amazon ECR credential helper, see [Registry Authentication](#).

1. Retrieve an authentication token and authenticate your Docker client to your registry.
Use the AWS CLI:

```
aws ecr get-login-password --region us-east-1 | docker login --username AWS --password-stdin 642660919026.dkr.ecr.us-east-1.amazonaws.com
```

Note: If you receive an error using the AWS CLI, make sure that you have the latest version of the AWS CLI and Docker installed.

2. Build your Docker image using the following command. For information on building a Docker file from scratch see the instructions [here](#). You can skip this step if your image is already built:

```
docker build -t my-docker-repo .
```

3. After the build completes, tag your image so you can push the image to this repository:

```
docker tag my-docker-repo:latest 642660919026.dkr.ecr.us-east-1.amazonaws.com/my-docker-repo:latest
```

4. Run the following command to push this image to your newly created AWS repository:

```
docker push 642660919026.dkr.ecr.us-east-1.amazonaws.com/my-docker-repo:latest
```

That's it, you have created repo successfully. Let us now create docker images and push it to the above repo in ECR.

Docker installation steps using default repository from Ubuntu

Update local packages by executing below command:

sudo apt updateInstall the below packages**sudo apt install gnupg2 pass -y***gnupg2 is tool for secure communication and data storage. It can be used to encrypt data and to create digital signatures.*Install docker**sudo apt install docker.io -y**Add Ubuntu user to Docker group**sudo usermod -aG docker \$USER**

We need to reload shell to have new group settings applied. Now you need to logout and log back in command line or execute the below command:

newgrp docker

The Docker service needs to be set up to run at startup. To do so, type in each command followed by enter:

sudo systemctl start docker**sudo systemctl enable docker****sudo systemctl status docker**

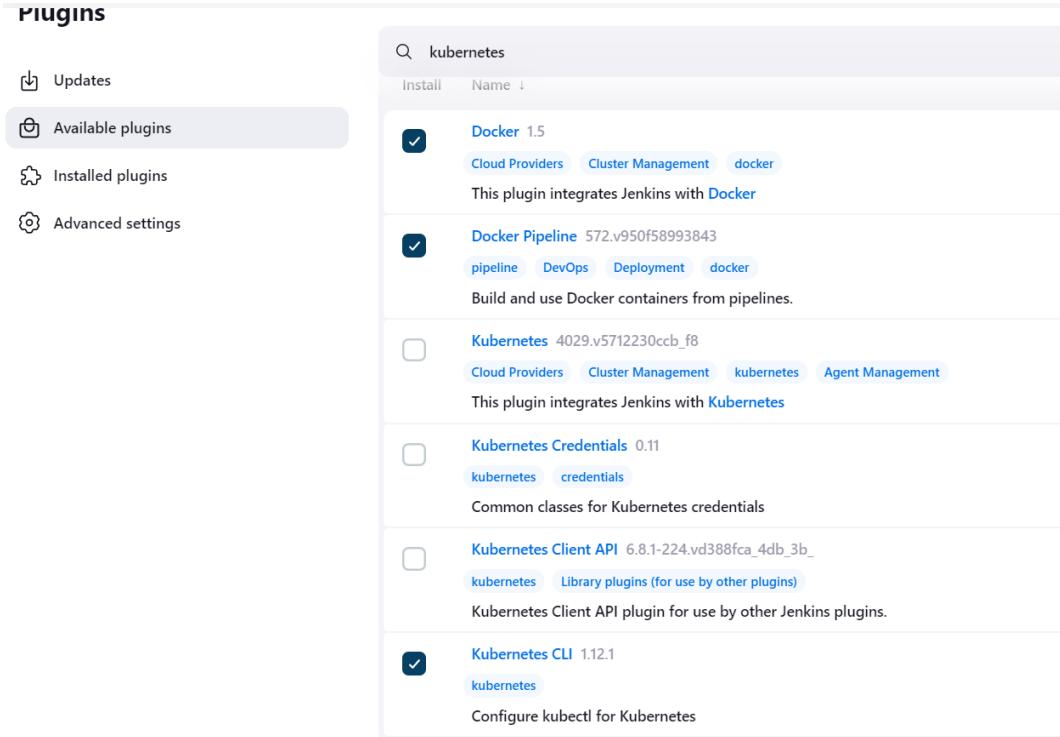
```
ubuntu@ip-172-31-24-229:~$ sudo systemctl status docker
● docker.service - Docker Application Container Engine
   Loaded: loaded (/lib/systemd/system/docker.service; enabled; vendor preset: enabled)
   Active: active (running) since Thu 2020-07-23 02:34:51 UTC; 12s ago
     Docs: https://docs.docker.com
      Main PID: 4299 (dockerd)
        Tasks: 8
       CGroup: /system.slice/docker.service
               └─4299 /usr/bin/dockerd -H fd:// --containerd=/run/containerd/containerd.sock

Jul 23 02:34:51 ip-172-31-24-229 dockerd[4299]: time="2020-07-23T02:34:51.427980999Z" level=info
Jul 23 02:34:51 ip-172-31-24-229 dockerd[4299]: time="2020-07-23T02:34:51.428072171Z" level=info
Jul 23 02:34:51 ip-172-31-24-229 dockerd[4299]: time="2020-07-23T02:34:51.428256519Z" level=info
Jul 23 02:34:51 ip-172-31-24-229 dockerd[4299]: time="2020-07-23T02:34:51.428600893Z" level=info
Jul 23 02:34:51 ip-172-31-24-229 dockerd[4299]: time="2020-07-23T02:34:51.630279701Z" level=info
Jul 23 02:34:51 ip-172-31-24-229 dockerd[4299]: time="2020-07-23T02:34:51.730563362Z" level=info
Jul 23 02:34:51 ip-172-31-24-229 dockerd[4299]: time="2020-07-23T02:34:51.850890001Z" level=info
Jul 23 02:34:51 ip-172-31-24-229 dockerd[4299]: time="2020-07-23T02:34:51.8508213536Z" level=info
Jul 23 02:34:51 ip-172-31-24-229 dockerd[4299]: time="2020-07-23T02:34:51.935841521Z" level=info
```

The above screenshot should confirm that Docker daemon is up and running.

Docker, Docker pipeline and Kubernetes CLI plug-ins are installed in Jenkins.Add jenkins user to Docker group**sudo usermod -a -G docker jenkins**Restart Jenkins service**sudo service jenkins restart**Reload system daemon files**sudo systemctl daemon-reload**Restart Docker service as well**sudo service docker stop****sudo service docker start**

Go to the Jenkins portal
Manage Jenkins > Plugins
Available plugins
Select the following:



The screenshot shows the Jenkins 'Available plugins' section with a search bar containing 'kubernetes'. The results are listed below:

- Docker 1.5** (checked) - This plugin integrates Jenkins with Docker.
- Docker Pipeline 572.v950f58993843** (checked) - Build and use Docker containers from pipelines.
- Kubernetes 4029.v5712230ccb_f8** (unchecked) - This plugin integrates Jenkins with Kubernetes.
- Kubernetes Credentials 0.11** (unchecked) - Common classes for Kubernetes credentials.
- Kubernetes Client API 6.8.1-224.vd388fca_4db_3b_** (unchecked) - Kubernetes Client API plugin for use by other Jenkins plugins.
- Kubernetes CLI 1.12.1** (checked) - Configure kubectl for Kubernetes.

Plugins

- Updates
- Available plugins
- Installed plugins
- Advanced settings
- Download progress

Download progress

| Preparation | |
|--------------------------------------|---|
| Cloud Statistics | ✓ Success |
| Authentication Tokens API | ✓ Success |
| Docker Commons | ✓ Success |
| Apache HttpComponents Client 5.x API | ✓ Success |
| Docker API | ✓ Success |
| Docker | ✓ Success |
| Docker Pipeline | ✓ Success |
| Kubernetes Client API | ✓ Success |
| Kubernetes Credentials | ✓ Success |
| Kubernetes CLI | ✓ Success |
| Loading plugin extensions | ... Running |

→ [Go back to the top page](#)
 (you can start using the installed plugins right away)

→ Restart Jenkins when installation is complete and no jobs are running

Note: do not opt to restart Jenkins

Step # 1 - Create Maven3 variable under Global tool configuration in Jenkins.
Manage Jenkins > Global tool configuration.

Make sure you create a **Maven3** variable under **Global tool configuration**.

Maven installations ^ / Edited

Add Maven

Maven
Name

MAVEN_HOME

Install automatically ?

Add Maven

Click Save

Step #2 - Create Credentials for connecting to Kubernetes Cluster using kubeconfig

Manage Jenkins > Credentials

Click on Add Credentials, use Kubernetes configuration from drop down.

Jenkins > Credentials > System > Global credentials (unrestricted)

[Back to credential domains](#)

[Add Credentials](#)

Global cre

Credentials that should be available:

| | |
|--|--------------------|
| | fa32f95a-2d3e-4c7b |
| | 0c5c8205-0068-4aa |

Icon: S M L

Use secret file from drop down.

System > Global credentials (unrestricted) >

Kind

- Username with password
- GitHub App
- OpenShift Username and Password
- SSH Username with private key
- Secret file**
- Secret text
- X.509 Client Certificate
- Certificate

Let's now go back to our Jenkins server to generate the Secret file.

Execute the below command to login as jenkins user.

sudo su - jenkins

You should see the nodes running in EKS cluster.

kubectl get nodes

```
ubuntu@ip-172-31-32-245:~$ kubectl get nodes
NAME           STATUS   ROLES      AGE     VERSION
ip-192-168-25-216.us-east-2.compute.internal   Ready    <none>    49m    v1.17.11-eks-cfdcc40
ip-192-168-51-213.us-east-2.compute.internal   Ready    <none>    49m    v1.17.11-eks-cfdcc40
```

Execute the below command to get kubeconfig info, copy the entire content of the file:

cat /var/lib/jenkins/.kube/config

Open your text editor or notepad, copy and paste the entire content and save in a file. We will upload this file.

Kind

Secret file

Scope

Global (Jenkins, nodes, items, all ch...)

File

Choose File No file chosen

ID

K8S

Description

OK

Enter ID as **K8S** and Choose File.

Kind
Secret file

Scope
Global (Jenkins, nodes, items, all children)

File
Choose File kuebconfig-jan19.txt

ID
K8S

Description
K8S description

OK

Upload the file and click OK to save.
Make note of the ID **K8S** for the pipeline script later.

Step # 3 - Create a pipeline in Jenkins.

Choose New Item > Pipeline

Enter an item name
Mykube
» Required field

Freestyle project
This is the central feature of Jenkins. Jenkins will build your project, combining any SCM with a build step or something other than software build.

Pipeline
Orchestrates long-running activities that can span multiple build agents. Suitable for building and/or organizing complex activities that do not easily fit in free-style job type.

Multi-configuration project
Suitable for projects that need a large number of different configurations, such as testing on multiple platforms, builds, etc.

Folder
Creates a container that stores nested items in it. Useful for grouping things together. Unlike a folder in your file system, it has its own separate namespace, so you can have multiple things of the same name as long as they are in different folders.

OK **Cancel** **Branch Pipeline**

Call the item **Mykube** and click **OK**.

Configure

Pipeline

Definition

Pipeline script

Script ?

```
1 * pipeline {  
2     agent any  
3  
4     stages {  
5         stage('Hello') {  
6             steps {  
7                 echo 'Hello World'  
8             }  
9         }  
10    }  
11}  
12
```

Hello World

Use Groovy Sandbox ?

Pipeline Syntax

Save Apply

This screenshot shows the Jenkins Pipeline configuration interface. On the left, there's a sidebar with 'Configure' and 'Pipeline' tabs, where 'Pipeline' is selected. The main area is titled 'Definition' and contains a 'Pipeline script' dropdown set to 'Pipeline script'. Below it is a code editor labeled 'Script' with the following Groovy code:

```
1 * pipeline {  
2     agent any  
3  
4     stages {  
5         stage('Hello') {  
6             steps {  
7                 echo 'Hello World'  
8             }  
9         }  
10    }  
11}  
12
```

The code editor has syntax highlighting for Groovy. A dropdown menu next to the code editor shows 'Hello World'. Below the code editor are two buttons: 'Use Groovy Sandbox' (checked) and 'Pipeline Syntax'. At the bottom are 'Save' and 'Apply' buttons.

Select Pipeline > Hello World

Script ?

```
1 * pipeline {  
2     agent any  
3  
4     stages {  
5         stage('Checkout') {  
6             steps {  
7                 echo 'Hello World'  
8             }  
9         }  
10    }  
11}  
12
```

This screenshot shows a portion of the Jenkins Pipeline configuration screen. It displays a Groovy script for a pipeline named 'Checkout'. The script defines a single stage named 'Checkout' which contains a single step that echoes 'Hello World'. The code editor has syntax highlighting for Groovy. The background of the code editor is light gray, while the specific code block is highlighted in a darker shade.

Click Pipeline Syntax

Sample Step

checkout: Check out from version control

checkout ?

SCM

Git

Repositories ?

Repository URL ?

`https://github.com/costas778/springboot-app`

Credentials ?

Select the above options. Ensure that the Repository URL has the URL pointing to the springboot-app as above.

The screenshot shows the Jenkins Pipeline Generator interface. At the top, there's a field labeled "Branch Specifier (blank for 'any') ?" containing the value "/main". Below it is a button "Add Branch". Under "Repository browser ?" is a dropdown set to "(Auto)". In the "Additional Behaviours" section, there are two checked checkboxes: "Include in polling? ?" and "Include in changelog? ?". At the bottom is a blue button labeled "Generate Pipeline Script". Below the button, the generated Jenkins pipeline script is displayed:

```
checkout scmGit(branches: [[name: '/main']], extensions: [], userRemoteConfigs: [[url: 'https://github.com/costas778/springboot-app']])
```

Click the **Generate Pipeline Script.**

Copy the string back into the pipeline script.

Script ?

```

1 * pipeline {
2     agent any
3
4     stages {
5         stage('Checkout') {
6             steps {
7                 checkout scmGit(branches: [[name: '/main']], extensions: [], userRemoteConfigs: [[url: 'https://github.com/costas778/springboot-app']])
```

A code editor window showing a Jenkins pipeline script. The script defines a pipeline with an 'agent any' block, followed by a 'stages' block containing a single 'Checkout' stage. The 'Checkout' stage uses the 'scmGit' step with specific parameters. A dropdown menu is open at the top right of the editor, showing the option "Hello World".

Save the script.

Step # 4 - Copy the pipeline code from below.

Make sure you change yellow highlighted values below as per your settings:

Your docker user id should be updated with the following:

642660919026.dkr.ecr.us-east-1.amazonaws.com/my-docker-repo

Your registry credentials ID from Jenkins we created earlier when we created the Secret file is as follows:

K8S

Finally, change the following **springboot-app** file. For myself this is located as follows:

<https://github.com/costas778/springboot-app>

Make changes in **eks-deploy-k8s.yaml** to pull Docker image from your AWS ECR repo.

main → [springboot-app / eks-deploy-k8s.yaml](#)

akannan1087 changed port no mapping

At 1 contributor

42 lines (42 sloc) | 806 Bytes

```

1 apiVersion: apps/v1
2 kind: Deployment
3 metadata:
4   labels:
5     app: springboot-app
6     name: springboot-app
7 spec:
8   replicas: 2
9   selector:
10    matchLabels:
11      app: springboot-app
12   template:
13     metadata:
14       labels:
15         app: springboot-app
16     spec:
17       containers:
18         - name: my-springboot-app
19           image: dkr.ecr.us-east-2.amazonaws.com/my-docker-repo
20           imagePullPolicy: Always
21         ports:
22           - containerPort: 8085
23

```

main → [springboot-app / eks-deploy-k8s.yaml](#)

akannan1087 changed port no mapping

At 1 contributor

42 lines (42 sloc) | 806 Bytes

```

1 apiVersion: apps/v1
2 kind: Deployment
3 metadata:
4   labels:
5     app: springboot-app
6     name: springboot-app
7 spec:
8   replicas: 2
9   selector:
10    matchLabels:
11      app: springboot-app
12   template:
13     metadata:
14       labels:
15         app: springboot-app
16     spec:
17       containers:
18         - name: my-springboot-app
19           image: 211223789150.dkr.ecr.us-east-2.amazonaws.com/my-docker-repo
20           imagePullPolicy: Always
21         ports:
22

```

Go to Spec. In image, update it with [642660919026.dkr.ecr.us-east-1.amazonaws.com/my-docker-repo](#)

```

ubuntu@Jenkins: ~
jenkins@Jenkins:~$ docker images
REPOSITORY                                     TAG      IMAGE ID      CREATED        SIZE
642660919026.dkr.ecr.us-east-1.amazonaws.com/my-docker-repo  latest   2c042c242c2d  3 minutes ago  285MB
lolhens/baseimage-openjre                      latest   7990e5390b8d  3 years ago   264MB
jenkins@Jenkins:~$
```

```

pipeline {
  tools {
    maven 'Maven3'
  }
}
```

```

        }
    agent any
    environment{
        registry = "account_id.dkr.ecr.us-east-2.amazonaws.com/my-docker-repo"
    }

    stages {
        stage('Cloning Git') {
            steps {
                checkout([$class: 'GitSCM', branches: [[name: '*/main']]],
doGenerateSubmoduleConfigurations: false, extensions: [], submoduleCfg: [], userRemoteConfigs:
[[credentialsId: '', url: 'https://github.com/<your GITHUB username>/springboot-app']]])
            }
        }
        stage ('Build') {
            steps {
                sh 'mvn clean install'
            }
        }
        // Building Docker images
        stage('Building image') {
            steps{
                script {
                    dockerImage = docker.build registry
                }
            }
        }
    }

    // Uploading Docker images into AWS ECR
    stage('Pushing to ECR') {
        steps{
            script {
                sh 'aws ecr get-login-password --region us-east-2 | docker login --username AWS --password-
stdin account_id.dkr.ecr.us-east-2.amazonaws.com'
                sh 'docker push account_id.dkr.ecr.us-east-2.amazonaws.com/my-docker-repo:latest'
            }
        }
    }

    stage('K8S Deploy') {
        steps{
            script {
                withKubeConfig([credentialsId: 'K8S', serverUrl: ""]){
                    sh ('kubectl apply -f eks-deploy-k8s.yaml')
                }
            }
        }
    }
}

```

Once the above script is updated you can overwrite the whole pipeline script.

Step # 5 - Build the pipeline

Once you create the pipeline and changes values per your configuration, click on **Build now:**



Step # 6 - Verify deployments to K8S

kubectl get pods

```
AK-DevOps-Coach:Downloads devopscoachings$ kubectl get pods
NAME           READY   STATUS    RESTARTS   AGE
my-dk-deployment-67984f8cd9-2vl47  1/1     Running   0          28m
```

kubectl get deployments

```
AK-DevOps-Coach:Downloads devopscoachings$ kubectl get deployments
NAME        READY   UP-TO-DATE   AVAILABLE   AGE
my-dk-deployment  1/1      1          1          27m
```

kubectl get services

```
AK-DevOps-Coach:Downloads devopscoachings$ kubectl get svc
NAME         TYPE      CLUSTER-IP       EXTERNAL-IP
kubernetes   ClusterIP  10.100.0.1      <none>
springboot-app   LoadBalancer  10.100.197.254  a7898d8a8449f48afb356e0d205069c2-41286874.us-east-2.elb.amazonaws.com
  80:32179/TCP  39s
```

If you see any errors after deploying the pods, you can check the pod logs.

kubectl logs <pod_name>

```
curl: (6) Could not resolve host: a7898d8a8449f48afb356e0d205069c2-41286874.us-east-2.elb.amazonaws.com
jenkins@Jenkins:~$ kubectl logs springboot-app-5cff7b788f-6sk7f
```

Steps # 7 - Access SpringBoot App in K8S cluster

```
Jenkins@Jenkins:~$ kubectl get deployment
NAME        READY   UP-TO-DATE   AVAILABLE   AGE
springboot-app  2/2      2          2          22s
jenkins@Jenkins:~$ kubectl get pods
NAME           READY   STATUS    RESTARTS   AGE
springboot-app-5cff7b788f-6sk7f  1/1     Running   0          30s
springboot-app-5cff7b788f-9cp79  1/1     Running   0          30s
jenkins@Jenkins:~$ kubectl get svc
NAME         TYPE      CLUSTER-IP       EXTERNAL-IP
  PORT(S)   AGE
kubernetes   ClusterIP  10.100.0.1      <none>
  443/TCP   21m
springboot-app   LoadBalancer  10.100.197.254  a7898d8a8449f48afb356e0d205069c2-41286874.us-east-2.elb.amazonaws.com
  80:32179/TCP  39s
jenkins@Jenkins:~$ 
jenkins@Jenkins:~$ curl a7898d8a8449f48afb356e0d205069c2-41286874.us-east-2.elb.amazonaws.com
curl: (6) Could not resolve host: a7898d8a8449f48afb356e0d205069c2-41286874.us-east-2.elb.amazonaws.com
jenkins@Jenkins:~$ kubectl
```

Once build is successful, go to browser and enter master or worker node public IP address along with port number mentioned above

http://loadbalancer_ip_address

You should see page like below:

My Awesome SpringBoot + Docker App

FirstName:

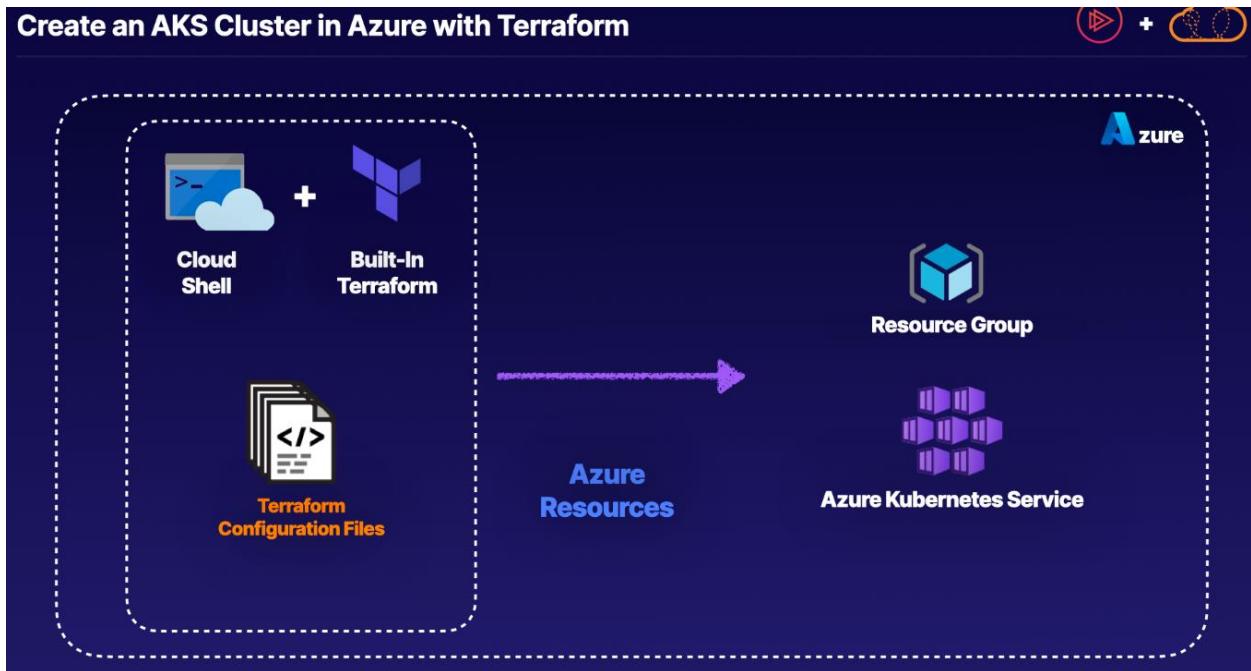
LastName:

`{{postResultMessage}}`

`Get All Customers`

`* {{cust.firstname + " " + cust.lastname}}`

Create an AKS Cluster in Azure with Terraform



Solution

A screenshot of the Azure portal showing a Resource Group named "664-3f7d28bd-create-an-aks-cluster-in-azure-with-t". The portal interface includes a search bar, navigation links (Overview, Activity log, Access control (IAM), Tags), and a toolbar with Create, Manage view, Delete resource group, Refresh, Export to CSV, Open query, and Assign tags. Below the toolbar, detailed information is provided: Subscription (move) : P8-Real Hands-On Labs, Subscription ID : 9734ed68-621d-47ed-babd-269110dbacb1, Deployments : 1 Succeeded, Location : East US, and Tags (edit) : Add tags.

Set Up the Cloud Shell and Lab Environment

In the Azure portal, click on the Cloud Shell icon (`>_`) at the top of the page, to the right of the search bar.

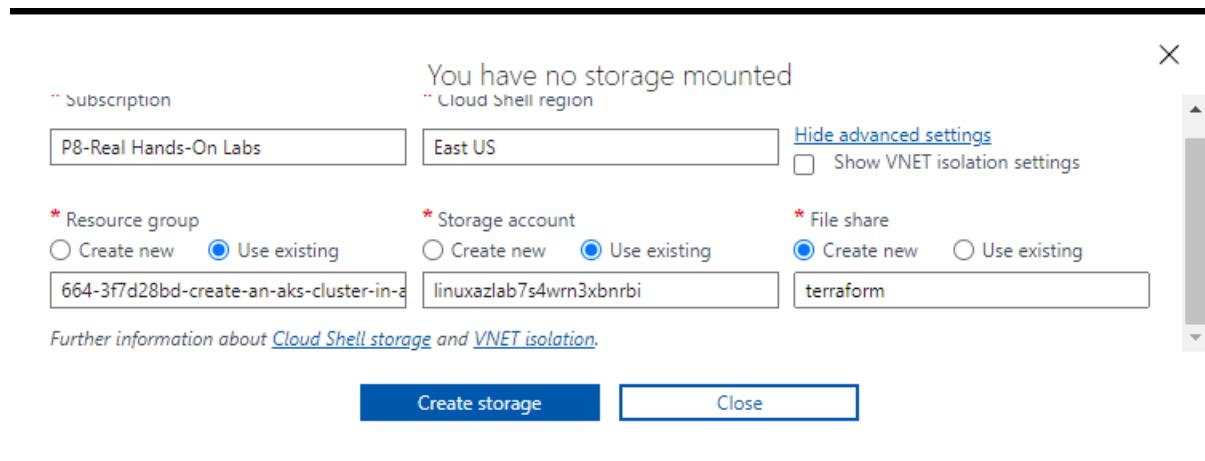
Select **Bash**.

Click **Show advanced settings**.

For the Cloud Shell region, select the **same region** as your **resource group location**.

For **Storage account**, choose **Use existing**.

Under **File share**, select **Create new** and type in the name of **terraform**.



Click **Create storage**. Your Cloud Shell should begin to configure.

Run the following command to pull down the lab setup script from the GitHub repo:

```
 wget https://raw.githubusercontent.com/ACloudGuru/advanced-terraform-with-  
azure/main/lab_aks_cluster/lab_7/setup.sh
```

```
Bash v | ? ⑤ ⑥ ⑦ ⑧ ⑨ ⑩

Requesting a Cloud Shell.Succeeded.
Connecting terminal...

Welcome to Azure Cloud Shell

Type "az" to use Azure CLI
Type "help" to learn about Cloud Shell

cloud [ ~ ]$ wget https://raw.githubusercontent.com/ACloudGuru/advanced-terraform-with-azure/main/lab_aks_cluster/lab_7_setup.sh
--2023-09-21 08:04:42--  https://raw.githubusercontent.com/ACloudGuru/advanced-terraform-with-azure/main/lab_aks_cluster/lab_7_setup.sh
Resolving raw.githubusercontent.com... 185.199.109.133, 185.199.110.133, 185.199.111.133, ...
Connecting to raw.githubusercontent.com[185.199.109.133]:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 218 [text/plain]
Saving to: 'lab_7_setup.sh'

lab_7_setup.sh          100%[=====] 218 --.-KB/s    in 0s

2023-09-21 08:04:42 (9.04 MB/s) - 'lab_7_setup.sh' saved [218/218]

cloud [ ~ ]$
```

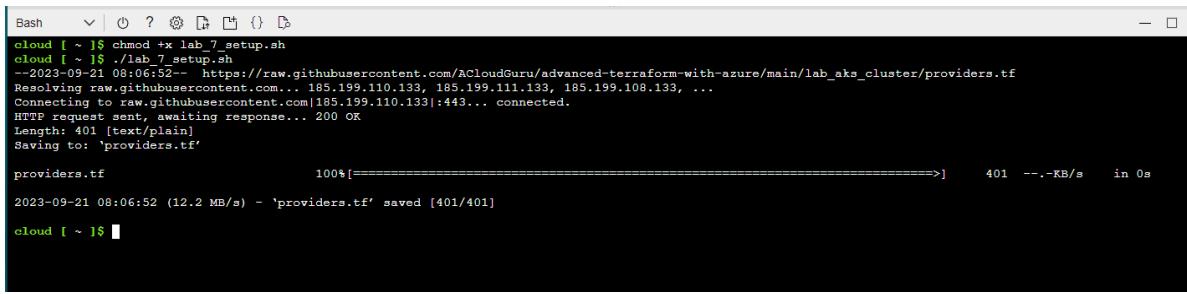
Run **ls** to list the contents.

```
cloud [ ~ ]$ ls  
clouddrive lab_7_setup.sh  
cloud [ ~ ]$ █
```

You should see the lab setup script listed, **lab_7_setup.sh**.

Run **chmod +x lab_7_setup.sh** to make it executable. . .

Run the script, `./lab_7_setup.sh`.



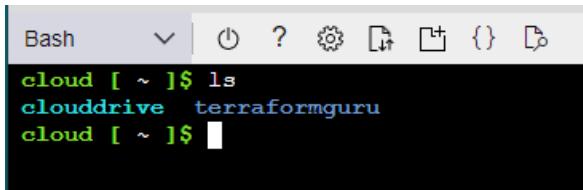
```
Bash      ~ | ⌂ ? ⚙️ 📁 ⓘ ⓘ
cloud [ ~ ]$ chmod +x lab_7_setup.sh
cloud [ ~ ]$ ./lab_7_setup.sh
--2023-09-21 08:06:52-- https://raw.githubusercontent.com/ACloudGuru/advanced-terraform-with-azure/main/lab_aks_cluster/providers.tf
Resolving raw.githubusercontent.com... 185.199.110.133, 185.199.111.133, 185.199.108.133, ...
Connecting to raw.githubusercontent.com[185.199.110.133]:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 401 [text/plain]
Saving to: 'providers.tf'

providers.tf          100%[=====]   401 --.-KB/s    in 0s

2023-09-21 08:06:52 (12.2 MB/s) - 'providers.tf' saved [401/401]

cloud [ ~ ]$
```

Run **ls** to list the contents.

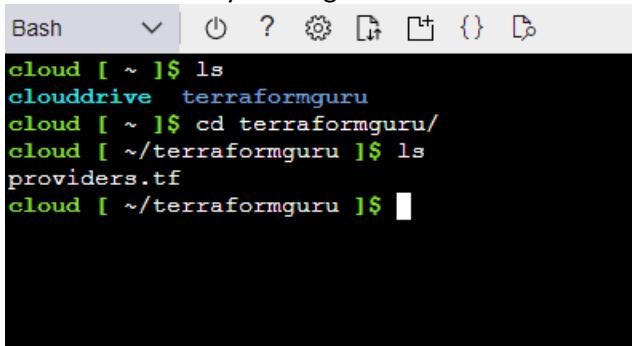


```
Bash      ~ | ⌂ ? ⚙️ 📁 ⓘ ⓘ
cloud [ ~ ]$ ls
clouddrive  terraformguru
cloud [ ~ ]$
```

You should see a **terraformguru** directory listed.

Run **cd terraformguru/** to change into that directory.

List the contents by running **ls**.



```
Bash      ~ | ⌂ ? ⚙️ 📁 ⓘ ⓘ
cloud [ ~ ]$ ls
clouddrive  terraformguru
cloud [ ~ ]$ cd terraformguru/
cloud [ ~/terraformguru ]$ ls
providers.tf
cloud [ ~/terraformguru ]$
```

You should see one configuration file listed: **providers.tf**.

Run **vim providers.tf** to look at the file.

```

Bash    ▾ | ⌂ ? ⌂ { } ⌂

terraform {

  required_version = ">=0.12"

  required_providers {
    azurerm = {
      source = "hashicorp/azurerm"
      version = "~>2.0"
    }
  }

  provider "azurerm" {
    features {}
    skip_provider_registration = true
  }

  # Import the Azure Resource Group
  resource "azurerm_resource_group" "k8s" {
    name      = "<RESOURCE_GROUP_NAME>"
    location = "<RESOURCE_GROUP_LOCATION>"
  }
}

```

Type **Esc** then **:q** to quit out of the file.

Import the Resource Group

Run **terraform init** to initialize the working directory.

```

Bash    ▾ | ⌂ ? ⌂ { } ⌂

Initializing the backend...

Initializing provider plugins...
- Finding hashicorp/azurerm versions matching "~> 2.0"...
- Installing hashicorp/azurerm v2.99.0...
- Installed hashicorp/azurerm v2.99.0 (signed by HashiCorp)

Terraform has created a lock file .terraform.lock.hcl to record the provider
selections it made above. Include this file in your version control repository
so that Terraform can guarantee to make the same selections by default when
you run "terraform init" in the future.

Terraform has been successfully initialized!

You may now begin working with Terraform. Try running "terraform plan" to see
any changes that are required for your infrastructure. All Terraform commands
should now work.

If you ever set or change modules or backend configuration for Terraform,
rerun this command to reinitialize your working directory. If you forget, other
commands will detect it and remind you to do so if necessary.
cloud [ ~/terraformguru ]$ 

```

Run **az group list** to look up the subscription ID.

```
Bash ~/?/cloud ~/terraformguru ]$ az group list
[
  {
    "id": "/subscriptions/9734ed68-621d-47ed-babd-269110dbacb1/resourceGroups/664-3f7d28bd-create-an-aks-cluster-in-azure-with-t",
    "location": "eastus",
    "managedBy": null,
    "name": "664-3f7d28bd-create-an-aks-cluster-in-azure-with-t",
    "properties": {
      "provisioningState": "Succeeded"
    },
    "tags": null,
    "type": "Microsoft.Resources/resourceGroups"
  }
]
cloud [ ~/terraformguru ]$
```

Copy the **subscription ID** to your clipboard. It should be located on the top line after "id":
Make sure to copy all of the characters in between the quotation marks.

/subscriptions/9734ed68-621d-47ed-babd-269110dbacb1/resourceGroups/664-3f7d28bd-create-an-aks-cluster-in-azure-with-t

Run the following command, making sure to paste in your copied subscription ID to replace <SUBSCRIPTION_ID>:

```
terraform import azurerm_resource_group.k8s <SUBSCRIPTION_ID>
terraform import azurerm_resource_group.k8s /subscriptions/9734ed68-621d-47ed-babd-269110dbacb1/resourceGroups/664-3f7d28bd-create-an-aks-cluster-in-azure-with-t
```

Note: It may take a minute to import your resource.

```
Bash ~/?/cloud ~/terraformguru ]$ terraform import azurerm_resource_group.k8s /subscriptions/9734ed68-621d-47ed-babd-269110dbacb1/resourceGroups/664-3f7d28bd-create-an-aks-cluster-in-azure-with-t...
azurerm_resource_group.k8s: Importing from ID "/subscriptions/9734ed68-621d-47ed-babd-269110dbacb1/resourceGroups/664-3f7d28bd-create-an-aks-cluster-in-azure-with-t"...
azurerm_resource_group.k8s: Import prepared!
azurerm_resource_group.k8s: Prepared azurerm_resource_group for import
azurerm_resource_group.k8s: Refreshing state... [id=/subscriptions/9734ed68-621d-47ed-babd-269110dbacb1/resourceGroups/664-3f7d28bd-create-an-aks-cluster-in-azure-with-t]
azurerm_resource_group.k8s: Import successful!

The resources that were imported are shown above. These resources are now in your Terraform state and will henceforth be managed by Terraform.

cloud [ ~/terraformguru ]$
```

Run **vim providers.tf** to edit the file.

Delete the comment hashes (#) in front of name and location.

Replace the placeholder <RESOURCE_GROUP_NAME> next to name.

Copy the resource group name located at the top left of the Azure portal, under Home.

Paste it into the file, to replace <RESOURCE_GROUP_NAME> making sure not to replace the quotation marks.

Replace the placeholder <RESOURCE_GROUP_LOCATION> next to location.

Copy the resource group location listed to the right of Location in the Azure portal.

(If you hover over it, a copy icon should appear that you can click to copy it to your clipboard.)

Paste it into the file, to replace <RESOURCE_GROUP_LOCATION> making sure not to replace the quotation marks.

Type Esc followed by :wq to save and quit the file.

```
Bash    ? { } 
cloud [ ~/terraformguru ]$ ls
providers.tf  terraform.tfstate
cloud [ ~/terraformguru ]$ vi providers.tf
cloud [ ~/terraformguru ]$ terraform fmt
providers.tf
cloud [ ~/terraformguru ]$ 
```

Run the following command to create an SSH key:

ssh-keygen -m PEM -t rsa -b 4096

Hit **Enter** to keep the defaults.

Hit **Enter** to leave the passphrase empty.

Hit **Enter** again to create your key pair.

```
Bash    ? { } 
cloud [ ~/terraformguru ]$ ssh-keygen -m PEM -t rsa -b 4096
Generating public/private rsa key pair.
Enter file in which to save the key (/home/cloud/.ssh/id_rsa):
Created directory '/home/cloud/.ssh'.
Enter passphrase (empty for no passphrase):
Enter same passphrase again:
Your identification has been saved in /home/cloud/.ssh/id_rsa
Your public key has been saved in /home/cloud/.ssh/id_rsa.pub
The key fingerprint is:
SHA256:5+adwHRxI5iUhX2iHYbDjc3jEwPMismI04BBzdiyLU cloud@cc-e0af4575-577b44cfcc-d5f21
The key's randomart image is:
+---[RSA 4096]---+
|...*+= o.=. |
|..o +.o .O.Bo |
| Eo o o.= O++ |
| o . +.oo=o=.. |
| . So ..o. |
| . . . |
| . o . |
| . o... |
| ooo. |
+---[SHA256]---+
cloud [ ~/terraformguru ]$ 
```

Add the AKS Config, Variables, and Outputs to the Configuration:

Run **vim aks.tf** to create your first configuration file.

Enter the following configuration:

```
resource "azurerm_kubernetes_cluster" "k8s" {
  name        = var.cluster_name
  location    = azurerm_resource_group.k8s.location
  resource_group_name = azurerm_resource_group.k8s.name
  dns_prefix   = var.dns_prefix

  linux_profile {
    admin_username = "ubuntu"
    ssh_key {
      key_data = file(var.ssh_public_key)
    }
  }
}
```

```

default_node_pool {
  name      = "agentpool"
  node_count = var.agent_count
  vm_size    = "Standard_D2s_v3"
  os_disk_size_gb = 30
}

service_principal {
  client_id   = var.aks_service_principal_app_id
  client_secret = var.aks_service_principal_client_secret
}

network_profile {
  load_balancer_sku = "Standard"
  network_plugin = "kubenet"
}

tags = {
  Environment = "Development"
}
}

```

Type **Esc** followed by **:wq** to save and quit the file.

Run **vim variables.tf** to create your next configuration file.

Enter the following configuration. Be sure to replace **<YOUR_RESOURCE_GROUP_LOCATION>** with the location of your resource group, and replace **<SERVICE_PRINCIPAL_APP_ID>** and **<SERVICE_PRINCIPAL_CLIENT_SECRET>** with the service principal IDs generated for this lab, which can be found in the lab credentials section.

```

variable "resource_group_location" {
  default = "<YOUR_RESOURCE_GROUP_LOCATION>"
}

variable "agent_count" {
  default = 3
}

variable "ssh_public_key" {
  default = "~/.ssh/id_rsa.pub"
}

variable "dns_prefix" {
  default = "k8sguru"
}

variable cluster_name {

```

```

    default = "k8sguru"
}

variable aks_service_principal_app_id {
    default = "<SERVICE_PRINCIPAL_APP_ID>"
}

variable aks_service_principal_client_secret {
    default = "<SERVICE_PRINCIPAL_CLIENT_SECRET>"
}

```

Type **Esc** followed by **:wq** to save and quit the file.

Run **vim output.tf** to create your final configuration file.

Enter the following configuration:

```

output "resource_group_name" {
    value = azurerm_resource_group.k8s.name
}

output "client_key" {
    value = azurerm_kubernetes_cluster.k8s.kube_config.0.client_key
}

output "client_certificate" {
    value = azurerm_kubernetes_cluster.k8s.kube_config.0.client_certificate
}

output "cluster_ca_certificate" {
    value = azurerm_kubernetes_cluster.k8s.kube_config.0.cluster_ca_certificate
}

output "cluster_username" {
    value = azurerm_kubernetes_cluster.k8s.kube_config.0.username
}

output "cluster_password" {
    value = azurerm_kubernetes_cluster.k8s.kube_config.0.password
}
output "kube_config" {
    value = azurerm_kubernetes_cluster.k8s.kube_config_raw
    sensitive = true
}
output "host" {
    value = azurerm_kubernetes_cluster.k8s.kube_config.0.host
}

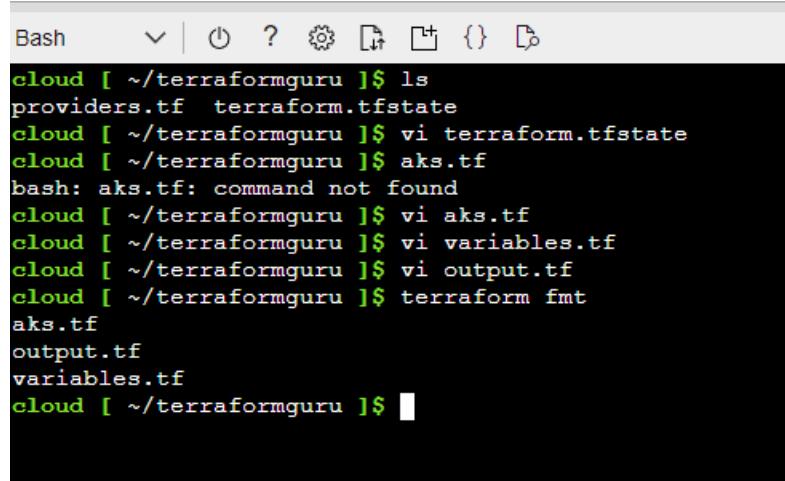
```

```
}
```

Type **Esc** followed by **:wq** to save and quit the file.

Deploy and Verify the Kubernetes Cluster is Running

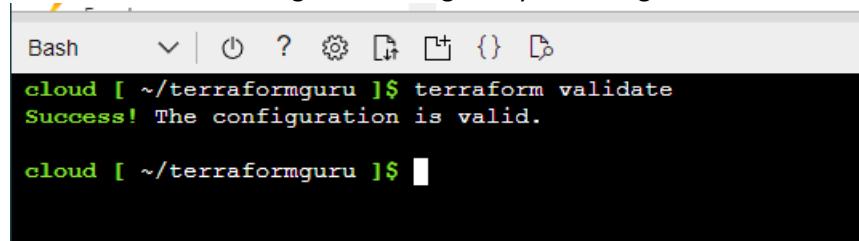
Run **terraform fmt** to check the formatting of your configuration files.
 Your aks.tf, output.tf, providers.tf, and variables.tf files should be listed.



```
Bash      ▾ | ⌁ ? ⚙️ ⌂ ⌃ ⌄ ⌅ ⌆
cloud [ ~/terraformguru ]$ ls
providers.tf  terraform.tfstate
cloud [ ~/terraformguru ]$ vi terraform.tfstate
cloud [ ~/terraformguru ]$ aks.tf
bash: aks.tf: command not found
cloud [ ~/terraformguru ]$ vi aks.tf
cloud [ ~/terraformguru ]$ vi variables.tf
cloud [ ~/terraformguru ]$ vi output.tf
cloud [ ~/terraformguru ]$ terraform fmt
aks.tf
output.tf
variables.tf
cloud [ ~/terraformguru ]$
```

Run **terraform validate** to validate the code in your configuration files.

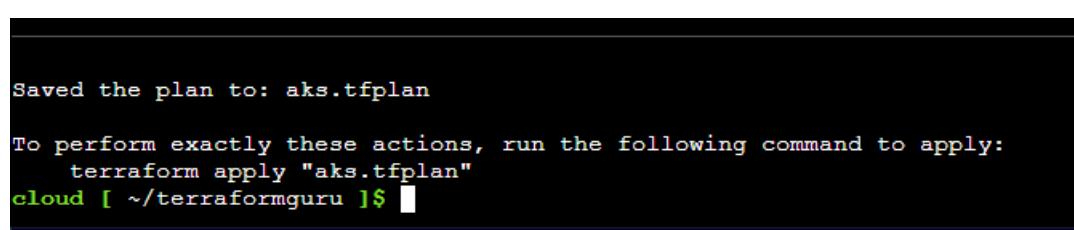
You should see a message confirming that your configuration is valid.



```
Bash      ▾ | ⌁ ? ⚙️ ⌂ ⌃ ⌄ ⌅ ⌆
cloud [ ~/terraformguru ]$ terraform validate
Success! The configuration is valid.

cloud [ ~/terraformguru ]$
```

Run **terraform plan -out aks.tfplan** to create your execution plan.



```
Saved the plan to: aks.tfplan

To perform exactly these actions, run the following command to apply:
  terraform apply "aks.tfplan"
cloud [ ~/terraformguru ]$
```

Run **terraform apply aks.tfplan** to execute your execution plan.

Note: It may take a couple of minutes to deploy your resources.

```

Bash    v | ⌁ ? ⌂ ⌄ ⌅ ⌆ ⌇ ⌈ ⌉ ⌊ ⌋
cloud [ ~/terraformguru ]$ ls
aks.tf  aks.tfpl  aks.tfplan  output.tf  providers.tf  terraform.tfstate  variables.tf
cloud [ ~/terraformguru ]$ terraform apply aks.tfplan
azurerm_kubernetes_cluster.k8s: Creating...
azurerm_kubernetes_cluster.k8s: Still creating... [10s elapsed]
azurerm_kubernetes_cluster.k8s: Still creating... [20s elapsed]
azurerm_kubernetes_cluster.k8s: Still creating... [30s elapsed]
azurerm_kubernetes_cluster.k8s: Still creating... [40s elapsed]
azurerm_kubernetes_cluster.k8s: Still creating... [50s elapsed]
azurerm_kubernetes_cluster.k8s: Still creating... [1m0s elapsed]
azurerm_kubernetes_cluster.k8s: Still creating... [1m10s elapsed]
azurerm_kubernetes_cluster.k8s: Still creating... [1m20s elapsed]
azurerm_kubernetes_cluster.k8s: Still creating... [1m30s elapsed]
azurerm_kubernetes_cluster.k8s: Still creating... [1m40s elapsed]
azurerm_kubernetes_cluster.k8s: Still creating... [1m50s elapsed]
azurerm_kubernetes_cluster.k8s: Still creating... [2m0s elapsed]
azurerm_kubernetes_cluster.k8s: Still creating... [2m10s elapsed]
azurerm_kubernetes_cluster.k8s: Still creating... [2m20s elapsed]

```

You will see a big block of text appear, which should mean that your cluster deployed successfully. You can scroll up to view the Apply complete message in green to confirm.

Scrolling down from the Apply complete message, you can view the client_certificate, client_key, cluster_ca_certificate, cluster_password, and cluster_username. Lastly, you should see the host address, kube_config, and resource_group_name.

```

azurerm_kubernetes_cluster.k8s: Still creating... [3m0s elapsed]
azurerm_kubernetes_cluster.k8s: Creation complete after 3m8s [id=/subscriptions/cluster-in-azure-with-t/providers/Microsoft.ContainerService/managedClusters/aks]
Apply complete! Resources: 1 added, 0 changed, 0 destroyed.

Outputs:

```

Run the following command to move your kube_config to a different file:

`echo "$(terraform output kube_config)" > ./azurek8s`

Run `cat ./azurek8s` to check the file.

You should see EOT at the beginning and end of the file, which will need to be removed.

Run `vim ./azurek8s` to edit the file.

Delete the <<EOT at the beginning and the EOT at the end of the file.

Type `Esc` followed by `:wq` to save and quit the file.

Run `export KUBECONFIG=./azurek8s` to create your environment variable.

Run `kubectl get nodes` to check if your nodes are running and healthy.

You should see your 3 nodes returned with a **STATUS** of Ready.

Go to the Kubernetes Services within the Azure portal to see the same 3 nodes with a **STATUS** of Ready as well.

k8sguru | Node pools

Kubernetes service

Microsoft Azure Search resources, services, and docs (G+)

Home > Kubernetes services > 664-3f7d28bd-create-an-aks-cluster-in-azure-with-t > k8sguru

Add node pool Refresh Start Stop Upgrade Kubernetes Update image Scale node pool Delete Troubleshoot Give feedback

Node pools Nodes

Node pools provide space for applications to run. Node pools of different types can be added to the cluster to handle a variety of workloads, existing node pools can be scaled and upgraded, or node pools that are no longer needed can be deleted. Each node pool will contain nodes backed by virtual machines. [Learn more about node pools](#)

| Node pool | Provisioning state | Power state | Node count | Mode | Kubernetes version | Node size | Operating system |
|-----------|--------------------|-------------|------------|--------|--------------------|-----------------|------------------|
| agentpool | Succeeded | Running | 3/3 ready | System | 1.26.6 | Standard_D2s_v3 | Linux |

Home > Kubernetes services > 664-3f7d28bd-create-an-aks-cluster-in-azure-with-t > k8sguru

k8sguru | Node pools

Kubernetes service

Microsoft Azure Search resources, services, and docs (G+)

Add node pool Refresh Show taints Troubleshoot Give feedback

about node pools

Filter by node pool name Enter the full node pool name

| Node | Status | CPU | Memory | Disk | Pods | Node pool | Kubernetes |
|-----------------------------------|--------|-----|--------|------|------|-----------|------------|
| aks-agentpool-29165910-vmss000000 | Ready | 10% | 20% | 72% | 9 | agentpool | 1.26.6 |
| aks-agentpool-29165910-vmss000001 | Ready | 11% | 17% | 70% | 5 | agentpool | 1.26.6 |
| aks-agentpool-29165910-vmss000002 | Ready | 9% | 20% | 73% | 8 | agentpool | 1.26.6 |

Microsoft Azure Search Kubernetes resources

Namespaces

Workloads Services and ingresses Storage Configuration Custom resources Run command

Create Delete Refresh Show labels Give feedback

Filter by namespace name Enter the full namespace name Add label filter

| Name | Status | Age |
|-----------------|--------|------------|
| kube-node-lease | Active | 16 minutes |
| kube-public | Active | 16 minutes |
| kube-system | Active | 16 minutes |
| default | Active | 16 minutes |

Addendum

```
cloud [ ~/terraformguru ]$ kubectl get namespaces
NAME      STATUS   AGE
default   Active  13m
kube-node-lease  Active  13m
kube-public  Active  13m
kube-system  Active  13m
cloud [ ~/terraformguru ]$ kubectl get pods
No resources found in default namespace.
cloud [ ~/terraformguru ]$
```

To check the current default namespace, run the following command:

kubectl config view | grep namespace

To set a new default namespace, run the following command:

kubectl config set-context --current --namespace=NAMESPACE_NAME
kubectl config set-context --current --namespace=kube-node-lease

To verify that the new default namespace has been set, run the following command:

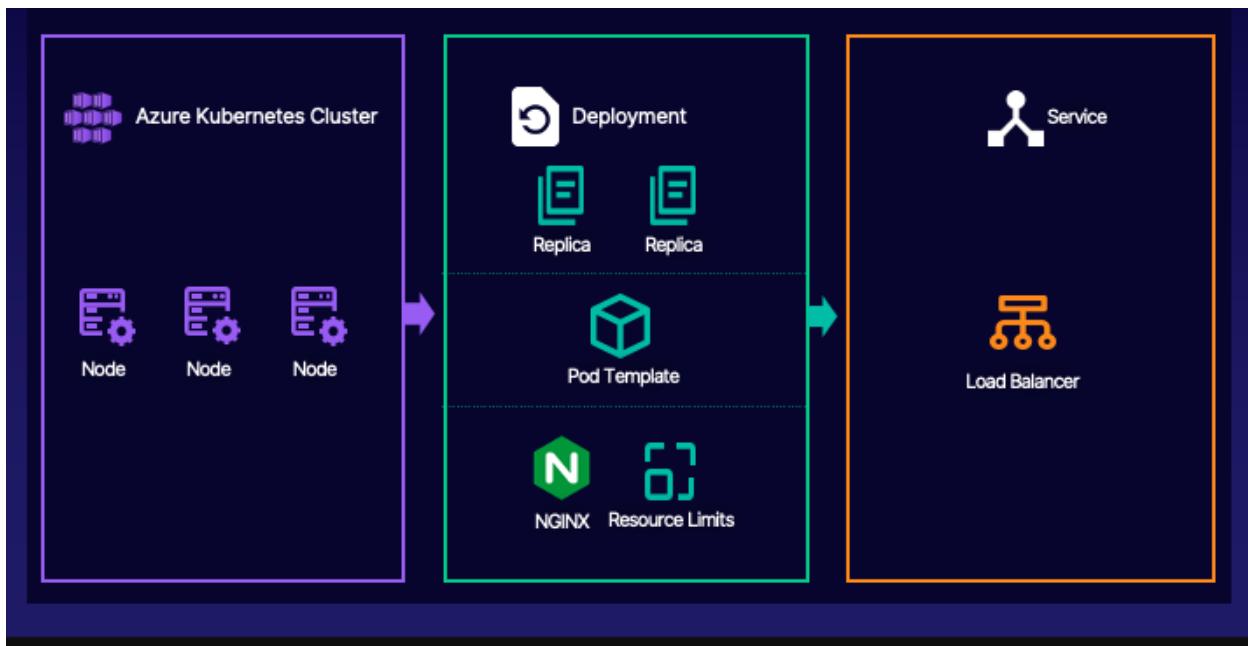
kubectl config view | grep namespace

Change Namespace for a Single Command:

If you want to use a different namespace for a single command, you can specify it using the "--namespace" flag. For example, to list all the pods in a different namespace, run the following command:

kubectl get pods --namespace=NAMESPACE_NAME
kubectl get pods --namespace=KUBE-NODE-LEASE

Deploying and Accessing an Application to an AKS Cluster



Introduction

We will work with the Azure Command Line Interface (CLI) to create an AKS cluster. We will then use the `kubectl` command line utility to deploy an application to that cluster, and then add a kubernetes service, so that the application is accessible over the internet.

Solution

Log in to the Azure portal using the credentials provided on the lab instructions page.

Create AKS Cluster

In the Azure Portal, click the Cloud Shell icon (`>_`) in the upper right.

Select **Bash**.

Click **Show advanced settings**.

Set the **Cloud Shell region** to the same location as the existing resource group.

For **Storage account**, select **Create new** and give it a globally unique name (e.g., "cloudshell" with a series of numbers at the end).

For **File share**, select **Create new** and give it a name of "**fileshare1**".

Click **Create storage**.

Locate the resource group name at the top of the webpage and copy it for later use:

491-cdc3bb61-deploying-and-accessing-an-applicatio

az group list

```
cloud [ ~]$ az group list
[
  {
    "id": "/subscriptions/0fce2870-d256-4119-b0a3-16293ac11bdc/resourceGroups/491-cdc3bb61-deploying-and-accessing-an-applicatio",
    "location": "eastus",
    "managedBy": null,
    "name": "491-cdc3bb61-deploying-and-accessing-an-applicatio",
    "properties": {
      "provisioningState": "Succeeded"
    },
    "tags": null,
    "type": "Microsoft.Resources/resourceGroups"
  }
]
cloud [ ~]$
```

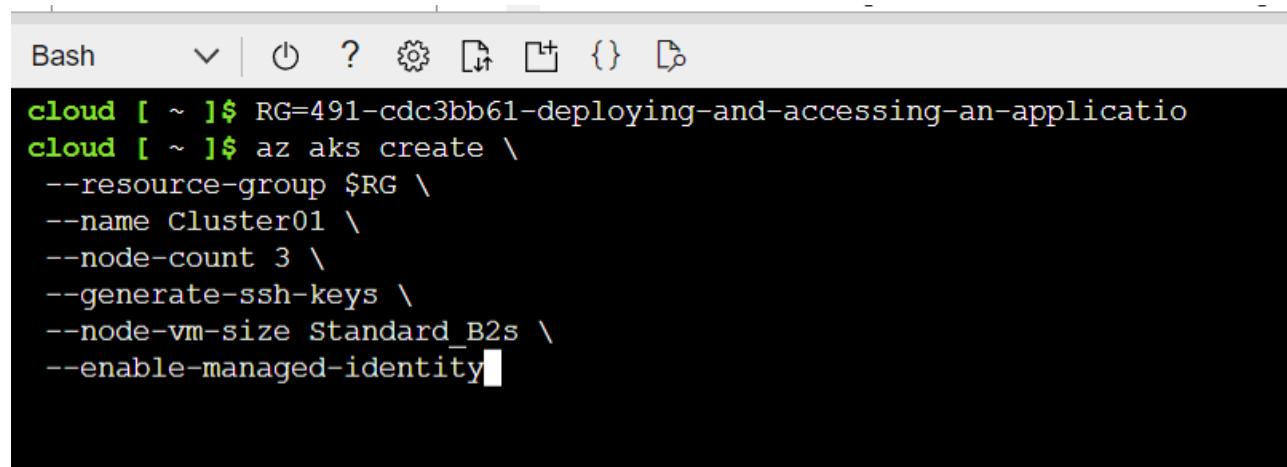
Create a variable for the resource group name:

```
RG=<RESOURCE_GROUP_NAME>
```

```
RG=491-cdc3bb61-deploying-and-accessing-an-applicatio
```

Create an AKS cluster:

```
az aks create \
--resource-group $RG \
--name Cluster01 \
--node-count 3 \
--generate-ssh-keys \
--node-vm-size Standard_B2s \
--enable-managed-identity
```



A screenshot of the Azure Cloud Shell interface. The title bar says "Bash". Below it is a toolbar with icons for copy, paste, refresh, and others. The main area shows a command-line session:

```
cloud [ ~ ]$ RG=491-cdc3bb61-deploying-and-accessing-an-applicatio
cloud [ ~ ]$ az aks create \
--resource-group $RG \
--name Cluster01 \
--node-count 3 \
--generate-ssh-keys \
--node-vm-size Standard_B2s \
--enable-managed-identity
```

```
--enable-managed-identity
SSH key files '/home/cloud/.ssh/id_rsa' and '/home/cloud/.ssh/id_rsa.pub' have been generated under ~/.ssh to allow SSH access to the VM. If using machines without permanent storage like Azure Cloud Shell without an attached file share, back up your keys to a safe location
[]\ Running ..
```

In the Azure Portal, under **All services**, navigate to **Kubernetes services**, and verify that the cluster was deployed successfully.

Back in the Cloud Shell, configure kubectl so that commands can be run against the cluster:

```
az aks get-credentials --name Cluster01 --resource-group $RG
```

```
Read more about the command in reference docs
cloud [ ~ ]$ RG=491-cdc3bb61-deploying-and-accessing-an-applicatio
cloud [ ~ ]$ az aks get-credentials --name Cluster01 --resource-group $RG
Merged "Cluster01" as current context in /home/cloud/.kube/config
cloud [ ~ ]$ []
```

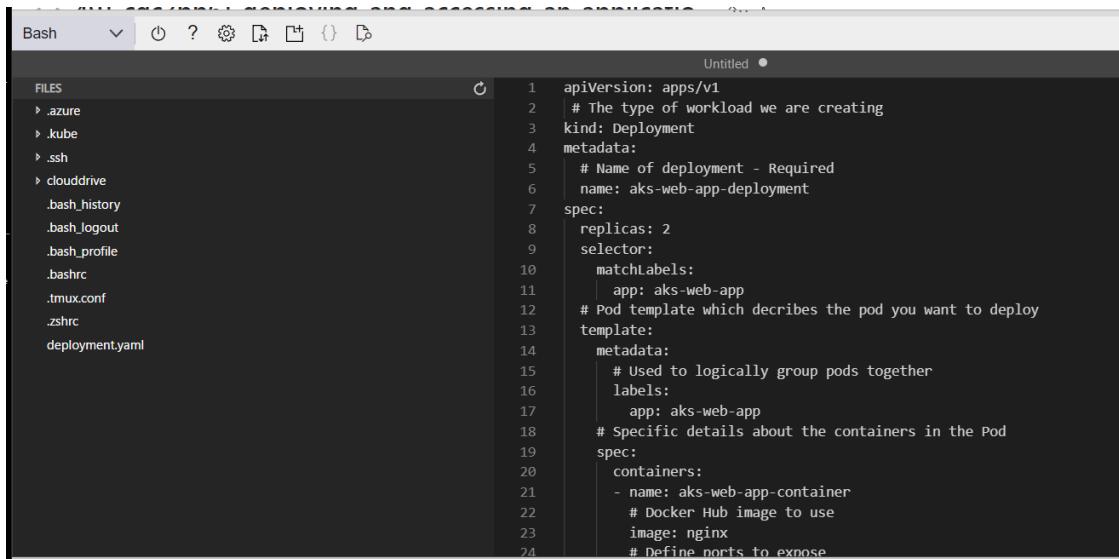
Deploy the Application

Create a new deployment manifest file:

touch deployment.yaml

Open the editor by clicking the curly braces {}, and open the **deployment.yaml** file.

Copy and paste the contents of the **deployment.yaml** file from the **GitHub** repo, into the new file.



```

apiVersion: apps/v1
# The type of workload we are creating
kind: Deployment
metadata:
  # Name of deployment - Required
  name: aks-web-app-deployment
spec:
  replicas: 2
  selector:
    matchLabels:
      app: aks-web-app
  # Pod template which describes the pod you want to deploy
  template:
    metadata:
      # Used to logically group pods together
      labels:
        app: aks-web-app
    # Specific details about the containers in the Pod
    spec:
      containers:
        - name: aks-web-app-container
          # Docker Hub image to use
          image: nginx
          # Define ports to expose

```

Click the menu icon in the top right (...), and click Close Editor.

Make sure to save the file. Use **ctrl + S**

Name the file **deployment.yaml**.

Deploy the application:

kubectl apply -f ./deployment.yaml

In the Azure Portal, under Kubernetes services, open **Cluster01**.

In the left-hand menu, under **Kubernetes Services**, click **Workloads**.

Home > Kubernetes services > Cluster01

| Name | Namespace | Ready | Up-to-date | Available |
|------------------------|-------------------|-------|------------|-----------|
| coredns | kube-system | ✓ 2/2 | 2 | 2 |
| coredns-autoscaler | kube-system | ✓ 1/1 | 1 | 1 |
| konnectivity-agent | kube-system | ✓ 2/2 | 2 | 2 |
| metrics-server | kube-system | ✓ 2/2 | 2 | 2 |
| azure-policy | kube-system | ✓ 1/1 | 1 | 1 |
| azure-policy-webhook | kube-system | ✓ 1/1 | 1 | 1 |
| gatekeeper-audit | gatekeeper-system | ✓ 1/1 | 1 | 1 |
| gatekeeper-controller | gatekeeper-system | ✓ 2/2 | 2 | 2 |
| aks-web-app-deployment | default | ✓ 2/2 | 2 | 2 |

In the **Deployments** tab, open the new deployment, and verify that the pods are running.

Access the Application Externally

Create a new file named **service.yaml**:

touch service.yaml

Open the editor by clicking the curly braces {}, and open the **service.yaml** file.

Copy and paste the contents of the **service.yaml** file from the **GitHub** repo, into the new file.

Click the menu icon in the top right (...), and click Close Editor.

Make sure to save the file. Use **ctrl + S**

Name the file **service.yaml**

Deploy the LoadBalancer service:

kubectl apply -f service.yaml

```

Requesting a Cloud Shell.Succeeded.
Connecting terminal...

cloud [ ~ ]$ kubectl apply -f service.yaml
service/aks-web-app-service created
cloud [ ~ ]$ █

```

Back in the **Azure Portal**, go to back to **Kubernetes services**. In **Cluster01**, click **Services and ingresses**.

Open the newly deployed **LoadBalancer**.

| <input type="checkbox"/> | Name | Namespace | Status | Type | Cluster IP | External IP | Ports |
|--------------------------|----------------------------|-------------------|-----------------|--------------|--------------|--|---------------|
| <input type="checkbox"/> | kubernetes | default | Ok | ClusterIP | 10.0.0.1 | | 443/TCP |
| <input type="checkbox"/> | kube-dns | kube-system | Ok | ClusterIP | 10.0.0.10 | | 53/UDP,53/TCP |
| <input type="checkbox"/> | metrics-server | kube-system | Ok | ClusterIP | 10.0.193.47 | | 443/TCP |
| <input type="checkbox"/> | azure-policy-webhook-s... | kube-system | Ok | ClusterIP | 10.0.218.42 | | 443/TCP |
| <input type="checkbox"/> | gatekeeper-webhook-se... | gatekeeper-system | Ok | ClusterIP | 10.0.199.195 | | 443/TCP |
| <input type="checkbox"/> | aks-web-app-service | default | Ok | LoadBalancer | 10.0.201.164 | 20.185.101.46 | 80:32656/TCP |

Click the External IP and verify access to the application.

▲ Not secure | <http://20.185.101.46>

Welcome to nginx!

If you see this page, the nginx web server is successfully installed and working. Further configuration is required.

For online documentation and support please refer to nginx.org. Commercial support is available at nginx.com.

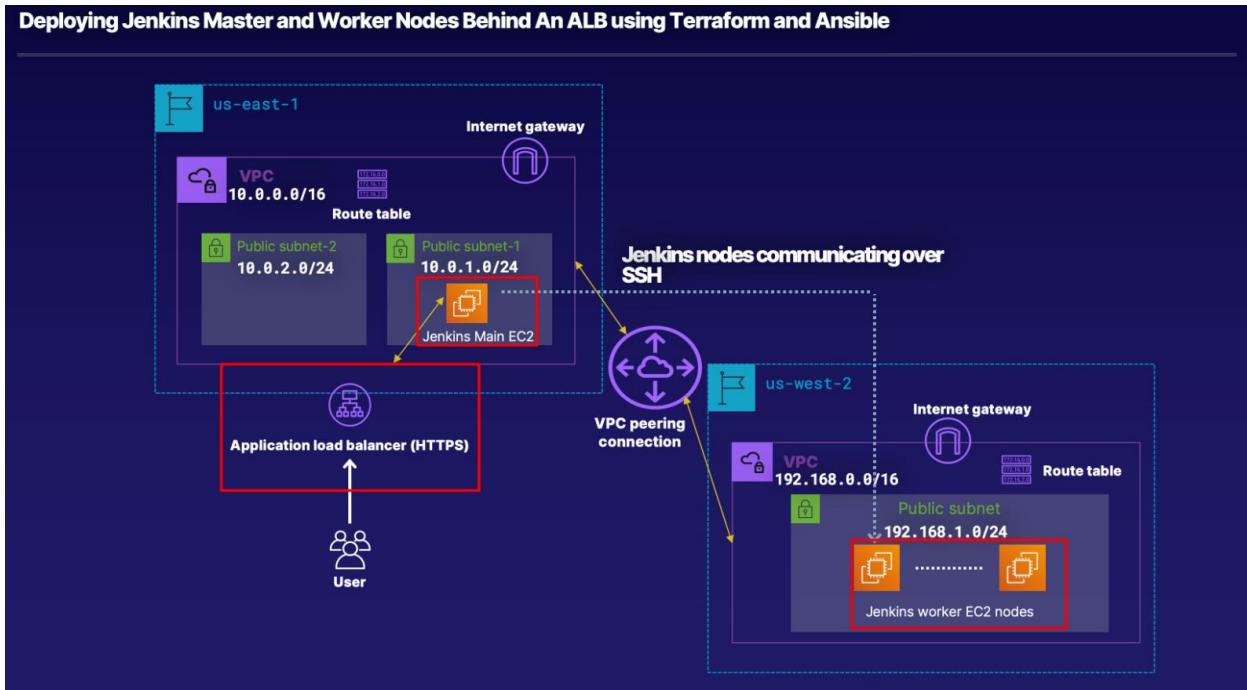
Thank you for using nginx.

Addendum

GitHub Repo files are located as follows:

<https://github.com/linuxacademy/content-aks-basics/tree/master/Lab-Deploying-and-accessing-an-application-to-an-AKS-cluster>

Deploying Jenkins Master and Worker Nodes in AWS Behind an ALB Using Terraform and Ansible



Solution

Pre-Steps:

1. Go and setup within an AWS environment (SANDBOXES are fine) a new EC2 instance within EC2:

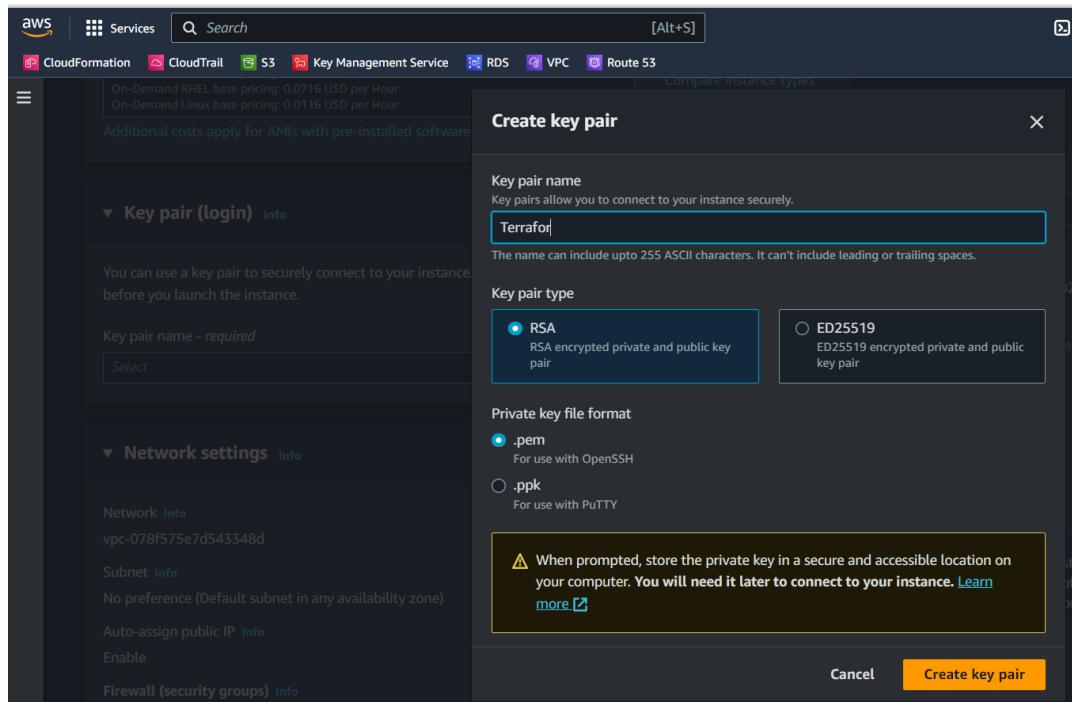
Click on **Launch Instance**

Choose one of the free tier Amazon Linux machines.

Ensure you also create a key pair. Call it "**Terraform**".

Click on **Create Key Pair**

Scroll to the bottom of the webpage and create the instance.



2.Launch a command prompt from your local system and login to the EC2 instance using the command below.

```

[ec2-user@ip-172-31-43-92 ~]
27/08/2023 10:19          1,674 pipelineex.pem
05/03/2023 09:17          1,678 SSHv2.pem
08/09/2023 23:49          1,678 Terraform.pem
29/03/2023 11:09          1,674 WebKey.pem
28/03/2023 11:51          1,678 WhizKey.pem
28/03/2023 13:15          1,674 WhizKey2.pem
30/03/2023 10:49          1,674 WhizKey30.pem
30/03/2023 11:30          1,674 WhizKey30a.pem
30/03/2023 11:50          1,678 WhizKey30b.pem
      21 File(s)           35,198 bytes
          0 Dir(s)   10,389,254,144 bytes free

C:\Users\User\Downloads>ssh ec2-user@100.24.48.169
ec2-user@100.24.48.169: Permission denied (publickey,gssapi-keyex,gssapi-with-mic).

C:\Users\User\Downloads>Terraform.pem ssh ec2-user@100.24.48.169

C:\Users\User\Downloads>ssh -i Terraform.pem ec2-user@100.24.48.169
,      #
~\_ ####_      Amazon Linux 2023
~~ \####\_
~~ \|##|
~~  \#/ __  https://aws.amazon.com/linux/amazon-linux-2023
~~  V~' '-'>
~~/
~~.-
~~ / \
~~ /m/
Last login: Fri Sep  8 20:50:51 2023 from 18.206.107.27
[ec2-user@ip-172-31-43-92 ~]$
```

ssh -i /path/to/your/Terraform.pem ec2-user@<Public IP of the EC2 Instance you just created>

Run the following **cat** command:

```
cat /etc/os-release
[cloud_user@ip-10-0-1-79 ~]$ cat /etc/os-release
NAME="Amazon Linux"
VERSION="2"
ID="amzn"
ID_LIKE="centos rhel fedora"
VERSION_ID="2"
PRETTY_NAME="Amazon Linux 2"
ANSI_COLOR="0;33"
CPE_NAME="cpe:2.3:o:amazon:amazon_linux:2"
HOME_URL="https://amazonlinux.com/"
[cloud_user@ip-10-0-1-79 ~]$
```

3. Check if Python is installed

Use the following command: **python3 –version**

NOTE: It is important to ensure that both **PIP** and **botocore boto3** libraries are installed on this system.

Please follow these steps:

```
sudo yum update
sudo yum install python3
sudo yum install python3-pip
pip install botocore boto3
python3 –version
```

NOTE: if **botocore boto3** libraries are missing than Jenkins will result in a Gateway 502 error!

4. Install Ansible using PIP

sudo pip3 install ansible

You may get the following warning:

WARNING: Running pip as the 'root' user can result in broken permissions and conflicting behaviour with the system package manager. It is recommended to use a virtual environment instead:

<https://pip.pypa.io/warnings/venv>

Ignore the above warning.

Issue the following to confirm the installation of ansible:

ansible –version

5. Install Git

sudo yum install git

Y

6. Authenticate to AWS using aws configure

Type **aws configure** and press <enter>.

Provide both the **Access Key ID** and **Secret Access key ID** went prompted.

Accept the **defaults** for the other prompts.

AWS Sandbox

Dig into Amazon Web Services console in a sandbox, so you can get your hands cloudy.

| | | |
|-------------------|--|---------------------------------|
| Username | cloud_user | <input type="button" value=""/> |
| Password | K&\$0U^yM3RX8Swvek5% | <input type="button" value=""/> |
| URL | https://949221891277.siginin.aws.amazon.com/cc | <input type="button" value=""/> |
| Access Key Id | AKIA52AQITTGW3GW2CFV | <input type="button" value=""/> |
| Secret Access Key | CteAU4DHV+JRTV9bvYy8hA+FdZ5HLsTru7KMyV | <input type="button" value=""/> |

NOTE: if you don't have both the Access key ID and Secret Access Key you can go to IAM within the AWS console and generate them.

Please review earlier projects for the process.

7. Install Terraform

Check if Terraform is installed on the system with the follow command:

terraform version

Note: The Terraform version should be **12.13**. If it's a different version there will be issues when processing the Terraform files. There will be different errors that will require editing of the code. To avoid such headaches simply ensure the version above is installed (even if you have to remove another version).

```
# Uninstall any existing Terraform version.
sudo rm /usr/local/bin/terraform
# Download Terraform v0.12.13 for Linux (adjust for your OS)
wget https://releases.hashicorp.com/terraform/0.12.13/terraform\_0.12.13\_linux\_amd64.zip
# Unzip the downloaded file
unzip terraform_0.12.13_linux_amd64.zip
# Move the Terraform binary to /usr/local/bin to make it globally accessible
sudo mv terraform /usr/local/bin/
# Verify the installation
terraform version
```

8. Download git files and deploy the Terraform code

Clone the GitHub Repo for Terraform Code

a. Use the git command to clone the GitHub repo which has the Terraform code for deploying the solution.

Execute the following command:

```
git clone https://github.com/ACloudGuru-Resources/content-deploying-to-aws-ansible-terraform.git
```

b. Generate a key pair using Ansible

```
cd content-deploying-to-aws-ansible-terraform/lab_jenkins_master_worker
```

Examine the contents of the directory you're in:

```
ls
```

Run the `gen_ssh_key.yaml` Ansible Playbook to generate a SSH Key Pair

```
ansible-playbook ansible_templates/gen_ssh_key.yaml
```

This Ansible Playbook will generate an SSH key pair for your user **cloud_user** which is required for deploying EC2 key pairs in our code.

Note: Alternatively, you may also run the following Linux command to do the same:

```
ssh-keygen -t rsa
```

Note: When this above command prompts for input, keep pressing enter until you're returned to the prompt.

Do not enter a passphrase.

c. Deploy the Terraform Code

```
terraform init
```

```
terraform fmt
```

```
terraform validate
```

```
terraform plan
```

```
terraform apply
```

Enter **yes** when prompted.

After terraform apply has run successfully, you can either use the AWS CLI on the Controller node to list and describe created resources or you can log in to the AWS Console to verify and investigate created resources.

After a successful terraform apply, you will get the DNS URL of the ALB.

Test it out to see if you can reach your Jenkins deployment.

Apply complete! Resources: 3 added, 0 changed, 1 destroyed.

Outputs:

```
Jenkins-Master-AMI-ID = ami-0e1c5d8c23330dee3
```

```
Jenkins-Master-Private-IP = 10.0.1.217
```

```
Jenkins-Worker-AMI-ID = ami-04288abc8d2000768
```

```
Jenkins-Worker-Private-IPs = {
```

```
    "i-0cd8f79141ddd30c4" = "192.168.1.137"
```

```
}
```

```
Jenkins-Worker-Public-IPs = {
```

```

    "i-0cd8f79141ddd30c4" = "18.237.211.240"
}
Loadbalancer-DNS-URL = jenkins-lb-22632709.us-east-1.elb.amazonaws.com

```

d. Use the following Jenkins credentials to login:

username: admin

password: password

The screenshot shows the Jenkins dashboard at the URL <http://jenkins-lb-1662211725.us-east-1.elb.amazonaws.com>. The page title is "Jenkins". The main content area features a "Welcome to Jenkins!" message and a "Start building your software project" button. On the left, there is a sidebar with various Jenkins management links: "New Item", "People", "Build History", "Project Relationship", "Check File Fingerprint", "Manage Jenkins", "My Views", "Job Config History", "Open Blue Ocean Homepage", and "Lockable Resources". Below the sidebar, a "Build Queue" section indicates "No builds in the queue." A search bar at the top right says "Search (CTRL+K)".

e. Finally, on the Terraform Controller node CLI, delete all resources which were created and ensure that it runs through successfully.

terraform destroy

Addendum1. If you fail to install the Python libraries above using PIP you will get the following errors:

```

aws_instance.jenkins-worker-oregon[0] (local-exec): [WARNING]: * Failed to parse /home/ec2-
user/content-deploying-to-aws-ansible-t
aws_instance.jenkins-worker-oregon[0] (local-exec):
erraform/lab_jenkins_master_worker/ansible_templates/inventory_aws/tf_aws_ec2.y
aws_instance.jenkins-worker-oregon[0] (local-exec): ml with auto plugin: Failed to import the required
Python library (botocore and
aws_instance.jenkins-worker-oregon[0] (local-exec): boto3) on ip-172-31-43-130.ec2.internal's Python
/usr/bin/python3. Please read
aws_instance.jenkins-worker-oregon[0] (local-exec): the module documentation and install it in the
appropriate location. If the
aws_instance.jenkins-worker-oregon[0] (local-exec): required library is installed, but Ansible is using the
wrong Python
aws_instance.jenkins-worker-oregon[0] (local-exec): interpreter, please consult the documentation on
ansible_python_interpreter
aws_instance.jenkins-worker-oregon[0] (local-exec): [WARNING]: * Failed to parse /home/ec2-
user/content-deploying-to-aws-ansible-t
aws_instance.jenkins-worker-oregon[0] (local-exec):
erraform/lab_jenkins_master_worker/ansible_templates/inventory_aws/tf_aws_ec2.y
aws_instance.jenkins-worker-oregon[0] (local-exec): ml with yaml plugin: Plugin configuration YAML file,
not YAML inventory
aws_instance.jenkins-worker-oregon[0] (local-exec): [WARNING]: * Failed to parse /home/ec2-
user/content-deploying-to-aws-ansible-t
aws_instance.jenkins-worker-oregon[0] (local-exec):
erraform/lab_jenkins_master_worker/ansible_templates/inventory_aws/tf_aws_ec2.y
aws_instance.jenkins-worker-oregon[0] (local-exec): ml with ini plugin: Invalid host pattern '---' supplied,
'---' is normally a
aws_instance.jenkins-worker-oregon[0] (local-exec): sign this is a YAML file.
aws_instance.jenkins-worker-oregon[0] (local-exec): [WARNING]: Unable to parse /home/ec2-
user/content-deploying-to-aws-ansible-terr
aws_instance.jenkins-worker-oregon[0] (local-exec):
aform/lab_jenkins_master_worker/ansible_templates/inventory_aws/tf_aws_ec2.yml
aws_instance.jenkins-worker-oregon[0] (local-exec): as an inventory source
aws_instance.jenkins-worker-oregon[0] (local-exec): [WARNING]: No inventory was parsed, only implicit
localhost is available
aws_instance.jenkins-worker-oregon[0] (local-exec): [WARNING]: provided hosts list is empty, only
localhost is available. Note that
aws_instance.jenkins-worker-oregon[0] (local-exec): the implicit localhost does not match 'all'
aws_instance.jenkins-worker-oregon[0] (local-exec): [WARNING]: Could not match supplied host
pattern, ignoring:
aws_instance.jenkins-worker-oregon[0] (local-exec): tag_Name_jenkins_worker_tf_1

```

The above was the Output at the end of **Terraform Apply**

It will generate a load balance string similar to the following.

DNS-URL = **jenkins-lb-22632709.us-east-1.elb.amazonaws.com**

However, when you load it into a web browser it will generate a Gateway 502 error!

2. If you fail to login with aws configure you will get the following during the Terraform Plan command.

```
ec2-user@ip-172-31-43-92 lab_jenkins_master_worker]$ terraform plan
```

Refreshing Terraform state in-memory prior to plan...

The refreshed state will be used to calculate this plan, but will not be persisted to local or remote state storage.

Error: error configuring Terraform AWS Provider: **no valid credential sources** for Terraform AWS Provider found.

Please see <https://registry.terraform.io/providers/hashicorp/aws> for more information about providing credentials.

Error: NoCredentialProviders: no valid providers in chain. Deprecated.

For verbose messaging see aws.Config.CredentialsChainVerboseErrors on network_setup.tf line 1, in provider "aws":

```
1: provider "aws" {
```

Error: error configuring Terraform AWS Provider: no valid credential sources for Terraform AWS Provider found.

Please see <https://registry.terraform.io/providers/hashicorp/aws> for more information about providing credentials.

Error: NoCredentialProviders: no valid providers in chain. Deprecated.

For verbose messaging see aws.Config.CredentialsChainVerboseErrors

on network_setup.tf line 7, in provider "aws":

```
7: provider "aws" {
```

3. If you have another version of Terraform than the one mentioned above you will get the following errors when you run Terraform Init:

Error: Invalid reference from destroy provisioner

```
|  
| on instances.tf line 60, in resource "aws_instance" "jenkins-worker-oregon":  
| 60: inline = [  
| 61:   "java -jar /home/ec2-user/jenkins-cli.jar -auth @/home/ec2-user/jenkins_auth -s  
http://${aws_instance.jenkins-master.private_ip}:8080 -auth @/home/ec2-user/jenkins_auth delete-  
node ${self.private_ip}"  
| 62: ]  
|
```

| Destroy-time provisioners and their connection configurations may only reference attributes of the related resource, via 'self',

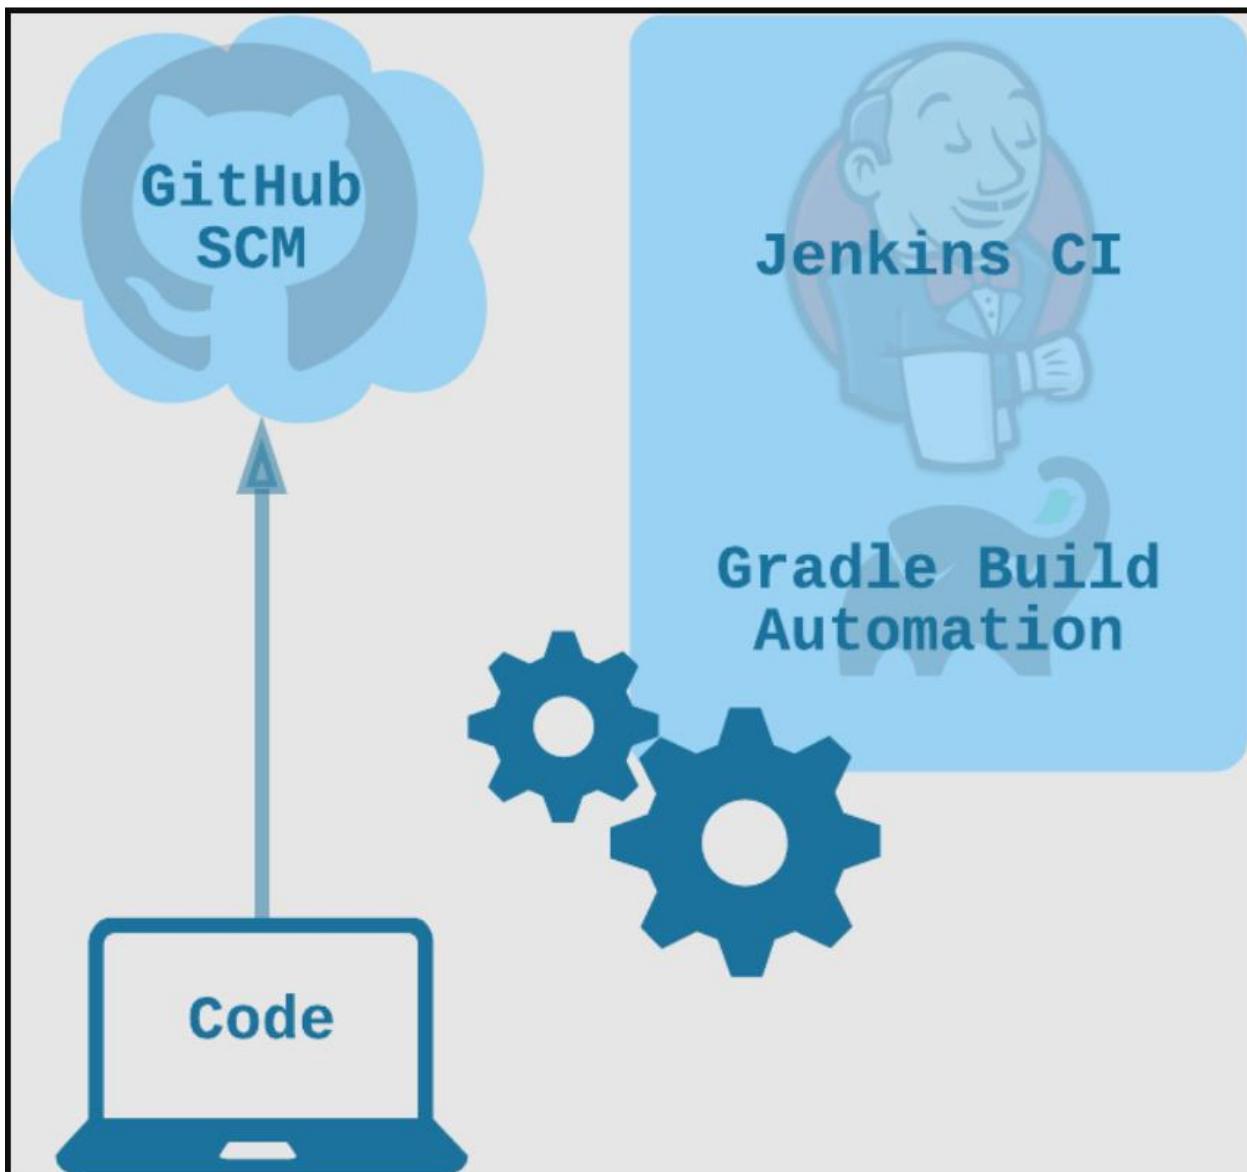
| 'count.index', or 'each.key'.

| References to other resources during the destroy phase can cause dependency cycles and interact poorly with create_before_destroy.

|

[cloud_user@9c438d91c5a44174a4a8c7a5cd9beb884c lab_jenkins_master_worker]\$

Deploying a Docker Container with Jenkins Pipelines



Pre-Steps:

1. Setup two (FREE) Centos 7.5 servers on the Azure portal.



On one we will Install Jenkins.

On the other we will setup our production server that we will deploy the Docker Container.

The screenshot shows the 'Create a virtual machine' wizard in the Microsoft Azure portal. It's on the 'Project details' step, where you select a subscription and resource group. Then it transitions to the 'Instance details' step, which includes fields for the virtual machine name (set to 'Jenkins'), region ('US East US'), availability options ('No infrastructure redundancy required'), security type ('Standard'), image ('CentOS 7.5 Free - x64 Gen1'), and VM architecture ('x64').

Give the Jenkins server a name (e.g., Jenkins)

Run with Azure Spot discount

Size Standard_DS3_v2 - 4 vcpus, 14 GiB memory (US\$213.89/month)

[See all sizes](#)

Info Item(s) availability based on policy assignment(s) for the selected scope.
vadergen2labs/Microsoft.Authorization/8ab9bdb77b304947a1347e88 ([Policy details](#))

[Request quota](#)
[Refresh quota](#)

Administrator account

Authentication type SSH public key Password

Info Azure now automatically generates an SSH key pair for you and allows you to store it for future use. It is a fast, simple, and secure way to connect to your virtual machine.

Username *

SSH public key source

Key pair name *

The options for both servers will likely be as above. It depends, in part, on your subscription. The key option to ensure you select **SSH public key source** radio button as the **authentication type**. This is because we intent to remote into the virtual machines from our local workstation.

Microsoft Azure

Home > 1-49cd24dc-playground-sandbox > Marketplace >

Create a virtual machine ...

Validation passed

Basics Disks Networking Management Monitoring Advanced Tags [Review + create](#)

Info Cost given below is an estimate and not the final price. Please use [Pricing calculator](#) for all your pricing needs.

Price

| | |
|--|---|
| CentOS 7.5 Free by Cognosys Terms of use Privacy policy | Not covered by credits <input type="checkbox"/> 0.0000 USD/hr |
| 1 X Standard DS3 v2 by Microsoft Terms of use Privacy policy | Subscription credits apply <input type="checkbox"/> 0.2930 USD/hr Pricing for other VM sizes |

TERMS

When the Validation passes as above click **create**.

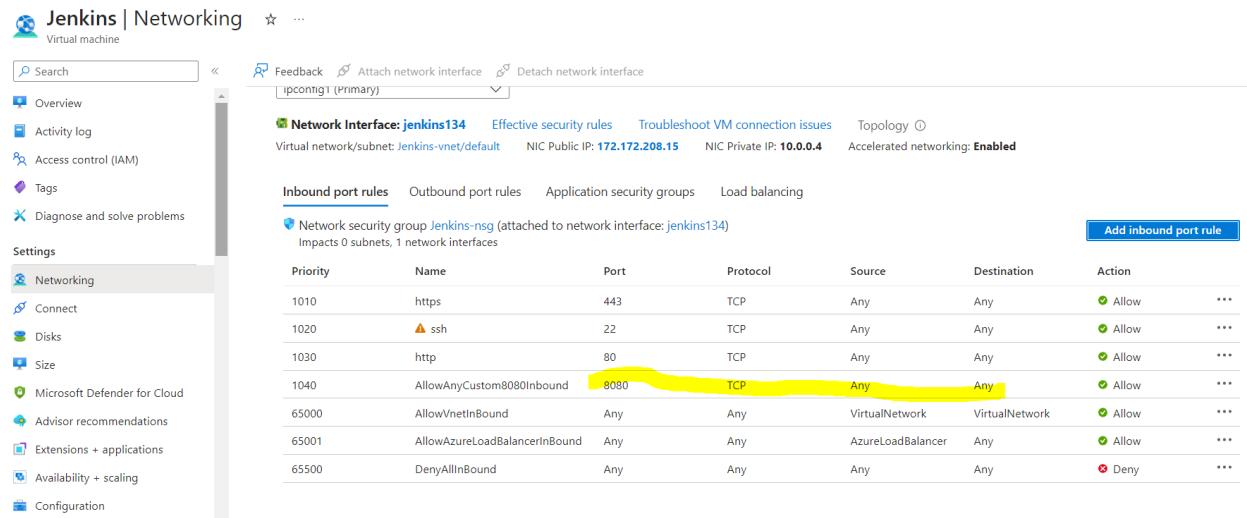
Generate new key pair

Info An SSH key pair contains both a public key and a private key. **Azure doesn't store the private key**. After the SSH key resource is created, you won't be able to download the private key again. [Learn more](#)

When you select the **Download private key and create resource** button this will generate a new key pair to help you securely login to the server remotely!

When the virtual machines are generated select each, one at a time, and create a Network interface inbound rule. This function is located by going to **Networking** under **Settings**.

The port rule you will create for each server is to enable port 8080 from any source are as shown below.



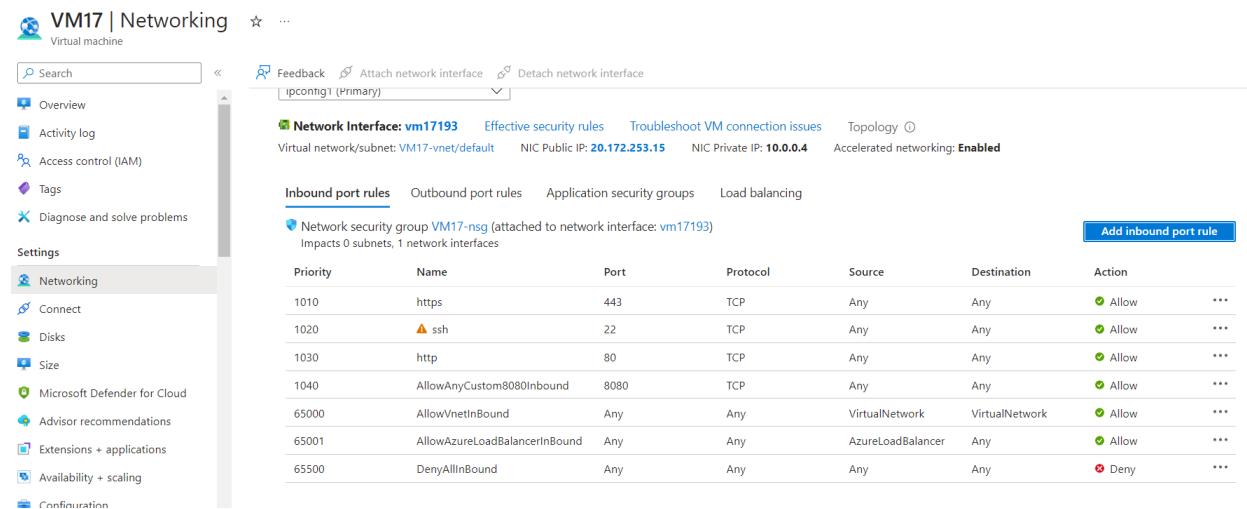
Jenkins | Networking

Network Interface: jenkins134

| Priority | Name | Port | Protocol | Source | Destination | Action |
|----------|-------------------------------|------|----------|-------------------|----------------|--------|
| 1010 | https | 443 | TCP | Any | Any | Allow |
| 1020 | ssh | 22 | TCP | Any | Any | Allow |
| 1030 | http | 80 | TCP | Any | Any | Allow |
| 1040 | AllowAnyCustom8080Inbound | 8080 | TCP | Any | Any | Allow |
| 65000 | AllowVnetInBound | Any | Any | VirtualNetwork | VirtualNetwork | Allow |
| 65001 | AllowAzureLoadBalancerInBound | Any | Any | AzureLoadBalancer | Any | Allow |
| 65500 | DenyAllInBound | Any | Any | Any | Any | Deny |

The port 8080 rule in the case of the Jenkins server is to allow it to communicate externally. Failure to do so will prevent you from loading the Jenkins web portal.

<Jenkins server public IP address>:8080



VM17 | Networking

Network Interface: vm17193

| Priority | Name | Port | Protocol | Source | Destination | Action |
|----------|-------------------------------|------|----------|-------------------|----------------|--------|
| 1010 | https | 443 | TCP | Any | Any | Allow |
| 1020 | ssh | 22 | TCP | Any | Any | Allow |
| 1030 | http | 80 | TCP | Any | Any | Allow |
| 1040 | AllowAnyCustom8080Inbound | 8080 | TCP | Any | Any | Allow |
| 65000 | AllowVnetInBound | Any | Any | VirtualNetwork | VirtualNetwork | Allow |
| 65001 | AllowAzureLoadBalancerInBound | Any | Any | AzureLoadBalancer | Any | Allow |
| 65500 | DenyAllInBound | Any | Any | Any | Any | Deny |

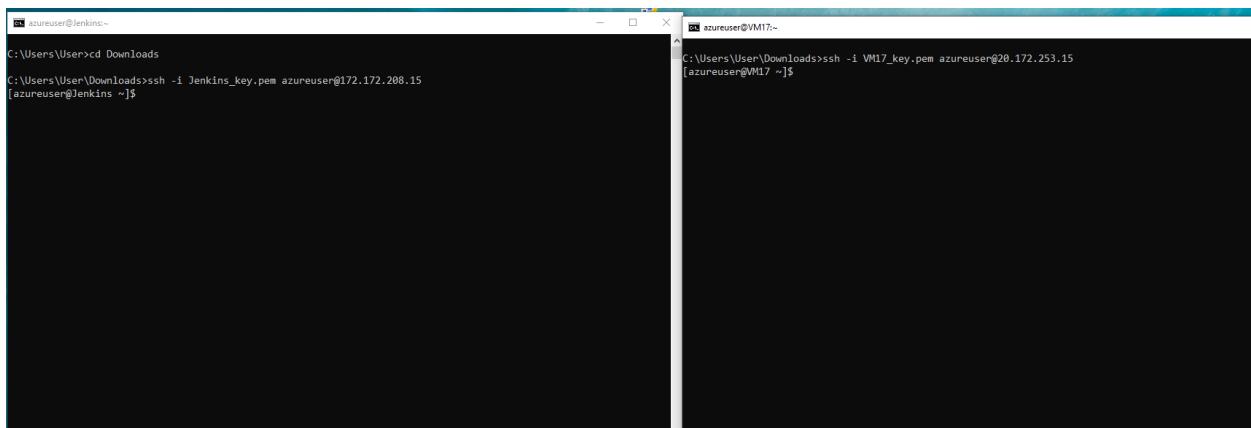
Follow the same process to create similar rule for the Production server. Failure, to do so will prevent the Train Schedule application loading.

<Production server public IP address>:8080

The screenshot shows the Microsoft Azure portal interface for a Jenkins virtual machine. On the left, the navigation menu includes options like Overview, Activity log, Access control (IAM), Tags, Diagnose and solve problems, Settings, Networking, Connect, Disks, Size, Microsoft Defender for Cloud, Advisor recommendations, Extensions + applications, Availability + scaling, and Configuration. Under the 'Connect' section, there's a 'Native SSH' option. The main content area displays the 'Native SSH' configuration page, which includes fields for Admin username (azureuser), Port (change) (22), and Just-in-time policy (Not configured for port 22). Below this, there are two sections: 'Recommended' (SSH using Azure CLI) and 'Most common' (Native SSH). The 'Native SSH' option is highlighted with a blue border. To the right, there are instructions for opening a local shell (Windows), copying and executing an SSH command (which includes a code editor with the command 'ssh -i ~/ssh/id_rsa.pem azureuser@172.172.208.15'), and other information.

Once the networking rule is completed, go to Connect under Settings. Locate the **Native SSH** option and click **Select**.

Follow the instructions to launch a SSH connection with the server.



Note: please ensure that the command above points to where the .PEM key pair file resides.

```

Host '44.212.64.255' is not in the trusted hosts file.
(sshd-25519 fingerprint sha1!! 18:50:bb:a7:54:fb:af:fe:44:10:bb:22:57:c8:26:4a:16:e1:8a:e9)
Do you want to continue connecting? (y/n) y
cloud_user@44.212.64.255's password:
[cloud_user@ip-10-0-1-162 ~]$ cat /etc/os-release
NAME="CentOS Linux"
VERSION="7 (Core)"
ID="centos"
ID_LIKE="rhel fedora"
VERSION_ID="7"
PRETTY_NAME="CentOS Linux 7 (Core)"
ANSI_COLOR="0;31"
CPE_NAME="cpe:/o:centos:centos:7"
HOME_URL="https://www.centos.org/"
BUG_REPORT_URL="https://bugs.centos.org/"

CENTOS_MANTISBT_PROJECT="CentOS-7"
CENTOS_MANTISBT_PROJECT_VERSION="7"
REDHAT_SUPPORT_PRODUCT="centos"
REDHAT_SUPPORT_PRODUCT_VERSION="7"

[cloud_user@ip-10-0-1-162 ~]$ █

```

The **cat /etc/os-release** commands should confirm the Linux OS version.

2. Commands to run on the Jenkins server using SSH

Install Jenkins

Update the server:

sudo yum update

Install Java and the epel-release package:

sudo yum install -y epel-release java-11-openjdk

Configure the Jenkins YUM repository:

sudo wget -O /etc/yum.repos.d/jenkins.repo <http://pkg.jenkins-ci.org/redhat-stable/jenkins.repo>

Install the Jenkins key:

sudo rpm --import <https://pkg.jenkins.io/redhat-stable/jenkins.io-2023.key>

Install Jenkins:

sudo yum -y install jenkins

Enable Jenkins:

sudo systemctl enable jenkins

Start Jenkins:

sudo systemctl start Jenkins

Check the status of Jenkins:

sudo systemctl status jenkins

Post Jenkins Install

In a browser, navigate to the Jenkins IP address:

JENKINS_PUBLIC_IP:**8080**

On the server, find the temporary password:

sudo cat /var/lib/jenkins/secrets/initialAdminPassword

Copy the temporary password.

In the browser, paste in the temporary admin password into the box provided.

Click **Install** suggested plugins.

Enter details for an "**admin**" user.

Click **Save and Continue**.

Click **Save and Finish**.

Click Start using Jenkins.

Install Docker Plugins

Manage Jenkins > Plugins

Go to **available** plugs. Search for Docker plugs.

Select all the Docker plugins.

Click on **Install**.

Select the option to reboot Jenkins after the install.

When the installation finishes go back to the SSH command prompt.

Install Git

sudo yum update

sudo yum install git

git --version

Install Docker

sudo yum update

sudo yum install -y yum-utils device-mapper-persistent-data lvm2

sudo yum-config-manager --add-repo <https://download.docker.com/linux/centos/docker-ce.repo>

sudo yum install -y docker-ce

sudo systemctl start docker

sudo systemctl enable docker

sudo docker --version

Set permissions for Docker

docker run hello-world

docker: Got permission denied while trying to connect to the Docker daemon socket at unix:///var/run/docker.sock: Post <http://%2Fvar%2Frun%2Fdocker.sock/v1.38/containers/create>: dial unix /var/run/docker.sock: connect: permission denied.

See 'docker run --help'.

sudo chmod 666 /var/run/docker.sock

docker run hello-world

Unable to find image 'hello-world:latest' locally

latest: Pulling from library/hello-world

1b930d010525: Pull complete

Digest: sha256:c3b4ada4687bbaa170745b3e4dd8ac3f194ca95b2d0518b417fb47e5879d9b5f

Status: Downloaded newer image for hello-world:latest

Hello from Docker!

This message shows that your installation appears to be working correctly.

To generate this message, Docker took the following steps:

1. The Docker client contacted the Docker daemon.
2. The Docker daemon pulled the "hello-world" image from the Docker Hub.
(amd64)
3. The Docker daemon created a new container from that image which runs the executable that produces the output you are currently reading.
4. The Docker daemon streamed that output to the Docker client, which sent it to your terminal.

To try something more ambitious, you can run an Ubuntu container with:

```
$ docker run -it ubuntu bash
```

Share images, automate workflows, and more with a free Docker ID:

<https://hub.docker.com/>

For more examples and ideas, visit:

<https://docs.docker.com/get-started/>

Install sshpass

Install sshpass

sudo yum install sshpass

3.Commands to run on the Production server using SSH

Update the server

sudo yum update

Install Java and the epel-release package

sudo yum install -y epel-release java-11-openjdk

Install Git

sudo yum update

sudo yum install git

git --version

Install Docker

sudo yum update

sudo yum install -y yum-utils device-mapper-persistent-data lvm2

sudo yum-config-manager --add-repo <https://download.docker.com/linux/centos/docker-ce.repo>

sudo yum install -y docker-ce

sudo systemctl start docker

sudo systemctl enable docker

sudo docker --version

Set permissions for Docker

When you run the following command on docker, by default, you will get the a permission denied error!

docker run hello-world

```
docker: Got permission denied while trying to connect to the Docker daemon socket at
unix:///var/run/docker.sock: Post http://%2Fvar%2Frun%2Fdocker.sock/v1.38/containers/create: dial
unix /var/run/docker.sock: connect: permission denied.
See 'docker run --help'.
```

sudo chmod 666 /var/run/docker.sock

docker run hello-world

```
Unable to find image 'hello-world:latest' locally
latest: Pulling from library/hello-world
1b930d010525: Pull complete
Digest: sha256:c3b4ada4687bbba170745b3e4dd8ac3f194ca95b2d0518b417fb47e5879d9b5f
Status: Downloaded newer image for hello-world:latest
```

Hello from Docker!

This message shows that your installation appears to be working correctly.

To generate this message, Docker took the following steps:

1. The Docker client contacted the Docker daemon.
2. The Docker daemon pulled the "hello-world" image from the Docker Hub.
(amd64)
3. The Docker daemon created a new container from that image which runs the executable that produces the output you are currently reading.
4. The Docker daemon streamed that output to the Docker client, which sent it to your terminal.

To try something more ambitious, you can run an Ubuntu container with:

\$ docker run -it ubuntu bash

Share images, automate workflows, and more with a free Docker ID:

<https://hub.docker.com/>

For more examples and ideas, visit:

<https://docs.docker.com/get-started/>

[Setup the Deploy account for Jenkins on the Production server.](#)

[azureuser@vm16 ~]\$ **whoami**

azureuser

[azureuser@vm16 ~]\$

Create a User:

sudo useradd -m deploy

Set a password:

sudo passwd deploy

Enter password

Confirm password

Note: we will use this later on the Jenkins portal.

List all user accounts

cat /etc/passwd

Host '54.210.24.6' is not in the trusted hosts file.

(ssh-ed25519 fingerprint sha1!! 5c:11:16:0b:b6:77:19:f9:aa:7c:64:01:02:c3:4d:97:e2:03:50:78)

Do you want to continue connecting? (y/n) y

cloud_user@54.210.24.6's password:

[cloud_user@ip-10-0-1-89 ~]\$ cat /etc/passwd

root:x:0:0:root:/root:/bin/bash

bin:x:1:1:bin:/bin/nologin

daemon:x:2:2:daemon:/sbin:/sbin/nologin

adm:x:3:4:adm:/var/adm:/sbin/nologin

lp:x:4:7:lp:/var/spool/lpd:/sbin/nologin

sync:x:5:0:sync:/sbin:/bin/sync

shutdown:x:6:0:shutdown:/sbin:/sbin/shutdown

halt:x:7:0:halt:/sbin:/sbin/halt

mail:x:8:12:mail:/var/spool/mail:/sbin/nologin

operator:x:11:0:operator:/root:/sbin/nologin

games:x:12:100:games:/usr/games:/sbin/nologin

ftp:x:14:50:FTP User:/var/ftp:/sbin/nologin

nobody:x:99:99:Nobody:/sbin/nologin

systemd-bus-proxy:x:999:998:systemd Bus Proxy:/sbin/nologin

systemd-network:x:192:192:systemd Network Management:/sbin/nologin

dbus:x:81:81:System message bus:/sbin/nologin

polkitd:x:998:997:User for polkitd:/sbin/nologin

tss:x:59:59:Account used by the trousers package to sandbox the tcscd daemon:/dev/null:/sbin/nologin

sshd:x:74:74:Privilege-separated SSH:/var/empty/sshd:/sbin/nologin

postfix:x:89:89::/var/spool/postfix:/sbin/nologin

chrony:x:997:995::/var/lib/chrony:/sbin/nologin

cloud_user:x:1000:1000:Cloud Assessment User:/home/cloud_user:/bin/bash

centos:x:1001:1001:Cloud User:/home/centos:/bin/bash

jenkins:x:996:993:Jenkins Automation Server:/var/lib/jenkins:/bin/false

deploy:x:1002:1002::/home/deploy:/bin/bash

[cloud_user@ip-10-0-1-89 ~]\$

Setup permissions for the Jenkins server so that it can SSH to the Production server.

When you run the following command you will get, by default, a permissions denied error as below

sshpass -p '<password for the deploy account on the production server>' -v ssh -o

StrictHostKeyChecking=no deploy@<production server Public IP Address>

```
[azureuser@Jenkins ~]$ sudo sshpass -p 'Harvee777$' -v ssh -o StrictHostKeyChecking=no
deploy@20.172.253.15
Warning: Permanently added '20.172.253.15' (ECDSA) to the list of known hosts.
Permission denied (publickey,gssapi-keyex,gssapi-with-mic).
[azureuser@Jenkins ~]$
```

The solution to this is to enable password authentication.

Go to the Production server and use the following commands:

sudo vi /etc/ssh/sshd_config

In the file, find the **PasswordAuthentication** line and make sure it ends with **yes**.
 Find the **ChallengeResponseAuthentication** option and disable it by adding **no**.
 If lines are commented out, remove the hash sign # to uncomment them.

Switch to command mode by pressing the Esc key.

Type : (colon). This will open the prompt bar in the bottom left corner of the window.
 Type wq after the colon and hit Enter. This will save in Vim the changes made to the file
 and exit you out of the editor.

Now restart SSH with the following command:

sudo systemctl restart sshd

Now return to the Jenkins server and issue the same command as before.

sshpass -p '<password for the deploy account on the production server>' -v ssh -o

StrictHostKeyChecking=no deploy@<production server Public IP Address>

```
[azureuser@Jenkins ~]$ sudo sshpass -p 'Harvee777$' -v ssh -o StrictHostKeyChecking=no deploy@172.172.155.55
SSHPASS searching for password prompt using match "assword"
SSHPASS read: deploy@172.172.155.55's password:
SSHPASS detected prompt. Sending password.
SSHPASS read:

[deploy@vm777 ~]$ logout
Connection to 172.172.155.55 closed.
[azureuser@Jenkins ~]$
```

It has the effect of creating an SSH session from the Jenkins server to the production server. Then logout to close the session.

The reason for the Jenkins to have these permissions is because it uses the above command to remote to the Production server then issue a docker command to pull the built docker container locally to the production server.

sshpass -p 'your_password_here' -v ssh -o StrictHostKeyChecking=no deploy@<production server public IP address> 'docker pull costas778/train-schedule'

NOTE: the password above is NOT one I will use for production purposes.

Deploying a Docker Container with a Jenkins Pipeline

We are, finally, in a position to deploy the Docker container using a Jenkins pipeline!

Go to **Manage Jenkins > Credentials**, then click **(global)**.

Click + Add Credentials

Create a global Jenkins credential called **deploy**

Kind: Username with password

Username: **deploy**

Password:

ID: **webserver_ID**

Description: **Webserver ID**

Use the same deploy password that you used for the production server.

Click **Create**

Create a global Jenkins credential for the Docker image registry (Docker Hub).

Click + Add Credentials:

Kind: Username with password

Username: <**DOCKER_HUB_USERNAME**>

Password: <**DOCKER_HUB_PASSWORD**>

ID: **docker_hub_login**

Description: **Docker Hub Login**

Click **Create**

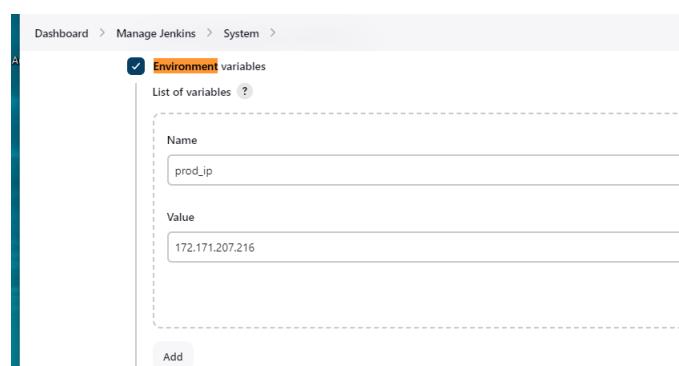
Note: You will need a Docker hub account to use Docker Hub as an image registry.

Configure a global property in Jenkins to store the production server IP by navigating to **Manage Jenkins > System** and adding an **environment variable** under **Global properties**.

Name: **prod_ip**

Value: <**PRODUCTION_SERVER_PUBLIC_IP_ADDRESS**>

Click **Save**



Go to GitHub

Make a personal fork from the following GitHub repository:

<https://github.com/ACloudGuru-Resources/cicd-pipeline-train-schedule-dockerdeploy>

https://github.com/<YOUR_GITHUB_USER>/cicd-pipeline-train-schedule-dockerdeploy

Generate a new **GitHub API key** to allow Jenkins to access the forked repo by navigating to **Profile > Settings > Developer Settings > Personal Access Tokens > Tokens (classic) > Generate new token (classic)**.

Note: **Jenkinsapp**

Permissions: **admin:repo_hook**

Copy the GitHub token.

Go back to the Jenkins portal.

Click **New Items**

In Jenkins, create a **Multibranch Pipeline** project named **train-schedule**.

Click **OK**

Under **Branch Sources**, click **Add** below **GitHub**, then choose **Jenkins** to configure your forked repository.

Kind: Username with password

Scope: Global

Username: <GITHUB_USERNAME>

Password: <YOUR_GITHUB_PERSONAL_ACCESS_TOKEN>

ID: **github_key**

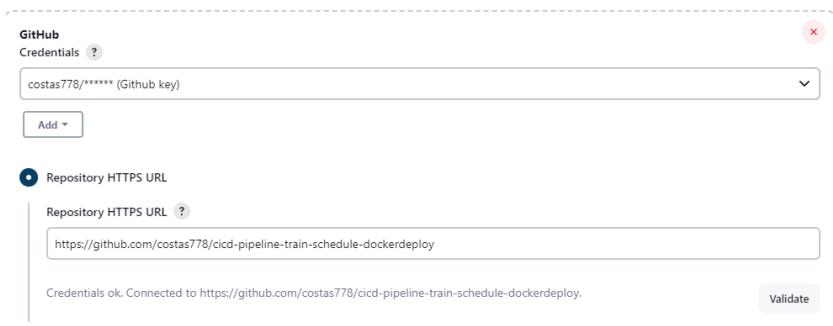
Description: **GitHub Key**

Click **Add**

Select the **GitHub Key** and paste your forked repository under **Repository HTTPS URL**:

https://github.com/<YOUR_GITHUB_USER>/cicd-pipeline-train-schedule-dockerdeploy

Click **Validate** to test your credentials.



Click **Save**.

Successfully deploy the train-schedule app to production as a Docker container using the Jenkins Pipeline.

Modify the **Jenkinsfile** in GitHub to build and push the Docker image to Docker Hub and commit the changes. Replace `willbla` from the example file with your own Docker Hub username.

```
pipeline {
    agent any
    stages {
        stage('Build') {
            steps {
                echo 'Running build automation'
                sh './gradlew build --no-daemon'
                archiveArtifacts artifacts: 'dist/trainSchedule.zip'
            }
        }
        stage('Build Docker Image') {
            when {
                branch 'master'
            }
            steps {
                script {
                    app = docker.build("<DOCKER_HUB_USERNAME>/train-schedule")
                    app.inside {
                        sh 'echo $(curl localhost:8080)'
                    }
                }
            }
        }
        stage('Push Docker Image') {
            when {
                branch 'master'
            }
            steps {
                script {
```

```
        docker.withRegistry('https://registry.hub.docker.com', 'docker_hub_login') {
            app.push("${env.BUILD_NUMBER}")
            app.push("latest")
        }
    }
}
}
}
```

In Jenkins, click **Build Now**.

Note: The initial build may take several minutes to complete.

In **Docker Hub**, under **Repositories**, select the **train-schedule** app.

Click the **Tags tab** to verify that the build was pushed successfully.

In GitHub, modify the **Jenkinsfile** to include a stage that pushes the build to the production server, and commit the changes.

```
stage ('DeployToProduction') {  
  
    when {  
  
        branch 'master'  
  
    }  
  
    steps {  
  
        input 'Deploy to Production'  
  
        milestone(1)  
  
        withCredentials ([usernamePassword(credentialsId: 'webserver_login', usernameVariable:  
'USERNAME', passwordVariable: 'USERPASS')]) {  
  
            script {
```

```

sh "sshpass -p '$USERPASS' -v ssh -o StrictHostKeyChecking=no $USERNAME@${env.prod_ip}

\"docker pull <DOCKER_HUB_USERNAME>/train-schedule:${env.BUILD_NUMBER}\"

try {

    sh "sshpass -p '$USERPASS' -v ssh -o StrictHostKeyChecking=no $USERNAME@${env.prod_ip}

\"docker stop train-schedule\"

    sh "sshpass -p '$USERPASS' -v ssh -o StrictHostKeyChecking=no $USERNAME@${env.prod_ip}

\"docker rm train-schedule\"

} catch (err) {

    echo: 'caught error: $err'

}

sh "sshpass -p '$USERPASS' -v ssh -o StrictHostKeyChecking=no $USERNAME@${env.prod_ip}

\"docker run --restart always --name train-schedule -p 8080:8080 -d <DOCKER_HUB_USERNAME>/train-
schedule:${env.BUILD_NUMBER}\"

}

}

}

}

```

Place your Docker hub username above and **commit** the changes.
In Jenkins, click **Build Now**.

Branch master

Full project name: train-schedule/master

Average stage times:

| Declarative: Checkout SCM | Build | Build Docker Image | Push Docker Image | DeployToProduction |
|---------------------------|---------|--------------------|-------------------|--------------------|
| 593ms | 1min 9s | 1min 46s | 7s | 26s |

Sept 15 11:50 No Changes

593ms 1min 9s 1min 46s 7s 26s (based on 10)

Permalinks

- Last build (#1), 4.9 sec ago
- Atom feed for all Atom feed for failures

Note: to proceed from Push Docker Image to **DeployToProduction** you will need to hover over Push Docker Image within the green area and allow it to proceed further.

Once the build is complete, using a web browser, verify that the application has been deployed successfully.

Branch master

Full project name: train-schedule/master

Average stage times:

| Declarative: Checkout SCM | Build | Build Docker Image | Push Docker Image | DeployToProduction |
|---------------------------|---------|--------------------|-------------------|------------------------|
| 593ms | 1min 9s | 1min 46s | 7s | 1min 38s (based on 10) |

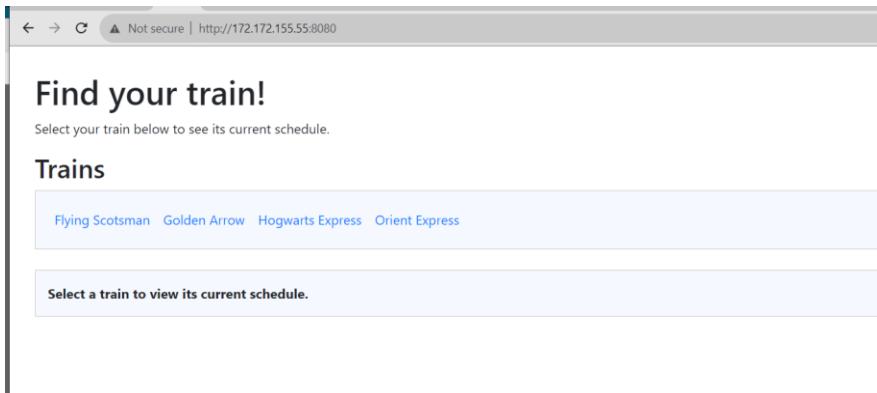
Sept 15 11:50 No Changes

593ms 1min 9s 1min 46s 7s 1min 38s (based on 10)

Permalinks

- Last build (#1), 4.9 sec ago
- Atom feed for all Atom feed for failures

Place the following in a web browser <PRODUCTION_SERVER_PUBLIC_IP_ADDRESS>:8080



If successful, you will see the above!

Addendum

1) Build Failure

Started by user Costas constantinou

08:44:29 Connecting to <https://api.github.com> using costas778/******** (Github key)

Obtained Jenkinsfile from 248d9bbd2a5c031c08db8c6b21a28f6336cc4173

[Pipeline] Start of Pipeline

[Pipeline] node

Running on Jenkins in /var/lib/jenkins/workspace/train-schedule_master

[Pipeline] {

[Pipeline] stage

[Pipeline] { (Declarative: Checkout SCM)

[Pipeline] checkout

The recommended git tool is: NONE

using credential github_key

> git rev-parse --resolve-git-dir /var/lib/jenkins/workspace/train-schedule_master/.git # timeout=10

Fetching changes from the remote Git repository

> git config remote.origin.url <https://github.com/costas778/cicd-pipeline-train-schedule-dockerdeploy.git> # timeout=10

Fetching without tags

Fetching upstream changes from <https://github.com/costas778/cicd-pipeline-train-schedule-dockerdeploy.git>

> git --version # timeout=10

> git --version # 'git version 1.8.3.1'

using GIT_ASKPASS to set credentials Github key

> git fetch --no-tags --progress <https://github.com/costas778/cicd-pipeline-train-schedule-dockerdeploy.git> +refs/heads/master:refs/remotes/origin/master # timeout=10

Checking out Revision 248d9bbd2a5c031c08db8c6b21a28f6336cc4173 (master)

> git config core.sparsecheckout # timeout=10

> git checkout -f 248d9bbd2a5c031c08db8c6b21a28f6336cc4173 # timeout=10

Commit message: "Update Jenkinsfile"

> git rev-list --no-walk 33c78dcc8603870d855e4caa9653c89b7848b739 # timeout=10

[Pipeline] }

```
[Pipeline] // stage
[Pipeline] withEnv
[Pipeline] {
[Pipeline] stage
[Pipeline] { (Build)
[Pipeline] echo
Running build automation
[Pipeline] sh
+ ./gradlew build --no-daemon
WARNING: An illegal reflective access operation has occurred
WARNING: Illegal reflective access by org.codehaus.groovy.vmplugin.v7.Java7$1
(file:/var/lib/jenkins/.gradle/wrapper/dists/gradle-6.0.1-bin/1lxlpkiy24sb18odw96cp4ojv/gradle-6.0.1/lib/groovy-all-1.3-2.5.8.jar) to constructor
java.lang.invoke.MethodHandles$Lookup(java.lang.Class,int)
WARNING: Please consider reporting this to the maintainers of
org.codehaus.groovy.vmplugin.v7.Java7$1
WARNING: Use --illegal-access=warn to enable warnings of further illegal reflective access operations
WARNING: All illegal access operations will be denied in a future release
> Task :nodeSetup UP-TO-DATE
> Task :npmSetup UP-TO-DATE
> Task :npmInstall UP-TO-DATE
> Task :npm_test
> cicd-pipeline-train-schedule-git@0.0.0 test /var/lib/jenkins/workspace/train-schedule_master
> mocha
Index Page
[0mGET / [32m200 [0m186.268 ms - 829[0m
    ✓ renders successfully (205ms)
```

Trains API

```
[0mGET /trains [32m200 [0m2.163 ms - 1093[0m
    ✓ returns data successfully
2 passing (227ms)
> Task :npm_build
> Task :zip
> Task :build
```

Deprecated Gradle features were used in this build, making it incompatible with Gradle 7.0.
Use '--warning-mode all' to show the individual deprecation warnings.

See

https://docs.gradle.org/6.0.1/userguide/command_line_interface.html#sec:command_line_warnings

```
BUILD SUCCESSFUL in 7s
6 actionable tasks: 3 executed, 3 up-to-date
[Pipeline] archiveArtifacts
Archiving artifacts
[Pipeline] }
[Pipeline] // stage
```

```

[Pipeline] stage
[Pipeline] { (Build Docker Image)
[Pipeline] script
[Pipeline] {
[Pipeline] }
[Pipeline] // script
[Pipeline] }
[Pipeline] // stage
[Pipeline] stage
[Pipeline] { (Push Docker Image)
Stage "Push Docker Image" skipped due to earlier failure(s)
[Pipeline] }
[Pipeline] // stage
[Pipeline] stage
[Pipeline] { (DeployToProduction)
Stage "DeployToProduction" skipped due to earlier failure(s)
[Pipeline] }
[Pipeline] // stage
[Pipeline] }
[Pipeline] // withEnv
[Pipeline] }
[Pipeline] // node
[Pipeline] End of Pipeline
Also: org.jenkinsci.plugins.workflow.actions.ErrorAction$ErrorId: 78a67a41-c4f8-405c-a52c-
f0f33958a0bf
groovy.lang.MissingPropertyException: No such property: docker for class: groovy.lang.Binding
at groovy.lang.Binding.getVariable(Binding.java:63)
at
org.jenkinsci.plugins.scriptsecurity.sandbox.groovy.SandboxInterceptor.onGetProperty(SandboxIn-
terceptor.java:285)
at org.kohsuke.groovy.sandbox.impl.Checker$7.call(Checker.java:375)
at org.kohsuke.groovy.sandbox.impl.Checker.checkedGetProperty(Checker.java:379)
at org.kohsuke.groovy.sandbox.impl.Checker.checkedGetProperty(Checker.java:355)
at org.kohsuke.groovy.sandbox.impl.Checker.checkedGetProperty(Checker.java:355)
at org.kohsuke.groovy.sandbox.impl.Checker.checkedGetProperty(Checker.java:355)
at com.cloudbees.groovy.cps.sandbox.SandboxInvoker.getProperty(SandboxInvoker.java:29)
at org.jenkinsci.plugins.workflow.cps.LoggingInvoker.getProperty(LoggingInvoker.java:121)
at com.cloudbees.groovy.cps.impl.PropertyAccessBlock.rawGet(PropertyAccessBlock.java:20)
at WorkflowScript.run(WorkflowScript:17)
at __cps.transform__(Native Method)
at
com.cloudbees.groovy.cps.impl.PropertyishBlock$ContinuationImpl.get(PropertyishBlock.java:73)
at com.cloudbees.groovy.cps.LValueBlock$GetAdapter.receive(LValueBlock.java:30)
at
com.cloudbees.groovy.cps.impl.PropertyishBlock$ContinuationImpl.fixName(PropertyishBlock.jav-
a:65)
at jdk.internal.reflect.GeneratedMethodAccessor386.invoke(Unknown Source)

```

```

at
java.base/jdk.internal.reflect.DelegatingMethodAccessorImpl.invoke(DelegatingMethodAccessorImpl.java:43)
at java.base/java.lang.reflect.Method.invoke(Method.java:566)
at
com.cloudbees.groovy.cps.impl.ContinuationPtr$ContinuationImpl.receive(ContinuationPtr.java:72)
at com.cloudbees.groovy.cps.impl.ConstantBlock.eval(ConstantBlock.java:21)
at com.cloudbees.groovy.cps.Next.step(Next.java:83)
at com.cloudbees.groovy.cps.Continuable$1.call(Continuable.java:152)
at com.cloudbees.groovy.cps.Continuable$1.call(Continuable.java:146)
at
org.codehaus.groovy.runtime.GroovyCategorySupport$ThreadCategoryInfo.use(GroovyCategorySupport.java:136)
at org.codehaus.groovy.runtime.GroovyCategorySupport.use(GroovyCategorySupport.java:275)
at com.cloudbees.groovy.cps.Continuable.run0(Continuable.java:146)
at
org.jenkinsci.plugins.workflow.cps.SandboxContinuable.access$001(SandboxContinuable.java:18)
at org.jenkinsci.plugins.workflow.cps.SandboxContinuable.run0(SandboxContinuable.java:51)
at org.jenkinsci.plugins.workflow.cps.CpsThread.runNextChunk(CpsThread.java:187)
at org.jenkinsci.plugins.workflow.cps.CpsThreadGroup.run(CpsThreadGroup.java:422)
at org.jenkinsci.plugins.workflow.cps.CpsThreadGroup$2.call(CpsThreadGroup.java:330)
at org.jenkinsci.plugins.workflow.cps.CpsThreadGroup$2.call(CpsThreadGroup.java:294)
at
org.jenkinsci.plugins.workflow.cps.CpsVmExecutorService$2.call(CpsVmExecutorService.java:97)
at java.base/java.util.concurrent.FutureTask.run(FutureTask.java:264)
at hudson.remoting.SingleLaneExecutorService$1.run(SingleLaneExecutorService.java:139)
at jenkins.util.ContextResettingExecutorService$1.run(ContextResettingExecutorService.java:28)
at jenkins.security.ImpersonatingExecutorService$1.run(ImpersonatingExecutorService.java:68)
at jenkins.util.ErrorLoggingExecutorService.lambda$wrap$0(ErrorLoggingExecutorService.java:51)
at java.base/java.util.concurrent.Executors$RunnableAdapter.call(Executors.java:515)
at java.base/java.util.concurrent.FutureTask.run(FutureTask.java:264)
at java.base/java.util.concurrent.ThreadPoolExecutor.runWorker(ThreadPoolExecutor.java:1128)
at java.base/java.util.concurrent.ThreadPoolExecutor$Worker.run(ThreadPoolExecutor.java:628)
at java.base/java.lang.Thread.run(Thread.java:829)

```

Could not update commit status, please check if your scan credentials belong to a member of the organization or a collaborator of the repository and repo:status scope is selected
GitHub has been notified of this commit's build result.

Finished: FAILURE

Solution:

Make sure that the **Docker plugins** are installed and configured correctly on your Jenkins server.

2) Failure to Deploy to Production (Stage)

```
Shell Script -- sshpass -p '${USERPASS}' -v ssh -o StrictHostKeyChecking=no
jenkinsad@172.171.207.216 "docker pull costas778/train-schedule:3" (self time 288ms)
```

Warning: A secret was passed to "sh" using Groovy String interpolation, which is insecure.

Affected argument(s) used the following variable(s): [USERPASS]

See <https://jenkins.io/redirect/groovy-string-interpolation> for details.

```
+ sshpass -p **** -v ssh -o StrictHostKeyChecking=no jenkinsad@172.171.207.216 'docker pull
costas778/train-schedule:3'
/var/lib/jenkins/workspace/train-schedule_master@tmp/durable-b344b4a1/script.sh: line 1: sshpass:
command not found
```

Solution:

Install SSH with the following command.

```
sudo yum install sshpass
```

3) Failure to Deploy To Production (Stage) after installing sshpass

Warning: A secret was passed to "sh" using Groovy String interpolation, which is insecure.

Affected argument(s) used the following variable(s): [USERPASS]

See <https://jenkins.io/redirect/groovy-string-interpolation> for details.

```
+ sshpass -p **** -v ssh -o StrictHostKeyChecking=no jenkinsad@172.171.207.216 'docker pull
costas778/train-schedule:4'
```

Warning: Permanently added '172.171.207.216' (ECDSA) to the list of known hosts.

Permission denied (publickey,gssapi-keyex,gssapi-with-mic).

Solution:

Go to sshd_config file on the production server and make the following changes:

```
sudo vi /etc/ssh/sshd_config
```

In the file, find the **PasswordAuthentication** line and make sure it ends with yes.

Find the **ChallengeResponseAuthentication** option and disable it by adding no.

If lines are commented out, remove the hash sign # to uncomment them.

Save the changes in the file then restart the SSH service for the changes to take effect.

```
sudo systemctl restart sshd
```

4) Build error - Failure to download Node with Gradle.

```
+ ./gradlew build --no-daemon
```

```
Downloading https://services.gradle.org/distributions/gradle-6.0.1-bin.zip
```

Unzipping /var/lib/jenkins/.gradle/wrapper/dists/gradle-6.0.1-bin/1lxlpkiy24sb18odw96cp4ojv/gradle-6.0.1-bin.zip to /var/lib/jenkins/.gradle/wrapper/dists/gradle-6.0.1-bin/1lxlpkiy24sb18odw96cp4ojv

Set executable permissions for: /var/lib/jenkins/.gradle/wrapper/dists/gradle-6.0.1-bin/1lxlpkiy24sb18odw96cp4ojv/gradle-6.0.1/bin/gradle

Welcome to Gradle 6.0.1!

Here are the highlights of this release:

- Substantial improvements in dependency management, including
 - Publishing Gradle Module Metadata in addition to pom.xml
 - Advanced control of transitive versions
 - Support for optional features and dependencies
 - Rules to tweak published metadata
- Support for Java 13
- Faster incremental Java and Groovy compilation
- New Zinc compiler for Scala
- VS2019 support
- Support for Gradle Enterprise plugin 3.0

For more details see <https://docs.gradle.org/6.0.1/release-notes.html>

WARNING: An illegal reflective access operation has occurred

WARNING: Illegal reflective access by org.codehaus.groovy.vmplugin.v7.Java7\$1
`(file:/var/lib/jenkins/.gradle/wrapper/dists/gradle-6.0.1-bin/1lxlpkiy24sb18odw96cp4ojv/gradle-6.0.1/lib/groovy-all-1.3-2.5.8.jar)` to constructor

`java.lang.invoke.MethodHandles$Lookup(java.lang.Class,int)`

WARNING: Please consider reporting this to the maintainers of

`org.codehaus.groovy.vmplugin.v7.Java7$1`

WARNING: Use --illegal-access=warn to enable warnings of further illegal reflective access operations

WARNING: All illegal access operations will be denied in a future release

> Task :nodeSetup

> Task :nodeSetup FAILED

FAILURE: Build failed with an exception.

* What went wrong:

Execution failed for task ':nodeSetup'.

> Could not resolve all files for configuration ':detachedConfiguration1'.

 > Could not download node-9.11.1-linux-x64.tar.gz (org.nodejs:node:9.11.1)

 > Could not get resource '<https://nodejs.org/dist/v9.11.1/node-v9.11.1-linux-x64.tar.gz>'.

 > Read timed out

* Try:

Run with --stacktrace option to get the stack trace. Run with --info or --debug option to get more log output. Run with --scan to get full insights.

* Get more help at <https://help.gradle.org>

Deprecated Gradle features were used in this build, making it incompatible with Gradle 7.0.

Use '--warning-mode all' to show the individual deprecation warnings.

See

https://docs.gradle.org/6.0.1/userguide/command_line_interface.html#sec:command_line_warnings

BUILD FAILED in 10m 17s

1 actionable task: 1 executed

Solution:

This is, likely, down to internet issues.

Select **Run Build** again.

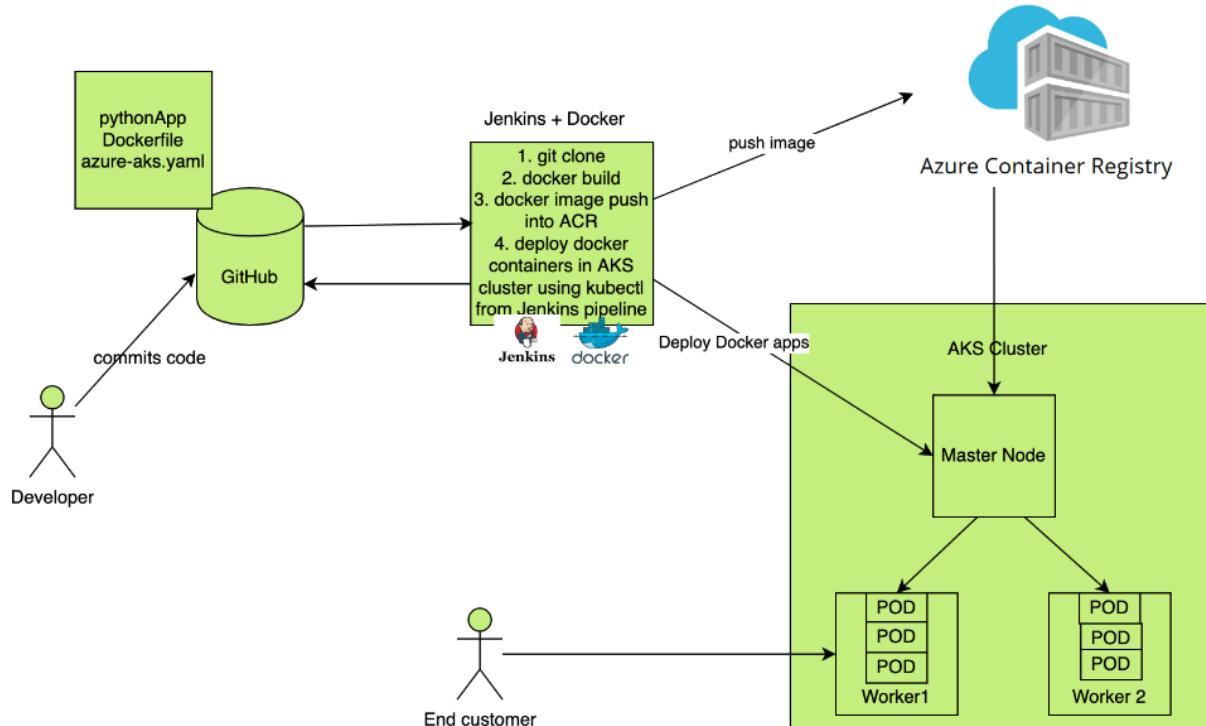
If this still fails then perform troubleshooting on your internet connection!

Alternatively, download the following manually and install it.

```
sudo wget https://nodejs.org/dist/v9.11.1/node-v9.11.1-linux-x64.tar.gz
https://nodejs.org/en/download
```

Deploy Python App into Kubernetes Cluster using Jenkins Pipeline | Containerize Python App and Deploy into AKS Cluster

Kubernetes(AKS) Deployment from Jenkins pipeline using Kubectl CLI



Pre-requisites:

Create an Ubuntu 22.0.4 LTS machine on the Azure Portal.

Remote into the Ubuntu machine using a keypair from your command prompt.

1. AKS Cluster is setup and running.

First check to see if Azure CLI is installed on your local machine.

Use the command, **az version** to confirm.

If you do not see a version of **azure-cli** then proceed to the following instructions:

Run the update first

sudo apt update

Second install with one command

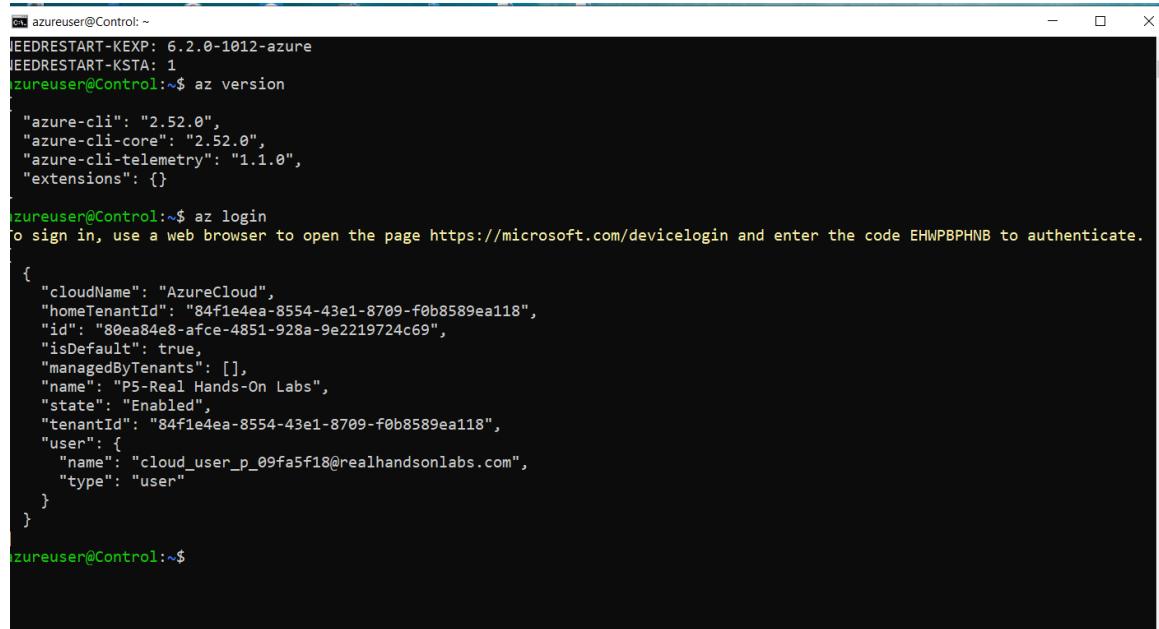
```
curl -sL https://aka.ms/InstallAzureCLIDeb | sudo bash
```

Check again to see if the system has a version of Azure CLI

az version

Second, insure you have an account setup in Azure cloud.

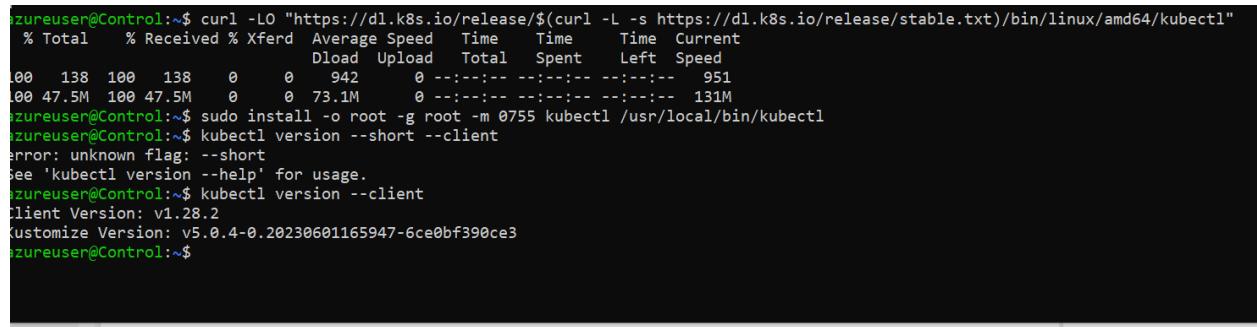
To sign in, use the **az login** command and follow the instructions in the command prompt.

az login


```
azureuser@Control:~$ az login
To sign in, use a web browser to open the page https://microsoft.com/devicelogin and enter the code EHWPBPHNB to authenticate.

{
  "cloudName": "AzureCloud",
  "homeTenantId": "84f1e4ea-8554-43e1-8709-f0b8589ea118",
  "id": "80ea84e8-afce-4851-928a-e2219724c69",
  "isDefault": true,
  "managedByTenants": [],
  "name": "P5-Real Hands-On Labs",
  "state": "Enabled",
  "tenantId": "84f1e4ea-8554-43e1-8709-f0b8589ea118",
  "user": {
    "name": "cloud_user_p_09fa5f18@realhandsonlabs.com",
    "type": "user"
  }
}

azureuser@Control:~$
```

Third, Install kubectl


```
azureuser@Control:~$ curl -LO "https://dl.k8s.io/release/$(curl -L -s https://dl.k8s.io/release/stable.txt)/bin/linux/amd64/kubectl"
  % Total    % Received % Xferd  Average Speed   Time   Time  Current
          Dload  Upload Total Spent   Left Speed
100  138  100  138    0     0  942      0 --:--:-- --:--:-- 951
100 47.5M  100 47.5M   0     0  73.1M      0 --:--:-- --:--:-- 131M
azureuser@Control:~$ sudo install -o root -g root -m 0755 kubectl /usr/local/bin/kubectl
azureuser@Control:~$ kubectl version --short --client
error: unknown flag: --short
See 'kubectl version --help' for usage.
azureuser@Control:~$ kubectl version --client
Client Version: v1.28.2
Kustomize Version: v5.0.4-0.20230601165947-6ce0bf390ce3
azureuser@Control:~$
```

Install kubectl binary with curl on Linux

```
curl -LO "https://dl.k8s.io/release/$(curl -L -s https://dl.k8s.io/release/stable.txt)/bin/linux/amd64/kubectl"
```

Install kubectl

```
sudo install -o root -g root -m 0755 kubectl /usr/local/bin/kubectl
```

Verify if kubectl got installed

kubectl version --short --client

```
ubuntu@ip-172-31-28-83:~$ kubectl version --short --client
Flag --short has been deprecated, and will be removed in the
become the default.
Client Version: v1.24.0
Kustomize Version: v4.5.4
ubuntu@ip-172-31-28-83:~$
```

Now let's setup the AKS cluster

Step 1

Make sure, once again, you are logged into the Azure portal first with
The following command:

az account show -o jsonc

```
azureuser@Control:~$ az account show -o jsonc
{
  "environmentName": "AzureCloud",
  "homeTenantId": "84f1e4ea-8554-43e1-8709-f0b8589ea118",
  "id": "80ea84e8-afce-4851-928a-9e2219724c69",
  "isDefault": true,
  "managedByTenants": [],
  "name": "PS-Real Hands-On Labs",
  "state": "Enabled",
  "tenantId": "84f1e4ea-8554-43e1-8709-f0b8589ea118",
  "user": {
    "name": "cloud_user_p_09fa5f18@realhandsonlabs.com",
    "type": "user"
  }
}azureuser@Control:~$
```

Step 2 - create a resource group first**az group create --name myResourceGroup --location southcentralus**Step 3 Create AKS cluster with 2 worker nodes**az aks create --resource-group myResourceGroup --name myAKSCluster --node-count 2****az aks create --resource-group 1-07ee85d4-playground-sandbox --name myAKSCluster --node-count 2**

NOTE: if you do not include the switch or flag, **--node-count 2** it will, by default, generate 2 nodes.

```
}
```

```
azureuser@Control:~$ az aks create --resource-group 1-07ee85d4-playground-sandbox --name myAKSCluster --node-count 2
■| Running ..
```

Output:

```
azureuser@Control:~$ az aks create --resource-group 1-07ee85d4-playground-sandbox --name
myAKSCluster --node-count 2
{
```

```
"aadProfile": null,
"addonProfiles": null,
"agentPoolProfiles": [
{
"availabilityZones": null,
"count": 2,
"creationData": null,
"currentOrchestratorVersion": "1.26.6",
"enableAutoScaling": false,
"enableEncryptionAtHost": false,
"enableFips": false,
"enableNodePublicIp": false,
"enableUltraSsd": false,
"gpuInstanceProfile": null,
"hostGroupId": null,
"kubeletConfig": null,
"kubeletDiskType": "OS",
"linuxOsConfig": null,
"maxCount": null,
"maxPods": 110,
"minCount": null,
"mode": "System",
"name": "nodepool1",
"nodeImageVersion": "AKSUbuntu-2204gen2containerd-202308.22.0",
"nodeLabels": null,
"nodePublicIpPrefixId": null,
"nodeTaints": null,
"orchestratorVersion": "1.26.6",
"osDiskSizeGb": 128,
"osDiskType": "Managed",
"osSku": "Ubuntu",
"osType": "Linux",
"podSubnetId": null,
"powerState": {
  "code": "Running"
},
"provisioningState": "Succeeded",
"proximityPlacementGroupId": null,
"scaleDownMode": null,
"scaleSetEvictionPolicy": null,
"scaleSetPriority": null,
"spotMaxPrice": null,
"tags": null,
"type": "VirtualMachineScaleSets",
"upgradeSettings": {
  "drainTimeoutInMinutes": null,
  "maxSurge": null
},
},
```

```

    "vmSize": "Standard_DS2_v2",
    "vnetSubnetId": null,
    "workloadRuntime": null
  },
],
"apiServerAccessProfile": null,
"autoScalerProfile": null,
"autoUpgradeProfile": {
  "nodeOsUpgradeChannel": "NodeImage",
  "upgradeChannel": null
},
"azureMonitorProfile": null,
"azurePortalFqdn": "myaksclust-1-07ee85d4-playg-80ea84-qkr1nwg3.portal.hcp.westus.azmk8s.io",
"currentKubernetesVersion": "1.26.6",
"disableLocalAccounts": false,
"diskEncryptionSetId": null,
"dnsPrefix": "myAKSClust-1-07ee85d4-playg-80ea84",
"enablePodSecurityPolicy": null,
"enableRbac": true,
"extendedLocation": null,
"fqdn": "myaksclust-1-07ee85d4-playg-80ea84-qkr1nwg3.hcp.westus.azmk8s.io",
"fqdnSubdomain": null,
"httpProxyConfig": null,
"id": "/subscriptions/80ea84e8-afce-4851-928a-9e2219724c69/resourcegroups/1-07ee85d4-
playground-sandbox/providers/Microsoft.ContainerService/managedClusters/myAKSCluster",
"identity": {
  "delegatedResources": null,
  "principalId": "26942b51-f827-4b6e-b0b0-6da6368b5e56",
  "tenantId": "84f1e4ea-8554-43e1-8709-f0b8589ea118",
  "type": "SystemAssigned",
  "userAssignedIdentities": null
},
"identityProfile": {
  "kubeletIdentity": {
    "clientId": "71c5b1ff-8c95-4cd8-b461-372eee4e2ffe",
    "objectId": "3cc16e6d-7e41-4097-a312-aadb5ae2adc0",
    "resourceId": "/subscriptions/80ea84e8-afce-4851-928a-9e2219724c69/resourcegroups/MC_1-
07ee85d4-playground-
sandbox_myAKSCluster_westus/providers/Microsoft.ManagedIdentity/userAssignedIdentities/myAKSCluster-agentpool"
  }
},
"kubernetesVersion": "1.26.6",
"linuxProfile": {
  "adminUsername": "azureuser",
  "ssh": {
    "publicKeys": [
      {
        "key": "-----BEGIN RSA PUBLIC KEY-----\n-----END RSA PUBLIC KEY-----"
      }
    ]
  }
}

```

```

    "keyData": "ssh-rsa
AAAAAB3NzaC1yc2EAAAQABAAQGrNWX7252YzD3Rr95Yq1NRommZLti9mmjcMEQYeHujWmy4
VyPV0eUdXd25rzzXk5CxG7JyTqvvmz+lubS/6L40blYll/Z4ckBrDRt9CEbJGDEYwb7qM9pNyoXtW0W+XfV4
Xpf/XIOG1OS/ei7DUs8mo8kNLwEyZXKr2ZvTpX51o9N7Ahvr81aiDLX6fcxyIL/x6zu57Xe1Zy2zqxhI/GfjdDQ4
HooejXndqqy21X1B8hoMT3jFWTfwTvQr5ItJmJswQpNmFa+WDV7166hQHroe7XBp6/GT9yECoP7Q331e
MKzUJ7X64JwEAUqmKhIzvmN6lzk+7OBau8X0cAnhMx"
    }
]
}
},
"location": "westus",
"maxAgentPools": 100,
"name": "myAKSCluster",
"networkProfile": {
  "dnsServiceIp": "10.0.0.10",
  "ipFamilies": [
    "IPv4"
  ],
  "loadBalancerProfile": {
    "allocatedOutboundPorts": null,
    "effectiveOutboundIPs": [
      {
        "id": "/subscriptions/80ea84e8-afce-4851-928a-9e2219724c69/resourceGroups/MC_1-07ee85d4-
playground-
sandbox_myAKSCluster_westus/providers/Microsoft.Network/publicIPAddresses/e8443d68-4b32-48fa-
bd74-71c057458ab5",
        "resourceGroup": "MC_1-07ee85d4-playground-sandbox_myAKSCluster_westus"
      }
    ],
    "enableMultipleStandardLoadBalancers": null,
    "idleTimeoutInMinutes": null,
    "managedOutboundIPs": {
      "count": 1,
      "countIpv6": null
    },
    "outboundIPs": null,
    "outboundIpPrefixes": null
  },
  "loadBalancerSku": "Standard",
  "natGatewayProfile": null,
  "networkDataplane": null,
  "networkMode": null,
  "networkPlugin": "kubernetes",
  "networkPluginMode": null,
  "networkPolicy": null,
  "outboundType": "loadBalancer",
  "podCidr": "10.244.0.0/16",
  "podCidrs": [
    ...
  ]
}

```

```

        "10.244.0.0/16"
    ],
    "serviceCidr": "10.0.0.0/16",
    "serviceCidrs": [
        "10.0.0.0/16"
    ]
},
"nodeResourceGroup": "MC_1-07ee85d4-playground-sandbox_myAKSCluster_westus",
"oidcIssuerProfile": {
    "enabled": false,
    "issuerUrl": null
},
"podIdentityProfile": null,
"powerState": {
    "code": "Running"
},
"privateFqdn": null,
"privateLinkResources": null,
"provisioningState": "Succeeded",
"publicNetworkAccess": null,
"resourceGroup": "1-07ee85d4-playground-sandbox",
"securityProfile": {
    "azureKeyVaultKms": null,
    "defender": null,
    "imageCleaner": null,
    "workloadIdentity": null
},
"servicePrincipalProfile": {
    "clientId": "msi",
    "secret": null
},
"sku": {
    "name": "Base",
    "tier": "Free"
},
"storageProfile": {
    "blobCsiDriver": null,
    "diskCsiDriver": {
        "enabled": true
    },
    "fileCsiDriver": {
        "enabled": true
    },
    "snapshotController": {
        "enabled": true
    }
},
"supportPlan": "KubernetesOfficial",

```

```

"systemData": null,
"tags": null,
"type": "Microsoft.ContainerService/ManagedClusters",
"upgradeSettings": null,
"windowsProfile": null,
"workloadAutoScalerProfile": {
  "keda": null,
  "verticalPodAutoscaler": null
}
}
}

```

Display Details of Cluster

```

az aks show --name myAKSCluster --resource-group myResourceGroup
az aks show --name myAKSCluster --resource-group 1-07ee85d4-playground-sandbox

```

The above command will display Cluster details.

| Node | Status | CPU | Memory | Disk | Pods | Kubernetes version |
|-----------------------------------|--------|-----|--------|------|------|--------------------|
| aks-nodepool1-35897272-vmss000000 | Ready | 10% | 28% | 16% | 10 | 1.26.6 |
| aks-nodepool1-35897272-vmss000001 | Ready | 10% | 30% | 16% | 12 | 1.26.6 |

The two nodes that we created can be seen above!

2. Jenkins Master and Slave is up and running.

Jenkins Master

Change Host Name to Jenkins

```
sudo hostnamectl set-hostname Jenkins
```

Perform update first

```
sudo apt update
```

Install Java 11

```
sudo apt install default-jdk -y
```

```
ubuntu@ip-172-31-44-89:~$ sudo apt install default-jdk -y
Reading package lists... Done
Building dependency tree
Reading state information... Done
The following additional packages will be installed:
at-spi2-core default-jdk-headless default-jre fonts-dejavu-extra libatk-bridge2.0-0
libatk-wrapper-java libatk-wrapper-java-jni libatk1.0-0 libatk1.0-data libatspi2.0-0
libdrm-amdgpu1 libdrm-intel1 libdrm-nouveau2 libdrm-radeon1 libfontenc1 libgif7 libgl1
libgl1-mesa-dri libglapi-mesa libglvnd0 libglx-mesa0 libglx0 libice-dev libice6
libllvm10 libpcicheck0 libpthread-stubs0-dev libsensors4 libsm-dev libsm6 libx11-dev
libx11-doc libx11-xcb1 libxau-dev libxaw7 libxcb-dri2-0 libxcb-dri3-0 libxcb-glx0
libxcb-present0 libxcb-shape0 libxcb-sync1 libxcb1-dev libxcomposite1 libxdamage1
```

Once install java, enter the below command.

Verify Java Version

java -version

```
ubuntu@ip-172-31-19-22:/opt$ java -version
openjdk version "11.0.6" 2020-01-14
OpenJDK Runtime Environment (Build 11.0.6+10-post-Ubuntu-1ubuntu118.04.1)
OpenJDK 64-Bit Server VM (build 11.0.6+10-post-Ubuntu-1ubuntu118.04.1, mixed mode, sharing)
ubuntu@ip-172-31-19-22:/opt$
```

Maven Installation

You can install Maven by executing below command:

sudo apt install maven -y

Now type the following:

mvn --version

```
ubuntu@ip-172-31-34-93:~$ mvn --version
Apache Maven 3.6.0
Maven home: /usr/share/maven
Java version: 11.0.7, vendor: Ubuntu, runtime: /usr/lib/jvm/java-11-openjdk-amd64
Default locale: en, platform encoding: UTF-8
OS name: "linux", version: "4.15.0-1065-aws", arch: "amd64", family: "unix"
ubuntu@ip-172-31-34-93:~$
```

Now let's start the Jenkins installation.

Jenkins Setup

Add a Repository key to the system

```
curl -fsSL https://pkg.jenkins.io/debian/jenkins.io-2023.key | sudo tee \
/usr/share/keyrings/jenkins-keyring.asc > /dev/null
```

```
ubuntu@ip-172-31-28-134:~$ curl -fsSL https://pkg.jenkins.io/debian/jenkins.io-2023.key | sudo tee \
/usr/share/keyrings/jenkins-keyring.asc > /dev/null
```

Append debian package repo address to the system

```
echo deb [signed-by=/usr/share/keyrings/jenkins-keyring.asc] \
https://pkg.jenkins.io/debian binary/ | sudo tee \
/etc/apt/sources.list.d/jenkins.list > /dev/null
```

```
ubuntu@ip-172-31-28-138:~$ echo deb [signed-by=/usr/share/keyrings/jenkins-keyring.gpg] \
https://pkg.jenkins.io/debian binary/ | sudo tee \
/etc/apt/sources.list.d/jenkins.list > /dev/null
```

Update Ubuntu package

sudo apt update

```
ubuntu@ip-172-31-37-211:~$ sudo apt-get update
Reading package lists... Done
Building dependency tree
Reading state information... Done
W: http://us-east-2.ec2.archive.ubuntu.com/ubuntu xenial InRelease
W: http://us-east-2.ec2.archive.ubuntu.com/ubuntu xenial-updates InRelease
W: http://us-east-2.ec2.archive.ubuntu.com/ubuntu xenial-backports
W: http://security.ubuntu.com/ubuntu xenial-security InRelease
W: http://ppa.launchpad.net/webupd8team/java/ubuntu xenial InRelease
ign:6 http://pkg.jenkins.io/debian-stable binary/ InRelease
get:7 http://pkg.jenkins.io/debian-stable binary/ Release [2,042 B]
get:8 http://pkg.jenkins.io/debian-stable binary/ Release.gpg [181 B]
get:9 http://pkg.jenkins.io/debian-stable binary/ Packages [13.2 kB]
Fetched 15.5 kB in 0s (17.8 kB/s)
Reading package lists... Done
```

Install Jenkins

sudo apt install jenkins -y

```
ubuntu@ip-172-31-37-211:~$ sudo apt-get install jenkins -y
Reading package lists... Done
Building dependency tree
Reading state information... Done
The following additional packages will be installed:
  daemon
The following NEW packages will be installed:
  daemon jenkins
0 upgraded, 2 newly installed, 0 to remove and 8 not upgraded.
Need to get 72.0 MB of archives.
After this operation, 75.1 MB of additional disk space will be used.
Get:1 http://us-east-2.ec2.archive.ubuntu.com/ubuntu xenial/universe amd64 daemon
Get:2 http://pkg.jenkins.io/debian-stable binary/ jenkins 2.121.3 [71.9 MB]
Fetched 72.0 MB in 5s (12.5 MB/s)
Selecting previously unselected package daemon.
(Reading database ... 51843 files and directories currently installed.)
Preparing to unpack .../daemon_0.6.4-1_amd64.deb ...
Unpacking daemon (0.6.4-1) ...
Selecting previously unselected package jenkins.
Preparing to unpack .../jenkins_2.121.3_all.deb ...
Unpacking jenkins (2.121.3) ...
Processing triggers for man-db (2.7.5-1) ...
Processing triggers for systemd (229-4ubuntu21.4) ...
Processing triggers for ureadahead (0.100.0-19) ...
Setting up daemon (0.6.4-1) ...
Setting up jenkins (2.121.3) ...
Processing triggers for systemd (229-4ubuntu21.4) ...
Processing triggers for ureadahead (0.100.0-19) ...
ubuntu@ip-172-31-37-211:~$
```

The above screenshot should confirm that Jenkins is successfully installed.

Allow Access in Jenkins in a web browser

Create an inbound rule within **Settings > Networking** for port **8080**.

Now go to browser and enter the virtual machines public IP address with port no **8080**.

http://Public IP address:8080

Unlock Jenkins

When this page loads, enter the below command in Git bash (Ubuntu console)

Unlock Jenkins

To ensure Jenkins is securely set up by the administrator, a password has been written to the log (not sure where to find it?) and this file on the server:

`/var/lib/jenkins/secrets/initialAdminPassword`

Please copy the password from either location and paste it below.

Administrator password

Get the initial password from the below file

`sudo cat /var/lib/jenkins/secrets/initialAdminPassword`

```
ubuntu@ip-172-31-37-211:~$ sudo cat /var/lib/jenkins/secrets/initialAdminPassword
df43c4644ce491298acefbac929bc5F
```

Now copy the password and paste it into the browser text box.

Then click on install suggested plug-ins.

Customize Jenkins

Plugins extend Jenkins with additional features to support many different needs.



Also create a username and password.

Enter everything as admin. At least username as **admin** password as **admin**

Create User

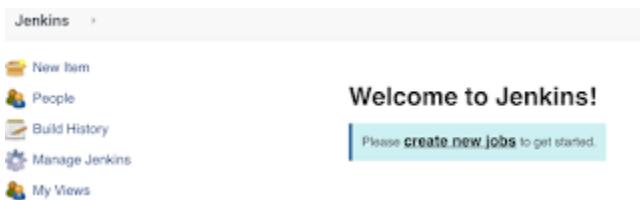
| | |
|-------------------|-----------------|
| Username: | admin |
| Password: | |
| Confirm password: | |
| Full name: | admin |
| E-mail address: | admin@admin.com |

Create User

Click on **Create User**.

Click on **Save and Finish**.

Click on **start using Jenkins**. Now you should see a screen like below:



Setup Jenkins slave. Then ensure it can run Docker builds.

Create a similar Ubuntu 22.4 LTS to be a slave.

Note: You need to create at least a t2.small Ubuntu 22.0.4 instance for this slave.
Only port 22 needs to be open

Configure it with the following:

Perform update first

sudo apt update

Install Java 11

sudo apt install default-jdk -y

```
ubuntu@ip-172-31-44-89:~$ sudo apt install default-jdk -y
Reading package lists... Done
Building dependency tree
Reading state information... Done
The following additional packages will be installed:
  at-spi2-core default-jdk-headless default-jre fonts-dejavu-extra libatk-bridge2.0-0
  libatk-wrapper-java libatk-wrapper-java-jni libatk1.0-0 libatk1.0-data libatspi2.0-0
  libdrm-amdgpu1 libdrm-intel1 libdrm-nouveau2 libdrm-radeon1 libfontenc1 libgif7 libgl1
  libgl1-mesa-dri libglapi-mesa libglvnd0 libglx-mesa0 libglx0 libice-dev libice6
  libllvm10 libpciauthentication0 libpthread-stubs0-dev libsensors4 libsm-dev libsm6 libx11-dev
  libx11-doc libx11-xcb1 libxau-dev libxaw7 libxcb-dri2-0 libxcb-dri3-0 libxcb-glx0
  libxcb-present0 libxcb-shape0 libxcb-sync1 libxcb1-dev libxcomposite1 libxdamage1
```

Once install java, enter the below command

Verify Java Version

java -version

```
ubuntu@ip-172-31-19-22:/opt$ java -version
openjdk version "11.0.6" 2020-01-14
OpenJDK Runtime Environment (Build 11.0.6+0-post-Ubuntu-1ubuntu118.04.1)
OpenJDK 64-Bit Server VM (Build 11.0.6+0-post-Ubuntu-1ubuntu118.04.1, mixed mode, sharing)
ubuntu@ip-172-31-19-22:/opt$
```

Maven Installation

You can install Maven by executing below command:

sudo apt install maven -y

you can type **mvn --version**

you should see the below output.

```
ubuntu@ip-172-31-34-93:~$ mvn -version
Apache Maven 3.6.0
Maven home: /usr/share/maven
Java version: 11.0.7, vendor: Ubuntu, runtime: /usr/lib/jvm/java-11-openjdk-amd64
Default locale: en, platform encoding: UTF-8
OS name: "linux", version: "4.15.0-1065-aws", arch: "amd64", family: "unix"
ubuntu@ip-172-31-34-93:~$
```

Steps for installing Docker

sudo apt-get update && sudo apt install docker.io -y

Add Jenkins to Docker Group

sudo usermod -aG docker jenkins

sudo newgrp docker

sudo systemctl daemon-reload

Restart Docker service

sudo systemctl start docker

sudo systemctl enable docker

sudo systemctl restart docker

sudo systemctl status docker

Test Permissions on Docker

docker run hello-world

permission denied while trying to connect to the Docker daemon socket at unix:///var/run/docker.sock:

Post "<http://%2Fvar%2Frun%2Fdocker.sock/v1.24/images/create?fromImage=tiangolo%2Fuwsgi-nginx-flask&tag=python3.6>": dial unix /var/run/docker.sock: connect: permission denied

Run the following:

sudo chmod 666 /var/run/docker.sock

Now run the following hello-world command again.

docker run hello-world

```

Hello from Docker!
This message shows that your installation appears to be working correctly.

To generate this message, Docker took the following steps:
1. The Docker client contacted the Docker daemon.
2. The Docker daemon pulled the "hello-world" image from the Docker Hub.
   (amd64)
3. The Docker daemon created a new container from that image which runs the
   executable that produces the output you are currently reading.
4. The Docker daemon streamed that output to the Docker client, which sent it
   to your terminal.

To try something more ambitious, you can run an Ubuntu container with:
$ docker run -it ubuntu bash

Share images, automate workflows, and more with a free Docker ID:
https://hub.docker.com/

For more examples and ideas, visit:
https://docs.docker.com/get-started/
azureuser@VM12:~$
```

Slave node configurationChange Host Name to Slave**sudo hostnamectl set-hostname Slave**Create User as Jenkins**sudo useradd -m jenkins****sudo -u jenkins mkdir /home/jenkins/.ssh**

```

ubuntu@ip-172-31-11-223:~$ sudo useradd -m jenkins
ubuntu@ip-172-31-11-223:~$ sudo -u jenkins mkdir /home/jenkins/.ssh
ubuntu@ip-172-31-11-223:~$
```

Now Login to Jenkins Master

Create SSH keys by executing below command:

ssh-keygen -t rsa -m PEM

Please overwrite your existing keys.

```

azureuser@Control:~$ sudo cat /var/lib/jenkins/secrets/initialAdminPassword
02fbe195b2b0433681388effdc07abce
azureuser@Control:~$ ssh-keygen -t rsa -m PEM
Generating public/private rsa key pair.
Enter file in which to save the key (/home/azureuser/.ssh/id_rsa):
/home/azureuser/.ssh/id_rsa already exists.
Overwrite (y/n)? y
Enter passphrase (empty for no passphrase):
Enter same passphrase again:
Your identification has been saved in /home/azureuser/.ssh/id_rsa
Your public key has been saved in /home/azureuser/.ssh/id_rsa.pub
The key fingerprint is:
SHA256:/TbxQmoycKBn7Wkjzz1NkhG9MVRKkSVWTRee48GRmU azureuser@Control
The key's randomart image is:
+---[RSA 3072]---+
|       .oooB=.oE |
|     .. =o.o.. |
|    . . . . o. . |
|   . . . . . . . |
|  . . . S.....o . |
| oo.= +. +... |
| o += = + ..o |
| .oo+o . . . |
|      ++o. |
+---[SHA256]---+
azureuser@Control:~$
```

Copy SSH Keys from Master to Slave

Execute the below command in Jenkins master EC2.

sudo cat ~/.ssh/id_rsa.pub

Copy the output of the above command:

```
ssh-rsa AAAAB3NzaC1yc2EAAAQABAAQgQDcKK7OR/+aMijfkUkqGXCOEqZMBY+069Dna38HIDtzAiKOYYqDuElcpbzOHSxjyrD6CPbaL6SC+Bx3zW2JIIWUVJZ3qsELjCzgLfKQLyvd9aWARzQcbcJGIOtMEx/kL2IUsb+sIFyDfI592nk0QpG5fYBC6Y8LrG9+FgQQRue8o7BC8aHvQuZ6o1c4k+aefc2CqQexonF9cDF4KMLnzYB+h3RRHUGMJ/ECd6OxmbprzQ4p1AeMcIV55OaJ7y34c/Pqq6Ilv2CF2ZpEd9dcROYW+hsNAZRHGEGdxmXIA5pmFuLh3DBVzUkVV0FL+/h5NUxfecY13BVTEfsQehCEdc1IE69oVUD3tsfxlt61IZrEdtkAcoJ9iLPY3WF8zLS1vMI7eeHDiYdSHSgFuZe2z/2D4peuY9/dKwc2nLqHSraD5vir6Xcz79YTx8+tA5r9aPpW9+hKV0FT9GiuzmyMjk8j/383wWcLqVTM2ZAkg44o+Y9bsK6Z4YJ1FY6qIY0= azureuser@Control
```

Now Login to Slave node and execute the below command

sudo -u jenkins vi /home/jenkins/.ssh/authorized_keys

This will be an empty file, now copy the public keys from **master** into above file.

Once you pasted the public keys in the above file in Slave, come out of the file by entering :wq!

Now go into master node**ssh jenkins@slave_node_ip****ssh jenkins@40.86.171.169**

```
ubuntu@ip-172-31-44-147:~$ ssh jenkins@172.31.11.223
The authenticity of host '172.31.11.223' can't be established.
ECDSA key fingerprint is SHA256:z3tpstU791wVTGNSbPgXfkScDv1Gj8S1vJI3Zwi5o/k.
Are you sure you want to continue connecting (yes/no)? yes
```

```
jenkins@Slave: ~
The authenticity of host '40.86.171.169 (40.86.171.169)' can't be established.
ED25519 key fingerprint is SHA256:OVLBA39AeEpXaall1yDE3SWpKQsjHutofNoGaOKMf/Bw.
This key is not known by any other names
Are you sure you want to continue connecting (yes/no/[fingerprint])? yes
Warning: Permanently added '40.86.171.169' (ED25519) to the list of known hosts.
Welcome to Ubuntu 22.04.3 LTS (GNU/Linux 6.2.0-1012-azure x86_64)

 * Documentation:  https://help.ubuntu.com
 * Management:     https://landscape.canonical.com
 * Support:        https://ubuntu.com/advantage

System information as of Wed Sep 20 10:56:48 UTC 2023

System load:  0.0615234375      Processes:           116
Usage of /:   6.5% of 28.89GB   Users logged in:    1
Memory usage: 4%                  IPv4 address for eth0: 10.0.0.5
Swap usage:   0%

Expanded Security Maintenance for Applications is not enabled.

0 updates can be applied immediately.

2 additional security updates can be applied with ESM Apps.
Learn more about enabling ESM Apps service at https://ubuntu.com/esm

The programs included with the Ubuntu system are free software;
the exact distribution terms for each program are described in the
individual files in /usr/share/doc/*/*copyright.

Ubuntu comes with ABSOLUTELY NO WARRANTY, to the extent permitted by
applicable law.

jenkins@Slave:~$
```

This is to make sure master is able to connect slave node. once you are successfully logged into slave, type **exit** to come out of slave.

```
Ubuntu comes with ABSOLUTELY NO WARRANTY, to the extent permitted by
applicable law.

$ exit
```

Register slave node in Jenkins:

Now to go to Jenkins Master.

Manage Jenkins > Manage nodes.

Node name: Slave
 Permanent Agent
Adds a plain, permanent other agent types apply
 Copy Existing Node
 Build Queue

Click on New Node.

Give name and check **permanent agent**.

Set no of executors as **1**.

Enter **/home/jenkins** as remote directory.

Name: Slave
 Description:
 # of executors: 1
 Remote root directory: /home/jenkins
 Labels:
 Usage: Use this node as much as possible
 Launch method: Launch agent agents via SSH
 Host: 172.31.11.223
 Credentials: jenkins (updated my private keys)
 Host Key Verification Strategy: Manually trusted key verification strategy

Select launch method as **Launch slaves nodes via SSH**.

Enter Slave node IP address as Host.

Click on Manage Jenkins > credentials.

Enter user name as **jenkins**.

Note: Make **jenkins** as lowercase as it is shown.

Kind as **SSH username with private key**.

Enter private key of master node directly by executing below command:



sudo cat ~/ssh/id_rsa

(Make sure you copy the whole key including the below without missing anything)

-----BEGIN RSA PRIVATE KEY-----

-----END RSA PRIVATE KEY-----

```
az@azureuser@Control: ~
azureuser@Control:~$ sudo cat ~/.ssh/id_rsa
-----BEGIN RSA PRIVATE KEY-----
MIIG5AIBAAKCAYEa3ciuzkfv2jIiX5FJKhlwtBKmTAwPt0vQ52t/B5Q7cwIitGK
g7HCKYWT0sY8qw+gj4Wi+kgvgcd81tISCF1FSWd6rBC4ws4C37SkC8r3fwlgEc
0Am3CRpTk5fH+c1FLG+rCbcgyBZefdp5NEKRu2AQumPC6xvfhYEEK7nvK0wQv
Gh70FM+qNKOJPmn3NgqkHsaJxfXAxeCjC582Afod0UR1BjCfxAnejsZm6a80Kd
QHjHCFeeTmie8t+HPz6qupSL9ghdmaRHFXETmFvobDQGURxhIHcZ15QoAzbh4d
wwC13FVdBS/v4eTVMX3nGndwVUxH7EHoQhHNZNROvaFVA97bH8ZbetSGaxHbZA
KCFYiz2N1hfMy0tbzJe3nhw4mMnUh0oBbmXts/9g+KXrmPf3ssHNpy6h0q2g+b4q
+13M+/WE8FPrQ0a/Wj6VvfoS1dBu/RorkWzsjCZP/9/N8FnC6lUzNmQjIOOKPmP
W7CumeGCdRW0qpWNAGMBAEoggGASgP/aOh0ccXjG0AMuf0dmVyfVfUssCWLsRHJ
vjG0nwa7xir413n+N+DW6T/6s2svdHeKAJEnaJBHn+Dr/NVfSs0CTSZ4Fkq1ItQGg
dwE2rKM6yGmNVz+84nmQuZGbDuWJYth5gEy06DPeGPWP5nOz0aLGNOCLde34zUwu
5p5XX4AMqDs5Twp80PZUAe5a2LngMrYV8UcBr7rz82UnBZvnavlVq0t2p/i/Atra
etR9sn+nAQ7+0RsJxvs2BKRVs+D2BVdt0XZh5S1377u+PFQlwEZKGJ5U4m74wq
UNefkLtl+SShyggn6zSua1lk8n0EdEBaQ/dEWox1zwL3H0WrmEnaUN1aKisQ20Xp
bxHRPNorwKRSKHYtqCs05748s2tR0VCTdipp6NPJFPSufeZIdKJ30zeGopbsu51
f+4iwMyb03PFWJniUkfMp0G+Yv2fppA637yD14yPs+oX3+AIlmfTs5kLL461RsF
rdJ7a2HeSw0X276YHx9Te10teaKdAoHBA02bsHe2fVFxFx8i3R/sLvwWix1/himSL
mbB7ouozBpig5anQNb1+0heCja3D+w3mFyH5oXM8R4pmxHYDv8ZPrEQ/ENBwgjqQ
bpizs1tin6Mqt2/x274f82voxiXnseJRLCPw44KzljM2tYSnes+Hod8/HvKG+ZRG
9BcbaNlwRYutHrgXSj2za6W2HTsak6h03tp3usMli2VY7qyWcPr/Uipy3IW0mgj
b/y8VuBeEkhiYsjylBHQeb+wx2HyoiT/owKbwDtMzqhowu7ofZt/teTna0+yL6e
QctnCyufhGvqiHx2tJ5DrZjqqPK0rLhwUeznc1hc9kVpZR0D77Zt8/f1Hb15ow
0SjBjDAcLuUn+swnswQ4dv+tnd5C/30Tt0f0HYDkDPP8heaEIoDoeLygZ8PFrC7b+
XmNm6ku5Szbs0jQ04c2CgqHz+qDm496AChvTmzPN8P7vAnET7sIAaSS3OffHi0
/T3BuKwD+vA6NdKbpj+2le7D+1TOdyK01svqQ8CgcEAi5V+oZUBaI1JyFwDaTOJ
UG1EBcuVZW0tMcQv/K11HpzGzm+GyN/yI-VNBtUvecRHsqo/H0C6FSITBpIJqohH
r/VufIRQ3KS1k/ddUj4p0cy1ReorURGsCqR/cowo+/HooWM4edDAfq2zhskUgt6
r++A2rD4Rw5VanAepw9+vJ5pjTpjtj2E66CEzuWjulhxwbrk0iay6qxshelQvox
4DBhV/A+95Z531s1HjGkT4cdAF4gs6Z90Z+msfts8dJaoHAYCpm20VLens9ffbV
26kPdebZMTt+3TVhgBKjMz+uzKkUr7PBSkLU4KPkvi1as9StBw9s15ekpk4tCqsP
q2so9jseG10dQSQStTr1+PF7F5VCXsoP+omIdM00hmunt+eM5b0zCpVPljoX1hNe
VLYUinujYxinUAalUKgpCkr/U31AKXdcCS2eLF/X1nR98DH00IUhVaHKm+ZfQ22
3I//Qx1/mpEZGp03ArYttHxy4qFYMeq/q1YUgb6XVmZkwicNaohBAKvhqwCst+bM
TZ2pMbTy2swGxoJVOAReyJ7HTWRhWvc3y8dQh0FhXM0mePmUSGwhYJD+VSbpqId
gn+cSsNdAqAIRCz3jfdwjVd1xJd2u3GZwVzViHzd7P+oPNINVtvPn1taaLuwGE
Xn0+qgyZP4mBmMoP/LMQODizzxbpAJinzpvVJgWjCG8VVGEDwJ3c6D0/LA9cA1wG
ipy+WuAKGRGIep8DH1Rtt28h9cWRtxS1EkppR3YRI2CD/eHe9wT14Q==

-----END RSA PRIVATE KEY-----
azureuser@Control:~$
```

Click **Save**.

Select **Host key verification strategy** as "manually trusted key verification strategy".

Launch agents via SSH

Host ?
40.86.171.169

Credentials ?
jenkins

Add ▾

Host Key Verification Strategy ?
Manually trusted key Verification Strategy

Require manual verification of initial connection ?

Advanced ▾ Edited

Availability ?
Keep this agent online as much as possible

Save

Click **Save**.

Click on launch agent..make sure it connects to agent node.

```
INFO: Both error and output logs will be printed to /home/jenkins/remoting
<===[JENKINS REMOTING CAPACITY]==>channel started
Remoting version: 3131.vf2b_b_798b_ce99
Launcher: SSHLauncher
Communication Protocol: Standard in/out
This is a Unix agent
WARNING: An illegal reflective access operation has occurred
WARNING: Illegal reflective access by jenkins.slaves.StandardOutputSwapper$ChannelSwapper to constructor java.io.FileDescriptor(int)
WARNING: Please consider reporting this to the maintainers of jenkins.slaves.StandardOutputSwapper$Chan
WARNING: Use --illegal-access=warn to enable warnings of further illegal reflective access operations
WARNING: All illegal access operations will be denied in a future release
Evacuated stdout
Agent successfully connected and online
```

| S | Name ↓ | Architecture | ... |
|---|--------|---------------|-----|
| | master | Linux (amd64) | |
| | Slave | Linux (amd64) | |

The screenshot shows the Jenkins 'Nodes' page. On the left, there's a sidebar with 'Nodes', 'Clouds', 'Node Monitoring', 'Build Queue' (with a note about builds in the queue), and 'Build Executor Status'. The main area is titled 'Nodes' and contains a table with columns: S, Name, Architecture, Clock Difference, Free Disk Space, Free Swap Space, Free Temp Space, Response Time, and a 'Configure' icon. Two nodes are listed: 'Built-In Node' (Linux amd64) and 'Slave' (Linux amd64). Both nodes show 'In sync' status, 24.75 GB and 26.99 GB free disk space, 0 B free swap space, 24.75 GB and 26.99 GB free temp space, and response times of 0ms and 77ms respectively.

Now you can kick start building the jobs, you will see Jenkins master runs jobs in slave nodes.

3. Docker, Docker pipeline and Kubectl CLI plug-ins are installed in Jenkins.

Managing Jenkins > Plugins

Available Plugins

Select and install the following:

Docker
Docker Pipeline
Kubectl CLI

The screenshot shows the Jenkins 'Manage Jenkins' > 'Plugins' page. A search bar at the top has 'kubernetes' typed into it. Below the search bar, there's an 'Install' button. Under the heading 'Available Plugins', two checkboxes are checked: 'Docker' and 'Docker Pipeline'. The 'Docker' plugin is described as integrating Jenkins with Docker, and the 'Docker Pipeline' plugin is described as building and using Docker containers from pipelines. Below this section, a list of other Kubernetes-related plugins is shown, each with a checkbox and a brief description. Most of these checkboxes are also checked.

| Plugin | Description | Last Updated |
|--|--|---------------|
| Kubernetes | build and use Docker containers from pipelines. | 21 days ago |
| Kubernetes Credentials | This plugin integrates Jenkins with Kubernetes | 21 days ago |
| Kubernetes Client API | Kubernetes Client API plugin for use by other Jenkins plugins. | 21 days ago |
| Kubernetes CLI | Configure kubectl for Kubernetes | 21 days ago |
| Kubernetes Credentials Provider | Provides a read only credentials store backed by Kubernetes. | 15 days ago |
| Kubernetes :: Pipeline :: DevOps Steps | pipeline kubernetes | 4 yr 7 mo ago |

→ [Go back to the top page](#)

(you can start using the installed plugins right away)

→ Restart Jenkins when installation is complete and no jobs are running

4. Azure container registry is setup in <https://portal.azure.com>

Run the below command to create your own private container registry using Azure Container Registry (ACR).

Make sure below green marked registry name is unique.

```
az acr create --resource-group myResourceGroup --name myacrrepo4321 --sku Standard --location southcentralus
```

```
az acr create --resource-group 1-07ee85d4-playground-sandbox --name costasrepo4321 --sku Standard --location westus
```

```
azureuser@Control ~
az acr create --resource-group 1-07ee85d4-playground-sandbox --name costasrepo4321 --sku Standard --location westus

{
  "adminUserEnabled": false,
  "anonymousPullEnabled": false,
  "creationDate": "2023-09-20T08:52:48.349911+00:00",
  "dataEndpointEnabled": false,
  "dataEndpointHostNames": [],
  "encryption": {
    "keyVaultProperties": null,
    "status": "disabled"
  },
  "id": "/subscriptions/80ea84e8-acfe-4851-928a-9e2219724c69/resourceGroups/1-07ee85d4-playground-sandbox/providers/Microsoft.ContainerRegistry/registries/costasrepo4321",
  "identity": null,
  "location": "westus",
  "loginServer": "costasrepo4321.azurecr.io",
  "name": "costasrepo4321",
  "networkRuleResources": "AzureServices"
}
```

Connect to the cluster

```
az aks get-credentials --resource-group myResourceGroup --name myAKSCluster --overwrite-existing
```

```
az aks get-credentials --resource-group 1-07ee85d4-playground-sandbox --name myAKSCluster --overwrite-existing
```

```
devopscoaching@DevOps-Coach-MacBook-Pro:~/Downloads$ az aks get-credentials --resource-group myResourceGroup --name myAKSCluster --overwrite-existing
Merged "myAKSCluster" as current context in /Users/devopscoaching/.kube/config
```

To verify the connection to your cluster, use the `kubectl get` command to return a list of the cluster nodes.

`kubectl get nodes`

```
AK-DevOps-01:~ devopscoaching$ kubectl get nodes
NAME           STATUS  ROLES   AGE    VERSION
aks-nodepool1-26195414-0  Ready   agent   13m   v1.17.11
aks-nodepool1-26195414-1  Ready   agent   13m   v1.17.11
```

```
# List all deployments in a specific namespace  
kubectl get deployments --all-namespaces=true
```

| AK-DevOps-01:- devopscoaching\$ kubectl get deployments --all-namespaces=true | | | | | |
|---|--------------------|-------|------------|-----------|-------|
| NAMESPACE | NAME | READY | UP-TO-DATE | AVAILABLE | AGE |
| kube-system | coredns | 2/2 | 2 | 2 | 6m25s |
| kube-system | coredns-autoscaler | 1/1 | 1 | 1 | 6m25s |
| kube-system | metrics-server | 1/1 | 1 | 1 | 6m25s |
| kube-system | omsagent-rs | 1/1 | 1 | 1 | 6m25s |
| kube-system | tunnelfront | 1/1 | 1 | 1 | 6m24s |

```
az azureuser@Control:~  
azureuser@Control:~$ az aks get-credentials --resource-group 1-07ee85d4-playground-sandbox --name myAKSCluster --overwrite-existing  
Merged "myAKSCluster" as current context in /home/azureuser/.kube/config  
azureuser@Control:~$ kubectl get nodes  
NAME STATUS ROLES AGE VERSION  
aks-nodepool1-35897272-vmss00000 Ready agent 15m v1.26.6  
aks-nodepool1-35897272-vmss00001 Ready agent 15m v1.26.6  
azureuser@Control:~$  
azureuser@Control:~$ kubectl get deployments --all-namespaces=true  
NAMESPACE NAME READY UP-TO-DATE AVAILABLE AGE  
gatekeeper-system gatekeeper-audit 1/1 1 1 7m25s  
gatekeeper-system gatekeeper-controller 2/2 2 2 7m25s  
kube-system azure-policy 1/1 1 1 7m26s  
kube-system azure-policy-webhook 1/1 1 1 7m25s  
kube-system coredns 2/2 2 2 16m  
kube-system coredns-autoscaler 1/1 1 1 16m  
kube-system konnectivity-agent 2/2 2 2 16m  
kube-system metrics-server 2/2 2 2 16m  
azureuser@Control:~$  
azureuser@Control:~$
```

Enter ID as K8S and choose enter directly and paste the above file content and save.

Global credentials (unrestricted)

Kind: Kubernetes configuration (kubeconfig)

Scope: Global (Jenkins, nodes, items, all child items, etc)

ID: K8S

Description: kube configuration

Kubeconfig: Enter directly
Content:

```
apiVersion: v1
clusters:
- cluster:
  certificate-authority-data:
LSQLS1CRUDUTIBDRVJUSUZ/JQFURS0tL80tCk1JSUN5RENDQWJD
```

 From a file on the Jenkins master
 From a file on the Kubernetes master node

OK

Step # 3 - Create a pipeline in Jenkins

Create a new pipeline job.

Enter an item name

myKubernetesDeployJob » Required field

Freestyle project

This is the central feature of Jenkins for building something other than software builds.

Pipeline

Orchestrates long-running activities by organizing complex activities that

Step # 4 - Copy the pipeline code from below

Make sure you change red highlighted values below:

Your docker user id should be updated.

your registry credentials ID from Jenkins from step # 1 should be copied

```
pipeline {
    agent {
        label 'mystave'
    }
    environment {
        //once you sign up for Docker hub, use that user_id here
        registry = "your_docker_hub_user_id/mypython-app"
        // update your credentials ID after creating credentials for connecting to Docker Hub
        registryCredential = 'dockerhub'
```

```

    dockerImage = "
}
stages {

    stage ('checkout') {
        steps {
            checkout([$class: 'GitSCM', branches: [[name: '/master']], doGenerateSubmoduleConfigurations: false, extensions: [], submoduleCfg: [], userRemoteConfigs: [[url: 'https://github.com/akannan1087/myPythonDockerRepo']]])
        }
    }

    stage ('Build docker image') {
        steps {
            script {
                dockerImage = docker.build registry
            }
        }
    }

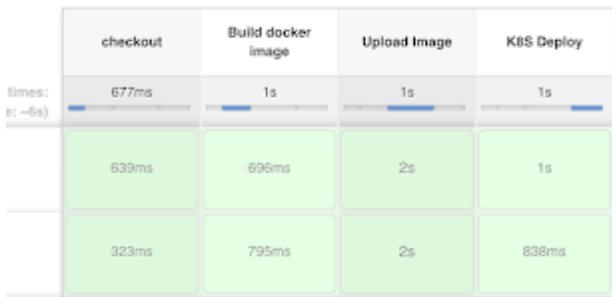
    // Uploading Docker images into Docker Hub
    stage('Upload Image') {
        steps{
            script {
                docker.withRegistry( "", registryCredential ) {
                    dockerImage.push()
                }
            }
        }
    }

    stage ('K8S Deploy') {
        steps {
            script {
                kubernetesDeploy(
                    configs: 'k8s-deployment.yaml',
                    kubeconfigId: 'K8S',
                    enableConfigSubstitution: true
                )
            }
        }
    }
}

```

[Step # 5 - Build the pipeline](#)

Once you create the pipeline and changes values per your Docker user id and credentials ID, click on



Step # 6 - Verify deployments to K8S

`kubectl get pods`

```
ubuntu@ip-172-31-2-128:~$ kubectl get pods
NAME           READY   STATUS    RESTARTS   AGE
ok1-controller-h8dck   1/1     Running   3          23h
my-php-deployment-8cd555cd6-b25bv   1/1     Running   0          4m2s
my-php-deployment-8cd555cd6-wbh1k   1/1     Running   0          4m2s
```

`kubectl get deployments`

```
ubuntu@ip-172-31-2-128:~$ kubectl get deployments
NAME           READY   UP-TO-DATE   AVAILABLE   AGE
my-php-deployment   2/2     2           2           7m27s
ubuntu@ip-172-31-2-128:~$
```

`kubectl get services`

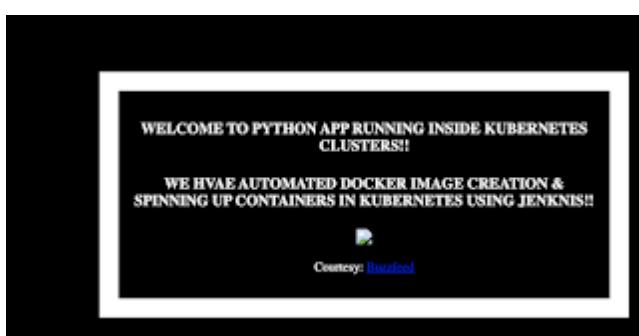
```
ubuntu@ip-172-31-2-128:~$ kubectl get services
NAME            TYPE      CLUSTER-IP   EXTERNAL-IP   PORT(S)        AGE
ok-angular-app   NodePort   10.103.50.84 <none>       8885:32004/TCP   23h
kubernetes       ClusterIP  10.96.0.1    <none>       443/TCP       26h
php-app-svc     LoadBalancer 10.97.39.217 <pending>    5000:30814/TCP   3m31s
ubuntu@ip-172-31-2-128:~$
```

Steps # 7 - Access Python App in K8S cluster

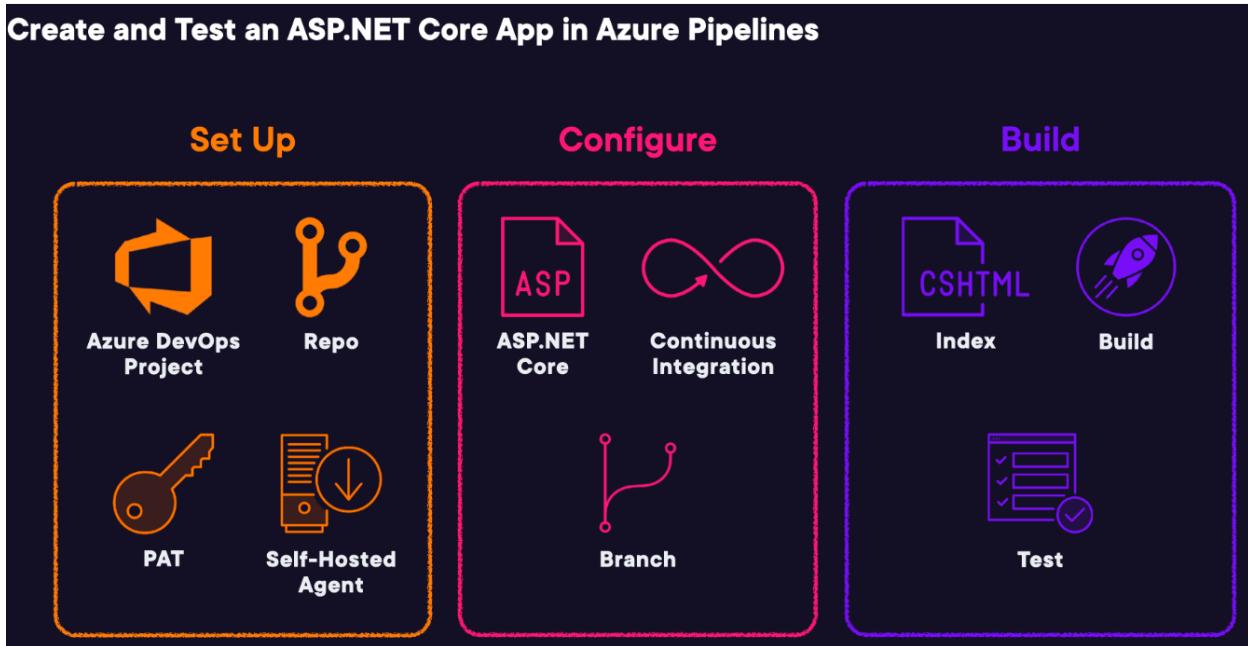
Once build is successful, go to browser and enter master or worker node public ip address along with port number mentioned above

`http://master_or_worker_node_public_ipaddress:port_no_from_above`

You should see page like below:



Create and Test an ASP.NET Core App in Azure Pipelines



Create the VM (and other resources using ARM)

Screenshot of the Azure portal showing a Resource Group named "1-b2909d3f-playground-sandbox". The portal interface includes a left sidebar with navigation links like Home, Overview, Activity log, Access control (IAM), Tags, Resource visualizer, Events, Deployments, Security, Deployment stacks, Policies, Properties, and Locks. The main content area displays the "Essentials" section for the resource group, showing details such as Subscription (move), Subscription ID, Tags, Deployments, Location (West US), and a Resources section with filtering and sorting options.

Go to the **Market Place** and locate an Ubuntu image

Microsoft Azure Search resources, services, and docs (G+)

Home > Marketplace > Create a virtual machine

Region * (US) West US

Availability options No infrastructure redundancy required

Security type Trusted launch virtual machines
Configure security features

Image * Ubuntu Server 20.04 LTS - x64 Gen2
See all images | Configure VM generation

VM architecture x64

Run with Azure Spot discount

Size * Standard_D2s_v3 - 2 vcpus, 8 GiB memory (US\$85.41/month)
See all sizes
Item(s) availability based on policy assignment(s) for the selected scope.
vadergen2labs/Microsoft.Authorization/8ab9bdb77b304947a1347e88 (Policy details)

Administrator account

Authentication type SSH public key
Password

Azure now automatically generates an SSH key pair for you and allows you to store it for future use. It is a fast, simple, and secure way to connect to your virtual machine.

Username * azureuser

SSH public key source Generate new key pair

Key pair name * linVM_key

Ensure that the **Region** is the same as the Resource group.

NOTE: Ubuntu 20.04 LTS was the only option.

Choose the **Size** Standard D2 with 2 CPUs as shown above.

For **Authentication** type, choose **SSH Public Key** and download the key pair locally to your system.
Go to end and create the ubuntu virtual machine.

NOTE: When you complete creating the VM all the other associated resources will also be created.

Navigate to the Virtual machine **linVM** within the resource group and select the link.

On the left hand menu, navigate to Settings and select **connect > Native SSH**

The screenshot shows the Azure portal interface for connecting to a VM. On the left, the 'Connect' option is selected under the 'Networking' section. The main area displays the 'Native SSH' configuration, which requires port 22 access and a public IP address. A command line example is shown:

```
ssh -i ~/ssh/id_rsa.pem azureuser@20.253.229.27
```

In a command prompt on your system type the following:

ssh -i <path to key pair>/<key pair name>.pem <username created earlier>@<public IP address of VM>

Once logged in, proceed to create the Pipeline within Azure DEVOPS.

Create the Azure DEVOPS Pipeline

In the main search bar, search for and select **Azure DevOps organizations**.

Scroll down and click **My Azure DevOps Organizations**.

Select your country/region from the dropdown list.

Click **Continue**.

Select **Create new organization > Continue > Continue** (complete the required CAPTCHA).

Set the following values to create a project:

Project name: **ASP_.NET**

Visibility: **Private**

Click + **Create project**.

Obtain and Push Project Code

In the Linux VM, pull code from the unit-test branch using the following URL:

git clone -b unit-test <https://github.com/ACloudGuru-Resources/content-az400-lab-resources.git>

```
nices_and_dacs (G+) 
azuser@linVM: ~

Enable ESM Apps to receive additional future security updates.
See https://ubuntu.com/esm or run: sudo pro status

The list of available updates is more than a week old.
To check for new updates run: sudo apt update

The programs included with the Ubuntu system are free software;
the exact distribution terms for each program are described in the
individual files in /usr/share/doc/*copyright.

Ubuntu comes with ABSOLUTELY NO WARRANTY, to the extent permitted by
applicable law.

To run a command as administrator (user "root"), use "sudo <command>".
See "man sudo_root" for details.

azuser@linVM:~$ git clone -b unit-test https://github.com/ACloudGuru-Resources/content-az400
Cloning into 'content-az400-lab-resources'...
remote: Enumerating objects: 4353, done.
remote: Counting objects: 100% (111/111), done.
remote: Compressing objects: 100% (104/104), done.
remote: Total 4353 (delta 29), reused 17 (delta 6), pack-reused 4242
Receiving objects: 100% (4353/4353), 356.10 MiB | 29.57 MiB/s, done.
Resolving deltas: 100% (1508/1508), done.
azuser@linVM:~$ █
```

List the contents of the directory with **ls**

Observe the folder is cloned to the VM.

Change into the directory shown:

cd content-az400-lab-resources/

View the remote origin:

git remote -v

It is currently the ACG GitHub repo.

Remove the remote origin:

git remote remove origin

```

az user@linVM: ~/content-az400-lab-resources
az user@linVM:~$ git clone -b unit-test https://github.com/ACloudGuru-Resources/content-az400-lab-resources.git
Cloning into 'content-az400-lab-resources'...
remote: Enumerating objects: 4353, done.
remote: Counting objects: 100% (111/111), done.
remote: Compressing objects: 100% (104/104), done.
remote: Total 4353 (delta 29), reused 17 (delta 6), pack-reused 4242
Receiving objects: 100% (4353/4353), 356.10 MiB | 29.57 MiB/s, done.
Resolving deltas: 100% (1508/1508), done.
az user@linVM:~$ ls
content-az400-lab-resources
az user@linVM:~$ cd content-az400-lab-resources/
az user@linVM:~/content-az400-lab-resources$ git remote -v
origin https://github.com/ACloudGuru-Resources/content-az400-lab-resources.git (fetch)
origin https://github.com/ACloudGuru-Resources/content-az400-lab-resources.git (push)
az user@linVM:~/content-az400-lab-resources$ git remote remove origin
az user@linVM:~/content-az400-lab-resources$ git remote -v
az user@linVM:~/content-az400-lab-resources$ git remote add origin https://clouduserp4b842ee7@dev.azure.com/clouduserp4b842ee7/ASP_Net/_git/ASP_Net
az user@linVM:~/content-az400-lab-resources$ git push -u origin --all
Password for 'https://clouduserp4b842ee7@dev.azure.com':

```

If you use the **git remote -v** command again you will find that nothing comes up!

Go back to **Azure DevOps**. From the left menu, click **Repos**.

Under **Push an existing repository** from command line, copy the commands.

The screenshot shows the Azure DevOps interface for managing repositories. On the left, there's a sidebar with options like 'Repos', 'Files', 'Commits', 'Pushes', 'Branches', 'Tags', 'Pull requests', 'Pipelines', and 'Test Plans'. The 'Repos' option is selected. In the main area, there's a 'Generate Git Credentials' button and a note about authentication. Below that, under 'Push an existing repository from command line', there are tabs for 'HTTPS' and 'SSH'. The 'HTTPS' tab is selected, showing the URL: `https://clouduserp9d19bec0@dev.azure.com/clouduserp9d19bec0/MyFirstProject/_git/MyFirstProject`. A terminal window at the bottom shows the command being run:

```

cloud_user@linVM:~/content-az400-lab-resources$ git remote -v
cloud_user@linVM:~/content-az400-lab-resources$ git remote add origin https://clouduserp9d19bec0@dev.azure.com/clouduserp9d19bec0/MyFirstProject/_git/MyFirstProject
cloud_user@linVM:~/content-az400-lab-resources$ git push -u origin --all
Password for 'https://clouduserp9d19bec0@dev.azure.com':

```

Return to the Linux VM and paste in the commands. You will be asked for a personal access token.

Create Personal Access Token

Back in Azure DevOps, in the upper right-hand corner, click the **User settings icon > Personal access tokens**.

Click **+ New Token** and set the following values:

Name: **linuxvm**
 Scopes: **Full access**
 Click **Create**.

Create a new personal access token X

Name

Organization

Expiration (UTC)
 (calendar icon)

Scopes
 Authorize the scope of access associated with this token
 Full access
 Custom defined

It will create a string like below:

das5yco2kl474qngkvp3fikjev6vtgxgymaverwa7kasm5nq3xfkq

Copy the **linuxvm** token and save it in a safe location.

Click **Close**.

```
azureuser@linVM: ~/content-az400-lab-resources
emote: Compressing objects: 100% (104/104), done.
emote: Total 4353 (delta 29), reused 17 (delta 6), pack-reused 4242
receiving objects: 100% (4353/4353), 356.10 MiB | 29.57 MiB/s, done.
resolving deltas: 100% (1508/1508), done.
azureuser@linVM:~$ ls
content-az400-lab-resources
azureuser@linVM:~/content-az400-lab-resources$ cd content-az400-lab-resources/
azureuser@linVM:~/content-az400-lab-resources$ git remote -v
origin https://github.com/ACloudGuru-Resources/content-az400-lab-resources.git (fetch)
origin https://github.com/ACloudGuru-Resources/content-az400-lab-resources.git (push)
azureuser@linVM:~/content-az400-lab-resources$ git remote remove origin
azureuser@linVM:~/content-az400-lab-resources$ git remote -v
azureuser@linVM:~/content-az400-lab-resources$ git remote add origin https://clouduserp4b842ee7@dev.azure.com/clouduserp4b842ee7/ASP_.NET/_git/ASP_.NET
azureuser@linVM:~/content-az400-lab-resources$ git push -u origin --all
assword for 'https://clouduserp4b842ee7@dev.azure.com':
numerating objects: 184, done.
ounting objects: 100% (184/184), done.
elta compression using up to 2 threads
ompressing objects: 100% (146/146), done.
riting objects: 100% (184/184), 1.20 MiB | 7.96 MiB/s, done.
otal 184 (delta 39), reused 150 (delta 25)
emote: Analyzing objects... (184/184) (212 ms)
emote: Validating commits... (8/8) done (1 ms)
emote: Storing packfile... done (97 ms)
emote: Storing index... done (44 ms)
o https://dev.azure.com/clouduserp4b842ee7/ASP_.NET/_git/ASP_.NET
* [new branch]      unit-test -> unit-test
ranch 'unit-test' set up to track remote branch 'unit-test' from 'origin'.
azureuser@linVM:~/content-az400-lab-resources$
```

Return to the Linux VM and paste in the access token as above.

If you now use `git remote -v` you will see the git hub code again.

In Azure DevOps, click the **Azure DevOps** icon in the upper-left corner.

Select your project.

From the left-hand navigation menu, click **Repos**. Verify that the code was pushed to Azure Repos.

| Name | Last change | Commits |
|------------------------------|--------------|---|
| AspNetCoreWebApplication | Apr 11, 2022 | ad3a4cd8 Updated package references Peter Ksiazek |
| AspNetCoreWebApplicationTest | Apr 11, 2022 | ad3a4cd8 Updated package references Peter Ksiazek |
| .gitignore | Oct 9, 2020 | 39c1b922 added .gitignore chadmcrowell |
| AspNetCoreWebApplication.sln | Oct 9, 2020 | 92a3ef16 unit test added Chad Crowell |
| azure-pipelines.yml | Oct 9, 2020 | 92a3ef16 unit test added Chad Crowell |
| README.md | May 31 | a540a49d Update README.md alexanderpotasnick |

Create a New Azure DevOps Pipeline

At the top of the page, click the **Azure DevOps** icon, then select the project.

On the bottom left, click **Project settings**.

Scroll down the left-hand pane and under **Pipelines**, select **Agent pools**.

Select the **Default** pool.

The screenshot shows the 'Project Settings' page for 'MyFirstProject'. The left sidebar has sections for General, Teams, Permissions, Notifications, Service hooks, Dashboards, Boards, Project configuration, Team configuration, GitHub connections, Pipelines, and Agent pools. The 'Agent pools' section is currently selected. It displays a list of agent pools:

| Name |
|-----------------|
| Azure Pipelines |
| Default |

Get the agent

X

Windows

macOS

Linux

x64

ARM

ARM64

RHEL6

System prerequisites

Configure your account

Configure your account by following the steps outlined [here](#).

Download the agent

[Download](#)

Create the agent

```
~/$ mkdir myagent && cd myagent
~/myagent$ tar zxvf ~/Downloads/vsts-agent-linux-x64-3.225.0.tar.gz
```

Configure the agent [Detailed instructions ↗](#)

```
~/myagent$ ./config.sh
```

Optionally run the agent interactively

If you didn't run as a service above:

```
~/myagent$ ./run.sh
```

Click New agent.Select the **Linux** tab.

Click the Copy icon, next to Download to copy the agent URL to your clipboard.

Paste it to a text file to use in the following steps.

<https://vstsagentpackage.azureedge.net/agent/3.225.0/vsts-agent-linux-x64-3.225.0.tar.gz>Keep the **Linux** tab and popup open.Go back to the Linux VM, and change directory to the root: **cd ~**Make and navigate to a new directory called Downloads:**mkdir Downloads && cd Downloads**Run **wget** against the agent URL you just copied:**wget <AGENT_URL>****wget <https://vstsagentpackage.azureedge.net/agent/3.225.0/vsts-agent-linux-x64-3.225.0.tar.gz>**

This will download the agent.

Navigate back to the root again by using **cd ~**.

Create and navigate to another directory called myagent:

mkdir myagent && cd myagent

Now copy the file downloaded from **downloads** to the **myagent** folder

example:

cp /home/cloud_user/Downloads/vsts-agent-linux-x64-3.225.0.tar.gz /home/cloud_user/myagent

NOTE: To work out the paths to use above use the command **pwd**.

Unzip the package you just downloaded. You can copy this same command from the popup in Azure DevOps to get the full package and version number:

tar zxvf <PACKAGE>

example:

tar zxvf vsts-agent-linux-x64-3.225.0.tar.gz

```
cloud_user@linVM:~/myagent$ ls
bin config.sh env.sh externals license.html run-docker.sh run.sh vsts-agent-linux-x64-3.225.0.tar.gz
cloud_user@linVM:~/myagent$
```

Review the contents of the directory (**ls**). Observe all of the files and folders for the agent.

Configure the Agent

Run the configuration script with the **config.sh** file:

```
az user@linVM: ~/myagent
az user@linVM:~/myagent$ az user@linVM:~/myagent$ ./config.sh

agent v3.225.0
(commit ef0cd4d)

>> End User License Agreements:

Building sources from a TFVC repository requires accepting the Team Explorer Everywhere End User License Agreement. This step is not required for building sources from Git repositories.

A copy of the Team Explorer Everywhere license agreement can be found at:
/home/azuser/myagent/license.html

Enter (Y/N) Accept the Team Explorer Everywhere license agreement now? (press enter for N) > y

>> Connect:

Enter server URL > https://dev.azure.com/clouduserp4b842ee7
Enter authentication type (press enter for PAT) >
Enter personal access token > *****
Connecting to server ...

>> Register Agent:

Enter agent pool (press enter for default) >
Enter agent name (press enter for linVM) >
Scanning for tool capabilities.
Connecting to the server.
Successfully added the agent
Testing agent connection.
Enter work folder (press enter for _work) >
2023-09-09 18:35:05Z: Settings Saved.
az user@linVM:~/myagent$
```

./config.sh

Type **y** to accept the license agreement.

When prompted for your server URL, go back to Azure DevOps and grab the URL from your browser. It will look something like <https://dev.azure.com/clouduser...>, followed by some numbers (i.e., copy everything up to and not including /ASP_.NET).

Paste this in the terminal.

Example:

<https://dev.azure.com/clouduserp4b842ee7>

NOTE: even adding / after the above URL will cause a failure and you will have start again!

Press **Enter** to accept the default authentication type of **PAT**.

For the personal access token option, copy and paste the personal access token you created and saved earlier in this lab.

example:

das5yco2kl474qngkvp3fikjev6vtgxgymaverwa7kasm5nq3xfkq

NOTE: this sandbox environment was deleted after creating this lab. Hence the above PAT is useless.

Press **Enter** to accept the default agent pool.

Press **Enter** to accept the default name.

Press **Enter** to accept the default work folder.

Run the agent interactively so that it is always listening for any jobs by using the following bash script.

./run.sh

```
azureuser@linVM:~/myagent$ ./run.sh
Scanning for tool capabilities.
Connecting to the server.
2023-09-09 18:38:03Z: Listening for Jobs
```

[Create the New Pipeline](#)

Go back to Azure DevOps, and in the left-hand navigation menu, click **Pipelines**.

Click **Create Pipeline**.

The screenshot shows the Azure DevOps Pipelines interface. On the left, a sidebar lists various project management and pipeline-related items: Overview, Boards, Repos, Pipelines (selected), Pipelines, Environments, Releases, Library, Task groups, Deployment groups, Test Plans, and Artifacts. At the top right, there are tabs for Connect, Select, Configure, and Review. Below these tabs, a section titled "Where is your code?" displays a list of source code repository options:

- Azure Repos Git (YAML) - Free private Git repositories, pull requests, and code search
- Bitbucket Cloud (YAML) - Hosted by Atlassian
- GitHub (YAML) - Home to the world's largest community of developers
- GitHub Enterprise Server (YAML) - The self-hosted version of GitHub Enterprise
- Other Git - Any generic Git repository
- Subversion - Centralized version control by Apache

At the bottom of the list, a note says "Use the classic editor to create a pipeline without YAML."

Click Use the **classic editor**.

Under Select a source, select **Azure Repos Git** and make the following selections:

Team project: **ASP_.NET**

Repository: **ASP_.NET**

Default branch: **unit-test**

Click **Continue**.

Select a source

Azure Repos Git GitHub GitHub Enterprise Server

Team project

ASP.NET

Repository

ASP.NET

Default branch for manual and scheduled builds

unit-test

Continue

Scroll down and select the **ASP.NET Core** template.

Click **Apply**.

On the right panel, for **Agent pool**, select **Default**.

For Project(s) to test, enter: ****/*[Tt]est/*.csproj**.

NOTE: Please avoid typos as this will cause the job to fail!

Next to Agent job 1, click the + icon to add a new task.

Search for and click Add for the **Use .NET Core** task.

NOTE: Please ensure you use this legacy version.

In the task list for Agent job 1, drag and move the **Use .NET Core task** to the top.

With the **Use .NET Core** task still selected, change the Version on the right to **3.1.x**.

At the top, click **Save & queue > Save > Save**.

The screenshot shows the Azure DevOps Pipeline Editor. The left pane displays the pipeline structure with a 'Pipeline' node at the top, followed by a 'Get sources' step (ASP.NET) and an 'Agent job 1' step. The 'Agent job 1' step contains several tasks: 'Use .NET Core sdk 3.1.x', 'Restore', 'Build', 'Test', and 'Publish'. Below these is a 'Publish Artifact' step. The right pane shows the detailed configuration for the 'Use .NET Core' task, including the task version (2.*), display name ('Use .NET Core sdk 3.1.x'), package to install ('SDK (contains runtime)'), version ('3.1.x'), and advanced options like 'Compatible Visual Studio version' and 'Path To Install .Net Core'.

Select the **Triggers** tab and check the box for **Enable continuous integration** on the right panel.
At the top, click **Save & queue** > **Save** > **Save**.

The screenshot shows the Azure DevOps Pipeline Editor with the 'Triggers' tab selected. The left pane shows a 'Continuous integration' section with an 'ASP.NET' trigger listed. The trigger configuration on the right includes the 'Enable continuous integration' checkbox checked, the 'Batch changes while a build is in progress' checkbox unchecked, and a 'Branch filters' section with a 'Type' dropdown set to 'Include' and a 'Branch specification' input field containing 'unit-test'. There are also 'Path filters' and 'Add' buttons.

Update the Branch

On the left menu, select **Repos**.

Select the **AspNetCoreWebApplication** folder on the right.

Navigate through the folders to **Views > Home**.

Open the **Index.cshtml** file.

To the right of the code pane, click **Edit**.

Make a minor change in the **<h1> Congratulations! </h1>**, add **You Did It!**.

<h1> Congratulations! You Did It!</h1>

At the top, click **Commit > Commit**.

Run the Pipeline and Verify Success

From the left navigation menu, select **Pipelines**.

Select the pipeline shown to open it. Then, select the currently running job.

The screenshot shows the Azure DevOps Pipelines interface. On the left, there's a sidebar with options like Overview, Boards, Repos, and Pipelines. The Pipelines option is selected and highlighted in grey. The main area shows a pipeline named "ASP.NET-ASP.NET Core-CI". A single run is listed under the "Runs" tab, with the status "Just now" and a duration of "20s". The run description is "#20230909.1 • Updated Index.cshtml". Below this, there's a section for "Recently run pipelines" which lists "MyFirstProject-ASP.NET Core-CI" with its last run details.

Pipelines

Recent All Runs

This screenshot shows the "Pipelines" page with the "Runs" tab selected. It displays a list of recently run pipelines. One pipeline, "MyFirstProject-ASP.NET Core-CI", is listed with its most recent run information: the run number "#20230909.1", the update description "• Updated Index.cshtml", and the stage "Individual CI for unit-test".

Select **Agent job 1** to view the progress.

The screenshot shows the Azure DevOps interface for a pipeline named 'ASP.NET-ASP.NET Core-Cl'. The left sidebar has 'Pipelines' selected. The main area shows a table of jobs. One job, 'Agent job 1', is expanded, showing its steps: 'Initialize job', 'Checkout ASP.NET@un...', 'Use .NET Core sdk 3.1.x', 'Restore', 'Build', 'Test', 'Publish', 'Publish Artifact', and 'Post-job: Checkout ASP...'. The 'Test' tab is selected, displaying a log of the test execution:

```

1 Starting: Test
2 =====
3 Task : .NET Core
4 Description : Build, test, package, or publish a dotnet application, or run a custom dotnet command
5 Version : 2.221.0
6 Author : Microsoft Corporation
7 Help : https://docs.microsoft.com/azure/devops/pipelines/tasks/build/dotnet-core-cli
8 =====
9 Info: .NET Core SDK/runtime 2.2 and 3.0 are now End of Life(EOL) and have been removed from all hosted agents. If you're using these
10 /home/azureuser/.myagent/_work/_tool/dotnet/dotnet test /home/azureuser/.myagent/_work/1/s/AspNetCoreWebApplicationTest/AspNetCoreWebAp
11 Test run for /home/azureuser/.myagent/_work/1/s/AspNetCoreWebApplicationTest/bin/Release/netcoreapp3.1/AspNetCoreWebApplicationTest.d
12 Microsoft (R) Test Execution Command Line Tool Version 16.7.1
13 Copyright (c) Microsoft Corporation. All rights reserved.
14
15 Starting test execution, please wait...
16
17 A total of 1 test files matched the specified pattern.
18 Results File: /home/azureuser/.myagent/_work/_temp/_linVM_2023-09-09_18_51_41.trx
19
20 Test Run Successful.
21 Total tests: 2
22 Passed: 2
23 Total time: 2.3819 Seconds

```

It will take a few minutes for the job to complete.
Once complete, click the back arrow above the job progress.

The screenshot shows the pipeline run details for job #20230909.1. The 'Tests' tab is selected. The summary shows the run was triggered by 'Cloud Student 4b842ee7'. It includes details about the repository, time started, related work items, and test coverage.

| Repository and version | Time started and elapsed | Related | Tests and coverage |
|------------------------|----------------------------|---|------------------------------------|
| ASP.NET unit-test | Today at 9:50 PM 1m 26s | 0 work items 1 published: 1 consumed | 100% passed Setup code coverage |

Select the **Tests** tab to verify success.

#20230909.1 • Updated Index.cshtml
ASP.NET-ASP.NET Core-CI

This run is being retained as one of 3 recent runs by unit-test (Branch). [View retention leases](#)

[Run new](#) [⋮](#)

[Summary](#) [Tests](#)

Summary

1 Run(s) Completed (1 Passed, 0 Failed)

| Total tests | Passed | Failed | Others |
|-------------|--------|--------|--------|
| 2 | 2 | 0 | 0 |

Pass percentage: 100% | Run duration: 2s 403ms | Tests not reported: 0

↑ 100% | ↑ +2s 403ms

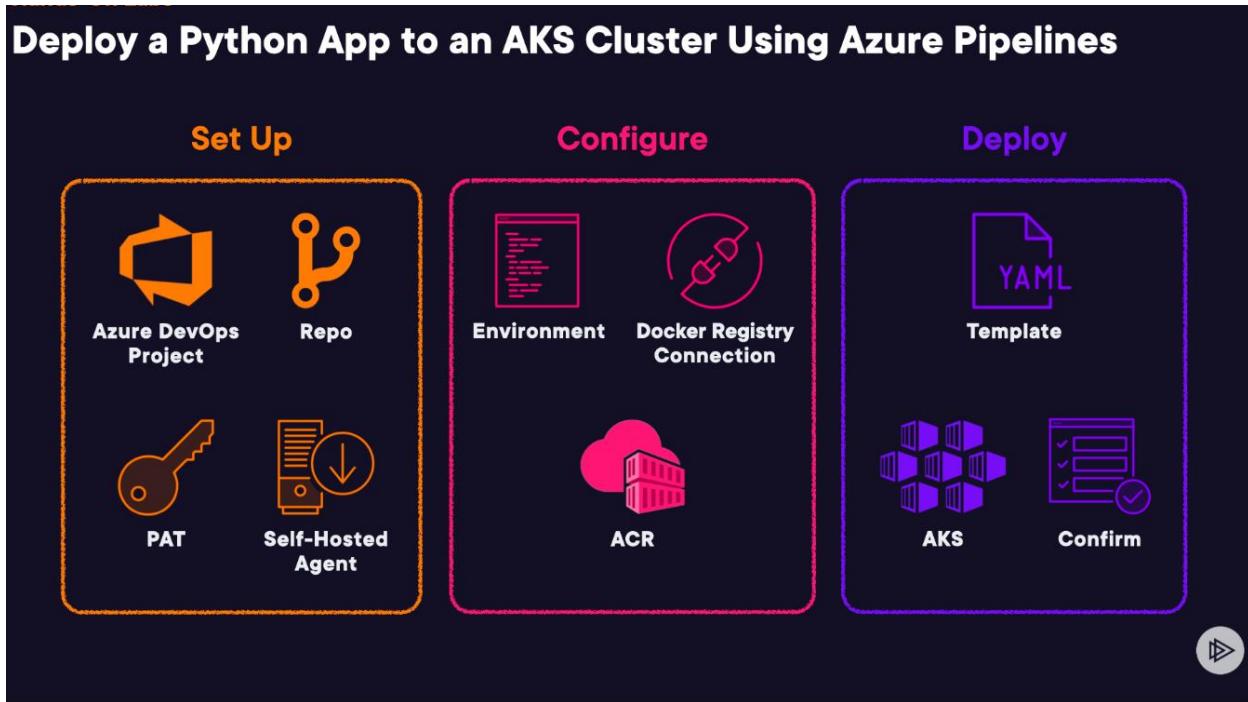
[Bug](#) [Link](#) [Filter by test or run name](#)

[Test run](#) [Column Options](#)

Tags: [Aborted \(+1\)](#)

 Hooray! There are no test failures.

Deploy a Python App to an AKS Cluster Using Azure Pipelines



Pre steps

1. Create a Linux virtual machine (Ubuntu)
2. Create a Container Registry
3. Create a Kubernetes service.
4. Setting up and configuring Docker on the Ubuntu server.

The screenshot shows the Azure portal interface for a resource group named "1-c7cc6262-playground-sandbox". The "Essentials" blade is open, displaying a list of resources:

| Type | Location |
|--------------------------------|----------|
| Container registry | East US |
| Metric alert rule | Global |
| Container registry replication | East US |
| Kubernetes service | East US |
| Metric alert rule | Global |
| Action group | Global |
| Virtual machine | East US |
| Public IP address | East US |

Login to the Azure portal sandbox and create the following resources we see above:

1. Create a Linux virtual machine (Ubuntu)

The screenshot shows the Microsoft Azure Marketplace interface. On the left, there's a sidebar with links like 'Get Started', 'Service Providers', 'Management' (which is currently selected), and 'My Marketplace'. The main area has a search bar at the top. Below it, there's a 'Create' dropdown menu. Underneath, there are several cards for different offerings. One card for 'Ubuntu Server 20.04 LTS' by Canonical is highlighted, showing its icon and name. Other cards include 'Virtual Machine' and 'Linux For The Cloud'.

2. Create a Container Registry

This screenshot shows the Microsoft Azure Marketplace interface again. The 'Management' link in the sidebar is still selected. A search bar at the top contains the text 'container registry'. Below the search bar, there's a checkbox labeled 'Azure services only'. The main content area displays search results for 'container registry', showing one result: 'Container Registry' by Microsoft, which is described as an 'Azure Service' for managing a Docker private registry.

3. Create a Kubernetes service.

Azure Kubernetes Service (AKS)
Microsoft
Azure Service
A managed cluster with a Kubernetes orchestrator for container deployments.

Now, SSH into the virtual machine with the keypair

Native SSH
Connecting using public IP address
52.146.88.8
Admin username: azureuser
Port (change): 22 [Check access](#) [Configure for this port](#)
Just-in-time policy: Not configured for port 22 [Configure for this port](#)

Recommended
Local machine | Azure portal
SSH using Azure CLI
Quickly connect in browser. Supports Azure AD authentication. Private key not required.
Using public IP address (52.146.88.8)
[Select](#)

Most common
Native SSH
No additional software needed. Private key connection. Best for those with existing SSH to VM with specified private key.
Using public IP address (52.146.88.8)
[Select](#)

Other Information
Using a Linux subsystem like WSL or Ubuntu on Terminal?
Copy your private key path to the Linux subsystem and ensure it has the correct read-only access.

ssh -i <key_pair>@<Public IP Address>

```
cat /etc/os-release
```

```
azureuser@VM12:~$ cat /etc/os-release
NAME="Ubuntu"
VERSION="20.04.6 LTS (Focal Fossa)"
ID=ubuntu
ID_LIKE=debian
PRETTY_NAME="Ubuntu 20.04.6 LTS"
VERSION_ID="20.04"
HOME_URL="https://www.ubuntu.com/"
SUPPORT_URL="https://help.ubuntu.com/"
BUG_REPORT_URL="https://bugs.launchpad.net/ubuntu/"
PRIVACY_POLICY_URL="https://www.ubuntu.com/legal/terms-and-policies/privacy-policy"
VERSION_CODENAME=focal
UBUNTU_CODENAME=focal
azureuser@VM12:~$
```

4. Setting up and configuring Docker on the Ubuntu server.

```
sudo apt update
```

```
sudo apt install -y \
    apt-transport-https \
    ca-certificates \
    curl \
    software-properties-common
```

```
curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo gpg --dearmor -o
/usr/share/keyrings/docker-archive-keyring.gpg
```

```
echo "deb [arch=amd64 signed-by=/usr/share/keyrings/docker-archive-keyring.gpg]
https://download.docker.com/linux/ubuntu $(lsb_release -cs) stable" | sudo tee
/etc/apt/sources.list.d/docker.list > /dev/null
```

```
sudo apt update
```

```
sudo apt install docker-ce
```

```
docker --version
```

```
sudo systemctl enable docker
sudo systemctl start docker
```

```
whoami
azureuser
```

```
sudo usermod -aG docker azureuser
sudo systemctl restart docker
```

docker run hello-world

permission denied while trying to connect to the Docker daemon socket at unix:///var/run/docker.sock:
Post "<http://var/run/docker.sock/v1.24/images/create?fromImage=tiangolo%2Fuwsgi-nginx-flask&tag=python3.6>": dial unix /var/run/docker.sock: connect: permission denied

Run the following:

sudo chmod 666 /var/run/docker.sock

Now run the following hello-world command again.

docker run hello-world

```
|Hello from Docker!
This message shows that your installation appears to be working correctly.

To generate this message, Docker took the following steps:
1. The Docker client contacted the Docker daemon.
2. The Docker daemon pulled the "hello-world" image from the Docker Hub.
   (amd64)
3. The Docker daemon created a new container from that image which runs the
   executable that produces the output you are currently reading.
4. The Docker daemon streamed that output to the Docker client, which sent it
   to your terminal.

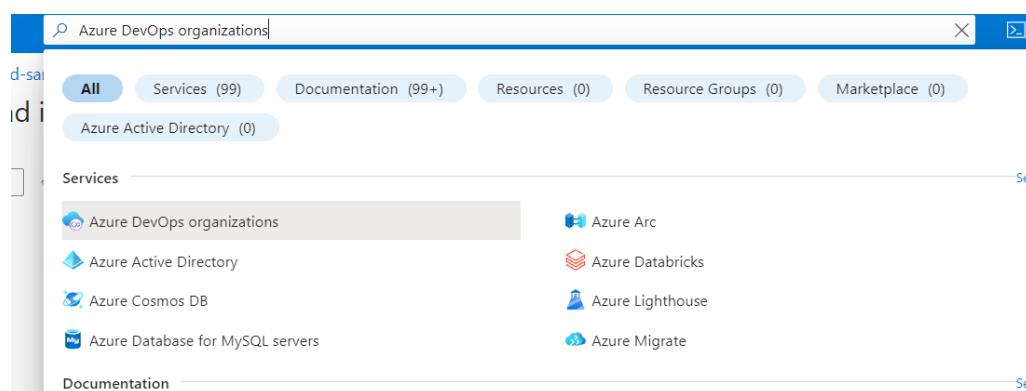
To try something more ambitious, you can run an Ubuntu container with:
$ docker run -it ubuntu bash

Share images, automate workflows, and more with a free Docker ID:
https://hub.docker.com/

For more examples and ideas, visit:
https://docs.docker.com/get-started/
azureuser@VM12:~$
```

Now we have completed the pre steps lets now proceed to setting up the Azure pipeline.

Create an Azure DevOps Organization



The screenshot shows the Azure DevOps portal's search interface. In the top search bar, the text "Azure DevOps organizations" is entered. Below the search bar, there are several navigation tabs: "All" (selected), "Services (99)", "Documentation (99+)", "Resources (0)", "Resource Groups (0)", and "Marketplace (0)". Under the "Services" tab, a list of services is displayed, with "Azure DevOps organizations" highlighted. Other listed services include Azure Arc, Azure Databricks, Azure Lighthouse, and Azure Migrate. At the bottom of the screen, there is a "Documentation" section.

In the main search bar, search for and select **Azure DevOps organizations**.

Scroll down and click **My Azure DevOps Organizations**.

The screenshot shows the Microsoft Azure portal interface. At the top, there's a header bar with a back/forward button, a search icon, and the URL <https://portal.azure.com/#view/AzureTfsExtension/OrganizationsTemp>. Below the header is a blue navigation bar with the Microsoft Azure logo and a search bar. A breadcrumb trail shows 'Home > Azure DevOps'. The main content area has a large orange banner at the top stating, 'We've made it easier to manage Azure DevOps billing and subscriptions. You can [set up billing](#) or [switch organizations](#)'. Below the banner, the text 'Azure DevOps' is displayed in large blue letters, followed by the subtext 'Plan smarter, collaborate better, and ship faster with a set of modern dev services'. At the bottom of the page, there are links: 'My Azure DevOps Organizations', 'Get started using Azure DevOps', and 'Billing management for Azure DevOps'.

Select your **country/region** from the dropdown list.

Select Create **new organization** > **Continue** > **Continue** (complete the required **CAPTCHA**).



Get started with Azure DevOps

Plan better, code together, ship faster with Azure DevOps

[Create new organization](#)

Set the following values to create a project:

Project name: **Dev**

Visibility: **Private**

Click + **Create project**.

Import Code and Set Up the Environment:

- In the sidebar menu, select **Repos**.

In the Import a repository section, click **Import**.

In the Clone URL field, enter the following repository:

<https://github.com/ACloudGuru-Resources/content-az400-lab-resources.git>

Click **Import**.

After the code import completes successfully, the Files page should open automatically.

The screenshot shows the Azure DevOps interface for a repository named 'aks'. The left sidebar has a 'Files' section selected. The main area displays a list of files with their last change date and committer information. A 'Create a pull request' button is visible at the top right.

| Name | Last change | Commits |
|---------------------|--------------|--|
| azure-vote | Oct 25, 2020 | a6de1eb8 added aks manifest Chad Crowell |
| manifests | 2h ago | e15850e5 Updated vote.yml Cloud Student 45b40fb2 |
| .gitignore | Oct 25, 2020 | a6de1eb8 added aks manifest Chad Crowell |
| azure-pipelines.yml | 2h ago | 0ef6d7fd Set up CI with Azure Pipelines Cloud Student 45b40fb2 |
| Dockerfile | Oct 25, 2020 | a6de1eb8 added aks manifest Chad Crowell |
| README.md | May 31 | 0b79db16 Update README.md alexanderpotasnick |

At the top of the page, use the dropdown to switch the branch from **DSC** to **aks**.

You should now see the files you need for the lab as above.

Create Personal Access Token

In the upper right-hand corner, click the **User settings icon > Personal access tokens**.

Click **+ New Token** and set the following values:

Name: pat1

Scopes: Full access

Click **Create**.

The dialog box shows a success message: "Success! You have successfully added a new personal access token. Copy the token now! Pat2 token". It includes a copied token field (fjtp7hxp6wu3hcyybqsgtzc) and a warning message: "Warning - Make sure you copy the above token now. We don't store it and you will not be able to see it again." A "Close" button is at the bottom.

Copy the **pat1** token and save it in a safe location.

Click **Close**.

Download Self-Hosted Agent

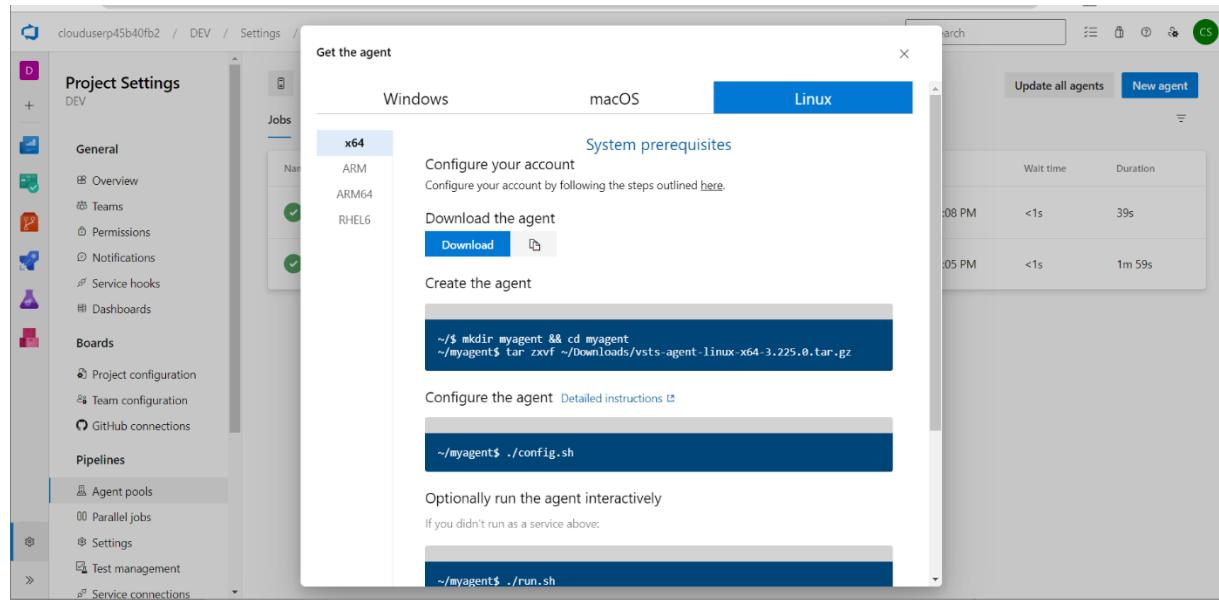
At the top of the page, click the **Azure DevOps icon**, then select the **project**.

On the bottom left, click **Project settings**.

Scroll down the left-hand pane and under **Pipelines**, select **Agent pools**.

Select the **Default pool**.

Click **New agent**.



Select the **Linux** tab.

Click the **Copy** icon, next to **Download** to copy the agent URL to your clipboard.
Paste it to a text file to use in the following steps. Keep this tab and popup open.

For example:

<https://vstsagentpackage.azureedge.net/agent/3.225.0/vsts-agent-linux-x64-3.225.0.tar.gz>

Setup the agent on the virtual machine

Ensure the directory is set to the root:

`cd ~`

Create and navigate to another directory called myagent:

`mkdir myagent && cd myagent`

Download the tar file within the agent directory.

wget <AGENT_URL>

wget <https://vstsagentpackage.azureedge.net/agent/3.225.0/vsts-agent-linux-x64-3.225.0.tar.gz>

```
azureuser@vm1:~/Downloads
Resolving vstsagentpackage.azureedge.net (vstsagentpackage.azureedge.net)... 72.21.81.200, 2606:2800:11f:17a5:191a:18d5:537:22f9
Connecting to vstsagentpackage.azureedge.net (vstsagentpackage.azureedge.net)|72.21.81.200|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 158318686 (151M) [application/octet-stream]
Saving to: 'vsts-agent-linux-x64-3.225.0.tar.gz'

vsts-agent-linux-x64-3.225.0. 100%[=====] 150.98M 192MB/s in 0.8s

2023-09-10 19:59:57 (192 MB/s) - 'vsts-agent-linux-x64-3.225.0.tar.gz' saved [158318686/158318686]

azureuser@vm1:~/myagent$ cd ~
azureuser@vm1:~$ mkdir Downloads && cd Downloads
azureuser@vm1:~/Downloads$ wget https://vstsagentpackage.azureedge.net/agent/3.225.0/vsts-agent-linux-x64-3.225.0.tar.gz
--2023-09-10 20:01:14- https://vstsagentpackage.azureedge.net/agent/3.225.0/vsts-agent-linux-x64-3.225.0.tar.gz
Resolving vstsagentpackage.azureedge.net (vstsagentpackage.azureedge.net)... 72.21.81.200, 2606:2800:11f:17a5:191a:18d5:537:22f9
Connecting to vstsagentpackage.azureedge.net (vstsagentpackage.azureedge.net)|72.21.81.200|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 158318686 (151M) [application/octet-stream]
Saving to: 'vsts-agent-linux-x64-3.225.0.tar.gz'

vsts-agent-linux-x64-3.225.0. 100%[=====] 150.98M 185MB/s in 0.8s

2023-09-10 20:01:15 (185 MB/s) - 'vsts-agent-linux-x64-3.225.0.tar.gz' saved [158318686/158318686]

azureuser@vm1:~/Downloads$
```

Unzip the package you just downloaded with the following command:

```
azureuser@VM12:~$ cd ./myagent/
azureuser@VM12:~/myagent$ ls
_diag bin env.sh license.html run.sh vsts-agent-linux-x64-3.225.0.tar.gz
_work config.sh externals run-docker.sh svc.sh
azureuser@VM12:~/myagent$
```

tar zxvf <PACKAGE>

You can also copy this same command from the popup in Azure DevOps to get the full package and version number:

```
azureuser@VM12:~$ cd ./myagent/
azureuser@VM12:~/myagent$ ls
_diag bin env.sh license.html run.sh vsts-agent-linux-x64-3.225.0.tar.gz
_work config.sh externals run-docker.sh svc.sh
azureuser@VM12:~/myagent$ tar zxvf vsts-agent-linux-x64-3.225.0.tar.gz
```

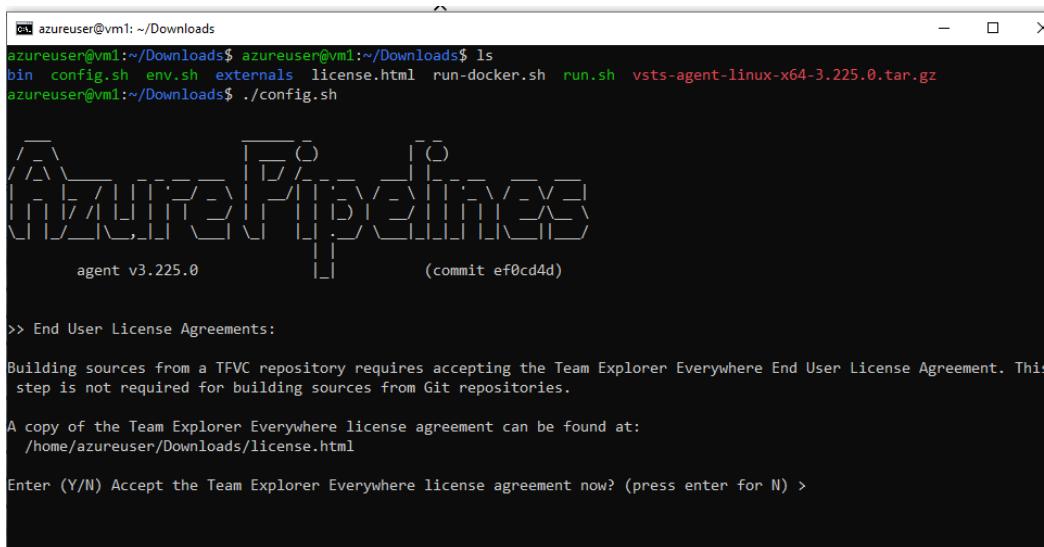
Review the contents of the directory (**ls**). Observe all the files and folders for the agent.

Configure the Agent

Run the configuration script with the following command:

./config.sh

Type **y** to accept the license agreement.



```
azureuser@vm1:~/Downloads$ ./config.sh
agent v3.225.0          (commit ef0cd4d)

>> End User License Agreements:

Building sources from a TFVC repository requires accepting the Team Explorer Everywhere End User License Agreement. This step is not required for building sources from Git repositories.

A copy of the Team Explorer Everywhere license agreement can be found at:
/home/azureuser/Downloads/license.html

Enter (Y/N) Accept the Team Explorer Everywhere license agreement now? (press enter for N) >
```

When prompted for your server URL, go back to Azure DevOps and grab the URL from your browser. It will look something like <https://dev.azure.com/clouduser...>, followed by some characters (i.e., copy everything up to but not including /dev).

For example:

<https://dev.azure.com/clouduserp45b40fb2>

Paste this in the terminal

Press **Enter** to accept the default authentication type of **PAT**.

For personal access token, copy and paste the personal access token you created and saved earlier.

```
az user@vm1: ~/Downloads
az user@vm1:~/Downloads$ ./config.sh
bin config.sh env.sh externals license.html run-docker.sh run.sh vsts-agent-linux-x64-3.225.0.tar.gz
az user@vm1:~/Downloads$ ./config.sh

agent v3.225.0          (commit ef0cd4d)

>> End User License Agreements:
Building sources from a TFVC repository requires accepting the Team Explorer Everywhere End User License Agreement. This step is not required for building sources from Git repositories.

A copy of the Team Explorer Everywhere license agreement can be found at:
/home/azuser/Downloads/license.html

Enter (Y/N) Accept the Team Explorer Everywhere license agreement now? (press enter for N) > y

>> Connect:
Enter server URL > https://dev.azure.com/clouduserpd2bb1751
Enter authentication type (press enter for PAT) >
Enter personal access token > *****
Connecting to server ...

>> Register Agent:
Enter agent pool (press enter for default) >
Enter agent name (press enter for vm1) >
Scanning for tool capabilities.
Connecting to the server.
Successfully added the agent
Testing agent connection.
Enter work folder (press enter for _work) >
2023-09-10 20:06:19Z: Settings Saved.
az user@vm1:~/Downloads$
```

Press Enter to accept the default **agent pool**.

Press Enter to accept the default **name**.

Press Enter to accept the default **work folder**.

Now run the agent interactively so that it is always listening for any jobs:

./run.sh

```
az user@vm1:~/Downloads$ ./run.sh
Scanning for tool capabilities.
Connecting to the server.
2023-09-10 20:07:13Z: Listening for Jobs
```

Now let's go back to the DevOps portal

Create Pipeline Environment

In the sidebar menu, select **Pipelines**, then select **Environments**.

Click **Create environment**.

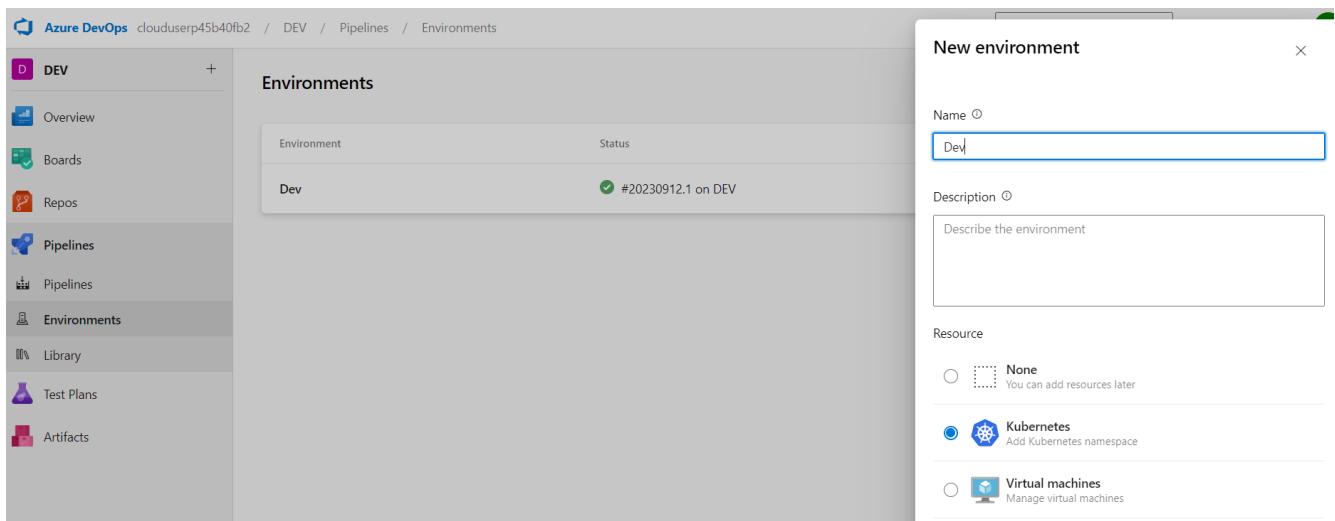
Fill in the environment details:

Name: Enter **dev**.

Resource: Select **Kubernetes**.

Click **Next**.

Ensure the resource details auto-populate, then click **Validate and create**.



Create a Service Connection and an ACR

In the bottom-left corner, click **Project Settings**.

This opens the settings in a new tab.

In the sidebar menu, select **Service connections**.

In the top right, click **New service connection**.

In the pane on the right, select **Docker Registry**, then click **Next**.

You'll need to gather some information from an **Azure Container Registry (ACR)** before you finish creating your service connection.

The screenshot shows the Azure DevOps interface for managing service connections. On the left, there's a sidebar with 'Project Settings' for the 'DEV' project, including sections for General, Boards, Pipelines, and more. The main area is titled 'Service connections' and lists existing connections like 'ACR' and 'Dev-K8-default-1694508846439'. A modal window titled 'New Docker Registry service connection' is open on the right, prompting for details: 'Docker Registry' (set to 'https://index.docker.io/v1/'), 'Docker ID' ('acrqnqworsxpnyk'), 'Docker Password' (redacted), 'Email (optional)' (empty), 'Service connection name' ('ACR'), and 'Description (optional)' (empty). Under 'Security', there's a checked checkbox for 'Grant access permission to all pipelines'. At the bottom of the modal are 'Back' and 'Save' buttons.

Navigate back to your Azure DevOps tab.

Navigate to **All resources** using the hamburger menu in the top-left corner.

Select the link for the **Azure container registry (ACR)**.

In the ACR's sidebar menu, select **Access keys** (under **Settings**).
You'll use these details for your service connection.

The screenshot shows the Microsoft Azure portal with the URL 'https://azuresdkforjs.dev.csharp.azure.com/'. The page displays the 'Access keys' section for a container registry named 'con1234567'. It shows two sets of access keys: 'password' and 'password2'. The 'password' key has a value of '9iT+eTuXkeH1jc+247LSWF9yy2dARzI5V5xWkVC5oq+ACR...'. There are 'Regenerate' and 'Delete' buttons next to each key. The left sidebar includes links for Overview, Activity log, Access control (IAM), Tags, Quick start, Events, Settings (with 'Access keys' selected), Encryption, Identity, Networking, and Microsoft Defender for Cloud.

Add the New Docker Registry service connection details:

Fill in the Docker Registry:

From the Access keys page, copy the **Login server**.

Navigate to the service connections tab and paste the login server into the **Docker Registry** field using the format https://<LOGIN_SERVER>.

Fill in the Docker ID:

Navigate to the ACR tab and copy the **Username**.

Navigate to the service connections tab and paste the username into the **Docker ID** field.

Fill in the Docker Password:

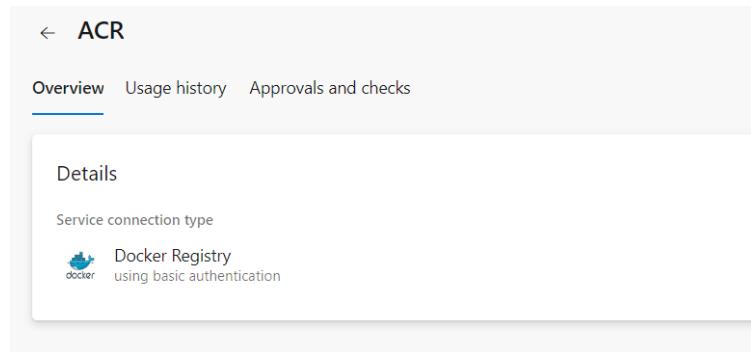
Navigate to the ACR tab and copy either **password**.

Navigate to the service connections tab and paste the password into the **Docker Password** field.

In the Service connection name field, enter **ACR**.

Below Security, check the **Grant access permission to all pipelines** checkbox.

Click **Save**.



Create the CI/CD Pipeline

Update the Deployment Image

From the service connections tab, use the sidebar menu to select **Repos**.

Select the **manifests** folder and open the **vote.yml** file.

In the top right, click **Edit**.

Update the deployment image:

Navigate to the **ACR** tab and copy the **Login server** again.

Navigate to the **vote.yaml** tab and replace the container image on line 59 with the Azure ACR DNS registry name, following the format `<LOGIN_SERVER>/azure-vote-front:v1`.

```

40 ---
41 apiVersion: apps/v1
42 kind: Deployment
43 metadata:
44   name: azure-vote-front
45 spec:
46   replicas: 1
47   selector:
48     matchLabels:
49       app: azure-vote-front
50   template:
51     metadata:
52       labels:
53         app: azure-vote-front
54     spec:
55       nodeSelector:
56         "beta.kubernetes.io/os": linux
57       containers:
58         - name: azure-vote-front
59           image: con1234567.azurecr.io/azure-vote-front:v1
60           resources:
61             requests:
62               cpu: 100m
63               memory: 128Mi
64             limits:
65               cpu: 250m
66               memory: 256Mi

```

Note: The ACR DNS name ends with `azurecr.io`.

In the top right, click **Commit**.

In the Commit pane, ensure the Branch name is set to `aks`, then click **Commit**.

Create the Pipeline

In the sidebar menu, select **Pipelines**.

Click **Create Pipeline**.

For the code location, select **Azure Repos Git**.

For the repository, select the dev repo.

Select **Existing Azure Pipelines YAML file**.

Fill in the YAML file details:

Branch: Select **aks**.

Path: Select **/azure-pipelines.yml**.

Click **Continue**.

Update the azure-pipelines.yml code:

Navigate to the **ACR** tab and copy the **Login server** again.

Navigate to the **pipeline** tab and replace the value for container Registry with your copied Azure ACR DNS registry name.

```

trigger:
- aks

variables:
  dockerRegistryServiceConnection: 'ACR'
  imageRepository: 'azure-unite-front'
  containerRegistry: 'ccr234567.azurecr.io'
  dockerfilePath: '**/Dockerfile'
  tag: '${{build.buildId}}'
  imagePullSecret: 'csacr5d93-auth'

stages:
- stage: Build
  displayName: Build stage
  jobs:
    - job: Build
      displayName: Build
  
```

In the top-right corner, click **Save and run**.

In the pane on the right, leave the default settings and click **Save and run**.

After the pipeline is triggered, select **Build stage** to monitor the build progress.

Note the banner in the terminal indicating that the pipeline needs additional permissions to continue the build.

Click **View** to the right of the banner.

Click **Permit** to the right of the requested permissions and confirm the selection.

The screenshot shows the Azure DevOps interface for a project named 'DEV'. The left sidebar has 'Pipelines' selected. A 'Build stage' table lists several tasks under the 'Build' job, each with a duration. To the right, a detailed log for the 'Post-job: Checkout DEV@aks to s' task is displayed, showing the command and its execution details.

```

1 Starting: Checkout DEV@aks to s
2 =====
3 Task      : Get sources
4 Description: Get sources from a repository. Supports Git, Tfvc, and SVN repositories.
5 Version   : 1.0.0
6 Author    : Microsoft
7 Help      : [More Information](https://go.microsoft.com/fwlink/?LinkId=798199)
8 =====
9 Cleaning any cached credential from repository: DEV (Git)
10 Finishing: Checkout DEV@aks to s

```

After the build has completed, select **Deploy to dev**.

The screenshot shows the Azure DevOps interface for the same project 'DEV'. The left sidebar has 'Pipelines' selected. A 'Build stage' table lists tasks under the 'Build' job. Below it, a 'Deploy to dev' section shows a table of tasks for the 'Deploy to dev' job. To the right, a detailed log for the 'Deploy to dev' task is displayed, showing the agent, start time, duration, and live console data.

```

1 Agent: VM12
2 Started: Just now
3 Duration: 35s
4 
5 Job live console data:
6 Starting: Deploy to dev
7 Async Command Start: DetectDockerContainer
8 Async Command End: DetectDockerContainer
9 Async Command Start: DetectDockerContainer
10 Async Command End: DetectDockerContainer
11 Finishing: Deploy to dev

```

Note the banner in the terminal indicating that the pipeline needs additional permissions to continue the deployment. If the banner does not display. Click the back arrow above the deployment to go back to the job summary. Then, click View in the yellow notification box.

Like before, allow all permissions. The deployment will take a few minutes to complete.

Access the AKS Cluster

Navigate to the ACR tab.

At the top, click **All resources**.

Select the link for the **Kubernetes service** cluster. (You can ignore any errors that appear as they are related to restrictions).

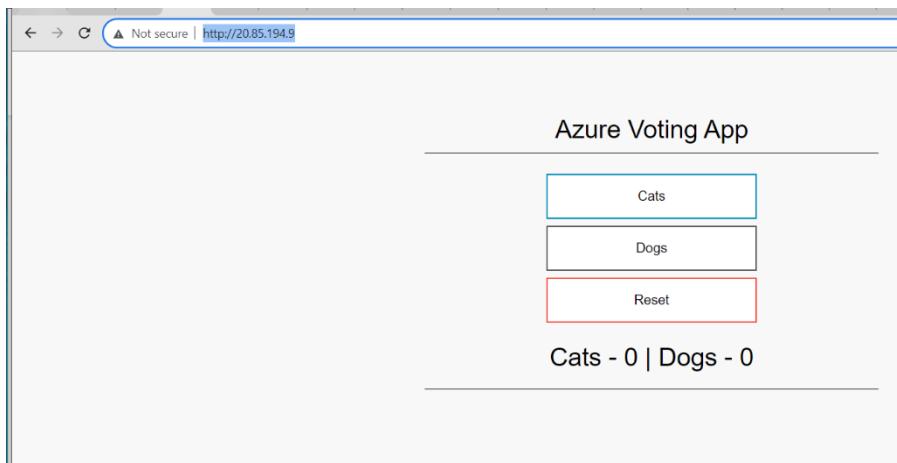
On the left menu, under **Kubernetes resources**, select **Services and ingresses**.

All services in the cluster are displayed. Observe the azure-vote-front service has an **External IP**

Click the link for the IP.

| Name | Namespace | Status | Type | Cluster IP | External IP | Ports | Age |
|------------------|-------------|--------|--------------|--------------|---------------|--------------|---------|
| kubernetes | default | Ok | ClusterIP | 10.0.0.1 | 443/TCP | 4 hours | ... |
| kube-dns | kube-system | Ok | ClusterIP | 10.0.0.10 | 53/UDP,53/TCP | 4 hours | ... |
| metrics-server | kube-system | Ok | ClusterIP | 10.0.159.21 | 443/TCP | 4 hours | ... |
| azure-vote-back | default | Ok | ClusterIP | 10.0.50.102 | 6379/TCP | 3 hours | ... |
| azure-vote-front | default | Ok | LoadBalancer | 10.0.161.109 | 20.85.194.9 | 80:32219/TCP | 3 hours |

The app opens in a new tab where you can click to vote on cats or dogs.



Addendum

Errors to avoid

1) The pipeline gives this error!

#20230910.1 • Set up CI with Azure Pipelines

Rerun failed jobs Run new :

This run will be cleaned up after 1 month based on your project settings.

Summary

Triggered by Cloud Student d2bb1751

View 40 changes

Repository and version

- Python_App
- aks 93ca53c3

Time started and elapsed

- Just now
- <1s

Related

- 0 work items
- 0 artifacts

Tests and coverage

- Get started

Errors 1

Job Deploy: Resource default does not exist in environment dev

20230910.1

View documentation for troubleshooting failed runs

Solution: Is your Docker installed and properly configured on your virtual machine? Please follow the instructions above in the pre steps!

2) The pipeline build gives this error!

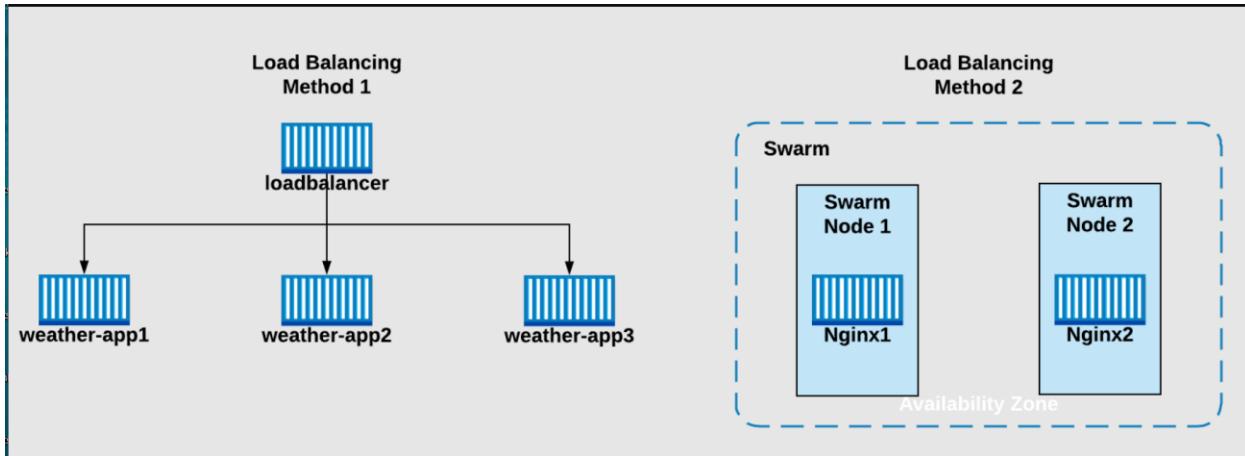
```
##[error]ERROR: permission denied while trying to connect to the Docker daemon socket at unix:///var/run/docker.sock: Get "http://%2Fvar%2Frun%2Fdocker.sock/_ping": dial unix /var/run/docker.sock: connect: permission denied
##[error]The process '/usr/bin/docker' failed with exit code 1
```

Solution: As indicated above, this is a permission issue related to your Docker on your virtual machine. Instructions for assigning permissions so that the DevOps pipeline can be

deployed to the virtual machine are above in pre steps.

Note: if the hello world test above fails with a similar message. Please do not proceed further with the Pre steps until this permission error is resolved. Otherwise, the pipeline Will fail!

Load Balancing Containers (Docker Swarm and Docker Compose)



Solution

This is as follows:

Prerequisites.

Before you begin, make sure you have:

Two CentOS 7 servers; one to be the master and the other to be the worker.

Master and Worker

```
[root@ip-10-0-1-4 ~]# cat /etc/os-release
```

```
NAME="CentOS Linux"
VERSION="7 (Core)"
ID="centos"
ID_LIKE="rhel fedora"
VERSION_ID="7"
PRETTY_NAME="CentOS Linux 7 (Core)"
ANSI_COLOR="0;31"
CPE_NAME="cpe:/o:centos:centos:7"
HOME_URL="https://www.centos.org/"
BUG_REPORT_URL="https://bugs.centos.org/"

CENTOS_MANTISBT_PROJECT="CentOS-7"
CENTOS_MANTISBT_PROJECT_VERSION="7"
REDHAT_SUPPORT_PRODUCT="centos"
REDHAT_SUPPORT_PRODUCT_VERSION="7"
```

You will need to have the following inbound port rules set.

Port **22** and **8080** need to be open.

SSH access to both servers with sudo privileges.

When setting up both server instances in the cloud, you will need to generate two .PEM key pairs to gain remote access to.

Installation and Configuration on the Master Node

Step 1: Update the System

SSH into the master server and update the system packages:

sudo yum update -y

Step 2: Install Docker

Install Docker on the master server:

sudo yum install -y docker

Step 3: Start and Enable Docker

Start and enable the Docker service:

sudo systemctl start docker

sudo systemctl enable docker

Step 4: Install Docker Compose

Install Docker Compose on the master node:

sudo curl -L "https://github.com/docker/compose/releases/latest/download/docker-compose-\$(uname -s)-\$(uname -m)" -o /usr/local/bin/docker-compose

sudo chmod +x /usr/local/bin/docker-compose

Installation and Configuration on the Worker Node

Step 1: Update the System

SSH into the master server and update the system packages:

sudo yum update -y

Step 2: Install Docker

Install Docker on the master server:

sudo yum install -y docker

Step 3: Start and Enable Docker

Start and enable the Docker service:

```
sudo systemctl start docker
sudo systemctl enable docker
```

Step 4: Install Docker Compose

Install Docker Compose on the master node:

```
sudo curl -L "https://github.com/docker/compose/releases/latest/download/docker-compose-
$(uname -s)-$(uname -m)" -o /usr/local/bin/docker-compose
```

```
sudo chmod +x /usr/local/bin/docker-compose
```

Work on the Master node

```
ssh cloud_user@PUBLIC_IP_ADDRESS
```

Become the root user:

```
sudo su -
```

Create a Docker Compose file

Look at the directories available to you.

```
ls
```

```
[root@ip-10-0-1-103 ~]# ls
anaconda-ks.cfg  lb-challenge  swarm-token.txt
[root@ip-10-0-1-103 ~]# pwd
/
```

Change to the lb-challenge directory:

```
cd lb-challenge
```

Use the VIM editor to create the **docker-compose.yml** file.

```
vi docker-compose.yml
```

Use **Ctrl + :**

type **set paste**

Use the letter I for insert

```
version: '3.2'
services:
  weather-app1:
    build: ./weather-app
    tty: true
    networks:
      - frontend
  weather-app2:
    build: ./weather-app
    tty: true
    networks:
      - frontend
  weather-app3:
    build: ./weather-app
    tty: true
    networks:
      - frontend

loadbalancer:
  build: ./load-balancer
  image: nginx
  tty: true
  ports:
    - '80:80'
  networks:
    - frontend

networks:
  frontend:
```

Now copy paste or type the above in the format as above.

When you have completed the above do the following to save and exit the file:

Ctrl + :

wq <enter>

Update nginx.conf

Change to the load-balancer directory:

cd load-balancer

Edit the nginx.conf file:

vi nginx.conf

The contents of your **nginx.conf** file should look like the following:

```
events { worker_connections 1024; }

http {
    upstream localhost {
        server weather-app1:3000;
        server weather-app2:3000;
        server weather-app3:3000;
    }
    server {
        listen 80;
        server_name localhost;
        location / {
            proxy_pass http://localhost;
            proxy_set_header Host $host;
        }
    }
}
```

Note: Please ensure that the format is as above. Otherwise the **nginx** file will not be processed!
Once done, save and exit with **wq <enter>**

Execute docker-compose up

Change directories:

```
cd ../
```

Execute the docker-compose up command:

```
docker-compose up --build -d
```

```
Compose does not use swarm mode to deploy services to multiple nodes in a swarm. All containers will be scheduled on the current node.
To deploy your application across the swarm, use `docker stack deploy`.

Creating network "lb-challenge_frontend" with the default driver
Building weather-app1
Step 1/10 : FROM node
latest: Pulling from library/node
167b8a53ca45: Download complete
b47a222d28fa: Download complete
debece5f9f3a9: Download complete
1d7ca7cd2e06: Downloading [=====] 29.14MB/211MB
94c7791033e8: Download complete
72ab0dfaf5cb: Downloading [=====] 7.796MB/47.86MB
3316ed2852d4: Download complete
ef5505406bea: Waiting
[...]
```

```
Step 10/10 : CMD ./bin/www
--> Using cache
--> 746f221ee7ca
Successfully built 746f221ee7ca
Successfully tagged lb-challenge_weather-app3:latest
Building loadbalancer
Step 1/4 : FROM nginx
latest: Pulling from library/nginx
a803e7c4b030: Pull complete
8b625c47d697: Pull complete
4d3239651a63: Pull complete
0f816efa513d: Pull complete
01d159b8db2f: Pull complete
5fb9a81470f3: Pull complete
9b1e1e7164db: Pull complete
Digest: sha256:32da30332506740a2f7c34d5dc70467b7f14ec67d912703568daff790ab3f
Status: Downloaded newer image for nginx:latest
--> 61395b4c586d
Step 2/4 : COPY nginx.conf /etc/nginx/nginx.conf
--> e9a8196303dc
Step 3/4 : EXPOSE 80
--> Running in 065c58de4883
Removing intermediate container 065c58de4883
--> 4ceaad870e19
Step 4/4 : CMD ["nginx", "-g", "daemon off;"]
--> Running in ad2713a8c142
Removing intermediate container ad2713a8c142
--> b72755fd3d7c
Successfully built b72755fd3d7c
Successfully tagged nginx:latest
Creating lb-challenge_weather-app3_1 ... done
Creating lb-challenge_weather-app1_1 ... done
Creating lb-challenge_loadbalancer_1 ... done
Creating lb-challenge_weather-app2_1 ... done
[root@ip-10-0-1-4 lb-challenge]# [...]
```

Now let's check our work:

docker ps

Copy the public IP address of the Swarm master and paste it into a new tab in your browser.

Create a Docker service using Docker Swarm

Change to the root directory:

cd ~/

View the contents of swarm-token.txt:

cat swarm-token.txt

```
[root@ip-10-0-1-4 lb-challenge]# cd ~/
[root@ip-10-0-1-4 ~]# cat swarm-token.txt
Swarm initialized: current node (yt05f9gtvto20vp2s209mjy5i) is now a manager.

To add a worker to this swarm, run the following command:

  docker swarm join --token SWMTKN-1-17j7h3wn62fkvit6zevjebaarb2o5imtrmdux7dwbonvnid4h1-3if43dsktfc7re2c8mezyl0i1 10.0.1.4:2377

To add a manager to this swarm, run 'docker swarm join-token manager' and follow the instructions.

[root@ip-10-0-1-4 ~]#
```

Copy the docker swarm join command from the previous step.

Note: You can also initiate this process external to the script by running the following on the Master:

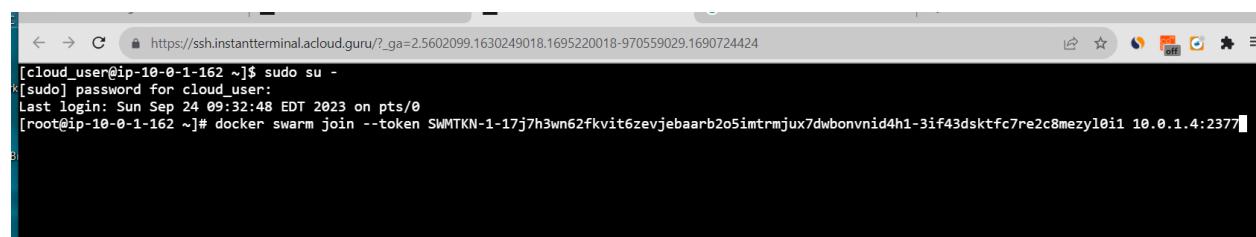
sudo docker swarm init

On Swarm worker node

Execute the command that was copied from the previous step.

i.e., the string docker swarm join –token <token> <IP address>:<port>

docker swarm join --token SWMTKN-1-17j7h3wn62fkvit6zevjebaarb2o5imtrmdux7dwbonvnid4h1-3if43dsktfc7re2c8mezyl0i1 10.0.1.4:2377



```
[cloud_user@ip-10-0-1-162 ~]$ sudo su -
[sudo] password for cloud_user:
Last login: Sun Sep 24 09:32:48 EDT 2023 on pts/0
[root@ip-10-0-1-162 ~]# docker swarm join --token SWMTKN-1-17j7h3wn62fkvit6zevjebaarb2o5imtrmjax7dwbonnid4h1-3if43dsktfc7re2c8mezy1oi1 10.0.1.4:2377
This node joined a swarm as a worker.
[root@ip-10-0-1-162 ~]#
```

On Swarm master

Create a Docker service by executing the following command:

```
docker service create --name nginx-app --publish published=8080,target=80 --replicas=2 nginx
```

```
[root@ip-10-0-1-4 ~]# docker service create --name nginx-app --publish published=8080,target=80 --replicas=2 nginx
[cp0sutcgldlw8mpb5ivjb4y2
overall progress: 2 out of 2 tasks
1/2: running [=====>]
2/2: running [=====>]
Verify: Service converged
[root@ip-10-0-1-4 ~]#
```

Note: the port 8080 used here. Please ensure that its open.

Verify this completed successfully by running the following command on both servers:

| CONTAINER ID | IMAGE | COMMAND | CREATED | STATUS | PORTS | NAMES |
|----------------------|---------------------------|--------------------------|--------------------|-------------------|-----------------------------------|-------------------|
| 691b98d167e5 | nginx:latest | "/docker-entrypoint..." | About a minute ago | Up About a minute | 80/tcp | nginx-app.1.1beqr |
| b2kp29b94bqzqyrlrg3q | nginx | "/docker-entrypoint..." | 8 minutes ago | Up 8 minutes | 0.0.0.0:80->80/tcp, :::80->80/tcp | lb-challenge_load |
| 7e3cd322282 | nginx | "/docker-entrypoint..." | 8 minutes ago | Up 8 minutes | 3000/tcp | lb-challenge_weat |
| balancer_1 | lb-challenge_weather-app2 | "docker-entrypoint.s..." | 8 minutes ago | Up 8 minutes | 3000/tcp | lb-challenge_weat |
| 98301bc56291 | lb-challenge_weather-app2 | "docker-entrypoint.s..." | 8 minutes ago | Up 8 minutes | 3000/tcp | lb-challenge_weat |
| her-app2_1 | lb-challenge_weather-app3 | "docker-entrypoint.s..." | 8 minutes ago | Up 8 minutes | 3000/tcp | lb-challenge_weat |
| 3ad6e3a29f83 | lb-challenge_weather-app3 | "docker-entrypoint.s..." | 8 minutes ago | Up 8 minutes | 3000/tcp | lb-challenge_weat |
| her-app3_1 | lb-challenge_weather-app1 | "docker-entrypoint.s..." | 8 minutes ago | Up 8 minutes | 3000/tcp | lb-challenge_weat |
| d3a92f585e7 | lb-challenge_weather-app1 | "docker-entrypoint.s..." | 8 minutes ago | Up 8 minutes | 3000/tcp | lb-challenge_weat |
| her-app1_1 | | | | | | |

docker ps

Note: If you wish to see all containers (whether running or shutdown) use the **docker ps -a** command.

Verify that the default nginx page loads in the browser:

PUBLIC_IP_ADDRESS:8080



Welcome to nginx!

If you see this page, the nginx web server is successfully installed and working. Further configuration is required.

For online documentation and support please refer to nginx.org.
Commercial support is available at nginx.com.

Thank you for using nginx.

Addendum

How to download files used for this project between your local machine and Centos.

```
[root@ip-10-0-1-103 ~]# ls
anaconda-ks.cfg  lb-challenge  swarm-token.txt
[root@ip-10-0-1-103 ~]# pwd
/
```

The files, due to the permissions, can only be worked on by someone with root account access. Hence you can only see the above with the root account.

You may, however, wish to download the files and folders to work on them from your own local machine.

To make matter more complicated, when using **WinSCP**, you are unable to login as a root user. Hence, you will not be able to download the files!

Rather than playing with both the permissions and commands like **sudo visudo** that can weaken security you can instead copy the files from the root profile to a non-root profile as follows:

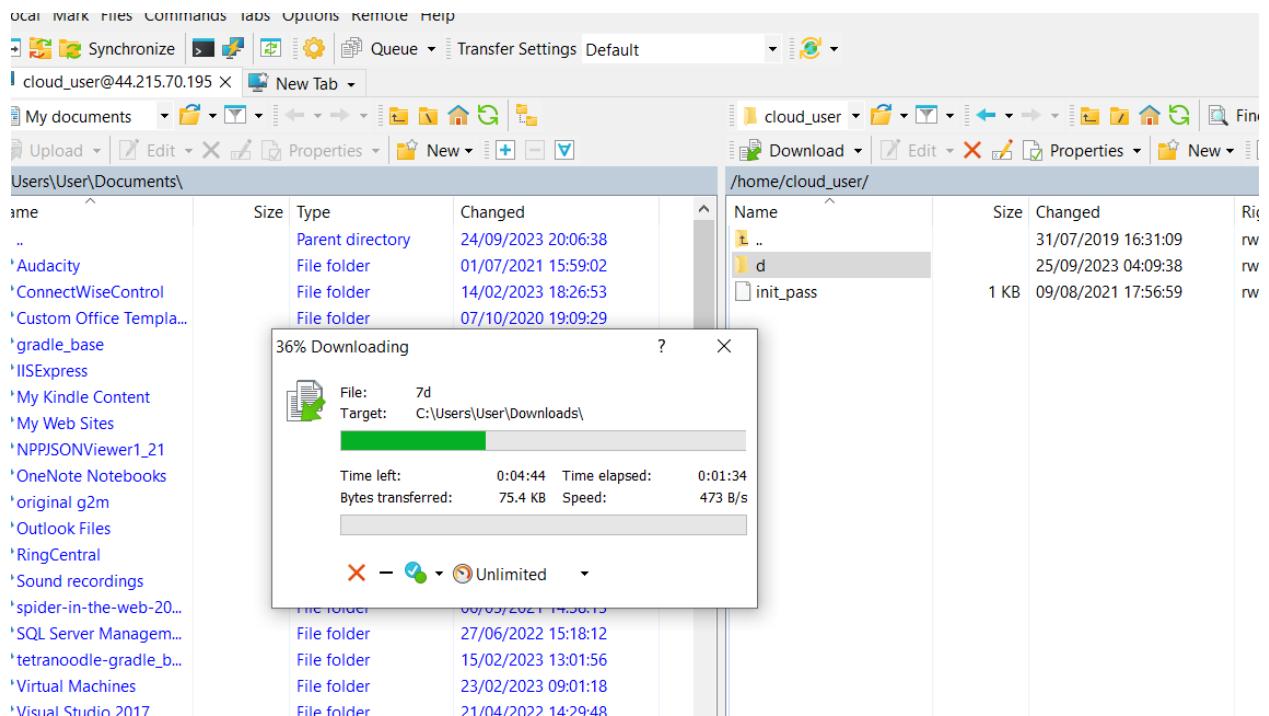
```
[root@ip-10-0-1-103 ~]# cp -r /root/lb-challenge/ /home/cloud_user/d
[root@ip-10-0-1-103 ~]# cp -r /root/anaconda-ks.cfg /home/cloud_user/d
[root@ip-10-0-1-103 ~]# cp -r /root/swarm-token.txt /home/cloud_user/d
[root@ip-10-0-1-103 ~]#
```

Note: **pwd** is helpful in constructing the folder paths.

cp == copy

-r == recursive command. Good when you have subfolders.

/path/source/ /path/destination

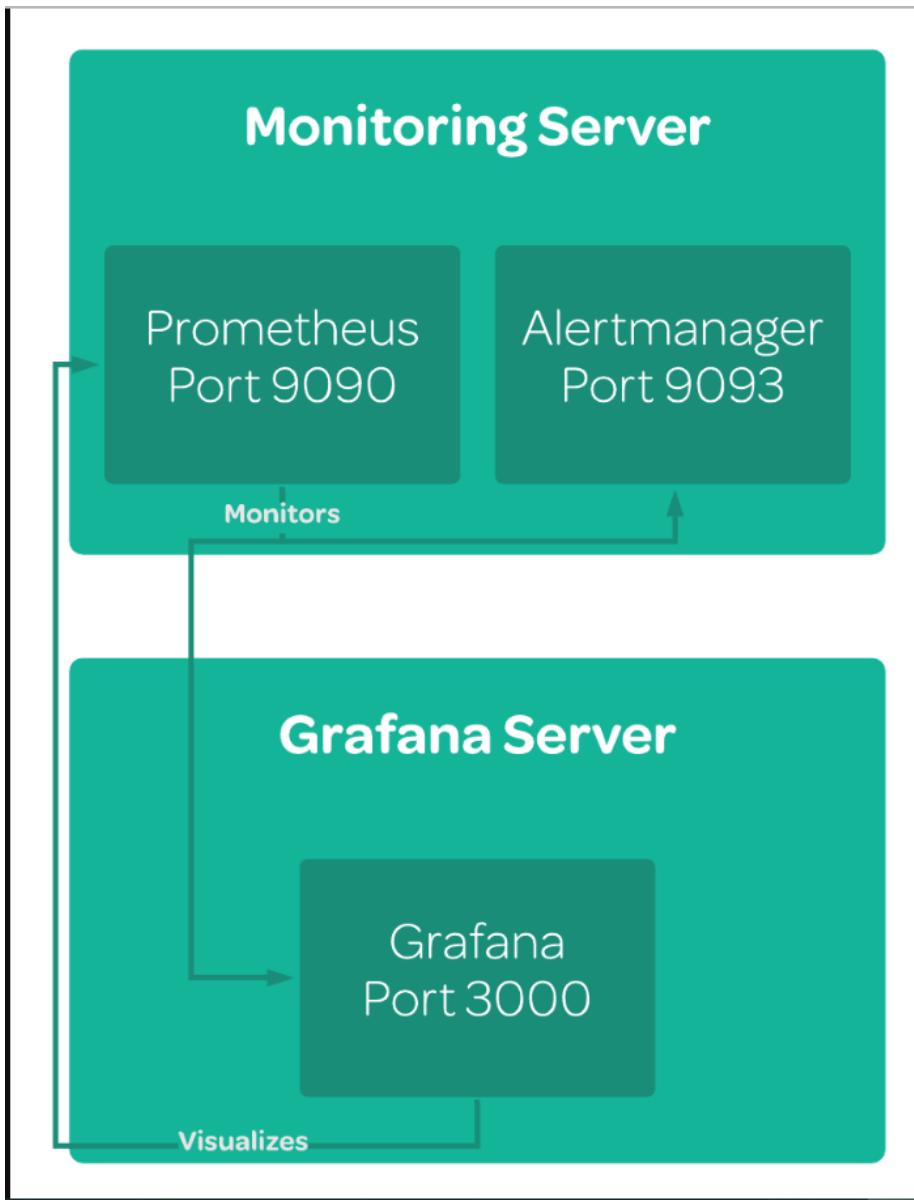


If you have any 'stubborn' files that won't download. No problem!

```
[root@ip-10-0-1-103 ~]# cat anaconda-ks.cfg
```

Use the cat command! Copy and paste the text locally.

Setting Up Prometheus and Adding Endpoints



Solution

Prerequisites:

You need two Linux servers.

One for installing **Prometheus** and the **Alert Manager**
The other for **Grafana**.

Call one **Monitoring** and the other **Grafana**.

As the following commands will, likely, run on any distro I am agnostic on which one you may choose.

For this project I will be using Centos 7

Logging In

NOTE: It's probably a good idea to have two or three terminals open, one for each node.

Install Docker and Kubernetes on All Servers

Most of these commands need to be run on each of the nodes. Pay attention though. Down at Step 10, we are going to do a little bit on just the master, and down at Step 15 we'll run something on just the nodes. There are notes down there, just be watching for them.

1 - Once we have logged in, we need to elevate privileges using sudo:

sudo su

2 - Disable SELinux:

setenforce 0

sed -i --follow-symlinks 's/SELINUX=enforcing/SELINUX=disabled/g' /etc/sysconfig/selinux

3 - Enable the br_netfilter module for cluster communication:

modprobe br_netfilter

echo '1' > /proc/sys/net/bridge/bridge-nf-call-iptables

4 - Ensure that the Docker dependencies are satisfied:

yum install -y yum-utils device-mapper-persistent-data lvm2

5 - Add the Docker repo and install Docker:

yum-config-manager --add-repo <https://download.docker.com/linux/centos/docker-ce.repo>

yum install -y docker-ce

6 - Set the cgroup driver for Docker to systemd, reload systemd, then enable and start Docker:

sed -i '/^ExecStart/ s/\$/ --exec-opt native.cgroupdriver=systemd/'

/usr/lib/systemd/system/docker.service

systemctl daemon-reload

systemctl enable docker --now

7 - Add the Kubernetes repo:

cat << EOF > /etc/yum.repos.d/kubernetes.repo

[kubernetes]

name=Kubernetes

baseurl=https://packages.cloud.google.com/yum/repos/kubernetes-el7-x86_64

enabled=1

gpgcheck=0

repo_gpgcheck=0

gpgkey=<https://packages.cloud.google.com/yum/doc/yum-key.gpg>

<https://packages.cloud.google.com/yum/doc/rpm-package-key.gpg>

EOF

8 - Install Kubernetes v1.14.0:

```
yum install -y kubelet-1.14.0-0 kubeadm-1.14.0-0 kubectl-1.14.0-0 kubernetes-cni-0.7.5
```

9 - Enable the kubelet service. The kubelet service will fail to start until the cluster is initialized, this is expected:

```
systemctl enable kubelet
```

Note: Complete the following section on the **MASTER ONLY!**

10 - Initialize the cluster using the IP range for Flannel:

```
kubeadm init --pod-network-cidr=10.244.0.0/16
```

11 - Copy the kubeadmn join command that is in the output. We will need this later.

NOTE: see **Addendum** for an example of the whole output from using the **kubeadm init** command.

12 - Exit sudo, copy the admin.conf to your home directory, and take ownership.

```
mkdir -p $HOME/.kube
```

```
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
```

```
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

13 - Deploy Flannel:

```
kubectl apply -f https://raw.githubusercontent.com/flannel-io/flannel/master/Documentation/kube-flannel-old.yaml
```

14 - Check the cluster state:

```
kubectl get pods --all-namespaces
```

Note: Complete the following steps on the **NODES ONLY!**

15 - Run the join command that you copied earlier, this requires running the command prefaced with sudo on the nodes (if we hadn't run sudo su to begin with).

Then we'll check the nodes from the **master**.

```
kubectl get nodes
```

[Create and Scale a Deployment Using kubectl](#)

Note: These commands will only be run on the **master node**.

16 - Create a simple deployment:

```
kubectl create deployment nginx --image=nginx
```

17 - Inspect the pod:

```
kubectl get pods
```

18 - Scale the deployment:

kubectl scale deployment nginx --replicas=4

19 - Inspect the pods. We should have four now:

kubectl get pods

Configure permissions for Docker on both Master and node/s

sudo chmod 666 /var/run/docker.sock

docker run hello-world

Also install the following on both Master and node/s

Install the Git package:

sudo yum install git -y

Install Python3 on both Master and node/s:

NOTE: It is important to ensure that both **PIP** and **botocore boto3** libraries are installed on this system.
Please follow these steps:

sudo yum update

sudo yum install python3

sudo yum install python3-pip

pip install botocore boto3

python3 –version

Install Java 11 on both Master and node/s

sudo apt install default-jdk -y

Please, ensure that the above ports are open externally and between each other.

Log into the monitoring server

ssh cloud_user@<MONITORING_SERVER_PUBLIC_IP_ADDRESS>

Note: When copying and pasting code into Vim as you will later, first enter

:set paste (and then i to enter insert mode) to avoid adding unnecessary spaces and hashes.

Set Up Prometheus

On the monitoring server, create a user for Prometheus:

sudo useradd --no-create-home --shell /bin/false prometheus

```
cloud_user@mon:~$ sudo useradd --no-create-home --shell /bin/false prometheus
[sudo] password for cloud_user:
```

Create the needed directories:

```
sudo mkdir /etc/prometheus
sudo mkdir /var/lib/Prometheus
```

```
cloud_user@mon:~$ sudo mkdir /etc/prometheus
cloud_user@mon:~$ sudo mkdir /var/lib/prometheus
```

Set the ownership of the /var/lib/prometheus directory:

```
sudo chown prometheus:prometheus /var/lib/Prometheus
```

```
cloud_user@mon:~$ sudo chown prometheus:prometheus /var/lib/prometheus
```

Download Prometheus:

```
cd /tmp/
wget https://github.com/prometheus/prometheus/releases/download/v2.7.1/prometheus-2.7.1.linux-amd64.tar.gz
```

2.7.1.linux-amd64.tar.gz

```
cloud_user@mon:/tmp$ wget https://github.com/prometheus/prometheus/releases/download/v2.7.1/prometheus-2.7.1.linux-amd64.tar.gz
--2023-09-26 07:58:37-- https://github.com/prometheus/prometheus/releases/download/v2.7.1/prometheus-2.7.1.linux-amd64.tar.gz
Resolving github.com (github.com)... 140.82.113.3
Connecting to github.com (github.com)|140.82.113.3|:443... connected.
HTTP request sent, awaiting response... 302 Found
Location: https://objects.githubusercontent.com/github-production-release-asset-2e65be/6838921/926b3400-2563-11e9-8088-63bcd1a7b246?X-Amz-Algorithm=AWS4-HMAC-SHA256&X-Amz-Credential=AKIAIWNJYAX4CSVEH53AX2F20230926X2fus-east-1%2Fs3%2Faws4_request&X-Amz-Date=20230926T075838Z&X-Amz-Expires=300&X-Amz-Signature=1d39eda3807ab9947a7b0249eea36747ce68187a46c14988f18337b500d87edb8&X-Amz-SignedHeaders=host&actor_id=0&key_id=0&repo_id=6838921&response-content-disposition=attachment%3B%20filename%3Dprometheus-2.7.1.linux-amd64.tar.gz&response-content-type=application%2Foctet-stream [following]
--2023-09-26 07:58:37-- https://objects.githubusercontent.com/github-production-release-asset-2e65be/6838921/926b3400-2563-11e9-8088-63bcd1a7b246?X-Amz-Algorithm=AWS4-HMAC-SHA256&X-Amz-Credential=AKIAIWNJYAX4CSVEH53AX2F20230926X2fus-east-1%2Fs3%2Faws4_request&X-Amz-Date=20230926T075838Z&X-Amz-Expires=300&X-Amz-Signature=1d39eda3807ab9947a7b0249eea36747ce68187a46c14988f18337b500d87edb8&X-Amz-SignedHeaders=host&actor_id=0&key_id=0&repo_id=6838921&response-content-disposition=attachment%3B%20filename%3Dprometheus-2.7.1.linux-amd64.tar.gz&response-content-type=application%2Foctet-stream
Resolving objects.githubusercontent.com (objects.githubusercontent.com)... 185.199.109.133, 185.199.110.133, 185.199.108.133, ...
Connecting to objects.githubusercontent.com (objects.githubusercontent.com)|185.199.109.133|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 38342763 (37M) [application/octet-stream]
Saving to: 'prometheus-2.7.1.linux-amd64.tar.gz'

prometheus-2.7.1.linux-amd64.tar.gz    100%[=====] 36.57M  181MB/s   in 0.2s
2023-09-26 07:58:37 (181 MB/s) - 'prometheus-2.7.1.linux-amd64.tar.gz' saved [38342763/38342763]
cloud_user@mon:/tmp$
```

Extract the files:

```
tar -xvf prometheus-2.7.1.linux-amd64.tar.gz
```

```
cloud_user@mon:~$ ls
cloud_user@mon:~$ cd /tmp/
cloud_user@mon:/tmp$ ls
prometheus-2.7.1.linux-amd64.tar.gz  systemd-private-0e2b4d0243044376a327d441faa9ad51-systemd-resolved.service-Crrxcz
snap-private-tmp                      systemd-private-0e2b4d0243044376a327d441faa9ad51-systemd-timesyncd.service-kQt5eu
cloud_user@mon:/tmp$ tar -xvf prometheus-2.7.1.linux-amd64.tar.gz
```

```
cloud_user@mon:~$ ls
cloud_user@mon:~$ cd /tmp/
cloud_user@mon:/tmp$ ls
prometheus-2.7.1.linux-amd64.tar.gz  systemd-private-0e2b4d0243044376a327d441faa9ad51-systemd-resolved.service-Crrxcz
snap-private-tmp                      systemd-private-0e2b4d0243044376a327d441faa9ad51-systemd-timesyncd.service-kQt5eu
cloud_user@mon:/tmp$ tar -xvf prometheus-2.7.1.linux-amd64.tar.gz
prometheus-2.7.1.linux-amd64/
prometheus-2.7.1.linux-amd64/LICENSE
prometheus-2.7.1.linux-amd64/prometheus.yml
prometheus-2.7.1.linux-amd64/consoles/
prometheus-2.7.1.linux-amd64/consoles/node-cpu.html
prometheus-2.7.1.linux-amd64/consoles/node.html
prometheus-2.7.1.linux-amd64/consoles/prometheus.html
prometheus-2.7.1.linux-amd64/consoles/prometheus-overview.html
prometheus-2.7.1.linux-amd64/consoles/node-overview.html
prometheus-2.7.1.linux-amd64/consoles/node-disk.html
prometheus-2.7.1.linux-amd64/consoles/index.html.example
prometheus-2.7.1.linux-amd64/console_libraries/
prometheus-2.7.1.linux-amd64/console_libraries/prom.lib
prometheus-2.7.1.linux-amd64/console_libraries/menu.lib
prometheus-2.7.1.linux-amd64/NOTICE
prometheus-2.7.1.linux-amd64/promtool
prometheus-2.7.1.linux-amd64/prometheus
cloud_user@mon:/tmp$
```

Move the configuration file and set the owner to the prometheus user:

```
cd prometheus-2.7.1.linux-amd64/
sudo mv console* /etc/prometheus
sudo mv prometheus.yml /etc/prometheus
sudo chown -R prometheus:prometheus /etc/prometheus
```

```
cloud_user@mon:/tmp$ cd prometheus-2.7.1.linux-amd64/
cloud_user@mon:/tmp/prometheus-2.7.1.linux-amd64$ sudo mv console* /etc/prometheus
[sudo] password for cloud_user:
cloud_user@mon:/tmp/prometheus-2.7.1.linux-amd64$ sudo mv prometheus.yml /etc/prometheus
cloud_user@mon:/tmp/prometheus-2.7.1.linux-amd64$ sudo chown -R prometheus:prometheus /etc/prometheus
cloud_user@mon:/tmp/prometheus-2.7.1.linux-amd64$
```

Move the binaries and set the owner:

```
sudo mv prometheus /usr/local/bin/
sudo mv promtool /usr/local/bin/
sudo chown prometheus:prometheus /usr/local/bin/prometheus
```

```
sudo chown prometheus:prometheus /usr/local/bin/promtool
```

Change back to your main directory, and then create the service file:

```
cloud_user@mon:/tmp/prometheus-2.7.1.linux-amd64$ sudo mv prometheus /usr/local/bin/
cloud_user@mon:/tmp/prometheus-2.7.1.linux-amd64$ sudo mv promtool /usr/local/bin/
cloud_user@mon:/tmp/prometheus-2.7.1.linux-amd64$ sudo chown prometheus:prometheus /usr/local/bin/prometheus
cloud_user@mon:/tmp/prometheus-2.7.1.linux-amd64$ sudo chown prometheus:prometheus /usr/local/bin/promtool
cloud_user@mon:/tmp/prometheus-2.7.1.linux-amd64$ █
```

```
cd
sudo vim /etc/systemd/system/prometheus.service
```

```
cloud_user@mon:/tmp/prometheus-2.7.1.linux-amd64$ cd
cloud_user@mon:~$ sudo vim /etc/systemd/system/prometheus.service
cloud_user@mon:~$ █
```

Add the following to the file:

```
[Unit]
Description=Prometheus
Wants=network-online.target
After=network-online.target
```

```
[Service]
User=prometheus
Group=prometheus
Type=simple
ExecStart=/usr/local/bin/prometheus \
  --config.file /etc/prometheus/prometheus.yml \
  --storage.tsdb.path /var/lib/prometheus/ \
  --web.console.templates=/etc/prometheus/consoles \
  --web.console.libraries=/etc/prometheus/console_libraries
```

```
[Install]
WantedBy=multi-user.target
```

Save and exit the file by pressing Escape followed by :wq.

```
[Unit]
Description=Prometheus
Wants=network-online.target
After=network-online.target

[Service]
User=prometheus
Group=prometheus
Type=simple
ExecStart=/usr/local/bin/prometheus \
    --config.file /etc/prometheus/prometheus.yml \
    --storage.tsdb.path /var/lib/prometheus/ \
    --web.console.templates=/etc/prometheus/consoles \
    --web.console.libraries=/etc/prometheus/console_libraries

[Install]
WantedBy=multi-user.target
```

Start Prometheus and make sure it automatically starts on boot:

```
sudo systemctl start prometheus  
sudo systemctl enable Prometheus
```

```
cloud_user@mon:~$ sudo systemctl start prometheus
cloud_user@mon:~$ sudo systemctl enable prometheus
Created symlink /etc/systemd/system/multi-user.target.wants/prometheus.service → /etc/systemd/system/prometheus.service.
cloud_user@mon:~$
```

Check Prometheus's status to ensure everything is working correctly:

sudo systemctl status prometheus

Set Up Alertmanager:

Still in the monitoring server, create the alertmanager system user:

```
sudo useradd --no-create-home --shell /bin/false alertmanager
```

Create the needed directories:

```
sudo mkdir /etc/alertmanager
```

```
cloud_user@mon:~$ sudo useradd --no-create-home --shell /bin/false alertmanager
cloud_user@mon:~$ sudo mkdir /etc/alertmanager
cloud_user@mon:~$ █
```

Download Alertmanager:

```
cd /tmp/
wget "https://github.com/prometheus/alertmanager/releases/download/v0.16.1/alertmanager-0.16.1.linux-amd64.tar.gz"
```

```
cloud_user@mon:~$ cd /tmp/
cloud_user@mon:/tmp$ wget "https://github.com/prometheus/alertmanager/releases/download/v0.16.1/alertmanager-0.16.1.linux-amd64.tar.gz"
--2023-09-26 08:23:36- https://github.com/prometheus/alertmanager/releases/download/v0.16.1/alertmanager-0.16.1.linux-amd64.tar.gz
Resolving github.com (github.com)... 140.82.113.4
Connecting to github.com (github.com)|140.82.113.4|:443... connected.
HTTP request sent, awaiting response... 302 Found
Location: https://objects.githubusercontent.com/github-production-release-asset-2e65be/11452538/5abdb580-257e-11e9-8be2-57c655a5e02d?X-Amz-Algorithm=AWS4-HMAC-SHA256&X-Amz-Credential=AKIAINWJYAX4CSVEH53A%2F20230926%2Fus-east-1%2Fs%2Faws4_request&X-Amz-Date=20230926T082337Z&X-Amz-Expires=300&X-Amz-Signature=db8f876fa7cc7a10f5e6fd09cf2e2b998b56900cd121658712eac2a542ece47c1&X-Amz-SignedHeaders=host&actor_id=0&key_id=0&repo_id=11452538&response-content-disposition=attachment%3B%20filename=alertmanager-0.16.1.linux-amd64.tar.gz&response-content-type=application%2Foctet-stream [following]
--2023-09-26 08:23:36- https://objects.githubusercontent.com/github-production-release-asset-2e65be/11452538/5abdb580-257e-11e9-8be2-57c655a5e02d?X-Amz-Algorithm=AWS4-HMAC-SHA256&X-Amz-Credential=AKIAINWJYAX4CSVEH53A%2F20230926%2Fus-east-1%2Fs%2Faws4_request&X-Amz-Date=20230926T082337Z&X-Amz-Expires=300&X-Amz-Signature=db8f876fe7cc7a10f5e6fd09cf2e2b998b56900cd121658712eac2a542ece47c1&X-Amz-SignedHeaders=host&actor_id=0&key_id=0&repo_id=11452538&response-content-disposition=attachment%3B%20filename=alertmanager-0.16.1.linux-amd64.tar.gz&response-content-type=application%2Foctet-stream
Resolving objects.githubusercontent.com (objects.githubusercontent.com)... 185.199.108.133, 185.199.111.133, 185.199.110.133, ...
Connecting to objects.githubusercontent.com (objects.githubusercontent.com)|185.199.108.133|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 19070056 (18M) [application/octet-stream]
Saving to: 'alertmanager-0.16.1.linux-amd64.tar.gz'

alertmanager-0.16.1.linux-amd64.tar.gz 100%[=====] 18.19M ---KB/s in 0.1s
2023-09-26 08:23:36 (169 MB/s) - 'alertmanager-0.16.1.linux-amd64.tar.gz' saved [19070056/19070056]

cloud_user@mon:/tmp$ █
```

Extract the files:

```
tar -xvf alertmanager-0.16.1.linux-amd64.tar.gz
```

```
cloud_user@mon:/tmp$ tar -xvf alertmanager-0.16.1.linux-amd64.tar.gz █
```

```
cloud_user@mon:/tmp$ tar -xvf alertmanager-0.16.1.linux-amd64.tar.gz
alertmanager-0.16.1.linux-amd64/
alertmanager-0.16.1.linux-amd64/NOTICE
alertmanager-0.16.1.linux-amd64/amtool
alertmanager-0.16.1.linux-amd64/LICENSE
alertmanager-0.16.1.linux-amd64/alertmanager
alertmanager-0.16.1.linux-amd64/alertmanager.yml
cloud_user@mon:/tmp$
```

`cd alertmanager-0.16.1.linux-amd64/`

```
cloud_user@mon:/tmp$ cd alertmanager-0.16.1.linux-amd64/
cloud_user@mon:/tmp/alertmanager-0.16.1.linux-amd64$
```

Move the binaries:

```
sudo mv alertmanager /usr/local/bin/
sudo mv amtool /usr/local/bin/
```

```
cloud_user@mon:/tmp/alertmanager-0.16.1.linux-amd64$ sudo mv alertmanager /usr/local/bin/
cloud_user@mon:/tmp/alertmanager-0.16.1.linux-amd64$ sudo mv amtool /usr/local/bin/
cloud_user@mon:/tmp/alertmanager-0.16.1.linux-amd64$
```

Set the ownership of the binaries:

```
sudo chown -R alertmanager:alertmanager /usr/local/bin/alertmanager
sudo chown -R alertmanager:alertmanager /usr/local/bin/amtool
```

```
cloud_user@mon:/tmp/alertmanager-0.16.1.linux-amd64$ sudo chown -R alertmanager:alertmanager /usr/local/bin/alertmanager
cloud_user@mon:/tmp/alertmanager-0.16.1.linux-amd64$ sudo chown -R alertmanager:alertmanager /usr/local/bin/amtool
cloud_user@mon:/tmp/alertmanager-0.16.1.linux-amd64$
```

Move the configuration file into the /etc/alertmanager directory:

```
sudo mv alertmanager.yml /etc/alertmanager/
```

```
cloud_user@mon:/tmp/alertmanager-0.16.1.linux-amd64$ sudo mv alertmanager.yml /etc/alertmanager/
cloud_user@mon:/tmp/alertmanager-0.16.1.linux-amd64$
```

Set the ownership of the /etc/alertmanager directory:

```
sudo chown -R alertmanager:alertmanager /etc/alertmanager/
```

```
cloud_user@mon:/tmp/alertmanager-0.16.1.linux-amd64$ sudo chown -R alertmanager:alertmanager /etc/alertmanager/
cloud_user@mon:/tmp/alertmanager-0.16.1.linux-amd64$
```

Create the alertmanager.service file for systemd:

```
sudo vim /etc/systemd/system/alertmanager.service
```

Add the following to the file:

```
[Unit]
Description=Alertmanager
Wants=network-online.target
After=network-online.target
```

```
[Service]
User=alertmanager
Group=alertmanager
Type=simple
WorkingDirectory=/etc/alertmanager/
ExecStart=/usr/local/bin/alertmanager \
--config.file=/etc/alertmanager/alertmanager.yml
```

```
[Install]
WantedBy=multi-user.target
```

Save and exit the file by pressing Escape followed by :wq.

```
[Unit]
Description=Alertmanager
Wants=network-online.target
After=network-online.target

[Service]
User=alertmanager
Group=alertmanager
Type=simple
WorkingDirectory=/etc/alertmanager/
ExecStart=/usr/local/bin/alertmanager \
--config.file=/etc/alertmanager/alertmanager.yml

[Install]
WantedBy=multi-user.target
~
~
~
```

Stop Prometheus, change back to your main directory, and then update the Prometheus configuration file to use Alertmanager:

```
sudo systemctl stop prometheus
cd
sudo vim /etc/prometheus/prometheus.yml
```

```
cloud_user@mon:/tmp/alertmanager-0.16.1.linux-amd64$ sudo systemctl stop prometheus
cloud_user@mon:/tmp/alertmanager-0.16.1.linux-amd64$ cd
cloud_user@mon:~/tmp/alertmanager-0.16.1.linux-amd64$ sudo vim /etc/prometheus/prometheus.yml
cloud_user@mon:~/tmp/alertmanager-0.16.1.linux-amd64$
```

Un-comment alertmanager, and change alertmanager to localhost (so it will look like the following):

```
alerting:
  alertmanagers:
    - static_configs:
      - targets:
        - localhost:9093
```

```
# Alertmanager configuration
alerting:
  alertmanagers:
    - static_configs:
      - targets:
          # - alertmanager:9093
```

```
# Alertmanager configuration
alerting:
  alertmanagers:
    - static_configs:
      - targets:
          - alertmanager:9093
```

Save and exit the file by pressing Escape followed by :wq.

Reload systemd, and then start the prometheus and alertmanager services:

```
sudo systemctl daemon-reload
sudo systemctl start prometheus
sudo systemctl start alertmanager
```

```
cloud_user@mon:~$ sudo systemctl daemon-reload
cloud_user@mon:~$ sudo systemctl start prometheus
cloud_user@mon:~$ sudo systemctl start alertmanager
cloud_user@mon:~$ █
```

Make sure alertmanager starts on boot:

```
sudo systemctl enable alertmanager
```

```
cloud_user@mon:~$ sudo systemctl enable alertmanager
Created symlink /etc/systemd/system/multi-user.target.wants/alertmanager.service → /etc/systemd/system/alertmanager.service.
cloud_user@mon:~$ █
```

Set Up Grafana

In a new terminal window, log in to the grafana server via SSH using the credentials provided on the lab page:

ssh cloud_user@<GRAFANA_SERVER_PUBLIC_IP_ADDRESS>

Install the prerequisite package:

sudo apt-get install libfontconfig

Y

```
cloud_user@grafana:~$ sudo apt-get install libfontconfig
[sudo] password for cloud_user:
Reading package lists... Done
Building dependency tree
Reading state information... Done
Note, selecting 'libfontconfig1' instead of 'libfontconfig'
The following packages were automatically installed and are no longer required:
  linux-aws-5.4.0-1060 linux-headers-5.4.0-1060-aws linux-image-5.4.0-1060-aws linux-modules-5.4.0-1060-aws
Use 'sudo apt autoremove' to remove them.
The following additional packages will be installed:
  fontconfig-config fonts-dejavu-core
The following NEW packages will be installed:
  fontconfig-config fonts-dejavu-core libfontconfig1
0 upgraded, 3 newly installed, 0 to remove and 26 not upgraded.
Need to get 1233 kB of archives.
After this operation, 4015 kB of additional disk space will be used.
Do you want to continue? [Y/n] ■
```

Download and install Grafana using the .deb package provided on the Grafana download page:

cd /tmp/

wget https://dl.grafana.com/oss/release/grafana_5.4.3_amd64.deb

```
cloud_user@grafana:~$ cd /tmp/
cloud_user@grafana:/tmp$ wget https://dl.grafana.com/oss/release/grafana_5.4.3_amd64.deb
--2023-09-26 08:45:29-  https://dl.grafana.com/oss/release/grafana_5.4.3_amd64.deb
Resolving dl.grafana.com (dl.grafana.com)... 146.75.38.217, 2a04:4e42:78::729
Connecting to dl.grafana.com (dl.grafana.com)|146.75.38.217|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 55405690 (53M) [application/x-debian-package]
Saving to: 'grafana_5.4.3_amd64.deb'

grafana_5.4.3_amd64.deb          100%[=====] 52.84M  166MB/s   in 0.3s
2023-09-26 08:45:30 (166 MB/s) - 'grafana_5.4.3_amd64.deb' saved [55405690/55405690]

cloud_user@grafana:/tmp$ sudo dpkg -i grafana_5.4.3_amd64.deb■
```

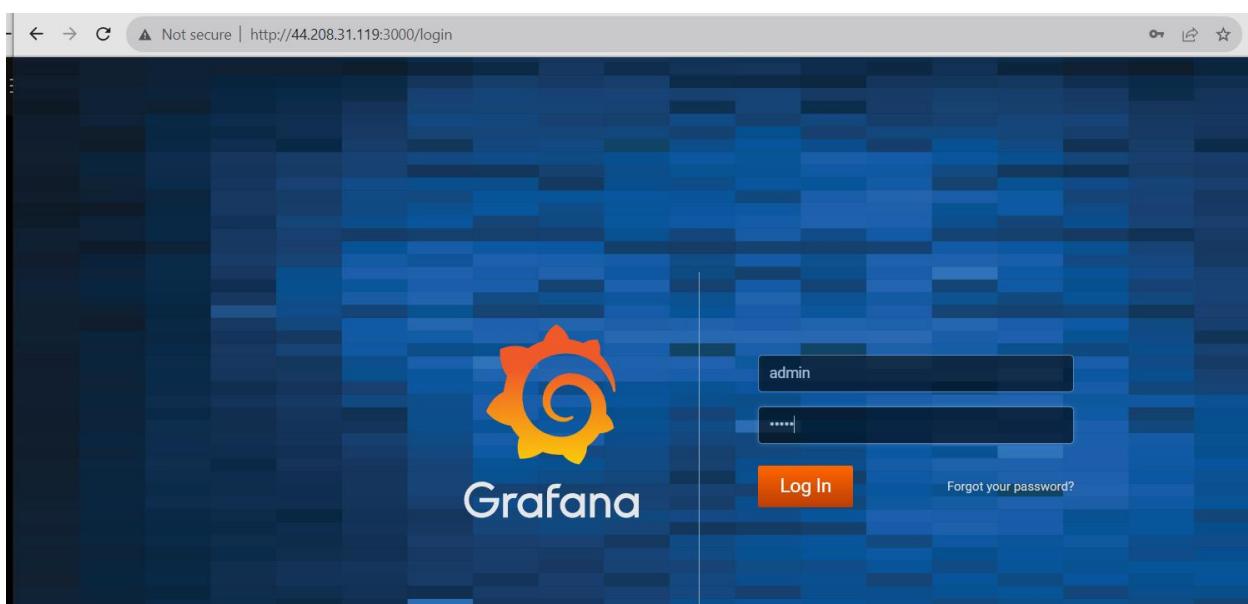
sudo dpkg -i grafana_5.4.3_amd64.deb

```
cloud_user@grafana:/tmp$ sudo dpkg -i grafana_5.4.3_amd64.deb
Selecting previously unselected package grafana.
(Reading database ... 149179 files and directories currently installed.)
Preparing to unpack grafana_5.4.3_amd64.deb ...
Unpacking grafana (5.4.3) ...
Setting up grafana (5.4.3) ...
Adding system user 'grafana' (UID 111) ...
Adding new user 'grafana' (UID 111) with group 'grafana' ...
Not creating home directory '/usr/share/grafana'.
## NOT starting on installation, please execute the following statements to configure grafana to start automatically using systemd
sudo /bin/systemctl daemon-reload
sudo /bin/systemctl enable grafana-server
## You can start grafana-server by executing
sudo /bin/systemctl start grafana-server
Processing triggers for systemd (237-3ubuntu10.57) ...
Processing triggers for ureadahead (0.100.0-21) ...
cloud_user@grafana:/tmp$
```

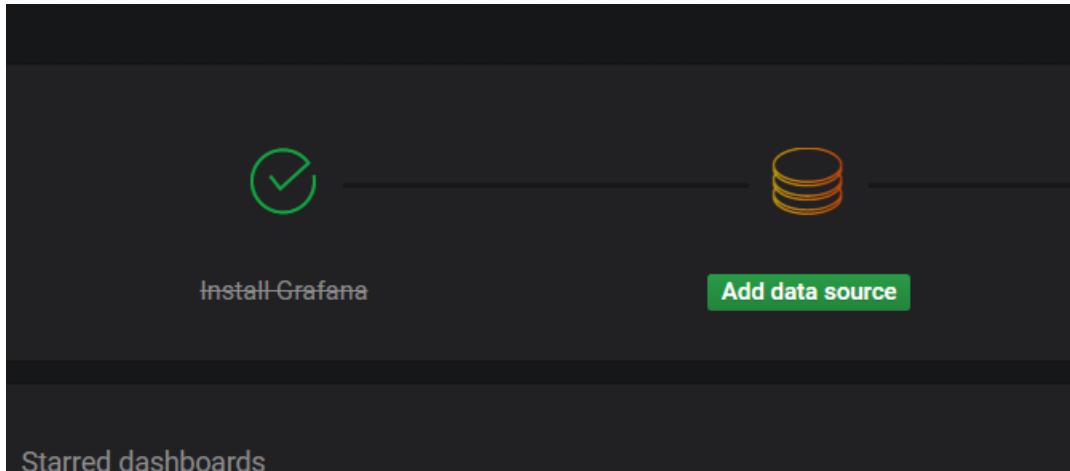
Start Grafana:**sudo systemctl start grafana-server**Ensure Grafana starts at boot:**sudo systemctl enable grafana-server**

```
cloud_user@grafana:/tmp$ sudo systemctl start grafana-server
cloud_user@grafana:/tmp$ sudo systemctl enable grafana-server
Synchronizing state of grafana-server.service with SysV service script with /lib/systemd/systemd-sysv-install.
Executing: /lib/systemd/systemd-sysv-install enable grafana-server
Created symlink /etc/systemd/system/multi-user.target.wants/grafana-server.service → /usr/lib/systemd/system/grafana-server.service.
cloud_user@grafana:/tmp$
```

Using the Grafana IP address listed on the lab page, access Grafana's web UI by navigating to **<GRAFANA_IP_ADDRESS>:3000** in a new browser tab.



Log in with the username **admin** and the password **admin**. Reset the password when prompted.

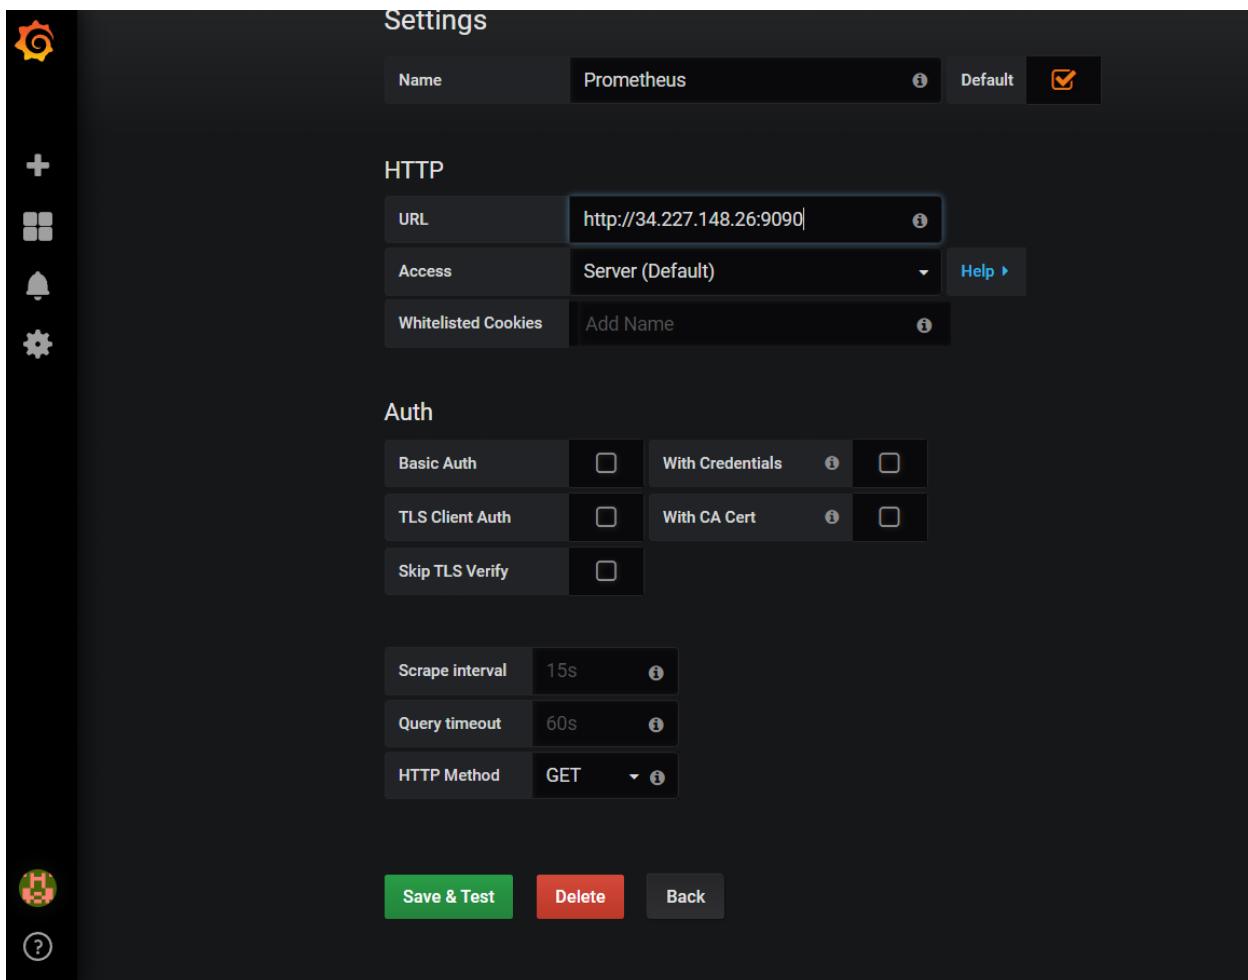
Add a Data Source

Click **Add data source** on the home page.

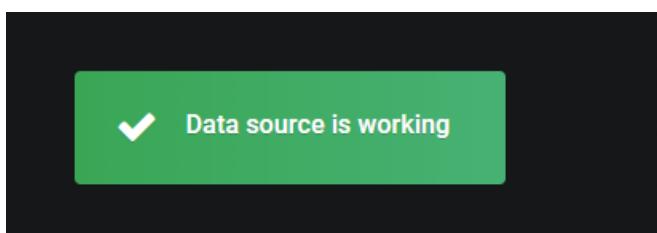
A screenshot of the Grafana Configuration page. The top navigation bar includes "Data Sources" (which is highlighted in orange), "Users", "Teams", "Plugins", "Preferences", and "API Keys". The main content area is titled "Choose data source type" and features a search bar with the placeholder "Filter by name or type". Below the search bar is a grid of nine data source options, each with an icon and a label: CloudWatch, Elasticsearch, Grafana Logging; Graphite, InfluxDB, Microsoft SQL Server; MySQL, OpenTSDB, PostgreSQL; Prometheus, Stackdriver, TestData DB.

Select **Prometheus**.

In the HTTP section, set the URL to http://<MONITORING_IP_ADDRESS>:9090 (using the public IP of the monitoring server provided on the lab page).



Click **Save & Test**.



Add Endpoints

Back in the **monitoring server** terminal, open the Prometheus configuration file:

```
sudo vim /etc/prometheus/prometheus.yml
```

At the end of the file, at the bottom of the scrape_configs section, add the Alertmanager endpoint (make sure it aligns with the - job_name: 'prometheus' line):

```
- job_name: 'alertmanager'
  static_configs:
    - targets: ['localhost:9093']
```

Beneath what you just added, add the Grafana endpoint (using the public IP address of the Grafana server):

```
- job_name: 'grafana'
  static_configs:
    - targets: ['<GRAFANA_IP_ADDRESS>:3000']
```

Save and exit the file by pressing Escape followed by :wq.

```
# The job name is added as a label `job=<
- job_name: 'prometheus'

# metrics_path defaults to '/metrics'
# scheme defaults to 'http'.

static_configs:
  - targets: ['localhost:9090']
- job_name: 'alertmanager'
  static_configs:
    - targets: ['localhost:9093']
- job_name: 'grafana'
  static_configs:
    - targets: ['44.208.31.119:3000']
```

Restart Prometheus:

sudo systemctl restart Prometheus

Check its status:

sudo systemctl status prometheus

```

cloud_user@mon:~$ sudo systemctl restart prometheus
cloud_user@mon:~$ sudo systemctl status prometheus
● prometheus.service - Prometheus
   Loaded: loaded (/etc/systemd/system/prometheus.service; enabled; vendor preset: enabled)
     Active: active (running) since Tue 2023-09-26 09:12:31 UTC; 22s ago
       Main PID: 32261 (prometheus)
         Tasks: 10 (limit: 1104)
        CGroup: /system.slice/prometheus --config.file /etc/prometheus/prometheus.yml --storage.tsdb.path /var/lib/prometheus/ --web.console.templates=/etc/
                  └─32261 /usr/local/bin/prometheus --config.file /etc/prometheus/prometheus.yml --storage.tsdb.path /var/lib/prometheus/ --web.console.templates=/etc/

Sep 26 09:12:31 mon prometheus[32261]: level=info ts=2023-09-26T09:12:31.631910232Z caller=main.go:304 host_details="(Linux 5.4.0-1092-aws #100~18.04.2-Ubuntu S
Sep 26 09:12:31 mon prometheus[32261]: level=info ts=2023-09-26T09:12:31.631952175Z caller=main.go:305 fd_limits="(soft=1024, hard=4096)"
Sep 26 09:12:31 mon prometheus[32261]: level=info ts=2023-09-26T09:12:31.631983667Z caller=main.go:306 vm_limits="(soft=unlimited, hard=unlimited)"
Sep 26 09:12:31 mon prometheus[32261]: level=info ts=2023-09-26T09:12:31.636602338Z caller=main.go:620 msg="Starting TSDB ... "
Sep 26 09:12:31 mon prometheus[32261]: level=info ts=2023-09-26T09:12:31.636671328Z caller=web.go:416 component=web msg="Start listening for connections" address=":9090"
Sep 26 09:12:31 mon prometheus[32261]: level=info ts=2023-09-26T09:12:31.661077185Z caller=main.go:635 msg="TSDB started"
Sep 26 09:12:31 mon prometheus[32261]: level=info ts=2023-09-26T09:12:31.661155465Z caller=main.go:695 msg="Loading configuration file" filename=/etc/prometheus
Sep 26 09:12:31 mon prometheus[32261]: level=info ts=2023-09-26T09:12:31.6621346Z caller=main.go:722 msg="Completed loading of configuration file" filename=/etc/
Sep 26 09:12:31 mon prometheus[32261]: level=info ts=2023-09-26T09:12:31.662173397Z caller=main.go:589 msg="Server is ready to receive web requests."
Sep 26 09:12:36 mon prometheus[32261]: level=error ts=2023-09-26T09:12:36.640107365Z caller=notifier.go:481 component=notifier alertmanager=http://alertmanager:9093
lines 1-18 (END)

```

Using the public IP address of the monitoring server, navigate to the Prometheus web UI in a new browser tab:

http://<MONITORING_IP_ADDRESS>:9090

The screenshot shows the Prometheus web interface. At the top, there's a navigation bar with links for 'Prometheus', 'Alerts', 'Graph', 'Status', and 'Help'. Below the navigation bar, there's a checkbox labeled 'Enable query history' which is unchecked. A large search bar is present with the placeholder 'Expression (press Shift+Enter for newlines)'. To the right of the search bar are two buttons: 'Execute' and a dropdown menu containing '- insert metric at cursor -'. Below the search bar, there are tabs for 'Graph' and 'Console', with 'Graph' being the active tab. Under the 'Graph' tab, there's a time range selector with 'Moment' in the center and arrows for navigating between time points. Below the time range is a section labeled 'Element' with the message 'no data'. At the bottom left of the interface, there's a blue button labeled 'Add Graph'.

Click **Status > Targets**.

The screenshot shows the Prometheus Targets page with three sections: alertmanager, grafana, and prometheus. Each section has a table with columns: Endpoint, State, Labels, Last Scrape, Scrape Duration, and Error. All endpoints are listed as UP.

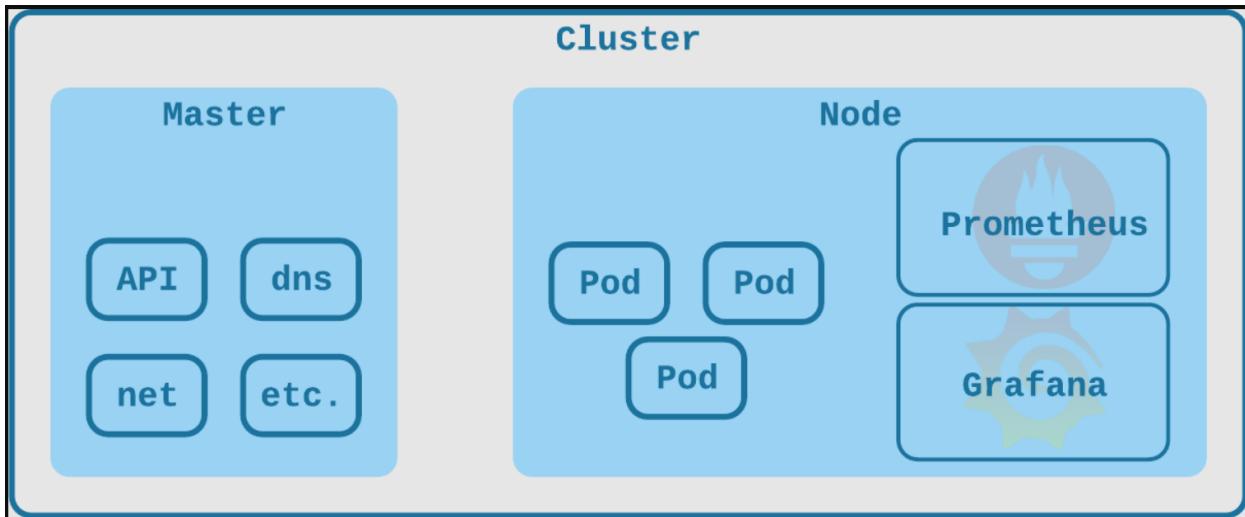
| Endpoint | State | Labels | Last Scrape | Scrape Duration | Error |
|-------------------------------|-------|--|-------------|-----------------|-------|
| http://localhost:9093/metrics | UP | instance="localhost:9093" job="alertmanager" | 10.458s ago | 3.023ms | |

| Endpoint | State | Labels | Last Scrape | Scrape Duration | Error |
|-----------------------------------|-------|---|-------------|-----------------|-------|
| http://44.208.31.119:3000/metrics | UP | instance="44.208.31.119:3000" job="grafana" | 7.493s ago | 5.42ms | |

| Endpoint | State | Labels | Last Scrape | Scrape Duration | Error |
|-------------------------------|-------|--|-------------|-----------------|-------|
| http://localhost:9090/metrics | UP | instance="localhost:9090" job="prometheus" | 3.753s ago | 4.961ms | |

Ensure all three endpoints are listed on the **Targets** page.

Monitoring in Kubernetes with Prometheus and Grafana



Creating a Kubernetes Cluster

Introduction

We will install and configure a Kubernetes cluster consisting of 1 master and 1 or 2 nodes. Once the installation and configuration are complete, we will have a 2-node Kubernetes cluster that uses Flannel as the network overlay.

We will be using Centos 7.x for all virtual machines.

Logging In

NOTE: It's probably a good idea to have two or three terminals open, one for each node.

Install Docker and Kubernetes on All Servers

Most of these commands need to be run on each of the nodes. Pay attention though. Down at Step 10, we are going to do a little bit on just the master, and down at Step 15 we'll run something on just the nodes. There are notes down there, just be watching for them.

1 - Once we have logged in, we need to elevate privileges using sudo:

```
sudo su
```

2 - Disable SELinux:

```
setenforce 0
```

```
sed -i --follow-symlinks 's/SELINUX=enforcing/SELINUX=disabled/g' /etc/sysconfig/selinux
```

3 - Enable the br_netfilter module for cluster communication:

```
modprobe br_netfilter
echo '1' > /proc/sys/net/bridge/bridge-nf-call-iptables
```

4 - Ensure that the Docker dependencies are satisfied:

```
yum install -y yum-utils device-mapper-persistent-data lvm2
```

5 - Add the Docker repo and install Docker:

```
yum-config-manager --add-repo https://download.docker.com/linux/centos/docker-ce.repo
yum install -y docker-ce
```

6 - Set the cgroup driver for Docker to systemd, reload systemd, then enable and start Docker:

```
sed -i '/^ExecStart/ s/$/ --exec-opt native.cgroupdriver=systemd/'
/usr/lib/systemd/system/docker.service
systemctl daemon-reload
systemctl enable docker --now
```

7 - Add the Kubernetes repo:

```
cat << EOF > /etc/yum.repos.d/kubernetes.repo
[kubernetes]
name=Kubernetes
baseurl=https://packages.cloud.google.com/yum/repos/kubernetes-el7-x86_64
enabled=1
gpgcheck=0
repo_gpgcheck=0
gpgkey=https://packages.cloud.google.com/yum/doc/yum-key.gpg
https://packages.cloud.google.com/yum/doc/rpm-package-key.gpg
EOF
```

8 - Install Kubernetes v1.14.0:

```
yum install -y kubelet-1.14.0-0 kubeadm-1.14.0-0 kubectl-1.14.0-0 kubernetes-cni-0.7.5
```

9 - Enable the kubelet service. The kubelet service will fail to start until the cluster is initialized, this is expected:

```
systemctl enable kubelet
```

Note: Complete the following section on the **MASTER ONLY!**

10 - Initialize the cluster using the IP range for Flannel:

```
kubeadm init --pod-network-cidr=10.244.0.0/16
```

11 - Copy the kubeadm join command that is in the output. We will need this later.

NOTE: see **Addendum** for an example of the whole output from using the **kubeadm init** command.

12 - Exit sudo, copy the admin.conf to your home directory, and take ownership.

```
mkdir -p $HOME/.kube
```

```
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
```

```
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

13 - Deploy Flannel:

```
kubectl apply -f https://raw.githubusercontent.com/flannel-io/flannel/master/Documentation/kube-flannel-old.yaml
```

14 - Check the cluster state:

```
kubectl get pods --all-namespaces
```

Note: Complete the following steps on the **NODES ONLY!**

15 - Run the join command that you copied earlier, this requires running the command prefaced with sudo on the nodes (if we hadn't run sudo su to begin with).

Then we'll check the nodes from the **master**.

kubectl get nodes

Create and Scale a Deployment Using kubectl

Note: These commands will only be run on the **master node**.

16 - Create a simple deployment:

kubectl create deployment nginx --image=nginx

17 - Inspect the pod:

kubectl get pods

18 - Scale the deployment:

kubectl scale deployment nginx --replicas=4

19 - Inspect the pods. We should have four now:

kubectl get pods

Configure permissions for Docker on both Master and node/s

sudo chmod 666 /var/run/docker.sock

docker run hello-world

Also install the following on both Master and node/s

Install the Git package:

sudo yum install git -y

Install Python3 on both Master and node/s:

NOTE: It is important to ensure that both **PIP** and **botocore boto3** libraries are installed on this system.

Please follow these steps:

sudo yum update

sudo yum install python3

sudo yum install python3-pip

pip install botocore boto3

python3 –version

Install Java 11 on both Master and node/s

```
sudo apt install default-jdk -y
```

Install components to use Gradle on both Master and node/s

```
sudo yum update
```

```
curl -o https://raw.githubusercontent.com/nvm-sh/nvm/v0.39.0/install.sh | bash
nvm install node
```

A task to perform on the Node

```
git clone https://github.com/ACloudGuru-Resources/cicd-pipeline-train-schedule-pipelines.git
```

```
cd
```

```
./gradlew build
```

```
./gradlew npm_start
```

Install Helm v2.8.2 on the Master

Helm is a package manager for Kubernetes!

Helm 3 can be installed many ways. We will install Helm 2.8.2 using the scripts option.

1.Download the script:

Download Helm v2.8.2 for your operating system from the Helm GitHub releases page.

Replace the URL with the appropriate one for your OS:

```
wget https://get.helm.sh/helm-v2.8.2-linux-amd64.tar.gz
tar -zxf helm-v2.8.2-linux-amd64.tar.gz
```

If you're using a different OS, adjust the URL accordingly.

2.Move Helm Binary to a Directory in Your PATH:

```
sudo mv linux-amd64/helm /usr/local/bin/
```

Move the Helm binary to a directory in your PATH so you can use it conveniently.

Initialize Helm with Tiller and the Specific Image:

```
helm init --stable-repo-url=https://charts.helm.sh/stable --wait --tiller-image
ghcr.io/helm/tiller:v2.8.2
```

```
cloud_user@ip-10-0-1-101:~$ helm init --stable-repo-url=https://charts.helm.sh/stable --wait --tiller-image ghcr.io/helm/tiller:v2.8.2
Creating /home/cloud_user/.helm
Creating /home/cloud_user/.helm/repository
Creating /home/cloud_user/.helm/repository/cache
Creating /home/cloud_user/.helm/repository/local
Creating /home/cloud_user/.helm/plugins
Creating /home/cloud_user/.helm/starters
Creating /home/cloud_user/.helm/cache/archive
Creating /home/cloud_user/.helm/repository/repositories.yaml
Adding stable repo with URL: https://charts.helm.sh/stable
Adding local repo with URL: http://127.0.0.1:8879/charts
$HELM_HOME has been configured at /home/cloud_user/.helm.

Tiller (the Helm server-side component) has been installed into your Kubernetes Cluster.

Please note: by default, Tiller is deployed with an insecure 'allow unauthenticated users' policy.
For more information on securing your installation see: https://docs.helm.sh/using_helm/#securing-your-helm-installation
Happy Helming!
cloud_user@ip-10-0-1-101:~$ █
```

Install Prometheus in the Kubernetes Cluster

To do this, make sure you have cloned the Kubernetes charts repo:

```
cd ~/
git clone https://github.com/kubernetes/charts
```

```
cloud_user@ip-10-0-1-101:~$ git clone https://github.com/kubernetes/charts
Cloning into 'charts'...
remote: Enumerating objects: 129385, done.
remote: Total 129385 (delta 0), reused 0 (delta 0), pack-reused 129385
Receiving objects: 100% (129385/129385), 284.57 MiB | 59.74 MiB/s, done.
Resolving deltas: 100% (87191/87191), done.
Checking connectivity... done.
cloud_user@ip-10-0-1-101:~$ █
```

```
cd charts
git checkout efdfcffe0b6973111ec6e5e83136ea74cdbe6527d
```

```
cloud_user@ip-10-0-1-101:~/charts$ git checkout efdfcffe0b6973111ec6e5e83136ea74cdbe6527d
Note: checking out 'efdfcffe0b6973111ec6e5e83136ea74cdbe6527d'.

You are in 'detached HEAD' state. You can look around, make experimental
changes and commit them, and you can discard any commits you make in this
state without impacting any branches by performing another checkout.

If you want to create a new branch to retain commits you create, you may
do so (now or later) by using -b with the checkout command again. Example:

  git checkout -b <new-branch-name>

HEAD is now at efdfcffe... [stable/prometheus] Ability to enable admin API (#5570)
cloud_user@ip-10-0-1-101:~/charts$ █
```

`cd ..`

Create a prometheus-values.yml file:

`vi prometheus-values.yml`

And paste in this content:

```
alertmanager:
  persistentVolume:
    enabled: false
server:
  persistentVolume:
    enabled: false
```

```
alertmanager:
  persistentVolume:
    enabled: false
server:
  persistentVolume:
    enabled: false
~
~
```

Save and close the file:

`:wq`

Use helm to install Prometheus with prometheus-values.yml:

```
helm install -f ~/prometheus-values.yml ~/charts/stable/prometheus --name prometheus --namespace prometheus
```

```
cloud_user@ip-10-0-1-101:~$ vi prometheus-values.yml
cloud_user@ip-10-0-1-101:~$ helm install -f ~/prometheus-values.yml ~/charts/stable/prometheus --name prometheus --namespace prometheus
```

NOTES:

The Prometheus server can be accessed via port 80 on the following DNS name from within your cluster:
prometheus-server.prometheus.svc.cluster.local

Get the Prometheus server URL by running these commands in the same shell:
`export POD_NAME=$(kubectl get pods --namespace prometheus -l "app=prometheus,component=server" -o jsonpath=".items[0].metadata.name")
 kubectl --namespace prometheus port-forward $POD_NAME 9090
 #####
 ##### WARNING: Persistence is disabled!!! You will lose your data when #####
 ##### the Server pod is terminated. #####
 #####`

The Prometheus alertmanager can be accessed via port 80 on the following DNS name from within your cluster:
prometheus-alertmanager.prometheus.svc.cluster.local

Get the Alertmanager URL by running these commands in the same shell:
`export POD_NAME=$(kubectl get pods --namespace prometheus -l "app=prometheus,component=alertmanager" -o jsonpath=".items[0].metadata.name")
 kubectl --namespace prometheus port-forward $POD_NAME 9093
 #####
 ##### WARNING: Persistence is disabled!!! You will lose your data when #####
 ##### the AlertManager pod is terminated. #####
 #####`

The Prometheus PushGateway can be accessed via port 9091 on the following DNS name from within your cluster:
prometheus-pushgateway.prometheus.svc.cluster.local

Get the PushGateway URL by running these commands in the same shell:
`export POD_NAME=$(kubectl get pods --namespace prometheus -l "app=prometheus,component=pushgateway" -o jsonpath=".items[0].metadata.name")
 kubectl --namespace prometheus port-forward $POD_NAME 9091`

For more information on running Prometheus, visit:
<https://prometheus.io/>

```
cloud_user@ip-10-0-1-101:~$
```

We can see which pods are running in this new namespace with the command:

kubectl get pods -n prometheus

```
cloud_user@ip-10-0-1-101:~$ kubectl get pods -n prometheus
NAME                               READY   STATUS    RESTARTS   AGE
prometheus-alertmanager-7b47d55bdf-7h4wd   2/2     Running   0          1m
prometheus-kube-state-metrics-6584885ccf-twdb2   1/1     Running   0          1m
prometheus-node-exporter-tk2td                1/1     Running   0          1m
prometheus-pushgateway-66c9fdb48f-d4xch      1/1     Running   0          1m
prometheus-server-65d5cc8544-b224f        2/2     Running   0          1m
cloud_user@ip-10-0-1-101:~$
```

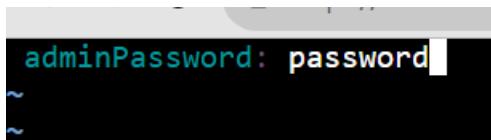
Install Grafana in the Kubernetes Cluster

Create a grafana-values.yml:

```
vi grafana-values.yml
```

Paste in this content (you will use this password to log in to Grafana):

```
adminPassword: password
```

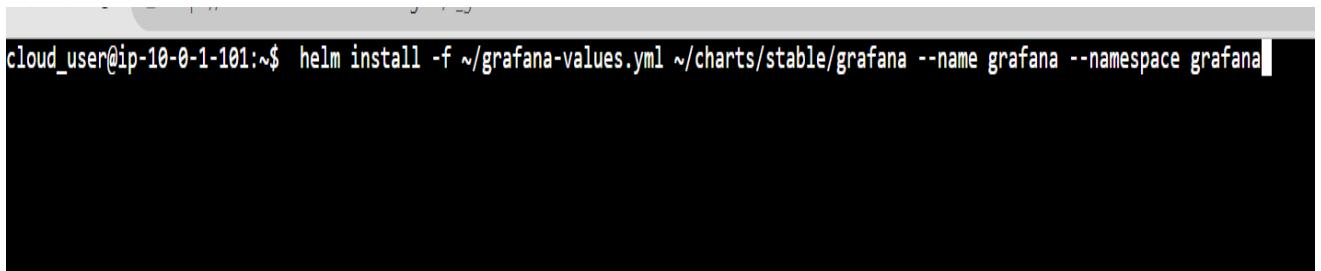
A screenshot of a terminal window. The text "adminPassword: password" is visible in the terminal area, indicating it has been typed or pasted into the file.

Save and close the file:

```
:wq
```

Use helm to install Grafana with grafana-values.yml:

```
helm install -f ~/grafana-values.yml ~/charts/stable/grafana --name grafana --namespace grafana
```

A screenshot of a terminal window. The command "helm install -f ~/grafana-values.yml ~/charts/stable/grafana --name grafana --namespace grafana" is visible at the top of the terminal window, indicating it has been typed or pasted into the terminal.

```

==> v1/ConfigMap
NAME          DATA AGE
grafana        1   0s
grafana-dashboards-json 0   0s

==> v1/Service
NAME      TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE
grafana  ClusterIP  10.104.74.80  <none>        80/TCP      0s

==> v1beta2/Deployment
NAME      DESIRED  CURRENT  UP-TO-DATE  AVAILABLE  AGE
grafana  1         1         1           0          0s

==> v1/Pod(related)
NAME          READY  STATUS      RESTARTS  AGE
grafana-d65c89cdc-778vg  0/1    ContainerCreating  0          0s

```

NOTES:

1. Get your 'admin' user password by running:

```
kubectl get secret --namespace grafana grafana -o jsonpath=".data.admin-password" | base64 --decode ; echo
```

2. The Grafana server can be accessed via port 80 on the following DNS name from within your cluster:

```
grafana.grafana.svc.cluster.local
```

Get the Grafana URL to visit by running these commands in the same shell:

```
export POD_NAME=$(kubectl get pods --namespace grafana -l "app=grafana,component=" -o jsonpath=".items[0].metadata.name")
kubectl --namespace grafana port-forward $POD_NAME 3000
```

3. Login with the password from step 1 and the username: admin

```
#####
##### WARNING: Persistence is disabled!!! You will lose your data when #####
#####           the Grafana pod is terminated.                         #####
#####
```

```
cloud_user@ip-10-0-1-101:~$ 
```

We can see which pods are running in this new namespace with the command:

```
kubectl get pods -n grafana
```

```
cloud_user@ip-10-0-1-101:~$ kubectl get pods -n grafana
NAME           READY   STATUS    RESTARTS   AGE
grafana-d65c89cdc-778vg   0/1     Running   0          1m
cloud_user@ip-10-0-1-101:~$ █
```

Deploy a NodePort Service to Provide External Access to Grafana

Make a file called grafana-ext.yml:

vi grafana-ext.yml

Paste in this content:

```
kind: Service
apiVersion: v1
metadata:
  namespace: grafana
  name: grafana-ext
spec:
  type: NodePort
  selector:
    app: grafana
  ports:
    - protocol: TCP
      port: 3000
      nodePort: 8081
```

```
kind: Service
apiVersion: v1
metadata:
  namespace: grafana
  name: grafana-ext
spec:
  type: NodePort
  selector:
    app: grafana
  ports:
    - protocol: TCP
      port: 3000
      nodePort: 8081█
```

Save and close the file:

:wq

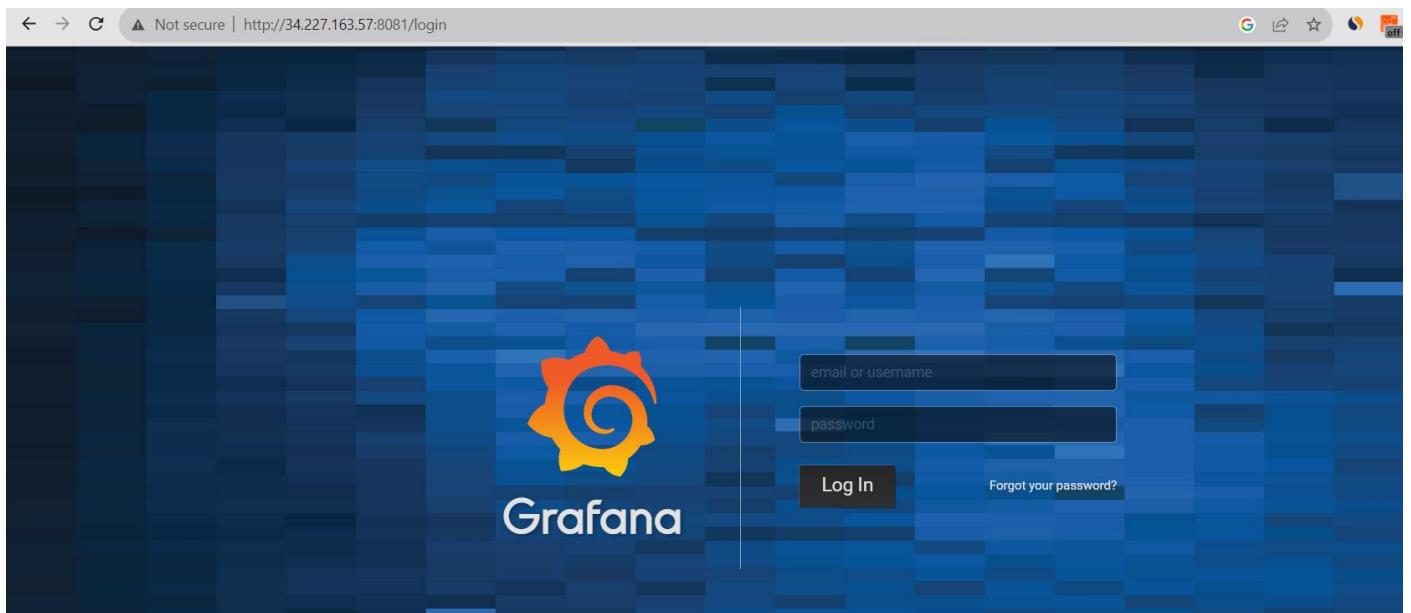
Deploy the service:

```
kubectl apply -f ~/grafana-ext.yml
```

```
cloud_user@ip-10-0-1-101:~$ vi grafana-ext.yml
cloud_user@ip-10-0-1-101:~$ kubectl apply -f ~/grafana-ext.yml
service "grafana-ext" created
cloud_user@ip-10-0-1-101:~$ █
```

Log in to Grafana using the Kubernetes Node Public IP

<Grafana Public IP address>:8081



Log in using the following credentials that were set earlier:

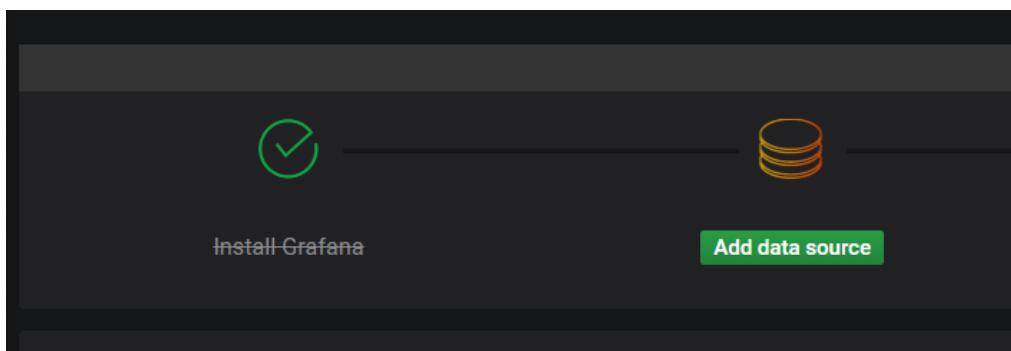
Username: admin

Password: password

Create the Monitoring Dashboards

Add a Datasource for Prometheus

Click on **Add data source**



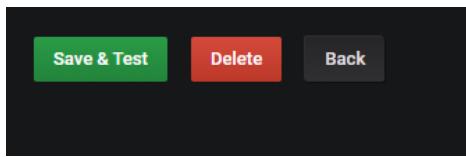
Name: **Kubernetes**

Type: **Prometheus**

URL: <http://prometheus-server.prometheus.svc.cluster.local>

Click **Save & Test**

The screenshot shows the 'Data Sources / Kubernetes' configuration page. At the top, there are tabs for 'Settings' (selected) and 'Dashboards'. The 'Name' field is set to 'Kubernetes' and is marked as 'Default'. The 'Type' is set to 'Prometheus'. Below this, the 'HTTP' section contains fields for 'URL' (set to 'http://localhost:9090') and 'Access' (set to 'proxy'). The 'Auth' section includes options for 'Basic Auth' and 'TLS Client Auth', both with checkboxes. A 'Skip TLS Verification (Insecure)' checkbox is also present. Under 'Advanced HTTP Settings', there are buttons for 'Whitelisted Cookies' and 'Add Name'. At the bottom, the 'Scrape interval' is set to '15s'.



Add the Kubernetes All Nodes Community Dashboard

Hover your mouse over the + in the left sidebar and click on **Import**.

In the Grafana.com Dashboard field, provide the ID **3131**.

Click outside of the field to load information about the dashboard.

In the Options section:

prometheus: **Kubernetes**

Click **Import**.

The screenshot shows the 'Import' dialog in Grafana. At the top, there's a logo with an orange arrow pointing up and the word 'Import'. Below it, the text 'Import dashboard from file or Grafana.com'.

Under the heading 'Importing Dashboard from [Grafana.com](#)', there are two cards:

- Published by: Bart Van Bos
- Updated on: 2017-09-06 14:14:13

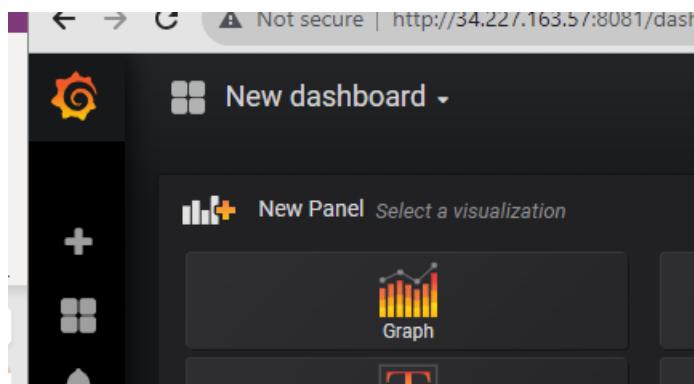
Below this, the 'Options' section contains two rows of configuration:

| Name | Kubernetes All Nodes | ✓ |
|------------|----------------------|---|
| prometheus | Kubernetes | ✓ |

At the bottom, there are two buttons: a green 'Import' button with a small icon and a grey 'Cancel' button.

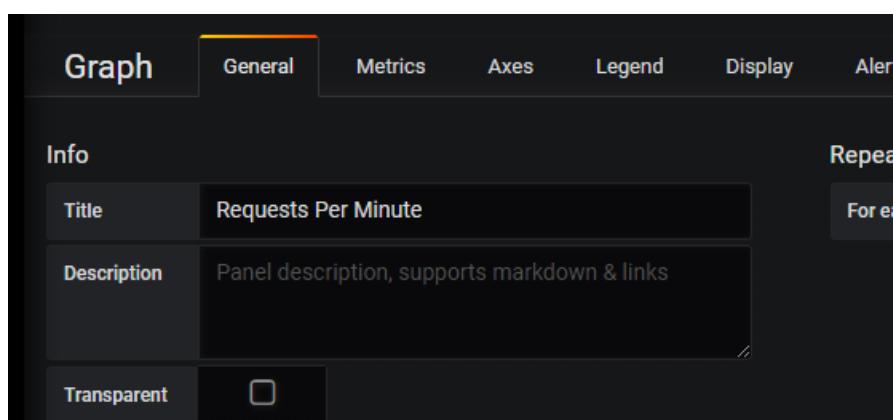
Hover your mouse over the + in the left sidebar and then click **Dashboard**.

Select the Graph panel.



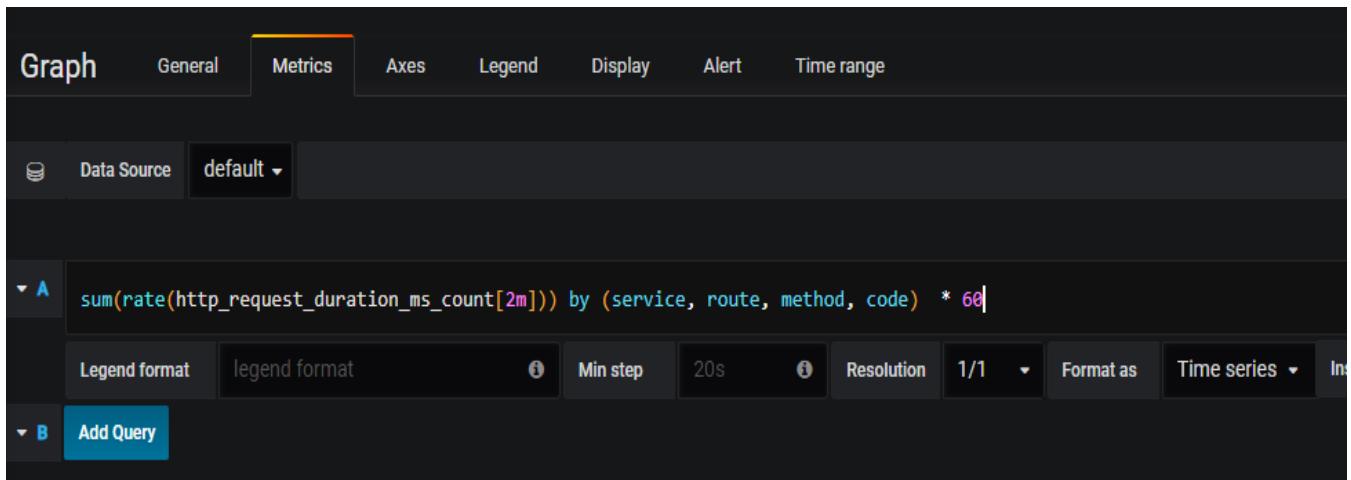
Hover over the Panel Title and click Edit.

In the General tab at the bottom of the screen, set the Title to Requests Per Minute.



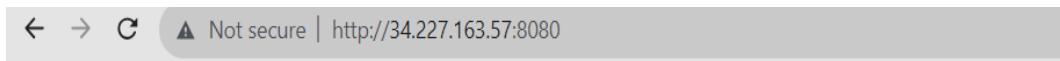
In the Metrics tab, paste in the following query:

```
sum(rate(http_request_duration_ms_count[2m])) by (service, route, method, code) * 60
```



With that in place, let's load our train-schedule app to give our graph some data:

<Grafana Public IP Address>:8080



Find your train!

Select your train below to see its current schedule.

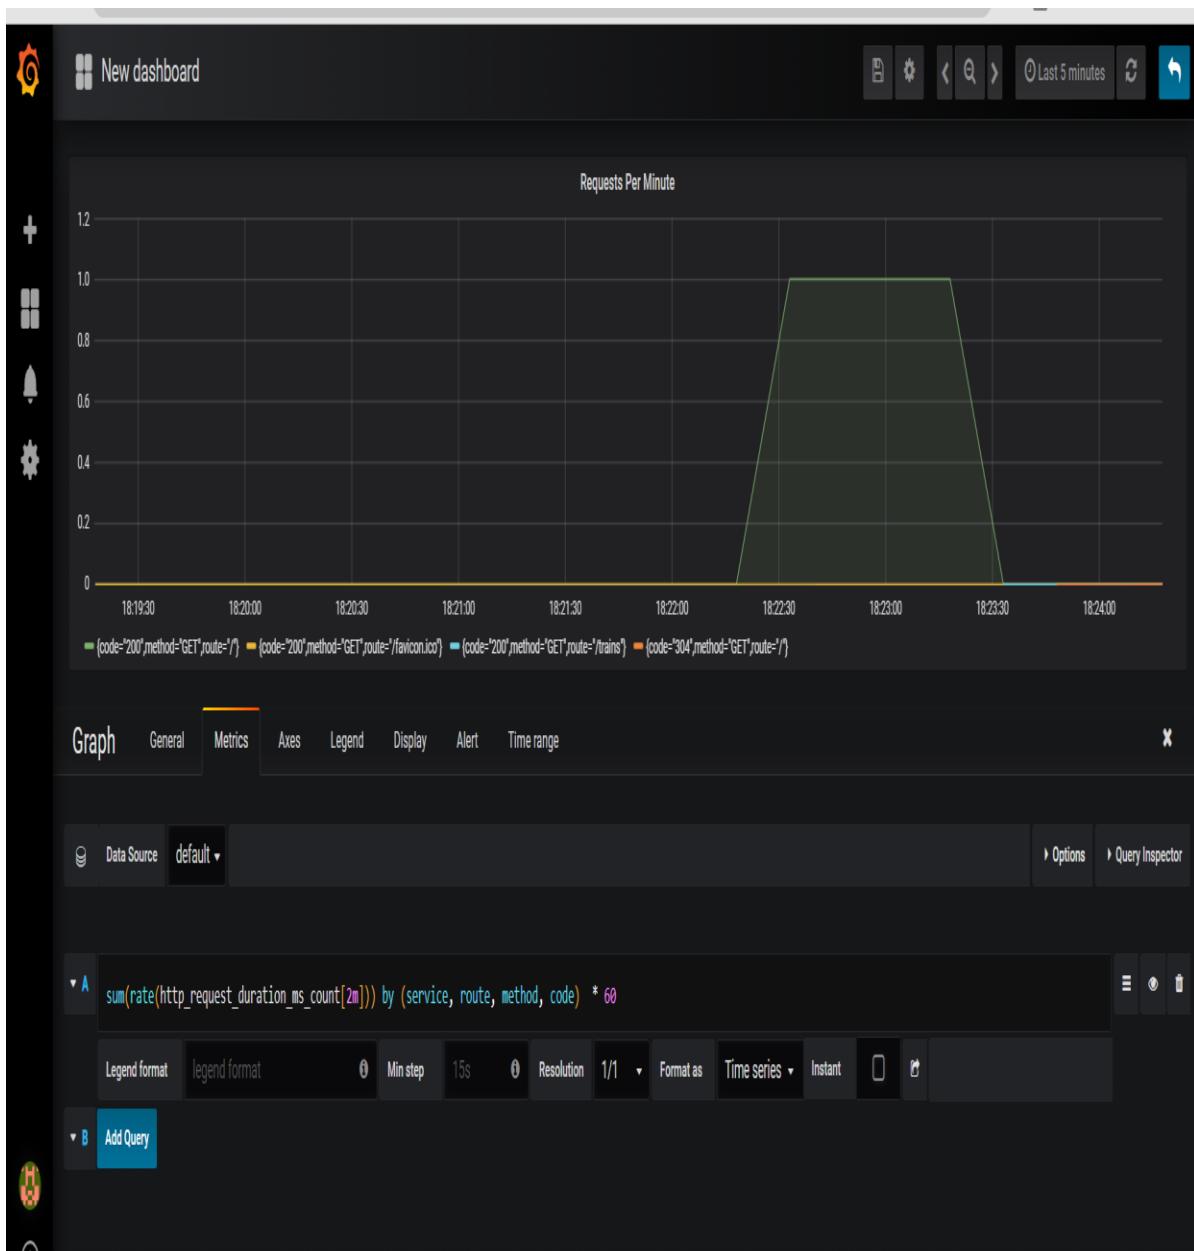
Trains

Flying Scotsman Golden Arrow Hogwarts Express Orient Express

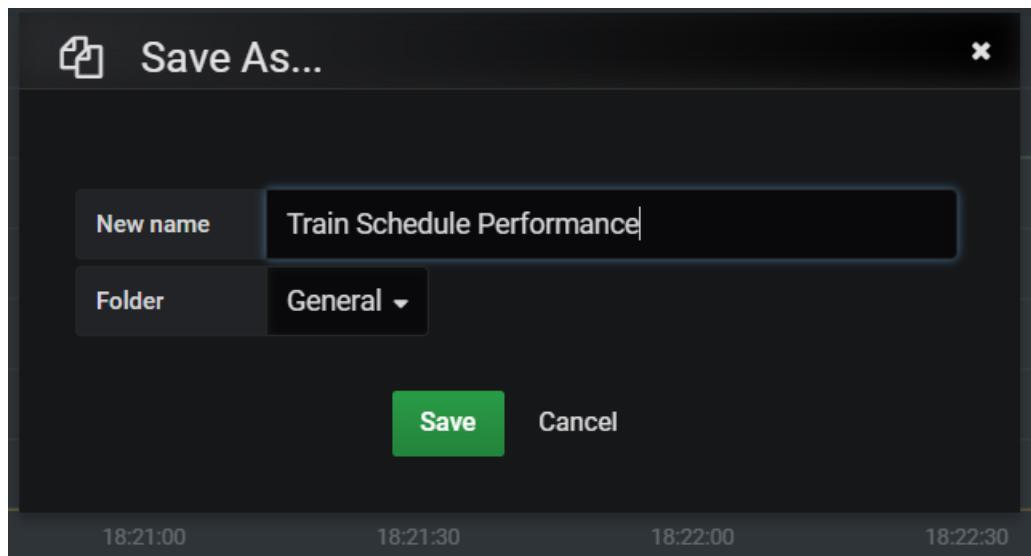
Select a train to view its current schedule.

Refresh the page a few times.

Now, navigate back to the Grafana dashboard tab in your browser.
In the top-right of the page, click Last 6 hours and change the time selection
to Last 5 minutes.



Click Back to dashboard in the top-right of the page.
Next, click Save dashboard, also in the top-right of the page.



Name the dashboard "**Train Schedule Performance**" and click **Save**.

||||||||||||||||||||||||||||||||||||||||||||||||
Addendum

Your Kubernetes control-plane has initialized successfully!

To start using your cluster, you need to run the following as a regular user:

```
mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

You should now deploy a pod network to the cluster.

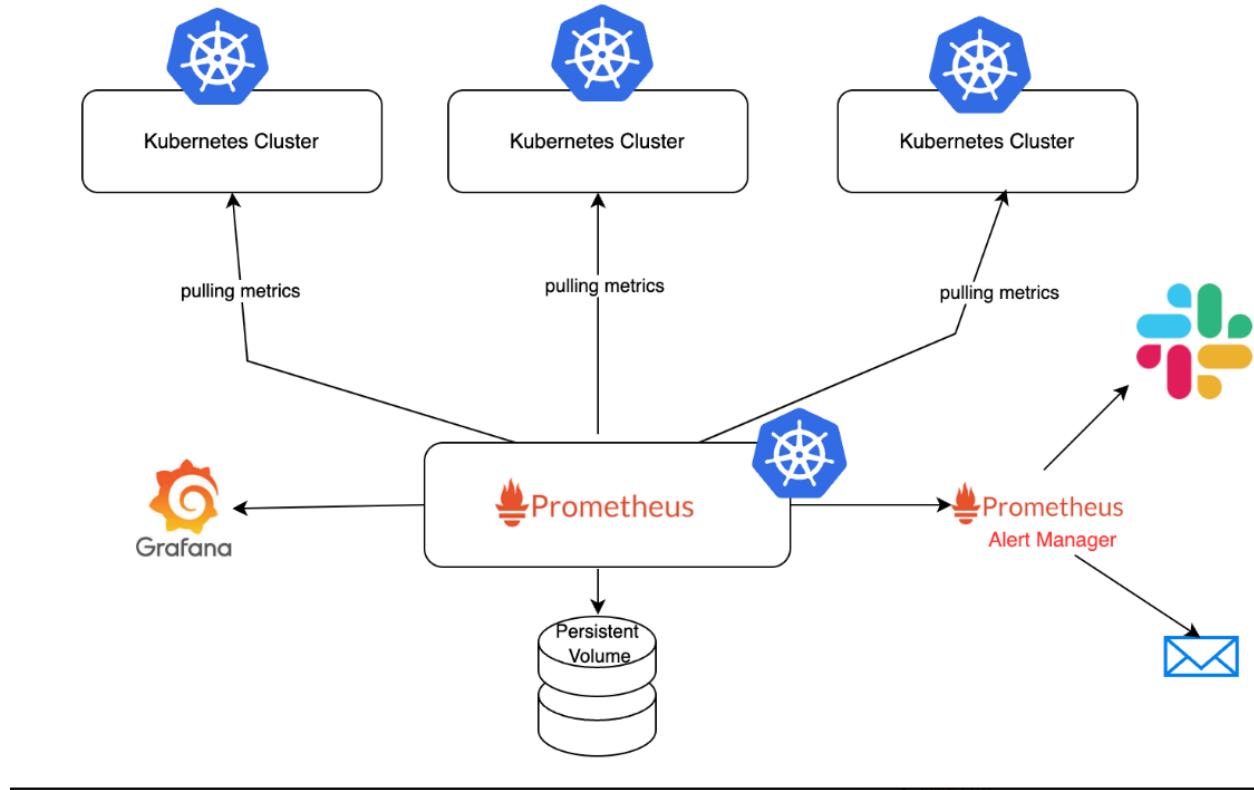
Run "kubectl apply -f [podnetwork].yaml" with one of the options listed at:
<https://kubernetes.io/docs/concepts/cluster-administration/addons/>

Then you can join any number of worker nodes by running the following on each as root:

```
kubeadm join 172.31.27.197:6443 --token 2kla1r.ncp3uo7myontulj7 \
--discovery-token-ca-cert-hash
sha256:228f75560859d6389ec4b6ac2b2b463f44bf42aa5f03c396596c8ec835e871ca
[root@9c438d91c5a44174a4a8c7a5cd9beb882c cloud_user]#
||||||||||||||||||||||||||||||||||||||||||||||||
```

How to setup monitoring on an EKS Kubernetes Cluster using Prometheus and Grafana

Prometheus and Grafana Setup for Monitoring Kubernetes Clusters



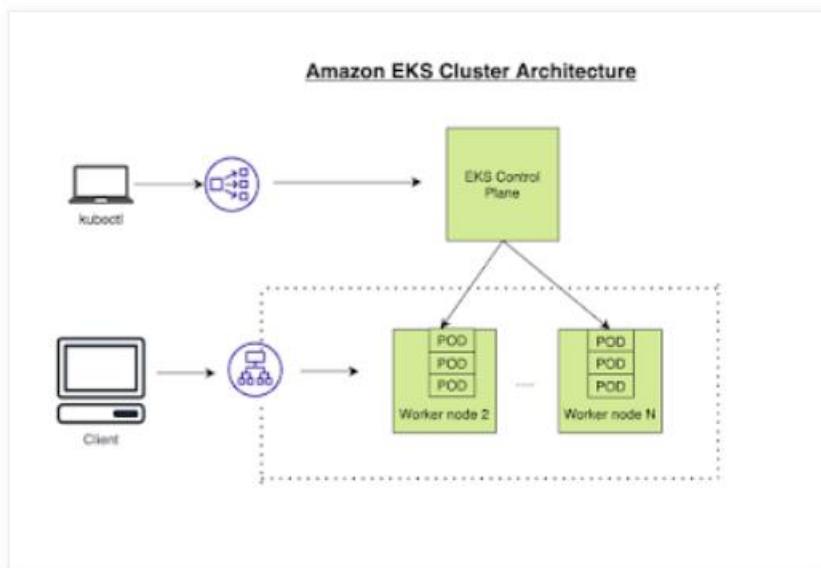
Pre-requisites:

- EC2 instance to access EKS cluster.
- EKS Cluster needs to be up and running.
- Install Helm3

How to create EKS cluster in AWS cloud using eksctl.

What is Amazon EKS

Amazon EKS is a fully managed container orchestration service. EKS allows you to quickly deploy a production ready Kubernetes cluster in AWS, deploy and manage containerized applications more easily with a fully managed Kubernetes service.

**Pre-requisites:****An EC2 instance (Ubuntu 22.4 LTS)**

Login to the AWS portal.

Go to **EC2****Launch Instance**

Launch instance

Name and tags [Info](#)

Name: VM [Add additional tags](#)

Application and OS Images (Amazon Machine Image) [Info](#)

An AMI is a template that contains the software configuration (operating system, application server, and applications) required to launch your instance. Search or Browse for AMIs if you don't see what you are looking for below.

Quick Start

| | | | | | | |
|--------------|-------|--------|---------|---------|------------|--|
| Amazon Linux | macOS | Ubuntu | Windows | Red Hat | SUSE Linux | Browse more AMIs Including AMIs from AWS, Marketplace and the Community |
|--------------|-------|--------|---------|---------|------------|--|

Amazon Machine Image (AMI)

| | |
|---|------------------------------------|
| Ubuntu Server 22.04 LTS (HVM), SSD Volume Type ami-053bb0d53c279ac90 (64-bit (x86)) / ami-0cbeebcd6dcdb0 (64-bit (Arm)) Virtualization: hvm ENA enabled: true Root device type: ebs | Free tier eligible |
|---|------------------------------------|

Select Ubuntu

Create a new key pair.

Key pair name
Key pairs allow you to connect to your instance securely.
VM

Key pair type
 RSA RSA encrypted private and public key pair
 ED25519 ED25519 encrypted private and public key pair

Private key file format
 .pem For use with OpenSSH
 .ppk For use with PuTTY

⚠️ When prompted, store the private key in a secure and accessible location on your computer. You will need it later to connect to your instance. [Learn more](#)

Cancel **Create key pair**

Click **Create key pair**.

Create the Instance.

Now select **View all instances**.

View all instances

Instances (1) [Info](#)

| <input type="checkbox"/> | Name | Instance ID | Instance state | Instance type | Status check | Alarm status | Available |
|--------------------------|------|---------------------|----------------------|---------------|---------------------------|--------------|-----------|
| <input type="checkbox"/> | VM | i-098a01231c18f85ab | Running | t2.micro | Initializing | No alarms | us-east-1 |

Select the instance on the left.

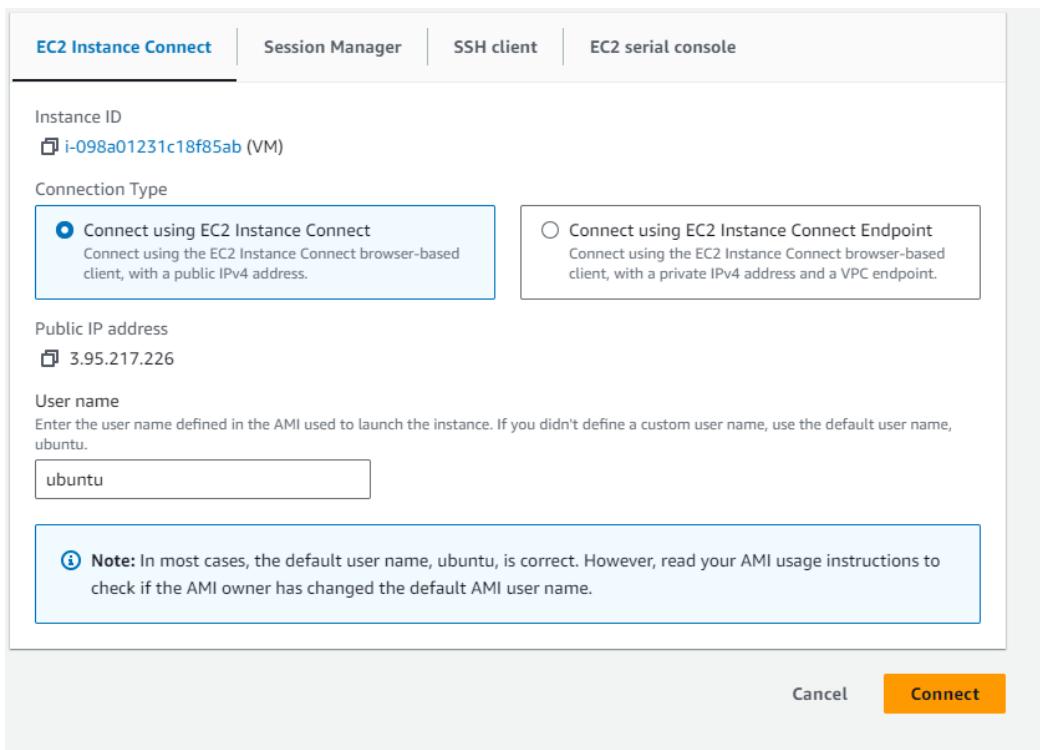
EC2 > Instances > i-098a01231c18f85ab

Instance summary for i-098a01231c18f85ab (VM) [Info](#)

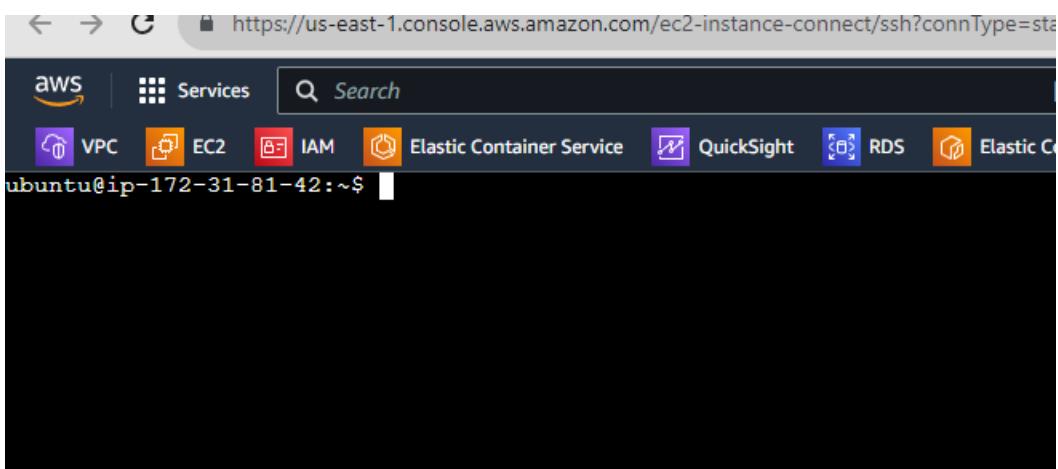
Updated less than a minute ago

Connect

Select the **Connect** button on the right.



Click the **Connect** button.



Install AWS CLI – Command line tools for working with AWS services, including Amazon EKS.

[Download the AWS CLI](#)

```
curl "https://awscli.amazonaws.com/awscli-exe-linux-x86_64.zip" -o "awscliv2.zip"
```

```
ubuntu@ip-172-31-81-42:~$ curl "https://awscli.amazonaws.com/awscli-exe-linux-x86_64.zip" -o "awscliv2.zip"
% Total    % Received % Xferd  Average Speed   Time     Time   Current
          Dload  Upload Total Spent   Left Speed
100 55.8M  100 55.8M    0     0  110M      0 --:--:-- --:--:-- 110M
```

Install the unzip tool

sudo apt install unzip

```
ubuntu@ip-172-31-81-42:~$ sudo apt install unzip
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
Suggested packages:
  zip
The following NEW packages will be installed:
  unzip
0 upgraded, 1 newly installed, 0 to remove and 0 not upgraded.
Need to get 174 kB of archives.
After this operation, 385 kB of additional disk space will be used.
Get:1 http://us-east-1.ec2.archive.ubuntu.com/ubuntu jammy-updates/main amd64 unzip amd64 6.0-26ubuntu3.1 [174 kB]
Fetched 174 kB in 0s (6351 kB/s)
Selecting previously unselected package unzip.
(Reading database ... 64295 files and directories currently installed.)
Preparing to unpack .../unzip_6.0-26ubuntu3.1_amd64.deb ...
Unpacking unzip (6.0-26ubuntu3.1) ...
Setting up unzip (6.0-26ubuntu3.1) ...
Processing triggers for man-db (2.10.2-1) ...
Scanning processes...
Scanning linux images...

Running kernel seems to be up-to-date.

No services need to be restarted.

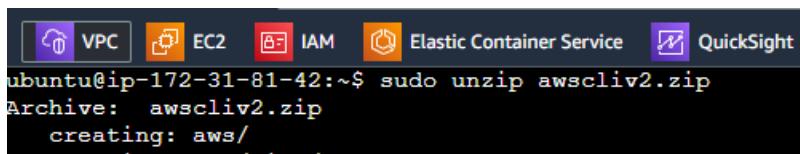
No containers need to be restarted.

No user sessions are running outdated binaries.

No VM guests are running outdated hypervisor (qemu) binaries on this host.
ubuntu@ip-172-31-81-42:~$
```

Unzip the AWS CLI

sudo unzip awscliv2.zip



Install the AWS CLI tool

sudo ./aws/install

```
ubuntu@ip-172-31-81-42:~$ sudo ./aws/install
You can now run: /usr/local/bin/aws --version
ubuntu@ip-172-31-81-42:~$ █
```

Check if the AWS CLI tool is installed

Check if the AWS CLI is installed

it should display the below output.

```
ubuntu@ip-172-31-81-42:~$ aws --version
aws-cli/2.13.22 Python/3.11.5 Linux/5.19.0-1025-aws exe/x86_64.ubuntu.22 prompt/off
ubuntu@ip-172-31-81-42:~$ 
```

Install eksctl – This is a command line tool for working with EKS clusters that automates many individual tasks.

The **eksctl** tool uses **CloudFormation** under the hood, creating one stack for the **EKS master control plane** and another stack for the **worker nodes**.

Download and extract the latest release of **eksctl** with the following command.

```
curl --silent --location
"https://github.com/weaveworks/eksctl/releases/latest/download/eksctl_$(uname -s)_amd64.tar.gz" | tar xz -C /tmp
```

```
ubuntu@ip-172-31-81-42:~$ curl --silent --location "https://github.com/weaveworks/eksctl/releases/latest/download/eksctl_$(uname -s)_amd64.tar.gz" | tar xz -C /tmp
ubuntu@ip-172-31-81-42:~$ 
```

Move the extracted binary to /usr/local/bin.

```
sudo mv /tmp/eksctl /usr/local/bin
```

```
ubuntu@ip-172-31-81-42:~$ sudo mv /tmp/eksctl /usr/local/bin
ubuntu@ip-172-31-81-42:~$ 
```

Check if the eksctl is running.

```
eksctl version
```

```
ubuntu@ip-172-31-81-42:~$ eksctl version
0.159.0
ubuntu@ip-172-31-81-42:~$ 
```

Install kubectl – This is a command line tool for working with Kubernetes clusters.

Kubernetes uses a command line utility called kubectl for communicating with the cluster API server. It is tool for controlling Kubernetes clusters. kubectl looks for a file named config in the \$HOME directory.

Downloading kubectl and placing it within a specific location.

```
sudo curl --silent --location -o /usr/local/bin/kubectl https://s3.us-west-2.amazonaws.com/amazon-eks/1.22.6/2022-03-09/bin/linux/amd64/kubectl
```

```
ubuntu@ip-172-31-81-42:~$ sudo curl --silent --location -o /usr/local/bin/kubectl https://s3.us-west-2.amazonaws.com/amazon-eks/1.22.6/2022-03-09/bin/linux/amd64/kub
ectl
ubuntu@ip-172-31-81-42:~$ █
```

Making the kubectl executable.

sudo chmod +x /usr/local/bin/kubectl

```
ubuntu@ip-172-31-81-42:~$ sudo chmod +x /usr/local/bin/kubectl
ubuntu@ip-172-31-81-42:~$ █
```

Verifying if kubectl got installed

kubectl version --short --client

```
ubuntu@ip-172-31-81-42:~$ kubectl version --short --client
Client Version: v1.22.6-eks-7d68063
ubuntu@ip-172-31-81-42:~$ █
```

NOTE: The **-short** flag doesn't work on all Linux distributions.

Install Helm 3

What is Helm and How to install Helm version 3?

Helm is a package manager for Kubernetes. Helm is the K8s equivalent of yum or apt. It accomplishes the same goals as Linux system package managers like APT or YUM: managing the installation of applications and dependencies behind the scenes and hiding the complexity from the user.

Why use Helm?

As the Kubernetes platform and ecosystem continued to expand, deploying one and only one Kubernetes configuration file (ie: a single YAML) was not the norm anymore. As number of K8S deployment files increased, how to manage those files? Helm solves that problem.

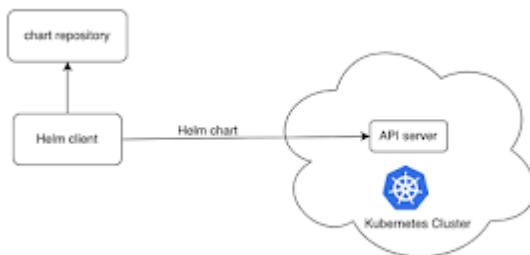
Helm Charts

Helm uses a packaging format called Charts. A Helm Chart is a collection of files that describe a set of Kubernetes resources. Helm Charts help you define, install, and upgrade even the most complex Kubernetes application. Charts are easy to create, version, share, and publish.

```
devopscaching@DevOpsCoachMBP springboot % tree mychart
mychart
└── Chart.yaml
    └── Charts
        └── templates
            ├── NOTES.txt
            ├── _helpers.tpl
            ├── deployment.yaml
            ├── hpa.yaml
            ├── ingress.yaml
            ├── service.yaml
            └── serviceaccount.yaml
            └── tests
                └── test-connection.yaml
    └── values.yaml

4 directories, 10 files
```

Helm Kubernetes Integration



In Helm 3 there is no tiller component (a feature of earlier versions of Helm). The Helm client directly interacts with the Kubernetes API for the helm chart deployment.

Helm 3 can be installed in many ways. We will, however, install Helm 3 using scripts option.

Download scripts

```
curl -fsSL -o get_helm.sh https://raw.githubusercontent.com/helm/helm/master/scripts/get-helm-3
```

```
ubuntu@ip-172-31-81-42:~$ curl -fsSL -o get_helm.sh https://raw.githubusercontent.com/helm/helm/master/scripts/get-helm-3
ubuntu@ip-172-31-81-42:~$
```

Provide permission

```
sudo chmod 700 get_helm.sh
```

```
ubuntu@ip-172-31-81-42:~$ sudo chmod 700 get_helm.sh
ubuntu@ip-172-31-81-42:~$
```

Execute script to install

```
sudo ./get_helm.sh
```

```
ubuntu@ip-172-31-81-42:~$ sudo ./get_helm.sh
Downloading https://get.helm.sh/helm-v3.13.0-linux-amd64.tar.gz
Verifying checksum... Done.
Preparing to install helm into /usr/local/bin
helm installed into /usr/local/bin/helm
ubuntu@ip-172-31-81-42:~$ █
```

Verify installation

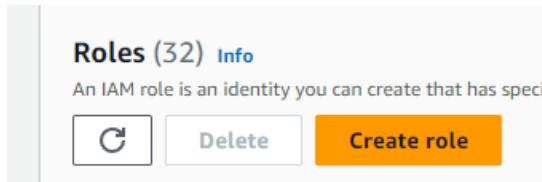
helm version --client

```
ubuntu@ip-172-31-81-42:~$ helm version --client
version.BuildInfo{Version:"v3.13.0", GitCommit:"825e86f6a7a38cef1112bfa606e4127a706749b1", GitTreeState:"clean", GoVersion:"go1.20.8"}
ubuntu@ip-172-31-81-42:~$ █
```

Create an IAM Role with Administrator Access

You need to create an **IAM** role with **AdministratorAccess** policy.

Go to AWS console, **IAM**, click on **Roles** and click **create role**



Select **AWS services**, Click **EC2**, Click on **Next permissions**.

Trusted entity type

- AWS service
Allow AWS services like EC2, Lambda, or others to perform actions in this account.
- AWS account
Allow entities in other AWS accounts belonging to you or a 3rd party to perform actions in this account.
- Web identity
Allows users federated by the specified external web identity provider to assume this role to perform actions in this account.
- SAML 2.0 federation
Allow users federated with SAML 2.0 from a corporate directory to perform actions in this account.
- Custom trust policy
Create a custom trust policy to enable others to perform actions in this account.

Use case
Allow an AWS service like EC2, Lambda, or others to perform actions in this account.

Service or use case

Choose a use case for the specified service.
Use case
 EC2
Allows EC2 instances to call AWS services on your behalf.

Now search for **AdministratorAccess** policy and click **Next**



Now give a role name and click **Create role**.

IAM > Roles > Create role

Step 1
[Select trusted entity](#)

Step 2
[Add permissions](#)

Step 3
[Name, review, and create](#)

Name, review, and create

Role details

Role name
Enter a meaningful name to identify this role.

Maximum 64 characters. Use alphanumeric and '+,-,@-' characters.

Description

Create role

Assign the role to EC2 instance

Go to **AWS** console, click on **EC2**, select the EC2 instance, Choose **Security**. Click on **Modify IAM role**.

The screenshot shows the AWS EC2 Instances page. There is one instance listed: 'VM' (Instance ID: i-098a01231c18f85ab, State: Running, Type: t2.micro). A context menu is open on the right side of the 'VM' row, with 'Modify IAM role' highlighted.

Choose the role you have created from the dropdown.

Select the role and click on **Apply**.

The screenshot shows the 'Modify IAM role' dialog. At the top, it shows the path: EC2 > Instances > i-098a01231c18f85ab > Modify IAM role. The main section is titled 'Modify IAM role' with a 'Info' link. It says 'Attach an IAM role to your instance.' Below that is a 'Instance ID' field containing 'i-098a01231c18f85ab (VM)'. Under 'IAM role', it says 'Select an IAM role to attach to your instance or create a new role if you haven't created any. The role you select replaces any roles that are currently attached to your instance.' A dropdown menu shows 'eks-admin-role' selected. To the right of the dropdown is a 'Create new IAM role' button. At the bottom are 'Cancel' and 'Update IAM role' buttons, with 'Update IAM role' being highlighted.

Instances > Attach/Replace IAM Role

Attach/Replace IAM Role

IAM role operation succeeded

Create EKS Cluster with two worker nodes using eksctl

```
eksctl create cluster --name demo-eks --region us-east-1 --nodegroup-name my-nodes --node-type t3.small --managed --nodes 2
```

the above command should create a EKS cluster in AWS, it might take 15 to 20 mins. The **eksctl** tool uses CloudFormation under the hood, creating one stack for the EKS master control plane and another stack for the worker nodes.

```
ubuntu@ip-172-31-81-42:~$ eksctl create cluster --name demo-eks --region us-east-1 --nodegroup-name my-nodes --node-type t3.small --managed --nodes 2
2023-09-28 08:53:06 [i] eksctl version 0.159.0
2023-09-28 08:53:06 [i] using region us-east-1
2023-09-28 08:53:06 [i] skipping us-east-1e from selection because it doesn't support the following instance type(s): t3.small
2023-09-28 08:53:06 [i] setting availability zones to [us-east-1d us-east-1c]
2023-09-28 08:53:06 [i] subnets for us-east-1d - public:192.168.0.0/19 private:192.168.64.0/19
2023-09-28 08:53:06 [i] subnets for us-east-1c - public:192.168.32.0/19 private:192.168.96.0/19
2023-09-28 08:53:06 [i] nodegroup "my-nodes" will use "" [AmazonLinux2/1.25]
2023-09-28 08:53:06 [i] using Kubernetes version 1.25
2023-09-28 08:53:06 [i] creating EKS cluster "demo-eks" in "us-east-1" region with managed nodes
2023-09-28 08:53:06 [i] will create 2 separate CloudFormation stacks for cluster itself and the initial managed nodegroup
2023-09-28 08:53:06 [i] if you encounter any issues, check CloudFormation console or try 'eksctl utils describe-stacks --region=us-east-1 --cluster=demo-eks'
2023-09-28 08:53:06 [i] Kubernetes API endpoint access will use default of (publicAccess=true, privateAccess=false) for cluster "demo-eks" in "us-east-1"
2023-09-28 08:53:06 [i] CloudWatch logging will not be enabled for cluster "demo-eks" in "us-east-1"
2023-09-28 08:53:06 [i] you can enable it with 'eksctl utils update-cluster-logging --enable-types=[SPECIFY-YOUR-LOG-TYPES-HERE] (e.g. all) --region=us-east-1 --cluster=demo-eks'
2023-09-28 08:53:06 [i]
2 sequential tasks: ( create cluster control plane "demo-eks",
  2 sequential sub-tasks: (
    wait for control plane to become ready,
    create managed nodegroup "my-nodes",
  )
)
2023-09-28 08:53:06 [i] building cluster stack "eksctl-demo-eks-cluster"
2023-09-28 08:53:06 [i] deploying stack "eksctl-demo-eks-cluster"
```

The screenshot shows the AWS CloudFormation console interface. On the left, there's a sidebar with options like 'Stacks', 'StackSets', 'Exports', 'Designer', 'Registry', 'Public extensions', and 'Activated extensions'. The main area is titled 'CloudFormation > Stacks' and shows a list of 'Stacks (2)'. The first stack, 'eksctl-demo-eks-cluster', has a status of 'CREATE_IN_PROGRESS' and was created on 2023-09-28 11:53:06 UTC+0300. Its description is 'EKS cluster [dedicated VPC: true, dedicated IAM: true] [created and managed by eksctl]'. The second stack, 'cfst-1449-3026bde225e07a2c93b07e789959fa23', has a status of 'CREATE_COMPLETE' and was created on 2023-09-28 10:44:09 UTC+0300.

This screenshot is from the same AWS CloudFormation console as the previous one, but it shows the state after the stacks have been fully created. Both 'eksctl-demo-eks-cluster' and 'cfst-1449-3026bde225e07a2c93b07e789959fa23' now have a status of 'CREATE_COMPLETE'.

The EKS cluster has been created.

```
eksctl get cluster --name demo-eks --region us-east-1
```

The above should confirm that EKS cluster is up and running.

```
ubuntu@ip-172-31-81-42:~$ eksctl get cluster --name demo-eks --region us-east-1
NAME      VERSION STATUS CREATED          VPC           SUBNETS
SECURITYGROUPS      PROVIDER
demo-eks      1.25   ACTIVE 2023-09-28T08:53:30Z  vpc-05fee4f22b2eb7f92  subnet-028e3389dd7a28274,subnet-059cb0cd9161466ae,subnet-0f5c03b32882d1e8f,subnet-0fd5d
41a07fbfc0de  sg-05c8b0d8016035fc7  EKS
ubuntu@ip-172-31-81-42:~$ █
```

Update Kube config by entering command below:

```
- aws eks update-kubeconfig --name demo-eks --region us-east-1
```

```
ubuntu@ip-172-31-81-42:~$ aws eks update-kubeconfig --name demo-eks --region us-east-1
Added new context arn:aws:eks:us-east-1:594182463744:cluster/demo-eks to /home/ubuntu/.kube/config
ubuntu@ip-172-31-81-42:~$ █
```

Connect to EKS cluster using kubectl commands

To view the list of worker nodes as part of EKS cluster.

kubectl get nodes

```
ubuntu@ip-172-31-81-42:~$ kubectl get nodes
NAME           STATUS    ROLES   AGE    VERSION
ip-192-168-28-73.ec2.internal  Ready    <none>  51s   v1.25.13-eks-43840fb
ip-192-168-52-210.ec2.internal  Ready    <none>  44s   v1.25.13-eks-43840fb
ubuntu@ip-172-31-81-42:~$ █
```

kubectl get ns

```
ubuntu@ip-172-31-81-42:~$ kubectl get ns
NAME        STATUS   AGE
default     Active   9m14s
kube-node-lease  Active   9m16s
kube-public   Active   9m16s
kube-system   Active   9m16s
ubuntu@ip-172-31-81-42:~$ █
```

[Deploy Nginx on a Kubernetes Cluster](#)

Let us run some apps to make sure they are deployed to Kubernetes cluster. The below command will create deployment:

```
kubectl create deployment nginx --image=nginx
```

```
ubuntu@ip-172-31-81-42:~$ kubectl create deployment nginx --image=nginx
deployment.apps/nginx created
ubuntu@ip-172-31-81-42:~$ █
```

[View Deployments](#)

```
kubectl get deployments
```

```
ubuntu@ip-172-31-81-42:~$ kubectl get deployments
NAME      READY   UP-TO-DATE   AVAILABLE   AGE
nginx    1/1     1           1           45s
ubuntu@ip-172-31-81-42:~$ █
```

[Errors during Cluster creation](#)

If you are having issues when creating a cluster, try to delete the cluster by executing the below command and re-create it.

```
eksctl delete cluster --name demo-eks --region us-east-1
```

```
[1-29 22:23:00] [E] An error(s) occurred and cluster hasn't been created properly, you may wish to check CloudFormation console
[1-29 22:23:00] [E] to cleanup resources, run "eksctl delete cluster --region=us-east-1 --name=demo-eks"
[1-29 22:23:00] [E] Creating CloudFormation stack "eksctl-demo-eks-cluster". AlreadyByTimestamp: Stack [eksctl-demo-eks-cluster] already exists under [AWS::CloudFormation::Stack], reason: ID: 6043851-4e40-470e-9225-8c330c2039
[1-29 22:23:00] [E] Failed to create Cluster "demo-eks"
```

or Login to **AWS console** --> **AWS Cloud formation** --> **delete the stack manually**.

You can also delete the cluster under AWS console --> Elastic Kubernetes Service --> Clusters

Click on Delete cluster

[Implementation steps](#)

We need to add the Helm Stable Charts for your local client. Execute the below command:

```
helm repo add stable https://charts.helm.sh/stable
```

```
ubuntu@ip-172-31-81-42:~$ helm repo add stable https://charts.helm.sh/stable
"stable" has been added to your repositories
ubuntu@ip-172-31-81-42:~$ █
```

[Add prometheus Helm repo](#)

```
helm repo add prometheus-community https://prometheus-community.github.io/helm-charts
```

```
ubuntu@ip-172-31-81-42:~$ helm repo add prometheus-community https://prometheus-community.github.io/helm-charts
"prometheus-community" has been added to your repositories
ubuntu@ip-172-31-81-42:~$
```

helm search repo prometheus-community

Prometheus and grafana helm chart moved to kube prometheus stack

| NAME | CHART VERSION | APP VERSION | DESCRIPTION |
|--|---------------|-------------|--|
| prometheus-community/alertmanager | 1.7.0 | v0.26.0 | The Alertmanager handles alerts sent by client ... |
| prometheus-community/alertmanager-snmp-notifier | 0.1.2 | v1.4.0 | The SNMP Notifier handles alerts coming from Pr... |
| prometheus-community/jiralert | 1.6.0 | v1.3.0 | A Helm chart for Kubernetes to install jiralert... |
| prometheus-community/kube-prometheus-stack | 51.2.0 | v0.68.0 | kube-prometheus-stack collects Kubernetes manif... |
| prometheus-community/kube-state-metrics | 5.13.0 | v2.10.0 | Install kube-state-metrics to generate and expo... |
| prometheus-community/prom-label-proxy | 0.6.0 | v0.7.0 | A proxy that enforces a given label in a given ... |
| prometheus-community/prometheus | 25.0.0 | v2.47.0 | Prometheus is a monitoring system and time seri... |
| prometheus-community/prometheus-adapter | 4.5.0 | v0.11.1 | A Helm chart for k8s prometheus adapter |
| prometheus-community/prometheus-blackbox-exporter | 8.4.0 | v0.24.0 | Prometheus Blackbox Exporter |
| prometheus-community/prometheus-cloudwatch-expo... | 0.25.2 | v0.15.4 | A Helm chart for prometheus cloudwatch-exporter |
| prometheus-community/prometheus-contrack-stats... | 0.5.7 | v0.4.15 | A Helm chart for contrack-stats-exporter |
| prometheus-community/prometheus-consul-exporter | 1.0.0 | v0.4.0 | A Helm chart for the Prometheus Consul Exporter |
| prometheus-community/prometheus-couchdb-exporter | 1.0.0 | v0.1.0 | A Helm chart to export the metrics from couchdb... |
| prometheus-community/prometheus-druid-exporter | 1.1.0 | v0.11.0 | Druid exporter to monitor druid metrics with Pr... |
| prometheus-community/prometheus-elasticsearch-e... | 5.3.0 | v1.6.0 | Elasticsearch stats exporter for Prometheus |
| prometheus-community/prometheus-fastly-exporter | 0.2.0 | v2.2.4 | A Helm chart for the Prometheus Fastly Exporter |
| prometheus-community/prometheus-ipmi-exporter | 0.1.0 | v0.1.6.1 | This is an IPMI exporter for Prometheus. |
| prometheus-community/prometheus-json-exporter | 0.7.1 | v0.5.0 | Install prometheus-json-exporter |
| prometheus-community/prometheus-kafka-exporter | 2.6.0 | v1.7.0 | A Helm chart to export the metrics from Kafka i... |
| prometheus-community/prometheus-modbus-exporter | 0.1.0 | v0.4.0 | A Helm chart for prometheus-modbus-exporter |
| prometheus-community/prometheus-mongodb-exporter | 3.4.0 | v0.39.0 | A Prometheus exporter for MongoDB metrics |
| prometheus-community/prometheus-mysql-exporter | 2.0.0 | v0.15.0 | A Helm chart for prometheus mysql exporter with... |
| prometheus-community/prometheus-nats-exporter | 2.13.0 | v0.12.0 | A Helm chart for prometheus-nats-exporter |
| prometheus-community/prometheus-nginx-exporter | 0.1.1 | v0.11.0 | A Helm chart for the Prometheus NGINX Exporter |
| prometheus-community/prometheus-node-exporter | 4.23.1 | v1.6.1 | A Helm chart for prometheus node-exporter |
| prometheus-community/prometheus-operator | 9.3.2 | v0.38.1 | DEPRECATED - This chart will be renamed. See ht... |
| prometheus-community/prometheus-operator-admiss... | 0.7.0 | v0.68.0 | Prometheus Operator Admission Webhook |
| prometheus-community/prometheus-operator-crds | 6.0.0 | v0.68.0 | A Helm chart that collects custom resource defi... |
| prometheus-community/prometheus-pgbouncer-exporter | 0.1.1 | v1.18.0 | A Helm chart for prometheus pgbouncer-exporter |
| prometheus-community/prometheus-pingdom-exporter | 2.5.0 | v20190610-1 | A Helm chart for Prometheus Pingdom Exporter |
| prometheus-community/prometheus-pingmesh-exporter | 0.3.0 | v1.1.0 | Prometheus Pingmesh Exporter |
| prometheus-community/prometheus-postgres-exporter | 5.1.0 | v0.14.0 | A Helm chart for prometheus postgres-exporter |
| prometheus-community/prometheus-pushgateway | 2.4.1 | v1.6.1 | A Helm chart for prometheus pushgateway |
| prometheus-community/prometheus-rabbitmq-exporter | 1.8.1 | v0.29.0 | Rabbitmq metrics exporter for prometheus |
| prometheus-community/prometheus-redis-exporter | 6.0.0 | v1.54.0 | Prometheus exporter for Redis metrics |
| prometheus-community/prometheus-smartctl-exporter | 0.6.0 | v0.11.0 | A Helm chart for Kubernetes |

Now create the Prometheus namespace

kubectl create namespace prometheus

```
ubuntu@ip-172-31-81-42:~$ kubectl create namespace prometheus
namespace/prometheus created
ubuntu@ip-172-31-81-42:~$
```

Install kube-prometheus-stack

Below is helm command to install kube-prometheus-stack. The helm repo kube-stack-prometheus (formerly prometheus-operator) comes with a grafana deployment embedded.

helm install stable prometheus-community/kube-prometheus-stack -n prometheus

```
ubuntu@ip-172-31-81-42:~$ helm install stable prometheus-community/kube-prometheus-stack -n prometheus
NAME: stable
LAST DEPLOYED: Thu Sep 28 09:15:30 2023
NAMESPACE: prometheus
STATUS: deployed
REVISION: 1
NOTES:
kube-prometheus-stack has been installed. Check its status by running:
  kubectl --namespace prometheus get pods -l "release=stable"

Visit https://github.com/prometheus-operator/kube-prometheus for instructions on how to create & configure Alertmanager and Prometheus instances using the Operator.
ubuntu@ip-172-31-81-42:~$
```

Lets check if prometheus and grafana pods are running already

kubectl get pods -n prometheus

```
ubuntu@ip-172-31-81-42:~$ kubectl get pods -n prometheus
NAME                               READY   STATUS    RESTARTS   AGE
alertmanager-stable-kube-prometheus-sta-alertmanager-0   2/2     Running   0          49s
prometheus-stable-kube-prometheus-sta-prometheus-0       2/2     Running   0          49s
stable-grafana-6c8fd88f56-zrcpr                         3/3     Running   0          54s
stable-kube-prometheus-sta-operator-75947cf97-xjmmv       1/1     Running   0          54s
stable-kube-state-metrics-7ccdc6767c-xj26t              1/1     Running   0          54s
stable-prometheus-node-exporter-w646x                   1/1     Running   0          54s
stable-prometheus-node-exporter-x6w9w                     1/1     Running   0          54s
ubuntu@ip-172-31-81-42:~$
```

kubectl get svc -n prometheus

```
ubuntu@ip-172-31-81-42:~$ kubectl get svc -n prometheus
NAME           TYPE      CLUSTER-IP   EXTERNAL-IP  PORT(S)        AGE
alertmanager-operated  ClusterIP  None         <none>       9093/TCP,9094/TCP,9094/UDP  90s
prometheus-operated  ClusterIP  None         <none>       9090/TCP     90s
stable-grafana      ClusterIP  10.100.186.94  <none>       80/TCP        95s
stable-kube-prometheus-sta-alertmanager  ClusterIP  10.100.254.197  <none>       9093/TCP,8080/TCP  95s
stable-kube-prometheus-sta-operator       ClusterIP  10.100.85.234  <none>       443/TCP       95s
stable-kube-prometheus-sta-prometheus   ClusterIP  10.100.229.18  <none>       9090/TCP,8080/TCP  95s
stable-kube-state-metrics               ClusterIP  10.100.245.2   <none>       8080/TCP     95s
stable-prometheus-node-exporter        ClusterIP  10.100.247.199  <none>       9100/TCP     95s
ubuntu@ip-172-31-81-42:~$
```

This confirms that prometheus and grafana have been installed successfully using Helm.

In order to make prometheus and grafana available outside the cluster, we need to use **LoadBalancer** or **NodePort** instead of **ClusterIP** within the service files.

Edit Prometheus Service

```
kubectl edit svc stable-kube-prometheus-sta-prometheus -n prometheus
```

```
ipFamilyPolicy: SingleStack
ports:
- name: http-web
  port: 9090
  protocol: TCP
  targetPort: 9090
- appProtocol: http
  name: reloader-web
  port: 8080
  protocol: TCP
  targetPort: reloader-web
selector:
  app.kubernetes.io/name: prometheus
  operator.prometheus.io/name: stable-kube-prometheus-sta-prometheus
sessionAffinity: None
status:
  loadBalancer: {}
```

Edit Grafana Service

```
kubectl edit svc stable-grafana -n prometheus
```

```
ipFamilyPolicy: SingleStack
ports:
- name: http-web
  port: 9090
  protocol: TCP
  targetPort: 9090
- appProtocol: http
  name: reloader-web
  port: 8080
  protocol: TCP
  targetPort: reloader-web
selector:
  app.kubernetes.io/name: prometheus
  operator.prometheus.io/name: stable-kube-prometheus-sta-prometheus
sessionAffinity: None
type: LoadBalancer
status:
  loadBalancer: {}
```

Verify if service is changed to **LoadBalancer** and also to get the **Load Balancer URL** for each application.

```
kubectl get svc -n prometheus
```

```
ubuntu@ip-172-31-81-42:~$ kubectl edit svc stable-grafana -n prometheus
service/stable-grafana edited
ubuntu@ip-172-31-81-42:~$ kubectl get svc -n prometheus
NAME           TYPE      CLUSTER-IP   EXTERNAL-IP          PORT(S)
AGE
alertmanager-operated   ClusterIP  None         <none>            9093/TCP,9094/TCP,90
94/UDP   13m
prometheus-operated    ClusterIP  None         <none>            9090/TCP
13m
stable-grafana        LoadBalancer 10.100.186.94  ad1a1ea7902e441d797edacd585fa999-733838404.us-east-1.elb.amazonaws.com  80:30905/TCP
13m
stable-kube-prometheus-sta-alertmanager ClusterIP 10.100.254.197 <none>            9093/TCP,8080/TCP
13m
stable-kube-prometheus-sta-operator       ClusterIP 10.100.85.234  <none>            443/TCP
13m
stable-kube-prometheus-sta-prometheus   LoadBalancer 10.100.229.18  a8ee684ef50a04fa99b6fa68a6a549e7-907360140.us-east-1.elb.amazonaws.com  9090:30489/TCP,8080:
30104/TCP 13m
stable-kube-state-metrics     ClusterIP 10.100.245.2   <none>            8080/TCP
13m
stable-prometheus-node-exporter ClusterIP 10.100.247.199 <none>            9100/TCP
13m
ubuntu@ip-172-31-81-42:~$ []
```

The screenshot shows the Prometheus web interface. At the top, there's a navigation bar with links for Prometheus, Alerts, Graph, Status, and Help. Below the navigation is a toolbar with checkboxes for "Use local time", "Enable query history", "Enable autocomplete" (which is checked), "Enable highlighting" (which is checked), and "Enable linter" (which is checked). A search bar contains the placeholder "Expression (press Shift+Enter for newlines)". Below the search bar are two tabs: "Table" (selected) and "Graph". Underneath the tabs is a "Evaluation time" selector with arrows for navigating between time ranges. A message "No data queried yet" is displayed below the evaluation time selector. At the bottom left is a blue "Add Panel" button.

Access Grafana UI in the browser

Get the URL after you run **kubectl get svc -n Prometheus** and place it in the browser.

The screenshot shows the Grafana login page. The address bar indicates the URL is "http://ad1a1ea7902e441d797edacd585fa999-733838404.us-east-1.elb.amazonaws.com/login". The main content features the Grafana logo (a stylized orange and yellow sun-like icon) and the text "Welcome to Grafana". Below this are two input fields: "Email or username" containing "email or username" and "Password" containing "password". To the right of the password field is a small eye icon for password visibility. At the bottom is a large blue "Log in" button. Below the "Log in" button is a link "Forgot your password?".

You should see the above.

You can now login with the following details:

UserName: admin

Password: prom-operator

The screenshot shows the Grafana home page. At the top, it says "Welcome to Grafana". Below that, there's a "Basic" section with a "TUTORIAL" card titled "DATA SOURCE AND DASHBOARDS" and "Grafana fundamentals". The card describes setting up Grafana and covers "Data source" and "Dashboards" steps. To the right, there are two "COMPLETE" cards: "Add your first data source" (with a database icon) and "Create your first dashboard" (with a grid icon). A sidebar on the left lists "Dashboards", "Starred dashboards", and "Recently viewed dashboards". On the right, there's a "Latest from the blog" section with a post about "client_java 1.0.0" dated Sep 27.

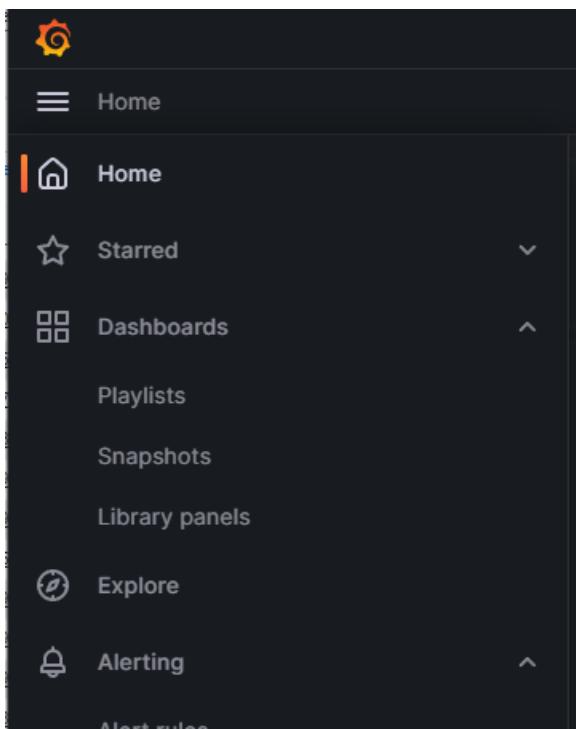
Create Dashboard in Grafana

In Grafana, we can create of various kinds of dashboards according to our needs.

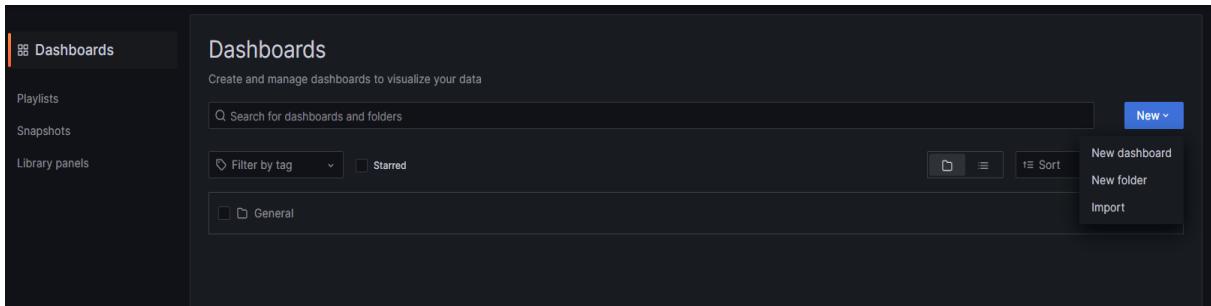
How to Create Kubernetes Monitoring Dashboard?

For creating a dashboard to monitor the cluster:

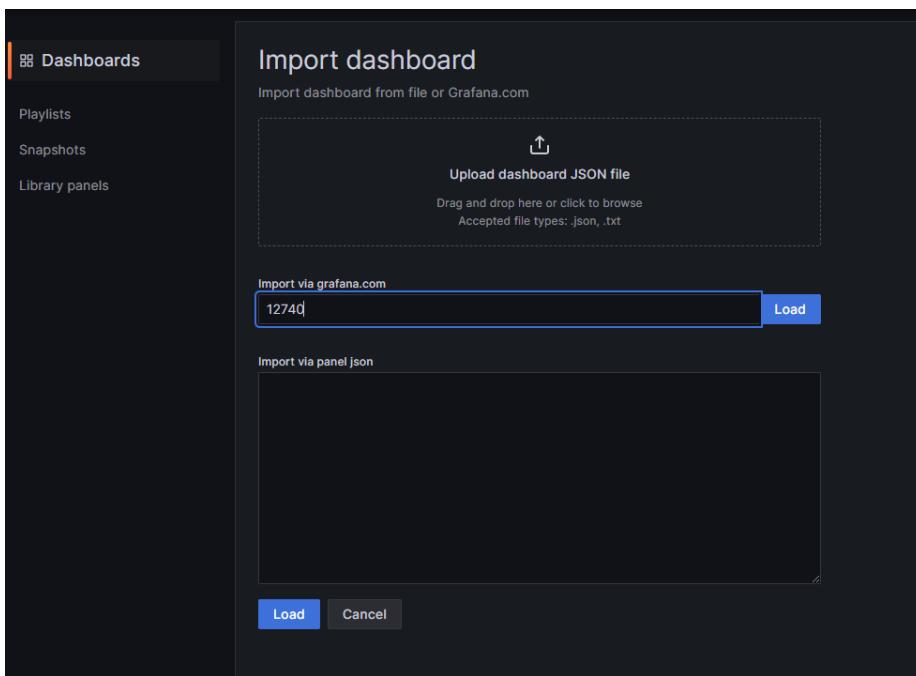
Click the hamburger at the top left than select **Dashboard**



Then on the right select **New > 'Import'**.

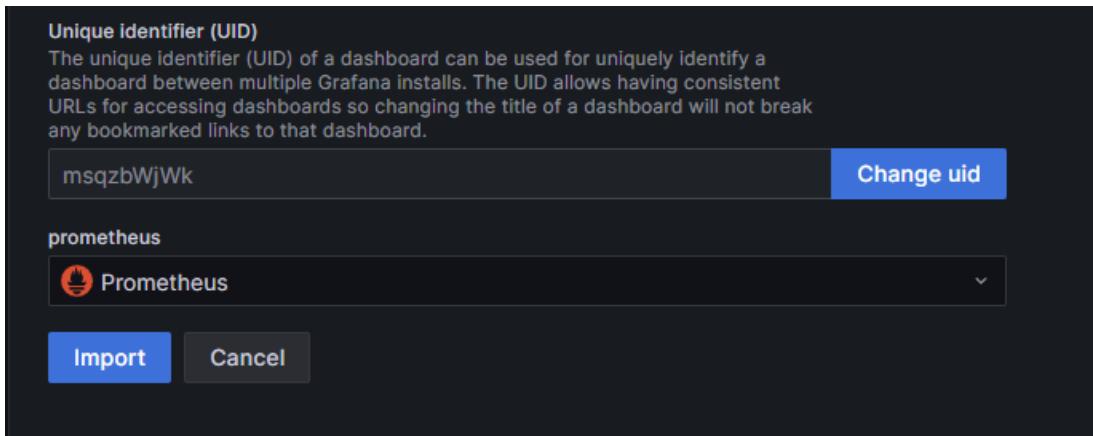


Enter **12740** dashboard id under [Grafana.com Dashboard](#).



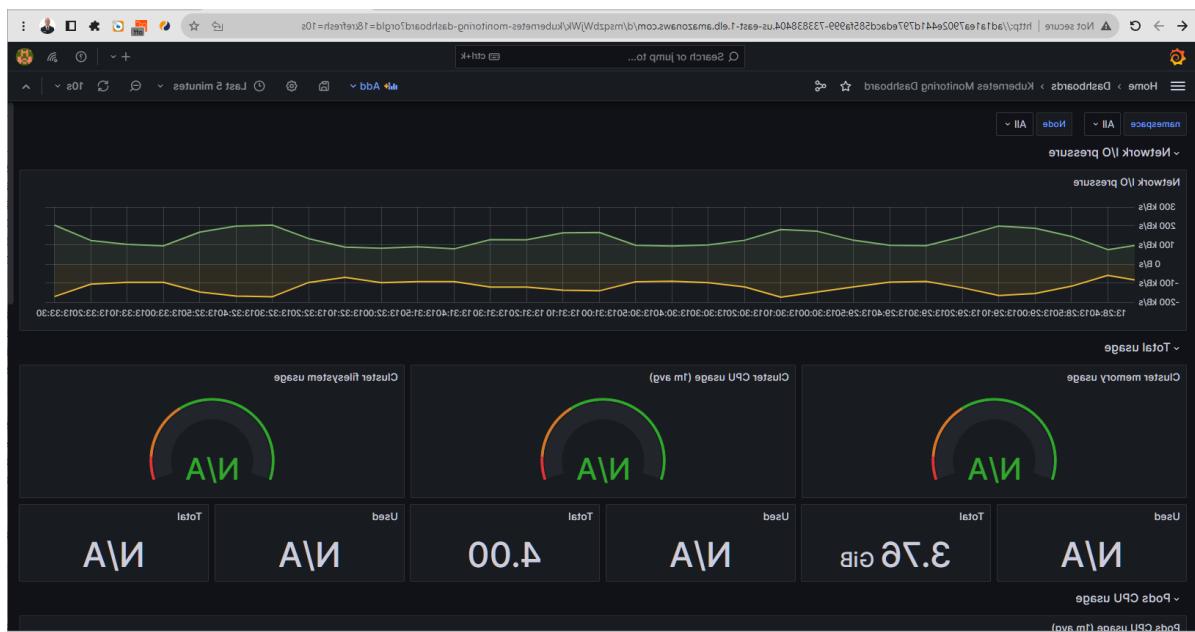
Click 'Load'.

Select 'Prometheus' as the endpoint under Prometheus data sources drop down.



Click 'Import'.

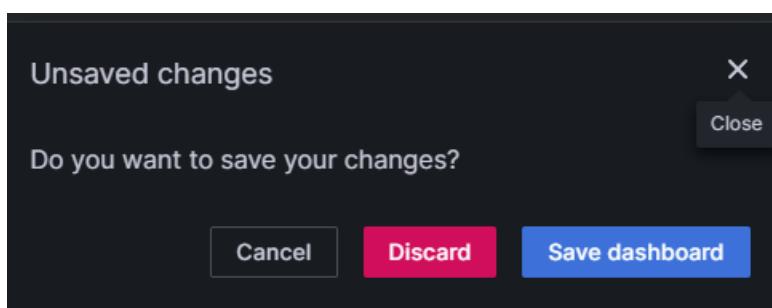
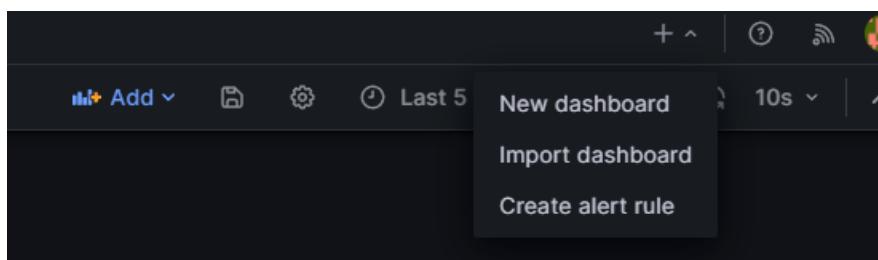
This will show the monitoring dashboard for all cluster nodes.



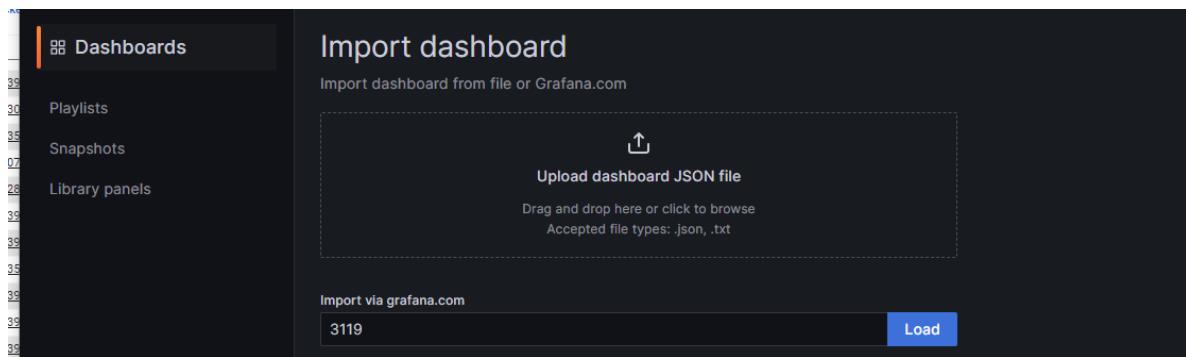
How to Create Kubernetes Cluster Monitoring Dashboard?

For creating a dashboard to monitor the cluster:

Click '+' button on the top right and select 'Import dashboard'.

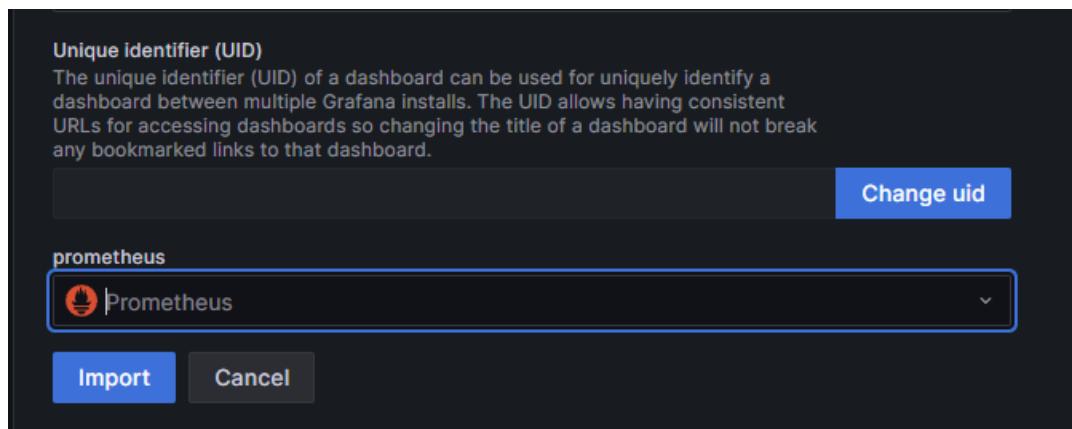


Enter **3119** dashboard id under Grafana.com Dashboard.



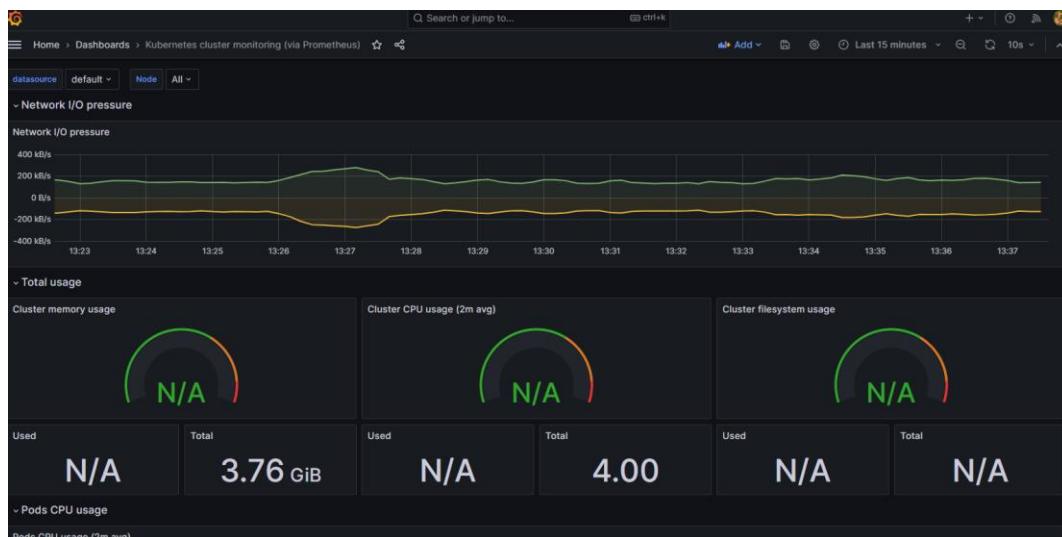
Click 'Load'.

Select 'Prometheus' as the endpoint under Prometheus data sources drop down.



Click 'Import'.

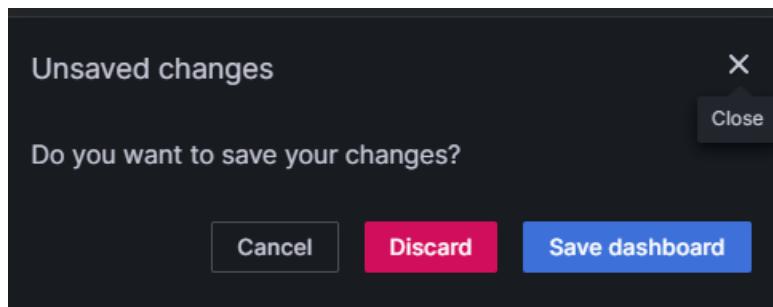
This will show a monitoring dashboard for all cluster nodes.



How to create a POD Monitoring Dashboard

For creating a dashboard to monitor the cluster:

Click '+' button on the top right panel and select 'Import dashboard'.

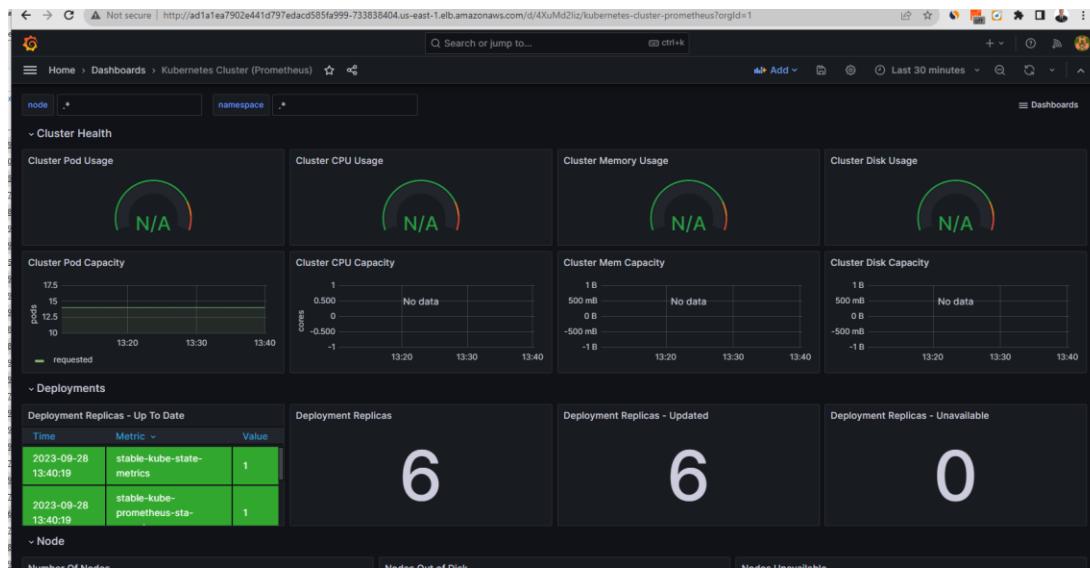


Enter **6417** dashboard id under Grafana.com **Dashboard**.

Click 'Load'.

Select '**Prometheus**' as the endpoint under Prometheus data sources drop down.
Click 'Import'.

This will show a monitoring dashboard for all cluster nodes.



Cleanup EKS Cluster

Use the below command to delete EKS cluster to avoid being charged by AWS.

eksctl delete cluster --name demo-eks --region us-east-1

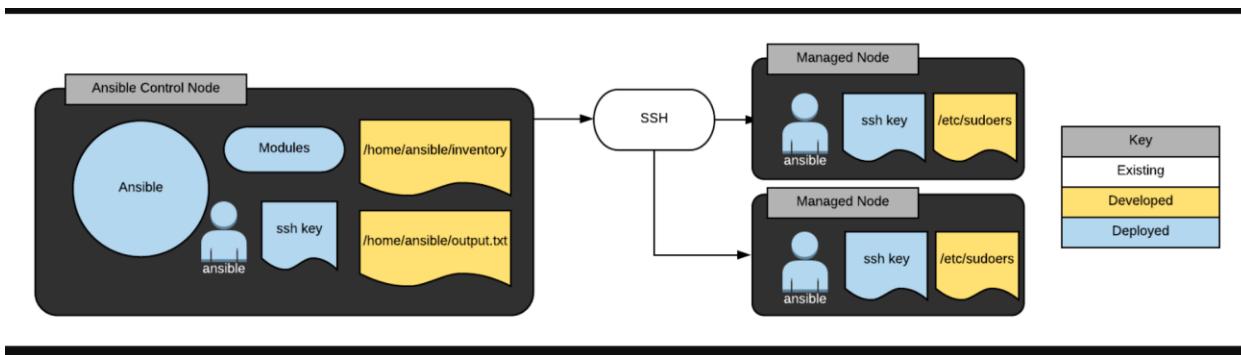
Alternatively, login to **AWS console --> AWS Cloudformation --> delete the stack manually**.

The screenshot shows the AWS CloudFormation console with the URL <https://us-east-1.console.aws.amazon.com/cloudformation/home?region=us-east-1#/stacks?filteringText=&filteringStatus=active&viewNested=true>. The left sidebar shows 'CloudFormation' with 'Stacks' selected. The main area displays 'Stacks (3)' with a table:

| Stack name | Status | Created time | Description |
|--|-----------------|------------------------------|--|
| eksctl-demo-eks-nodegroup-my-nodes | CREATE_COMPLETE | 2023-09-28 12:05:08 UTC+0300 | EKS Managed Nodes (SSH access: false) [created by eksctl] |
| eksctl-demo-eks-cluster | CREATE_COMPLETE | 2023-09-28 11:53:06 UTC+0300 | EKS cluster (dedicated VPC: true, dedicated IAM: true) [created and managed by eksctl] |
| cfst-1449-3026bde225e07a2c93b07e789959fa23 | CREATE_COMPLETE | 2023-09-28 10:44:09 UTC+0300 | - |

You can also delete the cluster under **AWS console --> Elastic Kubernetes Service --> Clusters**
Click on Delete cluster.

Getting Started with Ansible



Install Ansible on the Control Node:

Log in to the control node using ssh, cloud_user, and the provided public IP address and password:

```
ssh cloud_user@<PUBLIC IP>
```

To install Ansible on the control node:

```
sudo yum install ansible
```

```
[cloud_user@control1 ~]$ sudo yum install ansible
```

```
Installed:
  ansible.noarch 0:2.9.27-1.el7
```

```
Dependency Installed:
  python-cffi.x86_64 0:1.6.0-5.el7      python-enum34.noarch 0:1.0.4-1.el7      python-idna.noarch 0:2.4-1.el7      python-paramiko.noarch 0:2.1.1-9.el7
  python-ply.noarch 0:3.4-11.el7        python-pycparser.noarch 0:2.14-1.el7      python2-cryptography.x86_64 0:1.7.2-2.el7    python2-httplib2.noarch 0:0.18.1-3.el7
  python2-jmespath.noarch 0:0.9.4-2.el7  python2-pyasn1.noarch 0:0.1.9-7.el7      sshpass.x86_64 0:1.06-1.el7
```

```
Complete!
```

```
[cloud_user@control1 ~]$
```

Configure the ansible User on the Control Node:

Next, we'll configure the ansible user on the control node for ssh shared key access to managed nodes.

sudo su - ansible

```
[cloud_user@control1 ~]$ sudo su - ansible
[sudo] password for cloud_user:
[ansible@control1 ~]$ █
```

ssh-keygen

```
[ansible@control1 ~]$ ssh-keygen
Generating public/private rsa key pair.
Enter file in which to save the key (/home/ansible/.ssh/id_rsa):
Created directory '/home/ansible/.ssh'.
Enter passphrase (empty for no passphrase):
Enter same passphrase again:
Your identification has been saved in /home/ansible/.ssh/id_rsa.
Your public key has been saved in /home/ansible/.ssh/id_rsa.pub.
The key fingerprint is:
SHA256:hxr0J7gxHByy4BkPu3/DRKwQUNsioR6LNGEwgkGI2+Q ansible@control1
The key's randomart image is:
-----[RSA 2048]----+
|&B* . .
|B+o@ = .
|oOB + *
|=E+ = + .
|.o. . B S o
| . o * +
| . *
| ..
|
-----[SHA256]----+
[ansible@control1 ~]$ █
```

Note: Do not use a passphrase for the key pair.

Create a key pair for the ansible user on the control host, accepting the defaults when prompted:

Copy the public key to both node1 and node2:

ssh-copy-id node1

```
[ansible@control1 ~]$ ssh-copy-id node1
/bin/ssh-copy-id: INFO: Source of key(s) to be installed: "/home/ansible/.ssh/id_rsa.pub"
The authenticity of host 'node1 (10.0.1.85)' can't be established.
ECDSA key fingerprint is SHA256:kk5TX0feYjKZFmj56uyfNhkXTKQjMhVGgC4Lpsoek2k.
ECDSA key fingerprint is MD5:40:f0:88:b0:cb:e1:4d:66:62:ee:9f:97:fd:8a:1d:2c.
Are you sure you want to continue connecting (yes/no)? y
Please type 'yes' or 'no': y
Please type 'yes' or 'no': yes
/bin/ssh-copy-id: INFO: attempting to log in with the new key(s), to filter out any that are already installed
/bin/ssh-copy-id: INFO: 1 key(s) remain to be installed -- if you are prompted now it is to install the new keys
Password:

Number of key(s) added: 1

Now try logging into the machine, with: "ssh 'node1'"
and check to make sure that only the key(s) you wanted were added.

[ansible@control1 ~]$ ssh node1
[ansible@node1 ~]$ exit
logout
Connection to node1 closed.
[ansible@control1 ~]$
```

ssh-copy-id node2

```
[ansible@control1 ~]$ ssh-copy-id node2
/bin/ssh-copy-id: INFO: Source of key(s) to be installed: "/home/ansible/.ssh/id_rsa.pub"
The authenticity of host 'node2 (10.0.1.16)' can't be established.
ECDSA key fingerprint is SHA256:U4aqdBHcqJtTEsvVgUMsXI+hi4Z22Fkzp+1mUKuby4w.
ECDSA key fingerprint is MD5:ed:c2:60:f5:20:83:58:1b:76:86:15:29:3a:82:f1:eb.
Are you sure you want to continue connecting (yes/no)? yes
/bin/ssh-copy-id: INFO: attempting to log in with the new key(s), to filter out any that are already installed
/bin/ssh-copy-id: INFO: 1 key(s) remain to be installed -- if you are prompted now it is to install the new keys
Password:

Number of key(s) added: 1

Now try logging into the machine, with: "ssh 'node2'"
and check to make sure that only the key(s) you wanted were added.

[ansible@control1 ~]$ ssh node2
[ansible@node2 ~]$ exit
logout
Connection to node2 closed.
[ansible@control1 ~]$
```

Create a Simple Ansible Inventory

Next, we'll create a simple Ansible inventory on the control node in /home/ansible/inventory containing node1 and node2.

On the control host:

```
sudo su - ansible
touch /home/ansible/inventory
echo "node1" >> /home/ansible/inventory
```

```
echo "node2" >> /home/ansible/inventory
```

```
[ansible@control1 ~]$ touch /home/ansible/inventory
[ansible@control1 ~]$ echo "node1" >> /home/ansible/inventory
[ansible@control1 ~]$ echo "node2" >> /home/ansible/inventory
[ansible@control1 ~]$ 
```

Configure sudo Access for Ansible:

Now, we'll configure sudo access for Ansible on node1 and node2 such that Ansible may use sudo for any command with no password prompt.

Log in to node1 as cloud_user and edit the sudoers file to contain appropriate access for the ansible user:

```
[ansible@control1 ~]$ ssh cloud_user@node1
Password:
[cloud_user@node1 ~]$ sudo visudo
[sudo] password for cloud_user: 
```

ssh cloud_user@node1

sudo visudo

Add the following line to the file and save:

ansible ALL=(ALL) NOPASSWD: ALL

```
## Same thing without a password
# %wheel      ALL=(ALL)      NOPASSWD: ALL
ansible    ALL=(ALL)      NOPASSWD: ALL

## Allows members of the users group to mount and unmount the
## cdrom as root
# %users  ALL=/sbin/mount /mnt/cdrom, /sbin/umount /mnt/cdrom

## Allows members of the users group to shutdown this system
# %users  localhost=/sbin/shutdown -h now

## Read drop-in files from /etc/sudoers.d (the # here does not mean a comment)
#includeinclude /etc/sudoers.d
-- INSERT --
```

```
[ansible@control1 ~]$ ssh cloud_user@node1
Password:
[cloud_user@node1 ~]$ sudo visudo
[sudo] password for cloud_user:
[cloud_user@node1 ~]$ exit
logout
Connection to node1 closed.
[ansible@control1 ~]$ ssh cloud_user@node2
Password:
[cloud_user@node2 ~]$ sudo visudo
[sudo] password for cloud_user:
[cloud_user@node2 ~]$ logout
Connection to node2 closed.
[ansible@control1 ~]$ ]
```

Enter: **Logout**

Repeat these steps for node2, and then back out to the control node.

Verify Each Managed Node Is Accessible

Finally, we'll verify each managed node is able to be accessed by Ansible from the control node using the ping module.

Redirect the output of a successful command to /home/ansible/output.

To verify each node, run the following as the ansible user from the control host:

```
ansible -i /home/ansible/inventory node1 -m ping
```

```
[ansible@control1 ~]$ ansible -i /home/ansible/inventory node1 -m ping
node1 | SUCCESS => {
    "ansible_facts": {
        "discovered_interpreter_python": "/usr/bin/python"
    },
    "changed": false,
    "ping": "pong"
}
[ansible@control1 ~]$ █
```

```
ansible -i /home/ansible/inventory node2 -m ping
```

```
[ansible@control1 ~]$ ansible -i /home/ansible/inventory node2 -m ping
node2 | SUCCESS => {
    "ansible_facts": {
        "discovered_interpreter_python": "/usr/bin/python"
    },
    "changed": false,
    "ping": "pong"
}
[ansible@control1 ~]$ █
```

To redirect output of a successful command to /home/ansible/output:

```
ansible -i /home/ansible/inventory node1 -m ping > /home/ansible/output
```

```
[ansible@control1 ~]$ ansible -i /home/ansible/inventory node1 -m ping > /home/ansible/output
[ansible@control1 ~]$ █
```

```
[ansible@control1 ~]$ cat /home/ansible/output
node1 | SUCCESS => {
    "ansible_facts": {
        "discovered_interpreter_python": "/usr/bin/python"
    },
    "changed": false,
    "ping": "pong"
}
[ansible@control1 ~]$
```

cat /home/ansible/output

Addendum

cat /etc/os-release

```
[cloud_user@control1 ~]$ cat /etc/os-release
NAME="Red Hat Enterprise Linux Server"
VERSION="7.9 (Maipo)"
ID="rhel"
ID_LIKE="fedora"
VARIANT="Server"
VARIANT_ID="server"
VERSION_ID="7.9"
PRETTY_NAME="Red Hat Enterprise Linux Server 7.9 (Maipo)"
ANSI_COLOR="0;31"
CPE_NAME="cpe:/o:redhat:enterprise_linux:7.9:GA:server"
HOME_URL="https://www.redhat.com/"
BUG_REPORT_URL="https://bugzilla.redhat.com/"

REDHAT_BUGZILLA_PRODUCT="Red Hat Enterprise Linux 7"
REDHAT_BUGZILLA_PRODUCT_VERSION=7.9
REDHAT_SUPPORT_PRODUCT="Red Hat Enterprise Linux"
REDHAT_SUPPORT_PRODUCT_VERSION="7.9"
[cloud_user@control1 ~]$
```

```
cat /etc/hosts
```

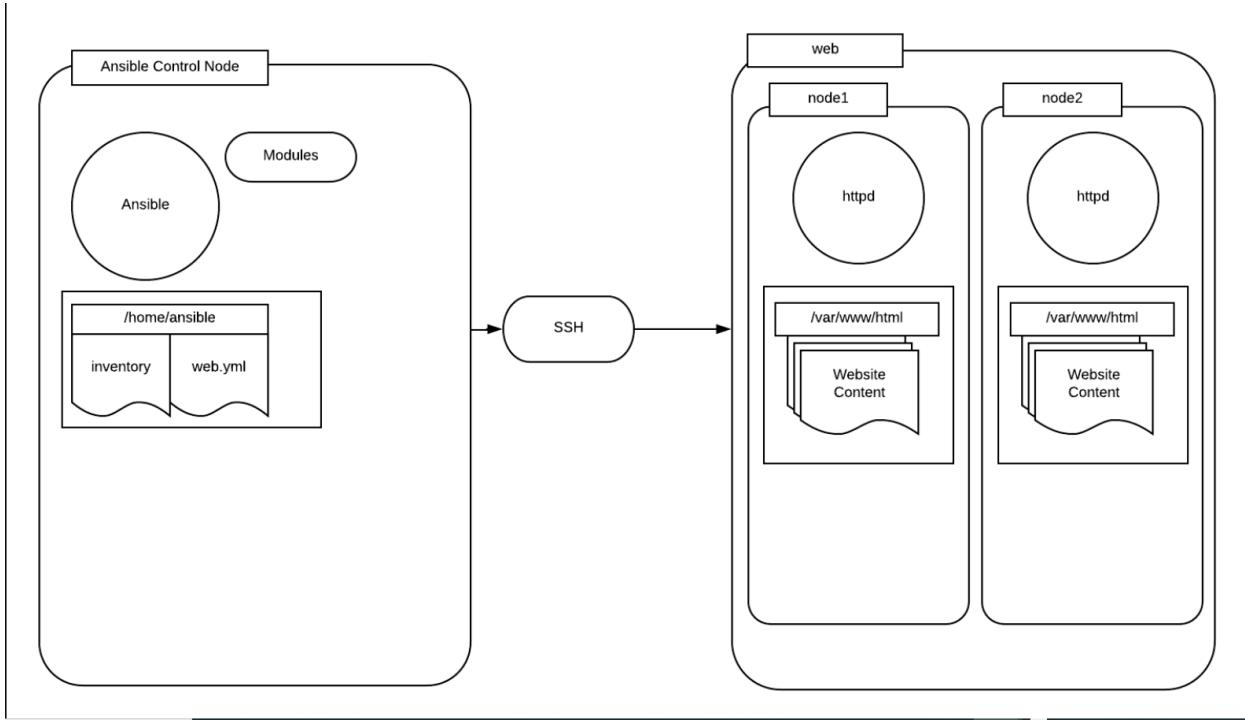
```
[cloud_user@control1 ~]$ cat /etc/hosts
127.0.0.1    localhost localhost.localdomain localhost4 localhost4.localdomain4
::1          localhost localhost.localdomain localhost6 localhost6.localdomain6

10.0.1.85 node1
10.0.1.16 node2
[cloud_user@control1 ~]$ █
```

```
getent passwd | awk -F: '{print $1}'
```

```
[cloud_user@control1 ~]$ getent passwd | awk -F: '{print $1}'
root
bin
daemon
adm
lp
sync
shutdown
halt
mail
operator
games
ftp
nobody
systemd-network
dbus
polkitd
sshd
postfix
chrony
ec2-user
linuxacademy
cloud_user
ansible
[cloud_user@control1 ~]$ █
```

Ansible Playbooks: The Basics



The Scenario

Our company has been increasing the deployment of small brochure-style websites for clients. The head of IT has decided that each client should have their own web server, for better client isolation, and has tasked us with creating concept automation to quickly deploy web-nodes with simple static website content.

We have been provided an Ansible control node (control1) and 2 test lab servers (node1 and node2) that have been preconfigured with the ansible user and key (see below).

We must create an Ansible inventory in **/home/ansible/inventory** containing a host group named **web**.

The **web** group should contain **node1** and **node2**.

Then we've got to design an Ansible playbook that will execute the following tasks on your configured inventory:

Install httpd

Start and enable the httpd service

Install a simple website provided on a repository server.

The playbook will be in /home/ansible/web.yml. The simple website may be accessed from <http://repo.example.com/website.tgz>.

Get Logged In

Once we're logged into the control1 server, become the ansible user (su - ansible) and we can get going.

```
[cloud_user@control1 ~]$ sudo su - ansible  
[sudo] password for cloud_user:  
[ansible@control1 ~]$ █
```

Create an inventory in /home/ansible/inventory That Contains a Host Group Named web. The web Group Should Contain node1 and node2

Use Vim to create the inventory file:

```
[ansible@control1]$ echo "[web]" >> /home/ansible/inventory  
[ansible@control1]$ echo "node1" >> /home/ansible/inventory  
[ansible@control1]$ echo "node2" >> /home/ansible/inventory
```

```
[ansible@control1 ~]$ echo "[web]" >> /home/ansible/inventory  
[ansible@control1 ~]$ echo "node1" >> /home/ansible/inventory  
[ansible@control1 ~]$ echo "node2" >> /home/ansible/inventory  
[ansible@control1 ~]$ █
```

It should contain this when we're done:

```
[web]  
node1  
node2
```

```
[ansible@control1 ~]$ cat /home/ansible/inventory  
[web]  
node1  
node2  
[ansible@control1 ~]$ █
```

Create a Playbook in /home/ansible/web.yml

```
[ansible@control1 ~]$ vi /home/ansible/web.yml
[ansible@control1 ~]$ 
```

Using Vim, we'll create our web.yml file with these contents:

```
---
- hosts: web
  become: yes
  tasks:
    - name: install httpd
      yum: name=httpd state=latest
    - name: start and enable httpd
      service: name=httpd state=started enabled=yes
    - name: retrieve website from repo
      get_url: url=http://repo.example.com/website.tgz dest=/tmp/website.tgz
    - name: install website
      unarchive: remote_src=yes src=/tmp/website.tgz dest=/var/www/html/
```

```
---
- hosts: web
  become: yes
  tasks:
    - name: install httpd
      yum: name=httpd state=latest
    - name: start and enable httpd
      service: name=httpd state=started enabled=yes
    - name: retrieve website from repo
      get_url: url=http://repo.example.com/website.tgz dest=/tmp/website.tgz
    - name: install website
      unarchive: remote_src=yes src=/tmp/website.tgz dest=/var/www/html/
~
~
```

Verify the Work by Executing the Playbook Using the Inventory

```
[ansible@control1]$ansible-playbook -i /home/ansible/inventory /home/ansible/web.yml
```

```
[ansible@control1 ~]$ ansible-playbook -i /home/ansible/inventory /home/ansible/web.yml

PLAY [web] ****
TASK [Gathering Facts] ****
ok: [node1]
ok: [node2]

TASK [install httpd] ****
changed: [node2]
changed: [node1]

TASK [start and enable httpd] ****
changed: [node2]
changed: [node1]

TASK [retrieve website from repo] ****
changed: [node2]
changed: [node1]

TASK [install website] ****
changed: [node2]
changed: [node1]

PLAY RECAP ****
node1                  : ok=5    changed=4    unreachable=0    failed=0    skipped=0    rescued=0    ignored=0
node2                  : ok=5    changed=4    unreachable=0    failed=0    skipped=0    rescued=0    ignored=0

[ansible@control1 ~]$ 
```

Addendum**cat /etc/ansible/hosts**

```
[ansible@control1 ~]$ cat /etc/ansible/hosts

## db-[99:101]-node.example.com

[nodes]
node1
node2

[ansible@control1 ~]$ 
```

Cat /etc/hosts

```
[ansible@control1 ~]$ cat /etc/hosts
127.0.0.1    localhost localhost.localdomain localhost4 localhost4.localdomain4
::1          localhost localhost.localdomain localhost6 localhost6.localdomain6

10.0.1.201 node1
10.0.1.45 node2
127.0.0.1      repo.example.com
[ansible@control1 ~]$ █
```

sudo visudo

```
## Read drop-in files from /etc/sudoers.d (the # here does not mean a comment)
#includedir /etc/sudoers.d
ec2-user      ALL=(ALL)      NOPASSWD: ALL
ansible       ALL=(ALL)      NOPASSWD: ALL
```

ssh node1

```
[ansible@control1 ~]$ ssh node1
Last login: Sat Sep 30 14:02:37 2023 from ip-10-0-1-120.ec2.internal
[ansible@node1 ~]$ sudo visudo
visudo: /etc/sudoers.tmp unchanged
[ansible@node1 ~]$ logout
Connection to node1 closed.
[ansible@control1 ~]$ █
```

```
getent passwd | awk -F: '{print $1}'
```

```
[ansible@control1 ~]$ getent passwd | awk -F: '{print $1}'  
root  
bin  
daemon  
adm  
lp  
sync  
shutdown  
halt  
mail  
operator  
games  
ftp  
nobody  
systemd-network  
dbus  
polkitd  
sshd  
postfix  
chrony  
ec2-user  
linuxacademy  
cloud_user  
ansible  
apache  
[ansible@control1 ~]$ █
```

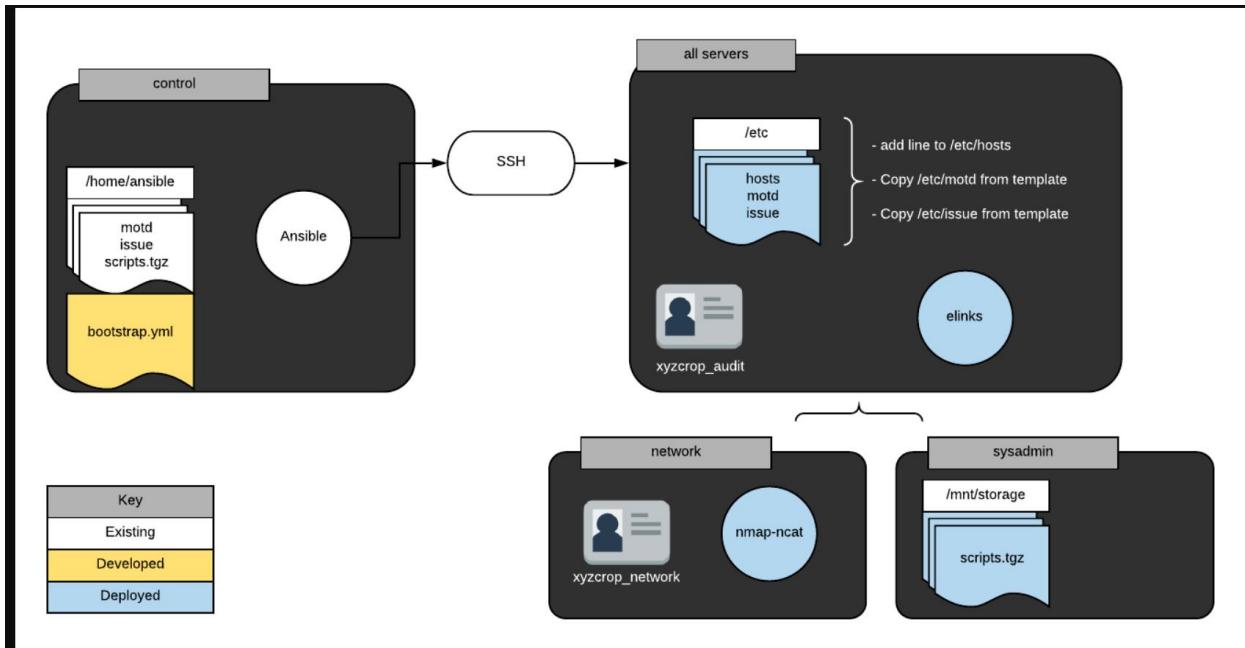
```
cat /etc/os-release
```

```
[ansible@control1 ~]$ cat /etc/os-release  
NAME="Red Hat Enterprise Linux Server"  
VERSION="7.9 (Maipo)"  
ID="rhel"  
ID_LIKE="fedora"  
VARIANT="Server"  
VARIANT_ID="server"  
VERSION_ID="7.9"  
PRETTY_NAME="Red Hat Enterprise Linux Server 7.9 (Maipo)"  
ANSI_COLOR="0;31"  
CPE_NAME="cpe:/o:redhat:enterprise_linux:7.9:GA:server"  
HOME_URL="https://www.redhat.com/"  
BUG_REPORT_URL="https://bugzilla.redhat.com/"  
  
REDHAT_BUGZILLA_PRODUCT="Red Hat Enterprise Linux 7"  
REDHAT_BUGZILLA_PRODUCT_VERSION=7.9  
REDHAT_SUPPORT_PRODUCT="Red Hat Enterprise Linux"  
REDHAT_SUPPORT_PRODUCT_VERSION="7.9"  
[ansible@control1 ~]$ █
```

Cat ~/.ssh/id_rsa.pub | cat ~/.ssh/id_rsa

```
[ansible@control1 ~]$ cat ~/.ssh/id_rsa.pub
ssh-rsa AAAAAB3NzaC1y2EAAAQABAAQCe6HRuRHqo1lBrIFys71sZJlgHufKghub1xUAnCEYemcYwoWT60FgDEuBaqCg4WQQU2k0DqEjdCtop9vbrhlcY50NbhwfXSF2WH2066kJqgh+TFQhRZh+KRWkwBPxG+IaP
8qYhoVW4XVxxTwxINExFbpmWwt+8SVAXN0rU/qBa0LtREq3ehz4rtsFltFzi7n5IhcIkklkjsofiuqq9GDBztxRrs12HkgL9WEVUOYy2+W4Tb2yqmK6IP+Ju187MueFcggFR9ShmOZxMw3x37Q/Z14CDMyPvE0/1/btH
m7JOGMr2rZYLoQxb05uWKAETZaX8r+m1GfLm01 ansible@control1
[ansible@control1 ~]$ cat ~/.ssh/id_rsa
-----BEGIN RSA PRIVATE KEY-----
MIIEogIBAAKCAQEAnneh0bkRG6BqNZQay8cr09bGSZYB7nyoIbm9cVAJwhGHpnGM
KFk+tBYAxLgwqo0FkK1NpNA6hI3QraKfb264ZXG0dDW4Vn10hd1h9tBupCaoIfk
xJUJWYfkVpMAT8RviGj/KmIaFVj0F1ccU1sSDRF2hW6ZlsLf/E1QfZdk1P6gWtC
7URKt3oc+K7bBzbRWYu510SIXCJJXibKH4rkKvRgwc7cua7NdhycC/VhFVDmMt
vluE29sqpiuID/ibnpQezLnhXIKxUoUjZjmCTft8d+0P2ZeAg0jMj7xDv5f27R5
uyThjK9q2NC6EMW60VMcB6E2W1I/K/pRny5jpQIDAQABAoIBAQCYf0/+xdzYRYpP
HUPYojM8tqcDUZrySQTyIcCcJ+zXuNHzaWfpLVWeiR/Wn1R2RVn1fV730g6s1
TEBiW9qz5N/N/ZRzxy9HT7byqgqPdZP/fatzbV/o0Dp2Beka2l1q1wAXbwjkeh1
6YJ4lqVJpuRBxv8mWR9r9E1A89zzP83CqAD8gXb8LwDtIGObwnwb6x0KP7kRKpzA
TSYeI6J7gJE8KCghLIGxWjt+N9BFx+x5xFHYwurc1nol/eSY3y6pvWqT7g9dtHe
NnLx7P+ZKrjtlawQnUZW31wsaNKzGJTK42+MTLU1MY38qHJ8laotakC8UBRA5c71
eYukGEdRAcGBANC8qspBriVA/5p1dzblMygVK9hsM0NvkErPNL+knUhhFaKZZ2GfZ
YpD047aIuMAuOJgD+Ws1EfGa9NzmqW86Us+x2fnM16jA3yvakZqWwqYPff9vz0d
IUBGX2A70jFzGasKy9VNS0dC40Kc+pPpZAejSKCavtH6NwgI1Fh/U5bAoGBAMQ
GldupCTrmalWkw3s5aJhv3+ZBQzd/f5tnzFzx2qXqroVdtHnmG9+T/2eORicgocVT
QTqd3MF7Xn3PhnWb8ZG7g9M1SF4IdHNotMPE9kuJyDkLOB4YaFx1QaeHaR/HES
T0i+PtcL0NTG4MzPuLo+TU+GvzSnnntBujaFbX/AoGBAMv983BMzgnYm7n8hzPw
C7495Q72LOjVLI8Is1U1TgBwC7w087uWYZ/Qz30aXsPCNzRna7vCd76VLCb2
+FAK0aFUCm2YYvL70hDugP/OjuIV0YibFuJDD1Mac7y5dwVn/KBFKoxHXEjoOCe
58YNDd+zZONPkt1i0WuHyQF9AoGAQChf1ge8E8Wa68rnH6xQreV4kyRi0Lr5Wdu
2S4VimuIsojV1QHTHwekahVcgTQRygQuFFcmcmHXRxxj/nmt5fjMV5ViEn1dQeC
MGwi05g3e7xkJF7u65wipokMg4w73IT+IxteUrdqV4xErpKVh1pa5mFNuM0mMsy
5SX01ScCfxw4cvdfNkql5boN1+VCm7fQqDbA8SK0PC60QYAO+vNx1GqcJojf1XX
4rkvRiOFxFil/WDXdr2ljxwJ1/3EBjhhHw3YHB9qSL4MtzuDBZOpPM6dTRWSLa
SqkGx2kdyO+uM0r0oSGrV7k1sovTIht6PjVcvde0BvkS3u8IqY0=
-----END RSA PRIVATE KEY-----
[ansible@control1 ~]$
```

Writing Your First Ansible Playbook



Playbooks are the heart of Ansible. They provide a means of executing numerous tasks

across any number of systems.

This exercise sets you on the path to Ansible mastery by allowing you to craft and execute an Ansible playbook.

Your supervisor has asked you to find a way to automate and audit basic system configuration for new servers in your environment. Given that Ansible is already configured at a basic level in your environment, the simplest solution would be to just write a playbook for bootstrapping your new hosts.

Create a playbook called **/home/ansible/bootstrap.yml** to fulfill the following bootstrap requirements:

All servers:

Edit /etc/hosts to include the following entry:

ansible.xyzcorp.com 169.168.0.1

Install elinks

Create the user xyzcorp_audit

Copy the files /home/ansible/motd and /home/ansible/issue to /etc/

Network servers:

Install nmap-ncat

Create the user xyzcorp_network**SysAdmin servers:**

Copy /home/ansible/scripts.tgz from the control node to /mnt/storage

The Ansible control node has been configured for you and each testing server has already been configured for use with Ansible. The default inventory has been configured to include the groups network and sysadmin.

Each group includes a sample host.

Solution

Begin by logging in to the Ansible control node:

ssh cloud_user@PUBLIC_IP_ADDRESS

Become the ansible user:

sudo -i -u ansible

```
[cloud_user@control ~]$ sudo -i -u ansible
[sudo] password for cloud_user:
```

Create a basic playbook for all servers in the Ansible inventory:

Create a basic playbook in **/home/ansible/bootstrap.yml** that completes the noted task required for all servers in the Ansible inventory.

Create the file **/home/ansible/bootstrap.yml** and add the following content:

vim /home/ansible/bootstrap.yml

```
---
- hosts: all
  become: yes
  tasks:
    - name: edit host file
      lineinfile:
        path: /etc/hosts
        line: "ansible.xyzcorp.com 169.168.0.1"
    - name: install elinks
      package:
        name: elinks
        state: latest
    - name: create audit user
      user:
        name: xyzcorp_audit
        state: present
    - name: update motd
      copy:
```

```
src: /home/ansible/motd
dest: /etc/motd
- name: update issue
copy:
  src: /home/ansible/issue
  dest: /etc/issue
```

```
---
- hosts: all
become: yes
tasks:
  - name: edit host file
    lineinfile:
      path: /etc/hosts
      line: "ansible.xyzcorp.com 169.168.0.1"
  - name: install elinks
    package:
      name: elinks
      state: latest
  - name: create audit user
    user:
      name: xyzcorp_audit
      state: present
  - name: update motd
    copy:
      src: /home/ansible/motd
      dest: /etc/motd
  - name: update issue
    copy:
      src: /home/ansible/issue
      dest: /etc/issue
```

~
~
~
~

Add a section to the playbook for the Network servers in the Ansible inventory.

Add a section to the playbook in **/home/ansible/bootstrap.yml** that completes the noted task required for the network servers in the Ansible inventory.

Create a play in the file **/home/ansible/bootstrap.yml** similar to the following:

```
---
- hosts: all
  become: yes
  tasks:
    - name: edit host file
      lineinfile:
        path: /etc/hosts
        line: "ansible.xyzcorp.com 169.168.0.1"
    - name: install elinks
      package:
        name: elinks
        state: latest
    - name: create audit user
      user:
        name: xyzcorp_audit
        state: present
    - name: update motd
      copy:
        src: /home/ansible/motd
        dest: /etc/motd
    - name: update issue
      copy:
        src: /home/ansible/issue
        dest: /etc/issue

- hosts: network
  become: yes
  tasks:
    - name: install netcat
      yum:
        name: nmap-ncat
        state: latest
    - name: create network user
      user:
        name: xyzcorp_network
        state: present
~
~
```

```
- hosts: network
  become: yes
  tasks:
    - name: install netcat
      yum:
        name: nmap-ncat
        state: latest
    - name: create network user
      user:
        name: xyzcorp_network
        state: present
```

Add a section to the playbook for the SysAdmin servers in the Ansible inventory.
 Add a section to the playbook in **/home/ansible/bootstrap.yml** that completes the noted task required for the sysadmin servers in the Ansible inventory.

Create a play in the file **/home/ansible/bootstrap.yml** similar to the following:

```

- hosts: network
  become: yes
  tasks:
    - name: install netcat
      yum:
        name: nmap-ncat
        state: latest
    - name: create network user
      user:
        name: xyzcorp_network
        state: present

- hosts: sysadmin
  become: yes
  tasks:
    - name: copy tarball
      copy:
        src: /home/ansible/scripts.tgz
        dest: /mnt/storage/■
-- INSERT (paste) --

```

```

- hosts: sysadmin
  become: yes
  tasks:
    - name: copy tarball
      copy:
        src: /home/ansible/scripts.tgz
        dest: /mnt/storage/

```

Execute playbook to verify your playbook works correctly.

Execute playbook **/home/ansible/bootstrap.yml** to verify your playbook works correctly.

Run **ansible-playbook /home/ansible/bootstrap.yml** from the control node.

```
[ansible@control ~]$ vim /home/ansible/bootstrap.yml
[ansible@control ~]$ ansible-playbook /home/ansible/bootstrap.yml
```

```
[ansible@control ~]$ ansible-playbook /home/ansible/bootstrap.yml

PLAY [all] ****
TASK [Gathering Facts] ****
ok: [node1]
ok: [localhost]
ok: [node2]

TASK [edit host file] ****
changed: [localhost]
changed: [node1]
changed: [node2]

TASK [install elinks] ****
changed: [localhost]
changed: [node1]
changed: [node2]

TASK [create audit user] ****
changed: [node1]
changed: [node2]
changed: [localhost]

TASK [update motd] ****
changed: [localhost]
changed: [node1]
changed: [node2]

TASK [update issue] ****
changed: [node1]
changed: [localhost]
changed: [node2]

PLAY [network] ****
TASK [Gathering Facts] ****
```

```
ok: [node1]

TASK [install netcat] ****
changed: [node1]

TASK [create network user] ****
changed: [node1]

PLAY [sysadmin] ****
TASK [Gathering Facts] ****
ok: [node2]

TASK [copy tarball] ****
changed: [node2]

PLAY RECAP ****
localhost      : ok=6    changed=5    unreachable=0    failed=0    skipped=0    rescued=0    ignored=0
node1         : ok=9    changed=7    unreachable=0    failed=0    skipped=0    rescued=0    ignored=0
node2         : ok=8    changed=6    unreachable=0    failed=0    skipped=0    rescued=0    ignored=0
```

Addendum**cat /etc/hosts**

```
[ansible@control ~]$ cat /etc/hosts
127.0.0.1 localhost localhost.localdomain localhost4 localhost4.localdomain4
::1 localhost localhost.localdomain localhost6 localhost6.localdomain6
10.0.1.92 node1
10.0.1.185 node2
[ansible@control ~]$
```

cat /etc/os-release

```
[ansible@control ~]$ cat /etc/os-release
NAME="CentOS Linux"
VERSION="7 (Core)"
ID="centos"
ID_LIKE="rhel fedora"
VERSION_ID="7"
PRETTY_NAME="CentOS Linux 7 (Core)"
ANSI_COLOR="0;31"
CPE_NAME="cpe:/o:centos:centos:7"
HOME_URL="https://www.centos.org/"
BUG_REPORT_URL="https://bugs.centos.org/"

CENTOS_MANTISBT_PROJECT="CentOS-7"
CENTOS_MANTISBT_PROJECT_VERSION="7"
REDHAT_SUPPORT_PRODUCT="centos"
REDHAT_SUPPORT_PRODUCT_VERSION="7"

[ansible@control ~]$
```

sudo visudo

```
## Read drop-in files from /etc/sudoers.d (the # here does not mean a comment)
#includedir /etc/sudoers.d
cloud_user ALL=(ALL)    ALL
ansible      ALL=(ALL)    NOPASSWD: ALL
```

sudo whoami

```
[ansible@control ~]$ sudo whoami
root
[ansible@control ~]$
```

cat /etc/ansible/hosts**or / and**

```
cat /home/ansible/inventory
```

```
[ansible@control ~]$ vim /home/ansible/bootstrap.yml
[ansible@control ~]$ ansible-playbook /home/ansible/bootstrap.yml
ERROR! We were unable to read either as JSON nor YAML, these are the errors we got from each:
JSON: No JSON object could be decoded

Syntax Error while loading YAML.
  mapping values are not allowed in this context

The error appears to be in '/home/ansible/bootstrap.yml': line 2, column 9, but may
be elsewhere in the file depending on the exact syntax problem.

The offending line appears to be:

---
- hosts: all
  ^ here
[ansible@control ~]$ █
```

```
---  
- hosts: all  
  become: yes  
  tasks:  
    - name: edit host file
```

How did I troubleshoot the above

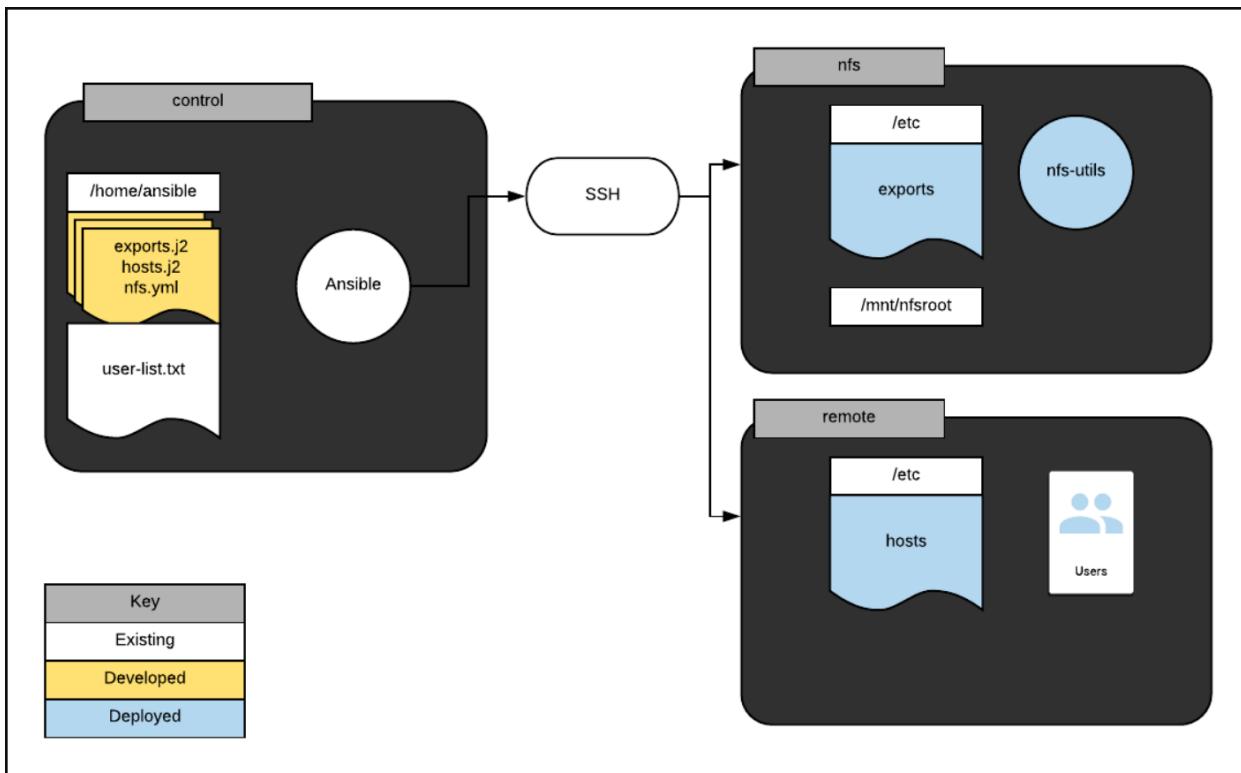
How did I solve the above?

First, I deleted the **- hosts: all** line and copy and pasted it again!

When that failed, I deleted the 3 dashes above --- and copy and pasted them again!

One key indicator for the --- having the wrong syntax was the colour (It was light white
Rather than a dull yellow).

Deploying Services Using Ansible



Introduction

The team has been working on automating the builds for common services using Ansible. You have been tasked with creating the build for a common NFS server. Write a playbook that satisfies the following requirements:

On nfs:

Make sure `nfs-utils` is installed.

Configure `/etc(exports` via an Ansible template stored in `/home/ansible/exports.j2`.

Deploy the template so that `/mnt/nfsroot` is exported with read and write to all hosts.

Note: Your template file should contain the following with the variable `{{ share_path }}` being defined within your playbook:

`{{ share_path }} *(rw)`

Note: The file `/etc(exports` on nfs should have the following content once deployed:

/mnt/nfsroot *(rw)

Create a handler that runs the command `export fs -a` if the file `/etc/exports` is modified in a playbook task.

You can assume all necessary firewall rules have been deployed.

On remote:

Configure `/etc/hosts` from a template file stored on control at `/home/ansible/etc.hosts.j2` with the following entries:

```
127.0.0.1 localhost,{{ ansible_hostname }}  
{{ nfs_ip }} {{ nfs_hostname }}
```

Note: You should populate the variables `{{ nfs_ip }}` and `{{ nfs_hostname }}` using magic variables in your playbook.

Create users from file stored on control at `/home/ansible/user-list.txt` only if the remote host has the file `/opt/user-agreement.txt`.

The Ansible control node has been configured for you and all other servers have already been configured for use with Ansible. The default inventory has been configured to include the group `remote` and server `nfs`.

Solution

Begin in by logging in to the Ansible control node as the `ansible` user.

ssh ansible@IP_ADDRESS

Create the necessary template files on the Ansible control node.

Create the files `exports.j2` in `/home/ansible` with the noted content:

vim /home/ansible/exports.j2

Provide the following content for this file:

`{{ share_path }} *(rw)`

Create the file `etc.hosts.j2` in `/home/ansible` with the noted content:

vim /home/ansible/etc.hosts.j2

Provide the following content for this file:

```
127.0.0.1      localhost {{ ansible_hostname }}  
{{ nfs_ip }}      {{ nfs_hostname }}
```

Create a playbook for the server 'nfs' in the Ansible inventory.

Create a playbook in **/home/ansible/nfs.yml** that meets the provided specification for the server nfs in the ansible inventory.

Edit the /home/ansible/nfs.yml file:

```
vim /home/ansible/nfs.yml
```

Add a playbook in **/home/ansible/nfs.yml** with the following content:

```
- hosts: nfs
  become: yes
  vars:
    share_path: /mnt/nfsroot
  tasks:
    - name: install nfs
      yum:
        name: nfs-utils
        state: latest
    - name: start and enable nfs-server
      service:
        name: nfs-server
        state: started
        enabled: yes
    - name: configure exports
      template:
        src: /home/ansible/exports.j2
        dest: /etc/exports
        notify: update nfs
  handlers:
    - name: update nfs exports
      command: exportfs -a
      listen: update nfs
~
```

```
- hosts: nfs
become: yes
vars:
  share_path: /mnt/nfsroot
tasks:
  - name: install nfs
    yum:
      name: nfs-utils
      state: latest
  - name: start and enable nfs-server
    service:
      name: nfs-server
      state: started
      enabled: yes
  - name: configure exports
    template:
      src: /home/ansible/exports.j2
      dest: /etc/exports
      notify: update nfs
handlers:
  - name: update nfs exports
    command: exportfs -a
    listen: update nfs
```

```
- hosts: nfs
become: yes
vars:
  share_path: /mnt/nfsroot
tasks:
  - name: install nfs
    yum:
      name: nfs-utils
      state: latest
  - name: start and enable nfs-server
    service:
      name: nfs-server
      state: started
      enabled: yes
  - name: configure exports
    template:
      src: /home/ansible/exports.j2
      dest: /etc/exports
      notify: update nfs
handlers:
  - name: update nfs exports
    command: exportfs -a
    listen: update nfs
```

Add a play for the remote host group.

Edit **/home/ansible/nfs.yml** such that changes are executed on the remote host group to meet the provided requirements.

Open the **/home/ansible/nfs.yml** file:

vim /home/ansible/nfs.yml

Edit the **/home/ansible/nfs.yml** file to contain the following content after the nfs play:

```
- hosts: remote
  become: yes
  vars:
    nfs_ip: "{{ hostvars['nfs']['ansible_default_ipv4']['address'] }}"
    nfs_hostname: "{{ hostvars['nfs']['ansible_hostname'] }}"
  vars_files:
    - /home/ansible/user-list.txt
  tasks:
    - name: configure hostsfile
      template:
        src: /home/ansible/etc.hosts.j2
        dest: /etc/hosts
    - name: get file status
      stat:
        path: /opt/user-agreement.txt
      register: filestat
    - name: debug info
      debug:
        var: filestat
    - name: create users
      user:
        name: "{{ item }}"
      when: filestat.stat.exists
      loop: "{{ users }}"
```

```
- hosts: remote
  become: yes
  vars:
    nfs_ip: "{{ hostvars['nfs']['ansible_default_ipv4']['address'] }}"
    nfs_hostname: "{{ hostvars['nfs']['ansible_hostname'] }}"
  vars_files:
    - /home/ansible/user-list.txt
  tasks:
    - name: configure hostsfile
      template:
        src: /home/ansible/etc.hosts.j2
        dest: /etc/hosts
    - name: get file status
      stat:
        path: /opt/user-agreement.txt
      register: filestat
    - name: debug info
```

```

debug:
  var: filestat
- name: create users
  user:
    name: "{{ item }}"
  when: filestat.stat.exists
  loop: "{{ users }}"

```

```

- hosts: remote
  become: yes
  vars:
    nfs_ip: "{{ hostvars['nfs']['ansible_default_ipv4']['address'] }}"
    nfs_hostname: "{{ hostvars['nfs']['ansible_hostname'] }}"
  vars_files:
    - /home/ansible/user-list.txt
  tasks:
    - name: configure hostsfile
      template:
        src: /home/ansible/etc.hosts.j2
        dest: /etc/hosts
    - name: get file status
      stat:
        path: /opt/user-agreement.txt
      register: filestat
    - name: debug info
      debug:
        var: filestat
    - name: create users
      user:
        name: "{{ item }}"
      when: filestat.stat.exists
      loop: "{{ users }}"
-- INSERT (paste) --

```

Execute playbook to verify your playbook works correctly.

Execute playbook /home/ansible/nfs.yml to verify your playbook works correctly.

Run the command below from the control node.

ansible-playbook /home/ansible/nfs.yml

```

        "mimetype": "text/plain",
        "mode": "0644",
        "mtime": 1696046537.6733913,
        "nlink": 1,
        "path": "/opt/user-agreement.txt",
        "pw_name": "root",
        "readable": true,
        "rgrp": true,
        "roth": true,
        "rusr": true,
        "size": 34,
        "uid": 0,
        "version": "18446744071586397626",
        "wgrp": false,
        "woth": false,
        "writable": true,
        "wusr": true,
        "xgrp": false,
        "xoth": false,
        "xusr": false
    }
}

TASK [create users] ****
changed: [node2] => (item=frank)
changed: [node2] => (item=judy)
changed: [node2] => (item=joe)
changed: [node2] => (item=sarah)
changed: [node2] => (item=sam)
changed: [node2] => (item=carry)

PLAY RECAP ****
nfs                  : ok=5     changed=4      unreachable=0      failed=0      skipped=0      rescued=0      ignored=0
node2                : ok=5     changed=2      unreachable=0      failed=0      skipped=0      rescued=0      ignored=0
[ansible@control ~]$ 

```

Cat /home/ansible/user-list.txt

```

[ansible@control ~]$ cat /home/ansible/user-list.txt
users:
- frank
- judy
- joe
- sarah
- sam
- carry
[ansible@control ~]$ 

```

Addendum

```
[ansible@control ~]$ vim /home/ansible/nfs.yml
[ansible@control ~]$ ansible-playbook /home/ansible/nfs.yml
ERROR! We were unable to read either as JSON nor YAML, these are the errors we got from each:
JSON: No JSON object could be decoded

Syntax Error while loading YAML.
  did not find expected '-' indicator

The error appears to be in '/home/ansible/nfs.yml': line 10, column 6, but may
be elsewhere in the file depending on the exact syntax problem.

The offending line appears to be:

    state: latest
  - name: start and enable nfs-server
    ^ here
[ansible@control ~]$
```

```
[ansible@control ~]$ vim /home/ansible/nfs.yml
[ansible@control ~]$ ansible-playbook /home/ansible/nfs.yml
ERROR! We were unable to read either as JSON nor YAML, these are the errors we got from each:
JSON: No JSON object could be decoded

Syntax Error while loading YAML.
  did not find expected '' indicator

The error appears to be in '/home/ansible/nfs.yml': line 20, column 5, but may
be elsewhere in the file depending on the exact syntax problem.

The offending line appears to be:

    notify: update nfs
  handlers:
    ^ here
[ansible@control ~]$
```

```
[ansible@control ~]$ ansible-playbook /home/ansible/nfs.yml
ERROR! We were unable to read either as JSON nor YAML, these are the errors we got from each:
JSON: No JSON object could be decoded

Syntax Error while loading YAML.
  did not find expected '-' indicator

The error appears to be in '/home/ansible/nfs.yml': line 25, column 2, but may
be elsewhere in the file depending on the exact syntax problem.

The offending line appears to be:

- hosts: remote
^ here
[ansible@control ~]$ █
```

```
- hosts: nfs
  become: yes
  vars:
    share_path: /mnt/nfsroot
  tasks:
    - name: install nfs
      yum:
        name: nfs-utils
        state: latest
    - name: start and enable nfs-server
      service:
        name: nfs-server
        state: started
        enabled: yes
    - name: configure exports
      template:
        src: /home/ansible/exports.j2
        dest: /etc/exports
        notify: update nfs
  handlers:
    - name: update nfs exports
      command: exportfs -a
      listen: update nfs
  . . .
```

```

- hosts: nfs
  become: yes
  vars:
    share_path: /mnt/nfsroot
  tasks:
    - name: install nfs
      yum:
        name: nfs-utils
        state: latest
    - name: start and enable nfs-server
      service:
        name: nfs-server
        state: started
        enabled: yes
    - name: configure exports
      template:
        src: /home/ansible/exports.j2
        dest: /etc/exports
      notify: update nfs
  handlers:
    - name: update nfs exports
      command: exportfs -a
      listen: update nfs

- hosts: remote
  become: yes
  vars:
    nfs_ip: "{{ hostvars['nfs']['ansible_default_ipv4']['address'] }}"
    nfs_hostname: "{{ hostvars['nfs']['ansible_hostname'] }}"
  vars_files:
    - /home/ansible/user-list.txt
  tasks:
    - name: configure hostsfile
      template:
        src: /home/ansible/etc.hosts.j2
        dest: /etc/hosts
-- INSERT --

```

The errors above were generated by syntax errors. Unfortunately, YAML is not agnostic about text not aligning like Terraform!

cat /etc/os-release

```

[ansible@control ~]$ cat /etc/os-release
NAME="CentOS Linux"
VERSION="7 (Core)"
ID="centos"
ID_LIKE="rhel fedora"
VERSION_ID="7"
PRETTY_NAME="CentOS Linux 7 (Core)"
ANSI_COLOR="0;31"
CPE_NAME="cpe:/o:centos:centos:7"
HOME_URL="https://www.centos.org/"
BUG_REPORT_URL="https://bugs.centos.org/"

CENTOS_MANTISBT_PROJECT="CentOS-7"
CENTOS_MANTISBT_PROJECT_VERSION="7"
REDHAT_SUPPORT_PRODUCT="centos"
REDHAT_SUPPORT_PRODUCT_VERSION="7"

[ansible@control ~]$ 

```

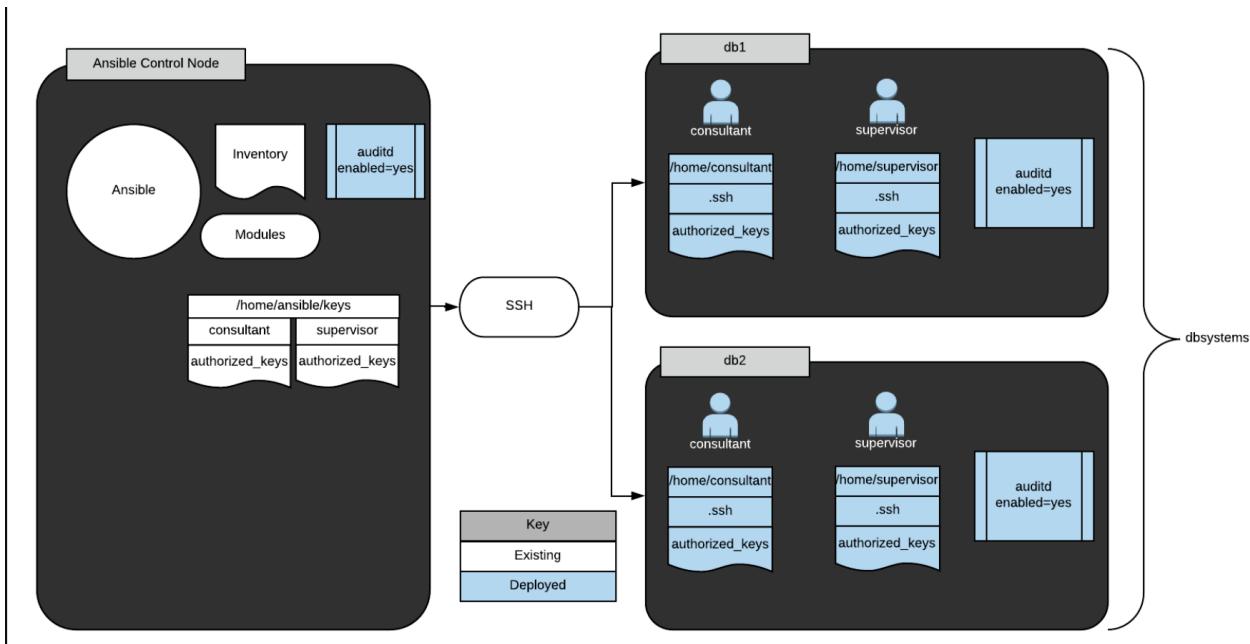
Sudo cat /etc/hosts (hosts is similar to the host.ini in windows)

```
[ansible@control ~]$ sudo cat /etc/hosts
127.0.0.1    localhost localhost.localdomain localhost4 localhost4.localdomain4
::1          localhost localhost.localdomain localhost6 localhost6.localdomain6
10.0.1.53    nfs node1
10.0.1.211   node2
[ansible@control ~]$
```

cat /etc/ansible/hosts

```
localhost
nfs
[remote]
node2
[ansible@control ~]$
```

Ad-Hoc Ansible Commands



The Scenario

Some consultants will be performing audits on a number of systems in our company's environment. We've got to create the user accounts listed in **/home/ansible/userlist.txt** and set up the provided public keys for their accounts. The security team has built a jump host for the consultants to access production systems and provided us with the full key-pair so we can set up and test the connection.

All hosts in dbsystems will need that public key installed so the consultants may use key-pair authentication to access the systems. We must also ensure the **auditd service** is enabled and running on all systems.

Important notes:

Ansible is already on the control node. If we connect to the server by clicking on the Public IP address in a web browser, we need to make sure we change to the ansible user, with the **sudo su - ansible** command.

The user ansible is present on all servers with appropriate shared keys for access to managed servers from the control node. We need to make sure we use this user to complete the commands.

The ansible user has the same password as cloud_user.

The default Ansible inventory has already been configured with the appropriate hosts and groups.

/etc/hosts entries are present on the control1 host for the managed servers.

Get Logged In

Once we're logged into the control1 server, become the ansible user (**su - ansible**) and we can get going.

```
[cloud_user@control1 ~]$ sudo su - ansible
[sudo] password for cloud_user:
[ansible@control1 ~]$ 
```

Create the User Accounts Noted in /home/ansible/userlist.txt

If we read the **userlist.txt** file in our home directory, we'll see consultant and supervisor.

```
[ansible@control1 ~]$ cat /home/ansible/userlist.txt
# Please create the the below users on all hosts in dbsystems
# You will find a pregenerated key pair and an authorized key file for each user in /home/ansible/keys

# User accounts to create:
consultant
supervisor
[ansible@control1 ~]$ 
```

Those are the two new user accounts we have to create:

```
[ansible@control1]$ ansible dbsystems -b -m user -a "name=consultant"
```

```
[ansible@control1 ~]$ ansible dbsystems -b -m user -a "name=consultant"
db1 | CHANGED => {
    "ansible_facts": {
        "discovered_interpreter_python": "/usr/bin/python"
    },
    "changed": true,
    "comment": "",
    "create_home": true,
    "group": 1004,
    "home": "/home/consultant",
    "name": "consultant",
    "shell": "/bin/bash",
    "state": "present",
    "system": false,
    "uid": 1004
}
db2 | CHANGED => {
    "ansible_facts": {
        "discovered_interpreter_python": "/usr/bin/python"
    },
    "changed": true,
    "comment": "",
    "create_home": true,
    "group": 1004,
    "home": "/home/consultant",
    "name": "consultant",
    "shell": "/bin/bash",
    "state": "present",
    "system": false,
    "uid": 1004
}
[ansible@control1 ~]$ 
```

```
[ansible@control1 ~]$ ansible dbsystems -b -m user -a "name=supervisor"
```

```
[ansible@control1 ~]$ ansible dbsystems -b -m user -a "name=supervisor"
db1 | CHANGED => {
    "ansible_facts": {
        "discovered_interpreter_python": "/usr/bin/python"
    },
    "changed": true,
    "comment": "",
    "create_home": true,
    "group": 1005,
    "home": "/home/supervisor",
    "name": "supervisor",
    "shell": "/bin/bash",
    "state": "present",
    "system": false,
    "uid": 1005
}
db2 | CHANGED => {
    "ansible_facts": {
        "discovered_interpreter_python": "/usr/bin/python"
    },
    "changed": true,
    "comment": "",
    "create_home": true,
    "group": 1005,
    "home": "/home/supervisor",
    "name": "supervisor",
    "shell": "/bin/bash",
    "state": "present",
    "system": false,
    "uid": 1005
}
[ansible@control1 ~]$
[ansible@control1 ~]$ █
```

Place Key Files in the Correct Location, /home/\$USER/.ssh/authorized_keys, on Hosts in dbsystems

```
[ansible@control1 ~]$ ansible dbsystems -b -m file -a "path=/home/consultant/.ssh state=directory owner=consultant group=consultant mode=0755"
```

```
[ansible@control1 ~]$ ansible dbsystems -b -m file -a "path=/home/consultant/.ssh state=directory owner=consultant group=consultant mode=0755"
db1 | CHANGED => {
    "ansible_facts": {
        "discovered_interpreter_python": "/usr/bin/python"
    },
    "changed": true,
    "gid": 1004,
    "group": "consultant",
    "mode": "0755",
    "owner": "consultant",
    "path": "/home/consultant/.ssh",
    "secontext": "unconfined_u:object_r:ssh_home_t:s0",
    "size": 6,
    "state": "directory",
    "uid": 1004
}
db2 | CHANGED => {
    "ansible_facts": {
        "discovered_interpreter_python": "/usr/bin/python"
    },
    "changed": true,
    "gid": 1004,
    "group": "consultant",
    "mode": "0755",
    "owner": "consultant",
    "path": "/home/consultant/.ssh",
    "secontext": "unconfined_u:object_r:ssh_home_t:s0",
    "size": 6,
    "state": "directory",
    "uid": 1004
}
[ansible@control1 ~]$ █
```

```
[ansible@control1]$ ansible dbsystems -b -m copy -a
"src=/home/ansible/keys/consultant/authorized_keys dest=/home/consultant/.ssh/authorized_keys
mode=0600 owner=consultant group=consultant"
```

```
[ansible@control1 ~]$ ansible dbsystems -b -m copy -a "src=/home/ansible/keys/consultant/authorized_keys dest=/home/consultant/.ssh/authorized_keys mode=0600 owner=co
nsultant group=consultant"
db1 | CHANGED => {
    "ansible_facts": {
        "discovered_interpreter_python": "/usr/bin/python"
    },
    "changed": true,
    "checksum": "a249d6a58d39b67f331b37c2ae8c24af42faab7",
    "dest": "/home/consultant/.ssh/authorized_keys",
    "gid": 1004,
    "group": "consultant",
    "md5sum": "6478a0a7abb698adbb0b9c6cb81bcb4f",
    "mode": "0600",
    "owner": "consultant",
    "secontext": "unconfined u:object_r:ssh_home_t:s0",
    "size": 395,
    "src": "/home/ansible/.ansible/tmp/ansible-tmp-1696086040.98-1764-124376452151507/source",
    "state": "file",
    "uid": 1004
}
db2 | CHANGED => {
    "ansible_facts": {
        "discovered_interpreter_python": "/usr/bin/python"
    },
    "changed": true,
    "checksum": "a249d6a58d39b67f331b37c2ae8c24af42faab7",
    "dest": "/home/consultant/.ssh/authorized_keys",
    "gid": 1004,
    "group": "consultant",
    "md5sum": "6478a0a7abb698adbb0b9c6cb81bcb4f",
    "mode": "0600",
    "owner": "consultant",
    "secontext": "unconfined u:object_r:ssh_home_t:s0",
    "size": 395,
    "src": "/home/ansible/.ansible/tmp/ansible-tmp-1696086041.03-1766-231417013313828/source",
    "state": "file",
    "uid": 1004
}
[ansible@control1 ~]$
```

```
[ansible@control1]$ ansible dbsystems -b -m file -a "path=/home/supervisor/.ssh state=directory
owner=supervisor group=supervisor mode=0755"
```

```
[ansible@control1]$ ansible dbsystems -b -m copy -a
"src=/home/ansible/keys/supervisor/authorized_keys dest=/home/supervisor/.ssh/authorized_keys
mode=0600 owner=supervisor group=supervisor"
```

Ensure auditd Is Enabled and Running on All Hosts

```
[ansible@control1]$ ansible all -b -m service -a "name=auditd state=started enabled=yes"
```

```
[ansible@control1 ~]$ ansible all -b -m service -a "name=auditd state=started enabled=yes"
```

Addendum

getent passwd | awk -F: '{print \$1}'

```
[ansible@control1 ~]$ getent passwd | awk -F: '{print $1}'
root
bin
daemon
adm
lp
sync
shutdown
halt
mail
operator
games
ftp
nobody
systemd-network
dbus
polkitd
sshd
postfix
chrony
ec2-user
linuxacademy
cloud_user
ansible
[ansible@control1 ~]$ █
```

cat /etc/hosts

```
[ansible@control1 ~]$ cat /etc/hosts
127.0.0.1    localhost localhost.localdomain localhost4 localhost4.localdomain4
::1          localhost localhost.localdomain localhost6 localhost6.localdomain6

10.0.1.224 db1 node1
10.0.1.229 db2 node2
[ansible@control1 ~]$ █
```

sudo visudo

```
## Read drop-in files from /etc/sudoers.d (the # here does not mean a comment)
#includedir /etc/sudoers.d
ec2-user      ALL=(ALL)      NOPASSWD: ALL
ansible      ALL=(ALL)      NOPASSWD: ALL
```

ssh node1

```
[ansible@control1 ~]$ ssh node1
Warning: Permanently added 'node1' (ECDSA) to the list of known hosts.
[ansible@db1 ~]$ ls
[ansible@db1 ~]$ sudo visudo
visudo: /etc/sudoers.tmp unchanged
[ansible@db1 ~]$ logout
Connection to node1 closed.
[ansible@control1 ~]$ █
```

cat ~/ssd/id_rsa.pub | cat ~/ssh/id_rsa

```
[ansible@control1 ~]$ cat ~/ssh/id_rsa.pub
ssh-rsa AAAAB3NzaC1yc2EAAQABAAAQAB0Q1BuUhXqZjZgPv0+xQwD/GNzOCB8mWGS/OTc1bPVVmQC6kZ+e8KXsx1dAZPUyPMySQCX+ZEQmJxZEMUsK3SFHpjLEQKhZjH8R9u8DxOodJ+IIRFKz/SY2AKPe/60KF
+QlXSiysb57kiMh3TDrx5zXU53Xcu7Gyz13SEt3ZV0mn91bt4lh5st1mrfXT2x07LGi19cU4u9/4UZAhlba/{Qt tvXhcY5URn0G5QHmw7BoYoEtKyQoX5SCvih9xIfaZF/zvXtGv3tSmIH3AaFD
adun1HE3XTMbggDCgv0AlfUth1ti6w6waWJpWf5C4K9yAb ansible@control1
[ansible@control1 ~]$ cat ~/ssh/id_rsa
-----BEGIN RSA PRIVATE KEY-----
MIIEpqIBAAKCAQEa5QIVYVGmSWYD70PsUMA/xjczzggfJ1hfzkJ3JWz1V2kAupGfn
vC17MZQGT1MqtTM-Kg1/mREj3iV3GRDPLrCt0nx6YyxECh8yR/EfbtA8TqHsfjJUR
ss+HNgcjv+j1hfkC1+crg7N051jd07JQ68ec11Ef7DMXLixess4t0h0/MjdGo
PqaQyFBsWVdjzpPZW7eJyebLdzkx109sd0x0xpXF0Lvf+FGBQ822v/0LbbyxYG0
VEZ9BrEB5sOwaGKBLSSkkF0kgr40fcSH2mRF8717Rr97Up1b9wABQ2nb9rxN10z
CWAIAwoK1dA3RVLVYzbvusOr2liT8H+QuCvcgGwIDAQABAoIBAQCKGLtia05ds/Pr
bMusja0nB5Y5jDSUHsgeoGIQcFV0iLMXU8ggf9DXP3zyGzKEvGQEvyr/fzPQx1qi
X65enP1vJmWY3TLnbAMDWFK0o/q4k0o3cvvuuueh7u7sU1phoiWtrKEYY+0ubj9
epiZEZfxmIevAay7MPHgwvlw/mBeX53XXYY810k8Lk479yhgcK110ZHU/UBeYdIRP
ekxD53egm3fcNdpYdTITCj1VWhWkEPwMMyY2qQZV2c2DthBnLWQBFevL7SgT
X80EcDx0p10iyM/hIMwdtYRLB8sL2Tu0hhHK1PHzYRY093jZfi0MpuesCsTp395sG
FtN2VU3hAoGAP5XvhIPoj0QunXltuOnZ21a4zvH9U/C4QRaTaZirRNap0kmlCxZ
wT7ddQa05v5MSME614ggzyneQxbhjqE9N4c33xUjFeAfAFe26k1c1DHu4+sfe
ou9pJPhkQfc86970ocZbVF12uiujoKfbvUQyd7Qlmn0M4o6ppQvV34JAoGBAOZr
vzwz5rXgVRmOfgCRs0gO1b1h6PPq41XGcmr50STPskrUv1PSV1t8VQEKY1JRx
175Y-BF+yrXiWhPH2NBd0ZSc1Wk85rBwEtzwFnwFD1/19w0+5G5sK6WFkxxi1
h1DvjhCFCyrlrDScAxTs0k3msAkX+RIPoI69yJYDAgGACb+z558/uRU9/6xogp73
PTOMvdBhmgJUL0k42oydw5gvHw89Yok2FTf0h1h1LTf1NnAaMPQ5Yz4X1/WMc
gkXCP0CdPDtsndnAiJ7hdKqtbzC8Ux1dryge0Npc+LnjfjE6nMCfprqi0JOXirqw
Enx9LzE6JdXbb1liy6BEHKCgYEaju1SpnGvrg1EAYjPDNM0J5C31igmaxMh
6442hsnPBlwsB23A6g6YLyLxsLMs4bGTFU9wvTfNcCX+KdAQTxvaQ909pfWr
+ob1Su0l+rX0to+uEMZf2+D4DYtjhle60ZMZVsVklG5xhLw3uMVLjBhal2F1
7C35bZEcgYEAF8fLi1atxR729scH8e/cpJFFknRE1tRnt/Q1udVrHqfn/fPsPt
ndKC/tyc+ZnAwBx1czEzh1v5r+S49iYqls9j0H0p2Y1TF8GKEn0x1ATA5ky5k
u0Z6LaKTdSzGaRM1khbeBiCmLovaIipqh5ialzz8Fr0jxPGuVIP0BGhk=
-----END RSA PRIVATE KEY-----
[ansible@control1 ~]$ █
```